

Started on	Saturday, 4 May 2024, 1:06 PM
State	Finished
Completed on	Saturday, 4 May 2024, 5:58 PM
Time taken	4 hours 52 mins
Grade	Not yet graded

Question **1**

Correct

Mark 1.00 out of 1.00

Match the left side use(or non-use) of a synchronization primitive with the best option on the right side.

This is the smallest primitive made available in software, using the hardware provided atomic instructions

spinlock



This tool is quite attractive in solving the main bounded buffer problem

semaphore



This tool is very useful for waiting for 'something'

condition variables



This tool is useful for event-wait scenarios

semaphore



This tool is more useful on multiprocessor systems

spinlock



Your answer is correct.

The correct answer is: This is the smallest primitive made available in software, using the hardware provided atomic instructions → spinlock, This tool is quite attractive in solving the main bounded buffer problem → semaphore, This tool is very useful for waiting for 'something' → condition variables, This tool is useful for event-wait scenarios → semaphore, This tool is more useful on multiprocessor systems → spinlock

Question **2**

Partially correct

Mark 0.43 out of 1.00

Select T/F for statements about Volume Managers.

Do pay attention to the use of the words physical partition and physical volume.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	A physical partition(made available as physical volume) can belong to only one volume group	✓
<input type="radio"/>	<input checked="" type="radio"/>	A volume group consists of multiple physical partitions	✗
<input type="radio"/>	<input checked="" type="radio"/>	A volume group can contain physical partitions(as physical volumes) from only one physical disk	✓
<input checked="" type="radio"/>	<input type="radio"/>	Extending the size of a logical volume does not automatically change the size of the file system	✗
<input type="radio"/>	<input checked="" type="radio"/>	A logical volume can span across two volume groups	✓
<input type="radio"/>	<input checked="" type="radio"/>	A volume group consists of multiple physical partitions	✗
<input type="radio"/>	<input checked="" type="radio"/>	A logical volume can be extended in size but upto the size of the maximum sized physical volume only	✗

A physical partition(made available as physical volume) can belong to only one volume group: True

A volume group consists of multiple physical partitions: False

A volume group can contain physical partitions(as physical volumes) from only one physical disk: False

Extending the size of a logical volume does not automatically change the size of the file system: True

A logical volume can span across two volume groups: False

A volume group consists of multiple physical partitions: False

A logical volume can be extended in size but upto the size of the maximum sized physical volume only: False

Question 3

Correct

Mark 0.50 out of 0.50

Match the parts of the code, with their description

```
void *thread1(void *arg) {
    while(run == 1) {
        pthread_mutex_lock(&lock);
        c++;
        pthread_mutex_unlock(&lock);
        c1++;
    }
}
```

pthread_mutex_unlock(&lock);	exit section	✓
c1++;	reminder	✓
c++;	critical section	✓
pthread_mutex_lock(&lock);	entry section	✓

Your answer is correct.

The correct answer is: pthread_mutex_unlock(&lock); → exit section, c1++; → reminder, c++; → critical section, pthread_mutex_lock(&lock); → entry section

Question 4

Partially correct

Mark 0.78 out of 1.00

Match the description of a memory management function with the name of the function that provides it, in xv6

Load contents from ELF into pages after allocating the pages first	loaduvm()	✗
Copy the code pages of a process	No such function	✓
setup the kernel part in the page table	setupkvm()	✓
Switch to user page table	switchuvm()	✓
Mark the page as in-accessible	clearpteu()	✓
Create a copy of the page table of a process	copyuvm()	✓
Switch to kernel page table	switchkvm()	✓
Load contents from ELF into existing pages	No such function	✗
Setup and load the user page table for initcode process	inituvm()	✓

The correct answer is: Load contents from ELF into pages after allocating the pages first → No such function, Copy the code pages of a process → No such function, setup the kernel part in the page table → setupkvm(), Switch to user page table → switchuvm(), Mark the page as in-accessible → clearpteu(), Create a copy of the page table of a process → copyuvm(), Switch to kernel page table → switchkvm(), Load contents from ELF into existing pages → loaduvm(), Setup and load the user page table for initcode process → inituvm()

Question **5**

Correct

Mark 1.00 out of 1.00

For the reference string

4 2 5 1 0 1 2 5 4 1 2

the number of page faults, including initial ones,
with LRU replacement and 2 frames are :

Answer:



4 -

4 2

5 2

5 1

0 1

-

2 1

2 5

4 5

4 1

2 1

The correct answer is: 10

Question 6

Correct

Mark 1.00 out of 1.00

Choice of the global or local replacement strategy is a subjective choice for kernel programmers. There are advantages and disadvantages on either side. Out of the following statements, that advocate either global or local replacement strategy, select those statements that have a logically CONSISTENT argument. (That is any statement that is logically correct about either global or local replacement)

Consistent **Inconsistent**

<input checked="" type="radio"/>	<input type="radio"/>	Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Global replacement may give highly variable per process completion time because number of page faults become un-predictable.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).	✓
<input checked="" type="radio"/>	<input type="radio"/>	Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.	✓

Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.: Consistent

Global replacement may give highly variable per process completion time because number of page faults become un-predictable.: Consistent

Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.: Consistent

Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).: Consistent

Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.: Consistent

Question **7**

Correct

Mark 1.00 out of 1.00

Select all the actions taken by iget()

- ☐ a. Returns the inode with inode-cache lock held
- ☐ b. Panics if inode does not exist in cache
- ☒ c. Returns the inode with reference count incremented ✓
- ☐ d. Returns a valid inode if not found in cache
- ☐ e. Returns the inode locked
- ☒ f. Returns an inode with given dev+inode-number from cache, if it exists in cache ✓
- ☒ g. Returns a free-inode , with dev+inode-number set, if not found in cache ✓

Your answer is correct.

The correct answers are: Returns an inode with given dev+inode-number from cache, if it exists in cache, Returns the inode with reference count incremented, Returns a free-inode , with dev+inode-number set, if not found in cache

Question 8

Partially correct

Mark 0.86 out of 1.00

Select T/F w.r.t physical disk handling in xv6 code

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	only direct blocks are supported	✗
<input type="radio"/>	<input checked="" type="radio"/>	device files are not supported	✓
<input checked="" type="radio"/>	<input type="radio"/>	The code supports IDE, and not SATA/SCSI	✓
<input checked="" type="radio"/>	<input type="radio"/>	disk driver handles only one buffer at a time	✓
<input checked="" type="radio"/>	<input type="radio"/>	only 2 disks are handled by default	✓
<input checked="" type="radio"/>	<input type="radio"/>	log is kept on the same device as the file system	✓
<input checked="" type="radio"/>	<input type="radio"/>	the superblock does not contain number of free blocks	✓

only direct blocks are supported: False

device files are not supported: False

The code supports IDE, and not SATA/SCSI: True

disk driver handles only one buffer at a time: True

only 2 disks are handled by default: True

log is kept on the same device as the file system: True

the superblock does not contain number of free blocks: True

Question 9

Correct

Mark 1.00 out of 1.00

Consider the following list of free chunks, in continuous memory management:

8k, 1k, 2k, 14k, 13k, 7k, 6k

Suppose there is a request for chunk of size 5k, then the free chunk selected under each of the following schemes will be

Best fit: 6k ✓

First fit: 8k ✓

Worst fit: 14k ✓

Question 10

Correct

Mark 1.00 out of 1.00

Select T/F for the disk block allocation scheme related statements

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	FAT uses linked allocation	✓
<input checked="" type="radio"/>	<input type="radio"/>	NVM storage devices are pushing for search for new block allocation schemes	✓
<input checked="" type="radio"/>	<input type="radio"/>	Continuous allocation leads to faster file access	✓
<input checked="" type="radio"/>	<input type="radio"/>	Continuous allocation allows to fetch any block with just one seek	✓
<input checked="" type="radio"/>	<input type="radio"/>	Continuous allocation may involve a costly search for free space	✓
<input checked="" type="radio"/>	<input type="radio"/>	The unix inode is based on indexed allocation	✓
<input checked="" type="radio"/>	<input type="radio"/>	Maximum file size limit is determined by disk block allocation scheme, as one of the factors.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Linked allocation does away with file size limitation to a large extent	✓

FAT uses linked allocation: True

NVM storage devices are pushing for search for new block allocation schemes: True

Continuous allocation leads to faster file access: True

Continuous allocation allows to fetch any block with just one seek: True

Continuous allocation may involve a costly search for free space: True

The unix inode is based on indexed allocation: True

Maximum file size limit is determined by disk block allocation scheme, as one of the factors.: True

Linked allocation does away with file size limitation to a large extent: True

Question 11

Correct

Mark 2.00 out of 2.00

Consider a demand-paging system with the following time-measured utilizations:

CPU utilization : 20%

Paging disk: 97.7%

Other I/O devices: 5%

For each of the following, indicate whether it will (or is likely to) improve CPU utilization (even if by a small amount). Explain your answers.

- a. Install a faster CPU : ✓
- b. Install a bigger paging disk. : ✓
- c. Increase the degree of multiprogramming. : ✓
- d. Decrease the degree of multiprogramming. : ✓
- e. Install more main memory.: ✓
- f. Install a faster hard disk or multiple controllers with multiple hard disks. : ✓
- g. Add prepaging to the page-fetch algorithms. : ✓
- h. Increase the page size. : ✓

Question 12

Complete

Marked out of 2.00

Which changes should be made to xv6 code to implement the vfork() system call ?

List your suggestions in a bullet-point fashion. Each point should mention

- (a) pseudo-code of new function to be added
- (b) prototype of any new function or new system call to be added
- (c) pseudo-code of changes to an existing function, describing lines to be removed, and lines to be added
- (d) precise declaration of new data structures to be added in C, or changes to the existing data structure
- (e) Name and a one-line description of new userland functionality to be added
- (f) Changes to Makefile
- (g) Any other change in a maximum of 20 words per change.

```
void  
vfork()
```

Question **13**

Correct

Mark 1.00 out of 1.00

For each function/code-point, select the status of segmentation setup in xv6

bootasm.S	<div>gdt setup with 3 entries, at start32 symbol of bootasm.S</div>	✓
kvmalloc() in main()	<div>gdt setup with 3 entries, at start32 symbol of bootasm.S</div>	✓
bootmain()	<div>gdt setup with 3 entries, at start32 symbol of bootasm.S</div>	✓
after startothers() in main()	<div>gdt setup with 5 entries (0 to 4) on all processors</div>	✓
entry.S	<div>gdt setup with 3 entries, at start32 symbol of bootasm.S</div>	✓
after seginit() in main()	<div>gdt setup with 5 entries (0 to 4) on one processor</div>	✓

Your answer is correct.

The correct answer is: bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor

Question 14

Correct

Mark 0.50 out of 0.50

You must have seen the error message "Segmentation fault, core dumped" very often.

With respect to this error message, mark the statements as True/False.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	On Linux, the message is printed only because the memory management scheme is segmentation	✓ No, it's just a term used, even if paging is used for memory management.
<input checked="" type="radio"/>	<input type="radio"/>	The process has definitely performed illegal memory access.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The core file can be analysed later using a debugger, to determine what went wrong.	✓ use gdb ./core ./executable-filename
<input type="radio"/>	<input checked="" type="radio"/>	The term "core" refers to the core code of the kernel.	✓ core means memory, all memory for the process.
<input checked="" type="radio"/>	<input type="radio"/>	The image of the process is stored in a file called "core", if the ulimit allows so.	✓ see ulimit -a
<input type="radio"/>	<input checked="" type="radio"/>	The illegal memory access was detected by the kernel and the process was punished by kernel.	✓ "detection" is done by CPU, not kernel.
<input checked="" type="radio"/>	<input type="radio"/>	On Linux, the process was sent a SIGSEGV signal and the default handler for the signal is "Term", so the process is terminated.	✓

On Linux, the message is printed only because the memory management scheme is segmentation: False

The process has definitely performed illegal memory access.: True

The core file can be analysed later using a debugger, to determine what went wrong.: True

The term "core" refers to the core code of the kernel.: False

The image of the process is stored in a file called "core", if the ulimit allows so.: True

The illegal memory access was detected by the kernel and the process was punished by kernel.: False

On Linux, the process was sent a SIGSEGV signal and the default handler for the signal is "Term", so the process is terminated.: True

Question 15

Correct

Mark 0.50 out of 0.50

exec() does this: `curproc->tf->eip = elf.entry`, but `userinit()` does this: `p->tf->eip = 0`; Select all the statements from below, that collectively explain this

- ☒ a. `exec()` loads from ELF file and the address of first instruction to be executed is given by 'entry' ✓
- ☐ b. the 'entry' in `initcode` is anyways 0
- ☐ c. the code of 'initcode' is loaded at physical address 0
- ☒ d. In `userinit()` the function `inituvm()` has mapped the code of 'initcode' to be starting at virtual address 0 ✓
- ☐ e. `elf.entry` is anyways 0, so both statements mean the same
- ☒ f. the `initcode` is created using `objcopy`, which discards all relocation information and symbols (like `entry`) ✓

The correct answers are: `exec()` loads from ELF file and the address of first instruction to be executed is given by 'entry', In `userinit()` the function `inituvm()` has mapped the code of 'initcode' to be starting at virtual address 0, the `initcode` is created using `objcopy`, which discards all relocation information and symbols (like `entry`)

Question 16

Correct

Mark 0.50 out of 0.50

Select the compiler's view of the process's address space, for each of the following MMU schemes:
(Assume that each scheme, e.g. paging/segmentation/etc is effectively utilised)

Relocation + Limit	one continuous chunk	✓
Segmentation	many continuous chunks of variable size	✓
Segmentation, then paging	many continuous chunks of variable size	✓
Paging	one continuous chunk	✓

Your answer is correct.

The correct answer is: Relocation + Limit → one continuous chunk, Segmentation → many continuous chunks of variable size, Segmentation, then paging → many continuous chunks of variable size, Paging → one continuous chunk

Question **17**

Correct

Mark 1.00 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- ☐ a. TLB hit ration has zero impact in effective memory access time in demand paging.
- ☒ b. The meaning of valid-invalid bit in page table is different in paging and demand-paging. ✓
- ☒ c. Demand paging always increases effective memory access time. ✓
- ☒ d. Demand paging requires additional hardware support, compared to paging. ✓
- ☒ e. Paging requires some hardware support in CPU ✓
- ☒ f. Calculations of number of bits for page number and offset are same in paging and demand paging. ✓
- ☐ g. With paging, it's possible to have user programs bigger than physical memory.
- ☒ h. Both demand paging and paging support shared memory pages. ✓
- ☒ i. With demand paging, it's possible to have user programs bigger than physical memory. ✓
- ☐ j. Paging requires NO hardware support in CPU

Your answer is correct.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question **18**

Correct

Mark 0.50 out of 0.50

Which of the following statements are correct about xv6 and multiprocessor support ?

- ☒ a. At any point in time, after main(), the kernel may be parallely executing on any of the processors. ✓
- ☒ b. Each processor on xv6 starts code execution in mpenter() when first processor sets "started" to 1. ✓
- ☒ c. In xv6 on x86, the first processor configures other processors in mpinit() ✓
- ☒ d. xv6 supports only SMP ✓
- ☒ e. xv6 supports a variable number of processors (upto 8) and adjusts it's data structures according to number of processors ✓

The correct answers are: xv6 supports only SMP, xv6 supports a variable number of processors (upto 8) and adjusts it's data structures according to number of processors, Each processor on xv6 starts code execution in mpenter() when first processor sets "started" to 1., In xv6 on x86, the first processor configures other processors in mpinit(), At any point in time, after main(), the kernel may be parallely executing on any of the processors.

Question **19**

Correct

Mark 0.50 out of 0.50

Rank the following storage systems from slowest (first) to fastest(last)

You can drag and drop the items below/above each other.

✓ Magnetic tapes

✓ Optical disk

✓ Hard-disk drives

✓ Nonvolatile memory

✓ Main memory

✓ Cache

✓ Registers

Your answer is correct.

The given semaphore implementation faces which problem?

Note: blocks means waits in a wait queue.

```
struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->sl));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <= 0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

signal(semaphore *s) {
    spinlock(&(s->sl));
    (s->val)++;
    x = dequeue(s->sl) and enqueue(readyq, x);
    spinunlock(&(s->sl));
}
```

- ☐ a. blocks holding a spinlock
- ☐ b. too much spinning, bounded wait not guaranteed
- ☐ c. deadlock
- ☒ d. not holding lock after unblock ✓

Your answer is correct.

The correct answer is: not holding lock after unblock

Question **21**

Correct

Mark 1.00 out of 1.00

Select the proper order of execution enforced by the given semaphore code

P1, P2, P3 are processes and S1, S2 are binary semaphores.

```
S1 = 0; S2 = 0;
```

P1:

```
Wait (S2);
```

```
Statement2;
```

P2:

```
Statement1;
```

```
Signal (S1);
```


P3:

```
Wait (S1)
```

```
Statement S3;
```

```
Signal (S2);
```

Answer:



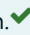
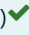
P2, P3, P1 

Question **22**

Correct

Mark 1.00 out of 1.00

Select all the correct statements, w.r.t. Copy on Write

- ☒ a. If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated 
- ☐ b. use of COW during fork() is useless if child called exit()
- ☒ c. COW helps us save memory 
- ☒ d. Fork() used COW technique to improve performance of new process creation. 
- ☒ e. Vfork() assumes that there will be no write, but rather exec() 
- ☐ f. use of COW during fork() is useless if exec() is called by the child

The correct answers are: Fork() used COW technique to improve performance of new process creation., If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated, COW helps us save memory, Vfork() assumes that there will be no write, but rather exec()

Question **23**

Correct

Mark 1.00 out of 1.00

Calculate the average waiting time using
FCFS scheduling
for the following workload

assuming that they arrive in this order during the first time unit:

Process Burst Time

P1	2
P2	6
P3	2
P4	3

Write only a number in the answer upto two decimal points.

Answer: ✓

P2 waits for 2 units

P3 waits for 2+6 units

P4 waits for 2 + 6 + 2 units of time

Total waiting = 2 + 2 + 6 + 2 + 6 + 2 = 20 units

Average waiting time = 20/4 = 5

The correct answer is: 5

Question **24**

Correct

Mark 1.00 out of 1.00

Shared memory is possible with which of the following memory management schemes ?

Select one or more:

- ☒ a. demand paging ✓
- ☒ b. paging ✓
- ☐ c. continuous memory management
- ☒ d. segmentation ✓

Your answer is correct.

The correct answers are: paging, segmentation, demand paging

In the code below assume that each function can be executed concurrently by many threads/processes. Ignore syntactical issues, and focus on the semantics.

This program is an example of

```
spinlock a, b; // assume initialized
thread1() {
    spinlock(b);
    //some code;
    spinlock(a);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
thread2() {
    spinlock(b);
    //some code;
    spinlock(a);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
```

- ☐ a. Self Deadlock
- ☐ b. Livelock
- ☐ c. Deadlock or livelock depending on actual race
- ☒ d. None of these ✓
- ☐ e. Deadlock

Your answer is correct.

The correct answer is: None of these

Question **26**

Correct

Mark 1.00 out of 1.00

For the reference string
3 4 3 5 2

using FIFO replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.
Select the correct statement.

Select one:

- ☒ a. Do not exhibit Balady's anomaly ✓
- ☐ b. Exhibit Balady's anomaly between 2 and 3 frames
- ☐ c. Exhibit Balady's anomaly between 3 and 4 frames

Your answer is correct.

The correct answer is: Do not exhibit Balady's anomaly

Question **27**

Correct

Mark 1.00 out of 1.00

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

11010010

Now, there is a request for a chunk of 45 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer:



The correct answer is: 11011110

Question **28**

Correct

Mark 0.50 out of 0.50

What does seginit() do?

- ☒ a. Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 3 ✓
- ☐ b. Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 0
- ☐ c. Nothing significant, just repetition of earlier GDT setup but with free frames list created now
- ☐ d. Nothing significant, just repetition of earlier GDT setup but with kernel page table allocated now
- ☐ e. Nothing significant, just repetition of earlier GDT setup but with 2-level paging setup done

The correct answer is: Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 3

Question **29**

Complete

Marked out of 2.00

Write all changes required to xv6 to add a buddy allocator.

Every change should be mentioned in terms of either of the following:

- (a) pseudo-code of new function to be added
- (b) prototype of any new function or new system call to be added
- (c) pseudo-code of changes to an existing function, describing lines to be removed, and lines to be added
- (d) **precise** declaration of new data structures to be added in C, or changes to the existing data structure
- (e) Name and a one-line description of new userland functionality to be added
- (f) Changes to Makefile
- (g) Any other change in a maximum of 20 words per change.

```
//Buddy structure
struct buddy
{ void *base;
  size_t size;
  struct buddy *next;
}

//free list array
struct buddy *freelist[MAX_ORDER];

//new functions
init_allocator(void *memory, size_t size)
void *allocator(size_t size)
void free(void *ptr)
void split_block(struct buddy *block, size_t order)
struct buddy *merge_blocks(struct buddy *block1, struct buddy *block2)

other changes would be to update memory management routines in xv6 to use buddy allocator functions for allocation and deallocation
```

Question **30**

Correct

Mark 1.00 out of 1.00

Select all the blocks that may need to be written back to disk (if updated, of-course), as "Yes", when an operation of deleting a file is carried out on ext2 file system.

An option has to be correct entirely to be marked "Yes"

Data blocks of the file	<div>No</div>	✓
Block bitmap(s) for all the blocks of the file	<div>Yes</div>	✓
Superblock	<div>Yes</div>	✓
One or more data bitmap blocks for the parent directory	<div>No</div>	✓
One or multiple data blocks of the parent directory	<div>No</div>	✓
Possibly one block bitmap corresponding to the parent directory	<div>Yes</div>	✓

Your answer is correct.

The correct answer is: Data blocks of the file → No, Block bitmap(s) for all the blocks of the file → Yes, Superblock → Yes, One or more data bitmap blocks for the parent directory → No, One or multiple data blocks of the parent directory → No, Possibly one block bitmap corresponding to the parent directory → Yes

Question **31**

Correct

Mark 1.00 out of 1.00

Map the function in xv6's file system code, to it's perceived logical layer.

dirlookup	directory	✓
stati	inode	✓
filestat()	file descriptor	✓
bmap	inode	✓
commit	logging	✓
bread	buffer cache	✓
ideintr	disk driver	✓
sys_chdir()	system call	✓
balloc	block allocation on disk	✓
ialloc	inode	✓
skipelem	pathname lookup	✓
namei	pathname lookup	✓

Your answer is correct.

The correct answer is: dirlookup → directory, stati → inode, filestat() → file descriptor, bmap → inode, commit → logging, bread → buffer cache, ideintr → disk driver, sys_chdir() → system call, balloc → block allocation on disk, ialloc → inode, skipelem → pathname lookup, namei → pathname lookup

Question **32**

Correct

Mark 0.50 out of 0.50

Assuming a 8- KB page size, what is the page numbers for the address 1093943 reference in decimal :
(give answer also in decimal)

Answer: ✓

The correct answer is: 134

Question **33**

Correct

Mark 1.00 out of 1.00

Mark the statements about named and un-named pipes as True or False

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	Named pipes can be used for communication between only "related" processes.	
<input checked="" type="radio"/>	<input type="radio"/>	Both types of pipes are an extension of the idea of "message passing".	
<input checked="" type="radio"/>	<input type="radio"/>	Both types of pipes provide FIFO communication.	
<input type="radio"/>	<input checked="" type="radio"/>	The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.	
<input checked="" type="radio"/>	<input type="radio"/>	Named pipe exists as a file	
<input checked="" type="radio"/>	<input type="radio"/>	Named pipes can exist beyond the life-time of processes using them.	
<input type="radio"/>	<input checked="" type="radio"/>	The pipe() system call can be used to create either a named or un-named pipe.	
<input checked="" type="radio"/>	<input type="radio"/>	Un-named pipes are inherited by a child process from parent.	
<input checked="" type="radio"/>	<input type="radio"/>	Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.	
<input type="radio"/>	<input checked="" type="radio"/>	A named pipe has a name decided by the kernel.	

Named pipes can be used for communication between only "related" processes.: False

Both types of pipes are an extension of the idea of "message passing".: True

Both types of pipes provide FIFO communication.: True

The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.: False

Named pipe exists as a file: True

Named pipes can exist beyond the life-time of processes using them.: True

The pipe() system call can be used to create either a named or un-named pipe.: False

Un-named pipes are inherited by a child process from parent.: True

Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.: True

A named pipe has a name decided by the kernel.: False

Question 34

Correct

Mark 2.00 out of 2.00

Suppose it is required to add the `chown()` (without notion of a group, just the notion of owner and others) system call to xv6. Select the changes, from the options given below, which are an absolute must to effect the addition of this system call.

Changes w.r.t. other system calls should be selected if those system calls are necessary for the implementation of the `chown()` system call.

Must	Not-Must		
<input type="radio"/>	<input checked="" type="radio"/>	Create an application program <code>su.c</code> which itself is a SUID program, and calls <code>setuid()</code> and <code>setuid()</code> with specified user-ID and then does <code>exec()</code> of a shell.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Add a uid field representing the owner, inside dinode	✓
<input type="radio"/>	<input checked="" type="radio"/>	Add few extra files belonging to different users, using <code>mkfs.c</code>	✓
<input checked="" type="radio"/>	<input type="radio"/>	Add the <code>chown()</code> system call which checks if the user with id "0" is calling this system call and then change the ownership of the on-disk and in-memory inode.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Add a mode field representing file permissions, inside dinode	✓
<input type="radio"/>	<input checked="" type="radio"/>	Add the username field to the on disk inode, that is dinode.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Add a <code>setuid()</code> , <code>seteuid()</code> system call, callable only by the user with ID or EUID equal to "0" to set the new user id of the process	✓
<input type="radio"/>	<input checked="" type="radio"/>	Mandatorily create a file like "passwd" which maps user-names to user-IDs and modify other utilities (like "ls") to show user-name instead of user-ID for the files.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Inherit the UID and EUID of the parent in <code>fork()</code>	✓
<input type="radio"/>	<input checked="" type="radio"/>	Modify <code>exec()</code> to check SUID bit on the executable file and set EUID of the process to UID of the executable	✓
<input type="radio"/>	<input checked="" type="radio"/>	Add a EUID field in the struct <code>proc</code> representing Effective user ID	✓
<input type="radio"/>	<input checked="" type="radio"/>	Add a UID field in the struct <code>proc</code> representing the USER-ID of the user who started this process	✓
<input checked="" type="radio"/>	<input type="radio"/>	Modify <code>mkfs.c</code> to add owner and permissions to each file created	✓

Create an application program su.c which itself is a SUID program, and calls setuid() and setuid() with specified user-ID and then does exec() of a shell.: Not-Must

Add a uid field representing the owner, inside dinode: Must

Add few extra files belonging to different users, using mkfs.c: Not-Must

Add the chown() system call which checks if the user with id "0" is calling this system call and then change the ownership of the on-disk and in-memory inode.: Must

Add a mode field representing file permissions, inside dinode: Must

Add the username field to the on disk inode, that is dinode.: Not-Must

Add a setuid(), seteuid() system call, callable only by the user with ID or EUID equal to "0" to set the new user id of the process: Not-Must

Mandatorily create a file like "passwd" which maps user-names to user-IDs and modify other utilities (like "ls") to show user-name instead of user-ID for the files.: Not-Must

Inherit the UID and EUID of the parent in fork(): Not-Must

Modify exec() to check SUID bit on the executable file and set EUID of the process to UID of the executable: Not-Must

Add a EUID field in the struct proc representing Effective user ID: Not-Must

Add a UID field in the struct proc representing the USER-ID of the user who started this process: Not-Must

Modify mkfs.c to add owner and permissions to each file created: Must

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
struct proc*  
myproc(void) {  
...  
    pushcli();  
    c = mycpu();  
    p = c->proc;  
    popcli();  
...  
}
```

Disable interrupts to avoid another process's pointer being returned



```
void  
yield(void)  
{  
...  
    release(&ptable.lock);  
}
```

Release the lock held by some another process



```
void  
acquire(struct spinlock *lk)  
{  
...  
    __sync_synchronize();  
}
```

Tell compiler not to reorder memory access beyond this line



```
void  
acquire(struct spinlock *lk)  
{  
    pushcli();  
...  
}
```

Disable interrupts to avoid deadlocks



```
void  
acquire(struct spinlock *lk)  
{  
...  
    getcallerpcs(&lk, lk->pcs);  
}
```

Traverse ebp chain to get sequence of instructions followed in functions calls



```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

    // The + in "+m" denotes a read-modify-write
    operand.
    asm volatile("lock; xchgl %0, %1" :
        "+m" (*addr), "=a" (result) :
        "1" (newval) :
        "cc");
    return result;
}
```

Atomic compare and swap instruction (to be expanded inline into code)



```
void
panic(char *s)
{
    ...
    panicked = 1;
```

Ensure that no printing happens on other processors



```
void
sleep(void *chan, struct spinlock *lk)
{
    ...
    if(lk != &ptable.lock){
        acquire(&ptable.lock);
        release(lk);
    }
```

Avoid a self-deadlock



Your answer is correct.

The correct answer is: `struct proc*`

```
myproc(void) {
    ...
    pushcli();
    c = mycpu();
    p = c->proc;
    popcli();
    ...
}
```

→ Disable interrupts to avoid another process's pointer being returned, `void`
`yield(void)`

```
{
    ...
    release(&ptable.lock);
}
```

→ Release the lock held by some another process, `void`

```
acquire(struct spinlock *lk)
{
    ...
    __sync_synchronize();
```

→ Tell compiler not to reorder memory access beyond this line, `void`

```
acquire(struct spinlock *lk)
{
    pushcli();
```

→ Disable interrupts to avoid deadlocks, `void`

```
acquire(struct spinlock *lk)
{
...
getcallerpcs(&lk, lk->pcs);
```

→ Traverse ebp chain to get sequence of instructions followed in functions calls, `static inline uint`

```
xchg(volatile uint *addr, uint newval)
```

```
{
    uint result;
```

// The + in "+m" denotes a read-modify-write operand.

```
asm volatile("lock; xchgl %0, %1" :
```

```
    "+m" (*addr), "=a" (result) :
```

```
    "1" (newval) :
```

```
    "cc");
```

```
return result;
```

```
} → Atomic compare and swap instruction (to be expanded inline into code), void
```

```
panic(char *s)
```

```
{
```

```
...
```

panicked = 1; → Ensure that no printing happens on other processors, `void`

```
sleep(void *chan, struct spinlock *lk)
```

```
{
```

```
...
```

```
if(lk != &ptable.lock){
```

```
    acquire(&ptable.lock);
```

```
    release(lk);
```

```
} → Avoid a self-deadlock
```

Question 36

Partially correct

Mark 1.80 out of 2.00

Following code claims to implement the command

```
/bin/ls -l | /usr/bin/head -3 | /usr/bin/tail -1
```

Fill in the blanks to make the code work.

Note: Do not include space in writing any option. x[1][2] should be written without any space, and so is the case with [1] or [2]. Pay attention to exact syntax and do not write any extra character like ';' or = etc.

```
int main(int argc, char *argv[]) {
    int pid1, pid2;
    int pfd[ 2 ][2];

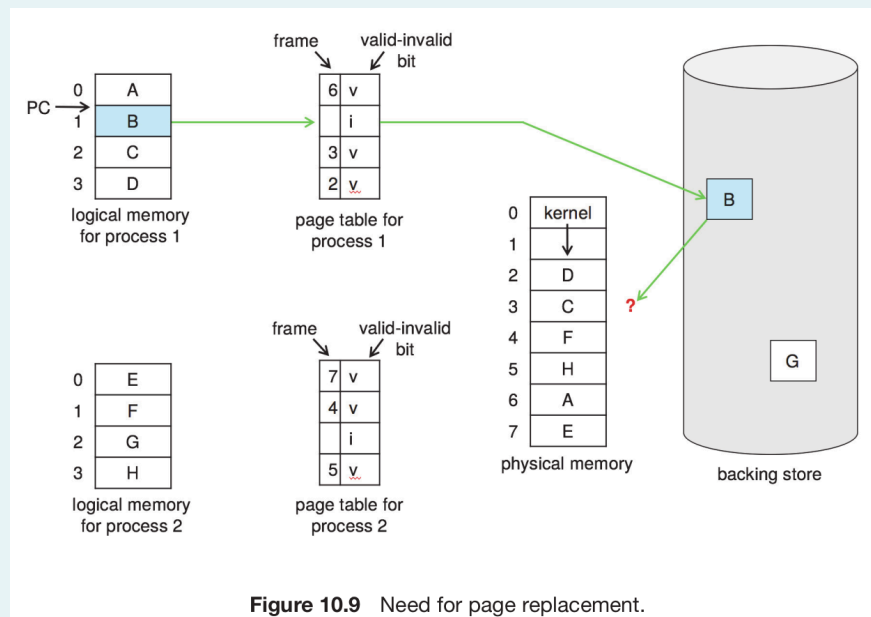
    pipe( pfd[0] );
    pid1 = fork() ;
    if(pid1 != 0) {
        close(pfd[0][ 0] );
        close( 1 );
        dup( pfd[0][1] );
        execl("/bin/ls", "/bin/ls", "-1 ", NULL);
    }
    pipe( pfd[1] );
    pid2 = fork();
    if(pid2 == 0) {
        close( pfd[1][1] );
        close(0);
        dup( pfd[0][0] );
        close(pfd[1][ 0] );
        close( 1 );
        dup( pfd[1][1] );
        execl("/usr/bin/head", "/usr/bin/head", "-3 ", NULL);
    } else {
        close(pfd[ 1][1] );
        close( 0 );
        dup( pfd[1][0] );
        close(pfd[ 0][0] );
        execl("/usr/bin/tail", "/usr/bin/tail", "-1 ", NULL);
    }
}
```

Question 37

Correct

Mark 1.00 out of 1.00

W.r.t the figure given below, mark the given statements as True or False.



True False

<input type="radio"/>	<input checked="" type="radio"/>	Global replacement means chose any of the frame from 0 to 7	✓
<input checked="" type="radio"/>	<input type="radio"/>	Kernel occupies two page frames	✓
<input checked="" type="radio"/>	<input type="radio"/>	Handling this scenario demands two disk I/Os	✓
<input checked="" type="radio"/>	<input type="radio"/>	The kernel's pages can not used for replacement if kernel is not pageable.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Page 1 of process 1 needs a replacement	✓
<input checked="" type="radio"/>	<input type="radio"/>	Local replacement means chose any of the frames 2, 3, 6	✓
<input type="radio"/>	<input checked="" type="radio"/>	Local replacement means chose any of the frame from 2 to 7	✓
<input checked="" type="radio"/>	<input type="radio"/>	Global replacement means chose any of the frame from 2 to 7	✓

Global replacement means chose any of the frame from 0 to 7: False

Kernel occupies two page frames: True

Handling this scenario demands two disk I/Os: True

The kernel's pages can not used for replacement if kernel is not pageable.: True

Page 1 of process 1 needs a replacement: True

Local replacement means chose any of the frames 2, 3, 6: True

Local replacement means chose any of the frame from 2 to 7: False

Global replacement means chose any of the frame from 2 to 7: True

Question **38**

Correct

Mark 0.50 out of 0.50

The approximate number of page frames created by kinit1 is

- ☐ a. 16
- ☐ b. 4
- ☐ c. 10
- ☒ d. 3000 ✓
- ☐ e. 1000
- ☐ f. 4000
- ☐ g. 2000

The correct answer is: 3000

Question **39**

Correct

Mark 1.00 out of 1.00

In xv6, The struct context is given as

```
struct context {  
    uint edi;  
    uint esi;  
    uint ebx;  
    uint ebp;  
    uint eip;  
};
```

Select all the reasons that explain why only these 5 registers are included in the struct context.

- ☒ a. esp is not saved in context, because context{} is on stack and it's address is always argument to switch() ✓
- ☐ b. esp is not saved in context, because it's not part of the context
- ☒ c. The segment registers are same across all contexts, hence they need not be saved ✓
- ☒ d. eax, ecx, edx are caller save, hence no need to save ✓
- ☐ e. xv6 tries to minimize the size of context to save memory space

Your answer is correct.

The correct answers are: The segment registers are same across all contexts, hence they need not be saved, eax, ecx, edx are caller save, hence no need to save, esp is not saved in context, because context{} is on stack and it's address is always argument to switch()

Question **40**

Partially correct

Mark 0.17 out of 1.00

Select correct statements about mounting

Select one or more:

- ☐ a. The existing name-space at the mount-point is no longer visible after mounting
- ☒ b. The mount point must be a directory ✓
- ☐ c. Mounting deletes all data at the mount-point
- ☐ d. Mounting is attaching a disk-partition with a filesystem on it, into another file system name-space
- ☐ e. The mount point can be a file as well
- ☐ f. On Linuxes mounting can be done only while booting the OS
- ☐ g. Even in operating systems with a pluggable kernel module for file systems, the code for mounting any particular file system must be already present in the operating system system kernel
- ☐ h. Mounting makes all disk partitions available as one name space
- ☐ i. In operating systems with a pluggable kernel module for file systems, the code for mounting a particular file system is provided by the module of that file system.
- ☐ j. It's possible to mount a partition on one computer, into namespace of another computer.

Your answer is partially correct.

You have correctly selected 1.

The correct answers are: Mounting is attaching a disk-partition with a filesystem on it, into another file system name-space, The mount point must be a directory, The existing name-space at the mount-point is no longer visible after mounting, Mounting makes all disk partitions available as one name space, In operating systems with a pluggable kernel module for file systems, the code for mounting a particular file system is provided by the module of that file system., It's possible to mount a partition on one computer, into namespace of another computer.

Question **41**

Correct

Mark 0.50 out of 0.50

Note: for this question you get full marks if you select all and only correct options, you get ZERO if at least one option is wrong or not selected.

Select all the correct statements about log structured file systems.

- ☒ a. log may be kept on same block device or another block device ✓
- ☒ b. file system recovery may end up losing data ✓
- ☒ c. even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery ✓
- ☐ d. file system recovery recovers all the lost data
- ☐ e. a transaction is said to be committed when all operations are written to file system

Your answer is correct.

The correct answers are: file system recovery may end up losing data, log may be kept on same block device or another block device, even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery

Question **42**

Correct

Mark 0.50 out of 0.50

Match the pair

Hashed page table	Linear search on collision done by OS (e.g. SPARC Solaris) typically	✓
Inverted Page table	Linear/Parallel search using page number in page table	✓
Hierarchical Paging	More memory access time per hierarchy	✓

Your answer is correct.

The correct answer is: Hashed page table → Linear search on collision done by OS (e.g. SPARC Solaris) typically, Inverted Page table → Linear/Parallel search using page number in page table, Hierarchical Paging → More memory access time per hierarchy

Question **43**

Partially correct

Mark 0.40 out of 0.50

Select all the complications added due to concurrent processing in operating systems.

Select one or more:

- ☒ a. Context of each process needs to be restored when it's scheduled ✓
- ☒ b. It requires a timer interrupt in hardware ✓
- ☒ c. Context of each process needs to be saved when it's interrupted by timer ✓
- ☐ d. OS needs to decide whether it's own code is interruptible or not
- ☒ e. Scheduling becomes more time consuming due to increased number of processes ✓

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: It requires a timer interrupt in hardware, Scheduling becomes more time consuming due to increased number of processes, OS needs to decide whether it's own code is interruptible or not, Context of each process needs to be saved when it's interrupted by timer, Context of each process needs to be restored when it's scheduled

Question 44

Correct

Mark 1.00 out of 1.00

Mark statements as T/F

All statements are in the context of preventing deadlocks.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	If a resource allocation graph contains a cycle then there is a guarantee of a deadlock	✓
<input checked="" type="radio"/>	<input type="radio"/>	Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens	✓
<input type="radio"/>	<input checked="" type="radio"/>	The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order	✓
<input checked="" type="radio"/>	<input type="radio"/>	Circular wait is avoided by enforcing a lock ordering	✓
<input checked="" type="radio"/>	<input type="radio"/>	If a resource allocation graph contains a cycle then there is a possibility of a deadlock	✓
<input checked="" type="radio"/>	<input type="radio"/>	Deadlock is not possible if any of these conditions is not met: Mutual exclusion, hold and wait, no pre-emption, circular wait.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The lock ordering to be followed to avoid circular wait is a protocol to be followed by programmers.	✓

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

Circular wait is avoided by enforcing a lock ordering: True

If a resource allocation graph contains a cycle then there is a possibility of a deadlock: True

Deadlock is not possible if any of these conditions is not met: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

The lock ordering to be followed to avoid circular wait is a protocol to be followed by programmers.: True

Question **45**

Correct

Mark 1.00 out of 1.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
void  
panic(char *s)  
{
```

```
...  
    panicked = 1;
```

Ensure that no printing happens on other processors



```
void  
acquire(struct spinlock *lk)  
{
```

```
...  
    getcallerpcs(&lk, lk->pcs);
```

Traverse ebp chain to get sequence of instructions followed in functions calls



```
void  
yield(void)  
{  
...  
    release(&ptable.lock);  
}
```

Release the lock held by some another process



Your answer is correct.

The correct answer is: void

panic(char *s)

```
{
```

```
...  
    panicked = 1; → Ensure that no printing happens on other processors, void
```

acquire(struct spinlock *lk)

```
{
```

```
...  
    getcallerpcs(&lk, lk->pcs); → Traverse ebp chain to get sequence of instructions followed in functions calls, void
```

yield(void)

```
{
```

```
...  
    release(&ptable.lock);
```

} → Release the lock held by some another process

Question **46**

Correct

Mark 1.00 out of 1.00

Mark the statements as True or False, w.r.t. mmap()

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	mmap() results in changes to page table of a process.	✓
<input type="radio"/>	<input checked="" type="radio"/>	mmap() results in changes to buffer-cache of the kernel.	✓
<input checked="" type="radio"/>	<input type="radio"/>	mmap() is a system call	✓
<input checked="" type="radio"/>	<input type="radio"/>	mmap() can be implemented on both demand paged and non-demand paged systems.	✓
<input type="radio"/>	<input checked="" type="radio"/>	MAP_FIXED guarantees that the mapping is always done at the specified address	✓
<input type="radio"/>	<input checked="" type="radio"/>	on failure mmap() returns NULL	✓
<input type="radio"/>	<input checked="" type="radio"/>	MAP_SHARED leads to a mapping that is copy-on-write	✓
<input checked="" type="radio"/>	<input type="radio"/>	on failure mmap() returns (void *)-1	✓
<input checked="" type="radio"/>	<input type="radio"/>	MAP_PRIVATE leads to a mapping that is copy-on-write	✓

mmap() results in changes to page table of a process.: True
mmap() results in changes to buffer-cache of the kernel.: False
mmap() is a system call: True
mmap() can be implemented on both demand paged and non-demand paged systems.: True
MAP_FIXED guarantees that the mapping is always done at the specified address: False
on failure mmap() returns NULL: False
MAP_SHARED leads to a mapping that is copy-on-write: False
on failure mmap() returns (void *)-1: True
MAP_PRIVATE leads to a mapping that is copy-on-write: True

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

\$ ls ./tmp/asdfksdf >/tmp/ddd 2>&1

Program 1

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    dup(fd);
    close(2);
    dup(fd);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Program 2

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    close(1);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(2);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Select all the correct statements about the programs

Select one or more:

- ☐ a. Both program 1 and 2 are incorrect
- ☐ b. Program 1 ensures 2>&1 and does not ensure > /tmp/ddd
- ☐ c. Program 2 makes sure that there is one file offset used for '2' and '1'
- ☒ d. Program 1 makes sure that there is one file offset used for '2' and '1' ✓
- ☐ e. Program 2 is correct for > /tmp/ddd but not for 2>&1
- ☒ f. Only Program 1 is correct ✓
- ☐ g. Program 1 does 1>&2
- ☐ h. Both programs are correct
- ☐ i. Program 2 ensures 2>&1 and does not ensure > /tmp/ddd
- ☐ j. Program 1 is correct for > /tmp/ddd but not for 2>&1
- ☐ k. Only Program 2 is correct
- ☐ l. Program 2 does 1>&2

Your answer is correct.

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'

Question **48**

Correct

Mark 1.00 out of 1.00

Given below is a sequence of reference bits on pages before the second chance algorithm runs. Before the algorithm runs, the counter is at the page marked (x). Write the sequence of reference bits after the second chance algorithm has executed once. In the answer write PRECISELY one space BETWEEN each number and do not mention (x).

0 0 1(x) 1 0 1 1

Answer: 0 0 0 0 1 1



The correct answer is: 0 0 0 0 1 1

Question **49**

Correct

Mark 0.50 out of 0.50

Map the block allocation scheme with the problem it suffers from

(Match pairs 1-1, match a scheme with the problem that it suffers from relatively the most, compared to others)

Indexed Allocation	Overhead of reading metadata blocks ▾	✓
Linked allocation	Too many seeks ▾	✓
Continuous allocation	need for compaction ▾	✓

Your answer is correct.


The correct answer is: Indexed Allocation → Overhead of reading metadata blocks, Linked allocation → Too many seeks, Continuous allocation → need for compaction

Question **50**


Correct

Mark 0.50 out of 0.50


Match the diagram with the threading model




Two-Level ☐ ✓



Many-One ☐ ✓




Many-Many ☐ ✓




One-One ☐ ✓

Your answer is correct.


The correct answer is:




→ Two-Level,



→ Many-One,



→ Many-Many,



→ One-One

Question **51**

Correct

Mark 0.50 out of 0.50

Suppose a program does a `scanf()` call.

Essentially the `scanf` does a `read()` system call.

This call will obviously "block" waiting for the user input.

In terms of OS data structures and execution of code, what does it mean?

Select one:

- ☐ a. OS code for `read()` will call scheduler
- ☐ b. OS code for `read()` will move the PCB of this process to a wait queue and return from the system call
- ☒ c. OS code for `read()` will move PCB of current process to a wait queue and call scheduler ✓
- ☐ d. `read()` will return and process will be taken to a wait queue
- ☐ e. `read()` returns and process calls `scheduler()`

Your answer is correct.

The correct answer is: OS code for `read()` will move PCB of current process to a wait queue and call scheduler

Mark the statements as True or False, w.r.t. thrashing

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Thrashing occurs when the total size of all processes's locality exceeds total memory size.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Thrashing occurs because some process is doing lot of disk I/O.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Thrashing can be limited if local replacement is used.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Thrashing is particular to demand paging systems, and does not apply to pure paging systems.	✓
<input type="radio"/>	<input checked="" type="radio"/>	mmap() solves the problem of thrashing.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Thrashing can occur even if entire memory is not in use.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The working set model is an attempt at approximating the locality of a process.	✓
<input checked="" type="radio"/>	<input type="radio"/>	During thrashing the CPU is under-utilised as most time is spent in I/O	✓

Thrashing occurs when the total size of all processes's locality exceeds total memory size.: True

Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.: True

Thrashing occurs because some process is doing lot of disk I/O.: False

Thrashing can be limited if local replacement is used.: True

Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.: False

Thrashing is particular to demand paging systems, and does not apply to pure paging systems.: True

mmap() solves the problem of thrashing.: False

Thrashing can occur even if entire memory is not in use.: False

The working set model is an attempt at approximating the locality of a process.: True

During thrashing the CPU is under-utilised as most time is spent in I/O: True

Question **53**

Partially correct

Mark 0.25 out of 1.00

Select the correct statements about directory entries in ext2:

- ☐ a. Maximum size of file name is 255
- ☒ b. It is possible that there is a hole between two directory entries. ✓
- ☐ c. When the first entry is removed, a compaction is done
- ☒ d. The size of a directory is = sum of sizes of all directory entries ✗
- ☒ e. The rec_len is always = 8 + name_len + number to round off name_len to multiple of 4 ✗
- ☒ f. The rec_len of the last entry is set to cover the directory's actual size. ✓
- ☐ g. The name_len is always a multiple of 4, by adding extra character to the actual name
- ☒ h. The name_len is always stored in multiple of 4 bytes, suffixing zeroes if required. ✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: It is possible that there is a hole between two directory entries., The rec_len of the last entry is set to cover the directory's actual size., The name_len is always stored in multiple of 4 bytes, suffixing zeroes if required., Maximum size of file name is 255

Question **54**

Correct

Mark 0.50 out of 0.50

If one thread opens a file with read privileges then, without opening the file

Select one:

- ☐ a. other threads in the another process can also read from that file
- ☐ b. none of these
- ☒ c. other threads in the same process can also read from that file ✓
- ☐ d. any other thread cannot read from that file

Your answer is correct.

The correct answer is: other threads in the same process can also read from that file

Question **55**

Correct

Mark 0.50 out of 0.50

Match the pairs of which action is taken by whom

- | | | |
|-----------|--|---|
| processor | <input type="text" value="detect memory violations"/> | ✓ |
| kernel | <input type="text" value="setup the MMU hardware"/> | ✓ |
| compiler | <input type="text" value="generate code using virtual addresses"/> | ✓ |

The correct answer is: processor → detect memory violations, kernel → setup the MMU hardware, compiler → generate code using virtual addresses

Question **56**

Correct

Mark 0.50 out of 0.50

which of the following, do you think, are valid concerns for making the kernel pageable?

- ☐ a. No data structure of kernel should be pageable
- ☒ b. The kernel's own page tables should not be pageable ✓
- ☒ c. The kernel must have some dedicated frames for it's own work ✓
- ☐ d. No part of kernel code should be pageable.
- ☒ e. The page fault handler should not be pageable ✓
- ☒ f. The disk driver and disk interrupt handler should not be pageable ✓

The correct answers are: The kernel's own page tables should not be pageable, The page fault handler should not be pageable, The kernel must have some dedicated frames for it's own work, The disk driver and disk interrupt handler should not be pageable

Question **57**

Incorrect

Mark 0.00 out of 1.00

Calculate the EAT in NANO-seconds (upto 2 decimal points) w.r.t. a page fault, given

Memory access time = 140 ns

Average page fault service time = 10 ms

Page fault rate = 0.9

Answer: ✗

The correct answer is: 9000014.00

For Virtual File System to work, which of the following changes are required to be done to an existing OS code (e.g. xv6)?

- ☒ a. The file system specific function pointers, for file system system-calls, need to be setup in the generic inode during lookup. ✓
- ☒ b. A mount() system call should be provided to mount a partition onto some directory in existing namespace rooted at "/" ✓
- ☒ c. Each open() needs to copy the function pointers from the inode of the parent directory into the inode of the child (if not already done), unless it's traversing a mount point. (This may be done as part of lookup() which is called by open()) ✓
- ☒ d. The operating system in-memory inode needs to be a generic-inode representing "inode" like data structure across multiple file systems. ✓
- ☒ e. The lookup() operation needs to check if it's crossing a mount point and call FS specific operations to read inodes/directories ✓
- ☒ f. The filesystem related system calls (e.g. read, write) need to invoke the file system specific functions (e.g. ext2_read, ext2_write, ntfs_read, ntfs_write) using function pointers. ✓
- ☒ g. Each file-system writer needs to provide the set of function pointers for VFS, and these function pointers need to be setup in generic inode of "/" of that file system during mount() ✓
- ☒ h. The generic inode needs to have a field representing if this inode is a mount point and also to refer/point to the root of the mounted file system's inode. ✓

The correct answers are: A mount() system call should be provided to mount a partition onto some directory in existing namespace rooted at "/", The filesystem related system calls (e.g. read, write) need to invoke the file system specific functions (e.g. ext2_read, ext2_write, ntfs_read, ntfs_write) using function pointers., The file system specific function pointers, for file system system-calls, need to be setup in the generic inode during lookup., The operating system in-memory inode needs to be a generic-inode representing "inode" like data structure across multiple file systems., The generic inode needs to have a field representing if this inode is a mount point and also to refer/point to the root of the mounted file system's inode., The lookup() operation needs to check if it's crossing a mount point and call FS specific operations to read inodes/directories, Each file-system writer needs to provide the set of function pointers for VFS, and these function pointers need to be setup in generic inode of "/" of that file system during mount(), Each open() needs to copy the function pointers from the inode of the parent directory into the inode of the child (if not already done), unless it's traversing a mount point. (This may be done as part of lookup() which is called by open())

Question **59**

Correct

Mark 0.50 out of 0.50

Predict the output of the program given here.

Assume that all the path names for the programs are correct. For example `"/usr/bin/echo"` will actually run echo command.

Assume that there is no mixing of printf output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good

output

should be written as good output

--

```
main() {  
    int i;  
    i = fork();  
    if(i == 0)  
        execl("/usr/bin/echo", "/usr/bin/echo", "hi", 0);  
    else  
        wait(0);  
    execl("/usr/bin/echo", "/usr/bin/echo", "one", 0);  
}
```

Answer: hi one



The correct answer is: hi one

Question **60**

Correct

Mark 0.50 out of 0.50

Map the technique with it's feature/problem

static loading	wastage of physical memory	✓
static linking	large executable file	✓
dynamic loading	allocate memory only if needed	✓
dynamic linking	small executable file	✓

The correct answer is: static loading → wastage of physical memory, static linking → large executable file, dynamic loading → allocate memory only if needed, dynamic linking → small executable file

Question 61

Partially correct

Mark 0.80 out of 1.00

Mark statements True/False w.r.t. change of states of a process.

Reference: The process state diagram (and your understanding of how kernel code works)




Figure 3.2 Diagram of process state.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	A RUNNING process becomes READY only on timer interrupt	✓
<input checked="" type="radio"/>	<input type="radio"/>	A process in RUNNING state becomes WAITING if resource is not available immediately	✓
<input type="radio"/>	<input checked="" type="radio"/>	Some forked processes may terminate normally without becoming ZOMBIE.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A READY process becomes RUNNING only when the scheduler selects it	✓
<input type="radio"/>	<input checked="" type="radio"/>	A process in RUNNING state becomes WAITING if the scheduler decides that it should WAIT for its turn	✗

A RUNNING process becomes READY only on timer interrupt: True

A process in RUNNING state becomes WAITING if resource is not available immediately: True

Some forked processes may terminate normally without becoming ZOMBIE.: False

A READY process becomes RUNNING only when the scheduler selects it: True

A process in RUNNING state becomes WAITING if the scheduler decides that it should WAIT for its turn: False

Question 62

Correct

Mark 2.00 out of 2.00

The following processes are being scheduled using a pre-emptive, priority-based, round-robin scheduling algorithm.

<u>Process</u>	<u>Priority</u>	<u>Burst</u>	<u>Arrival</u>
P_1	8	15	0
P_2	3	20	0
P_3	4	20	20
P_4	4	20	25
P_5	5	5	45
P_6	5	15	55

Each process is assigned a numerical priority, with a higher number indicating a higher relative priority. The scheduler will execute the currently highest priority process for its full duration, unless it gets pre-empted by newly arriving higher priority process. For processes with the same priority, a round-robin scheduler will be used with a time quantum of 10 units. If a process is pre-empted by a higher-priority process, the pre-empted process is placed at the front of the queue of same priority processes (so that its turn continues when higher priority process is over).

The order in which the processes get scheduled is (write your answer without a space, e.g. P1,P2,P3,P4,P5) :



The turn-around time for the process P2 is:



The turn-around time for the process P6 is:



The process P5 finishes at time unit:



0-15 P1

15-20 P2

20-40 P3

40-45 P4

45-50 P5

50-55 P4

55-70 P6

70-80 P4

80-95 P2

--

P2 turnaround time = 95 - 15 = 80

Question **63**

Correct

Mark 1.00 out of 1.00

Order the following events, related to page fault handling, in correct order

1. ✓ MMU detects that a page table entry is marked "invalid"
2. ✓ Page fault interrupt is generated
3. ✓ Page fault handler in kernel starts executing
4. ✓ Page fault handler detects that it's a page fault and not illegal memory access
5. ✓ Empty frame is found
6. ✓ Disk read is issued
7. ✓ Page faulting process is made to wait in a queue
8. ✓ Other processes scheduled by scheduler
9. ✓ Disk Interrupt occurs
10. ✓ Disk interrupt handler runs
11. ✓ Page table of page faulted process is updated
12. ✓ Page faulted process is moved to ready-queue

Your answer is correct.

[◀ Homework questions: Basics of MM, xv6 booting](#)

Jump to...



[Link to Ubuntu VM for this course \(VDI File, 15+GB\) ►](#)