

<b>Started on</b>	Saturday, 16 March 2024, 1:32 PM
<b>State</b>	Finished
<b>Completed on</b>	Saturday, 16 March 2024, 3:50 PM
<b>Time taken</b>	2 hours 17 mins
<b>Grade</b>	<b>11.42</b> out of 15.00 ( <b>76.13%</b> )

Question **1**

Correct

Mark 0.50 out of 0.50

Select the correct statements about hard and soft links

Select one or more:

- ☐ a. Hard links can span across partitions while soft links can't
- ☐ b. Deleting a soft link deletes both the link and the actual file
- ☐ c. Deleting a soft link deletes only the actual file
- ☐ d. Soft links increase the link count of the actual file inode
- ☒ e. Hard links enforce separation of filename from it's metadata in on-disk data structures. ✓
- ☒ f. Hard links increase the link count of the actual file inode ✓
- ☐ g. Deleting a hard link always deletes the file
- ☒ h. Deleting a hard link deletes the file, only if link count was 1 ✓
- ☐ i. Soft link shares the inode of actual file
- ☒ j. Hard links share the inode ✓
- ☒ k. Soft links can span across partitions while hard links can't ✓
- ☒ l. Deleting a soft link deletes the link, not the actual file ✓

Your answer is correct.

The correct answers are: Soft links can span across partitions while hard links can't, Hard links increase the link count of the actual file inode, Deleting a soft link deletes the link, not the actual file, Deleting a hard link deletes the file, only if link count was 1, Hard links share the inode, Hard links enforce separation of filename from it's metadata in on-disk data structures.

Question **2**

Incorrect

Mark 0.00 out of 0.50

In the code below assume that each function can be executed concurrently by many threads/processes.  
Ignore syntactical issues, and focus on the semantics.

This program is an example of

```
spinlock a, b; // assume initialized
thread1() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
thread2() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
```

- ☐ a. Self Deadlock
- ☒ b. Deadlock or livelock depending on actual race ❌
- ☐ c. Deadlock
- ☐ d. None of these
- ☐ e. Livelock

Your answer is incorrect.

The correct answer is: None of these

Question **3**

Correct

Mark 0.50 out of 0.50

The variable \$stack in entry.S is

- ☒ a. a memory region allocated as a part of entry.S ✔️
- ☐ b. located at the value given by %esp as setup by bootmain()
- ☐ c. located at less than 0x7c00
- ☐ d. located at 0x7c00
- ☐ e. located at 0

The correct answer is: a memory region allocated as a part of entry.S

Question **4**

Correct

Mark 0.50 out of 0.50

Which of the following is done by mappages()?

- ☐ a. allocate page frame if required
- ☒ b. create page table mappings to the range given by "pa" and "pa + size" ✓
- ☐ c. allocate page directory if required
- ☒ d. create page table mappings for the range given by "va" and "va + size" ✓
- ☒ e. allocate page table if required ✓

The correct answers are: create page table mappings for the range given by "va" and "va + size", allocate page table if required, create page table mappings to the range given by "pa" and "pa + size"

Question **5**

Partially correct

Mark 0.38 out of 0.50

The kernel ELF file contains these headers

Program Header:

```
LOAD off 0x00001000 vaddr 0x80100000 paddr 0x00100000 align 2**12
      filesz 0x00007aab memsz 0x00007aab flags r-x
LOAD off 0x00009000 vaddr 0x80108000 paddr 0x00108000 align 2**12
      filesz 0x00002516 memsz 0x0000d4a8 flags rw-
STACK off 0x00000000 vaddr 0x00000000 paddr 0x00000000 align 2**4
      filesz 0x00000000 memsz 0x00000000 flags rwx
```

mark the statemetns as True/False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	First header is for the code/text	✓
<input checked="" type="radio"/>	<input type="radio"/>	Third header is for stack	✓
<input type="radio"/>	<input checked="" type="radio"/>	in bootmain() the third header leads to allocation of no-memory.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Second header is for Data/Globals	✓

First header is for the code/text: True

Third header is for stack: True

in bootmain() the third header leads to allocation of no-memory.: True

Second header is for Data/Globals: True

Question **6**

Partially correct

Mark 0.25 out of 0.50

Mark statements about deadlocks as True or false

True	False			
<input checked="" type="radio"/>	<input type="radio"/>	A deadlock necessarily requires a cycle in the resource allocation graph	✓	
<input checked="" type="radio"/>	<input type="radio"/>	A deadlock is possible only if all the 4 conditions of mutual exclusion, cyclic wait, hold and wait, and no preemption are satisfied	✓	
<input checked="" type="radio"/>	<input type="radio"/>	Cycle in the resource allocation graph does not necessarily mean a deadlock	✓	
<input type="radio"/>	<input checked="" type="radio"/>	Deadlocks are the same as livelocks	✗	
<input type="radio"/>	<input checked="" type="radio"/>	A deadlock must involve at least two processes	✗	
<input checked="" type="radio"/>	<input type="radio"/>	Deadlocks are not possible if there is no race	✗	

A deadlock necessarily requires a cycle in the resource allocation graph: True

A deadlock is possible only if all the 4 conditions of mutual exclusion, cyclic wait, hold and wait, and no preemption are satisfied: True

Cycle in the resource allocation graph does not necessarily mean a deadlock: True

Deadlocks are the same as livelocks: False

A deadlock must involve at least two processes: False

Deadlocks are not possible if there is no race: True

Question **7**  
Partially correct  
Mark 0.43 out of 0.50

Given below are statements about concurrency and parallelism

Select T/F

A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time.

True	False			
<input type="radio"/>	<input checked="" type="radio"/>	Parallel systems allow more than one task to progress while concurrent systems do not.	✓	
<input type="radio"/>	<input checked="" type="radio"/>	A concurrent system allows more than one task to progress while a parallel system does not.	✓	
<input checked="" type="radio"/>	<input type="radio"/>	It is possible to have concurrency without parallelism	✓	
<input checked="" type="radio"/>	<input type="radio"/>	A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time.	✓	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Both concurrency and parallelism are the same.	✗	
<input type="radio"/>	<input checked="" type="radio"/>	It is possible to have parallelism without concurrency	✓	
<input type="radio"/>	<input checked="" type="radio"/>	It is not possible to have concurrency without parallelism.	✓	

Parallel systems allow more than one task to progress while concurrent systems do not.: False  
A concurrent system allows more than one task to progress while a parallel system does not.: False  
It is possible to have concurrency without parallelism: True  
A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time.: True  
Both concurrency and parallelism are the same.: False  
It is possible to have parallelism without concurrency: False  
It is not possible to have concurrency without parallelism.: False

Question **8**  
Incorrect  
Mark 0.00 out of 0.50

In an ext2 file system, if the block size is 4KB and partition size is 32 GB, then the number of block groups will be:

Answer:  ✗

The correct answer is: 256.00

## Question 9

Correct

Mark 0.50 out of 0.50

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

True	False			
<input type="radio"/>	<input checked="" type="radio"/>	String arguments are first copied to trapframe and then from trapframe to kernel's other variables.	✓	
<input checked="" type="radio"/>	<input type="radio"/>	The functions like argint(), argstr() make the system call arguments available in the kernel.	✓	
<input checked="" type="radio"/>	<input type="radio"/>	The arguments to system call originally reside on process stack.	✓	
<input checked="" type="radio"/>	<input type="radio"/>	The arguments are accessed in the kernel code using esp on the trapframe.	✓	
<input type="radio"/>	<input checked="" type="radio"/>	Integer arguments are stored in eax, ebx, ecx, etc. registers	✓	
<input type="radio"/>	<input checked="" type="radio"/>	The arguments to system call are copied to kernel stack in trapasm.S	✓	
<input checked="" type="radio"/>	<input type="radio"/>	String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer	✓	
<input checked="" type="radio"/>	<input type="radio"/>	Integer arguments are copied from user memory to kernel memory using argint()	✓	

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False

The functions like argint(), argstr() make the system call arguments available in the kernel.: True

The arguments to system call originally reside on process stack.: True

The arguments are accessed in the kernel code using esp on the trapframe.: True

Integer arguments are stored in eax, ebx, ecx, etc. registers: False

The arguments to system call are copied to kernel stack in trapasm.S: False

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

Integer arguments are copied from user memory to kernel memory using argint(): True

## Question 10

Correct

Mark 0.50 out of 0.50

The variable 'end' used as argument to kinit1 has the value

- ☐ a. 80110000
- ☐ b. 80000000
- ☒ c. 801154a8 ✓
- ☐ d. 80102da0
- ☐ e. 8010a48c
- ☐ f. 81000000

The correct answer is: 801154a8

Question **11**

Partially correct

Mark 0.04 out of 0.50

Map ext2 data structure features with their purpose

Used  
directories  
count in  
group  
descriptor

Choose...

Free  
blocks  
count in  
superblock  
and group  
descriptor

for statistical information on directories created

Combining  
file type  
and access  
rights in  
one  
variable

aligns all memory accesses on word boundary, improving performance

Inode  
bitmap is  
one block

Choose...

Many  
copies of  
Superblock

All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to

File Name  
is padded

Choose...

Inode  
table

Inodes are stored continuously because often OS issues many inode read requests at a time

Mount  
count in  
superblock

Choose...

A group

Try to keep all the data of a directory and it's file close together in a group

Inode  
table  
location in  
Group  
Descriptor

All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to

Block  
bitmap is  
one block

Choose...

rec\_len  
field in  
directory  
entry

Choose...

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: **Used directories count in group descriptor** → attempt is made to evenly spread the first-level directories, this count is used there, **Free blocks count in superblock and group descriptor** → Redundancy to help fsck restore consistency, **Combining file type and access rights in one variable** → saves 1 byte of space, **Inode bitmap is one block** → limits total number of files that can belong to a group, **Many copies of Superblock** → Redundancy to ensure the most crucial data structure is not lost, File Name is padded → aligns all memory accesses on word boundary, improving performance, **Inode table** → All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk, **Mount count in superblock** → to enforce file check after certain amount of mounts at boot time, **A group** → Try to keep all the data of a directory and it's file close together in a group, **Inode table location in Group Descriptor** → Obvious, as it's per group and not per file-system, **Block bitmap is one block** → Limits the size of a block group, thus improving on purpose of a group, **rec\_len field in directory entry** → allows holes and linking of entries in directory

Question **12**

Partially correct

Mark 0.38 out of 0.50

Select the correct statements about `sched()` and `scheduler()` in xv6 code

- ☒ a. `scheduler()` switches to the selected process's context ✓
- ☒ b. `sched()` switches to the scheduler's context ✓
- ☐ c. When either `sched()` or `scheduler()` is called, it results in a context switch
- ☒ d. After call to `swtch()` in `scheduler()`, the control moves to code in `sched()` ✓
- ☒ e. `sched()` and `scheduler()` are co-routines ✓
- ☒ f. After call to `swtch()` in `sched()`, the control moves to code in `scheduler()` ✓
- ☒ g. Each call to `sched()` or `scheduler()` involves change of one stack inside `swtch()` ✓
- ☐ h. When either `sched()` or `scheduler()` is called, it does not return immediately to caller

Your answer is partially correct.

You have correctly selected 6.

The correct answers are: `sched()` and `scheduler()` are co-routines, When either `sched()` or `scheduler()` is called, it does not return immediately to caller, When either `sched()` or `scheduler()` is called, it results in a context switch, `sched()` switches to the scheduler's context, `scheduler()` switches to the selected process's context, After call to `swtch()` in `scheduler()`, the control moves to code in `sched()`, After call to `swtch()` in `sched()`, the control moves to code in `scheduler()`, Each call to `sched()` or `scheduler()` involves change of one stack inside `swtch()`



Consider this program.  
Some statements are identified using the // comment at the end.  
Assume that = is an atomic operation.

```
#include <stdio.h>
#include <pthread.h>
long c = 0, c1 = 0, c2 = 0, run = 1;
void *thread1(void *arg) {
    while(run == 1) { //E
        c = 10; //A
        c1 = c2 + 5; //B
    }
}
void *thread2(void *arg) {
    while(run == 1) { //F
        c = 20; //C
        c2 = c1 + 3; //D
    }
}
int main() {
    pthread_t th1, th2;
    pthread_create(&th1, NULL, thread1, NULL);
    pthread_create(&th2, NULL, thread2, NULL);
    sleep(2);
    run = 0;
    fprintf(stdout, "c = %ld c1+c2 = %ld c1 = %ld c2 = %ld \n", c, c1+c2, c1, c2);
    fflush(stdout);
}
```

Which statements are part of the critical Section?

Yes		No		
<input type="radio"/>	<input checked="" type="radio"/>	E	✓	
<input type="radio"/>	<input checked="" type="radio"/>	F	✓	
<input checked="" type="radio"/>	<input type="radio"/>	A	✗	
<input checked="" type="radio"/>	<input type="radio"/>	C	✗	
<input checked="" type="radio"/>	<input type="radio"/>	B	✓	
<input checked="" type="radio"/>	<input type="radio"/>	D	✓	

- E: No
- F: No
- A: No
- C: No
- B: Yes
- D: Yes

Question **14**

Incorrect

Mark 0.00 out of 0.50

Will this code work for a spinlock() operation? The intention here is to call compare-and-swap() only if the lock is not held (the if condition checks for the same).

```
void spinlock(int *lock) {
{
    while (true) {
        if (*lock == 0) {
            /* lock appears to be available */
            if (!compare_and_swap(lock, 0, 1))
                break
        }
    }
}
```

- ☒ a. No, because in the case of both processes succeeding in the "if" condition, both may end up acquiring the lock. ❌
- ☐ b. Yes, because there is no race to update the lock variable
- ☐ c. No, because this breaks the atomicity requirement of compare-and-test.
- ☐ d. Yes, because no matter in which order the if-check and compare-and-swap run in multiple processes, only one process will succeed in compare-and-swap() and others will keep looping in while-loop.

Your answer is incorrect.

The correct answer is: Yes, because no matter in which order the if-check and compare-and-swap run in multiple processes, only one process will succeed in compare-and-swap() and others will keep looping in while-loop.

Question **15**

Incorrect

Mark 0.00 out of 0.50

Why V2P\_WO is used in entry.S and not V2P ?

- ☒ a. Because the processor can not do a type casting at run time ❌
- ☐ b. Because entry.S is an assembly code file and assemblers do not know about data types and type casting.
- ☐ c. It's a mistake. They could have used the same macro in both places.
- ☐ d. The two macros are different. They lead to different calculations.
- ☐ e. The typecasting has the effect of creating virtual address, while without typecast we get physical address.

Your answer is incorrect.

The correct answer is: Because entry.S is an assembly code file and assemblers do not know about data types and type casting.

Question **16**

Correct

Mark 0.50 out of 0.50

Select the correct statements about paging (not demand paging) mechanism

Select one or more:

- ☒ a. OS creates the page table for every process ✓
- ☐ b. An invalid entry on a page means, either it was illegal memory reference or the page was not present in memory.
- ☐ c. User process can update it's own PTBR
- ☐ d. User process can update it's own page table entries
- ☒ e. An invalid entry on a page means, it was an illegal memory reference ✓
- ☒ f. Page table is accessed by the MMU as part of execution of an instruction ✓
- ☒ g. The PTBR is loaded by the OS ✓
- ☐ h. Page table is accessed by the OS as part of execution of an instruction

Your answer is correct.

The correct answers are: OS creates the page table for every process, The PTBR is loaded by the OS, Page table is accessed by the MMU as part of execution of an instruction, An invalid entry on a page means, it was an illegal memory reference

Question **17**

Correct

Mark 0.50 out of 0.50

Match the code with it's functionality

**S = 5**

**Wait(S)**

**Critical Section**

**Signal(S)**

S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statemetn2;

Signal(S1);

P3:

Wait(S1);

Statement S3;

**S = 0**

**P1:**

**Statement1;**

**Signal(S)**

**P2:**

**Wait(S)**

**Statment2;**

**S = 1**

**Wait(S)**

**Critical Section**

**Signal(S);**

Counting semaphore



Execution order P2, P1, P3



Execution order P1, then P2



Binary Semaphore for mutual exclusion



Your answer is correct.

The correct answer is: **S = 5**

**Wait(S)**

**Critical Section**

**Signal(S)** → Counting semaphore, S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statemetn2;

Signal(S1);

P3:

Wait(S1);

Statement S3; → Execution order P2, P1, P3, **S = 0**

P1:

Statement1;

Signal(S)

P2:

Wait(S)

Statment2; → Execution order P1, then P2, **S = 1**

**Wait(S)**

**Critical Section**

**Signal(S);** → Binary Semaphore for mutual exclusion

Question **18**

Partially correct

Mark 0.38 out of 0.50

Suppose a file is to be created in an ext2 file system, in an existing directory `/a/b/`. Select from below, the list of blocks that may need modification.

Select one or more:

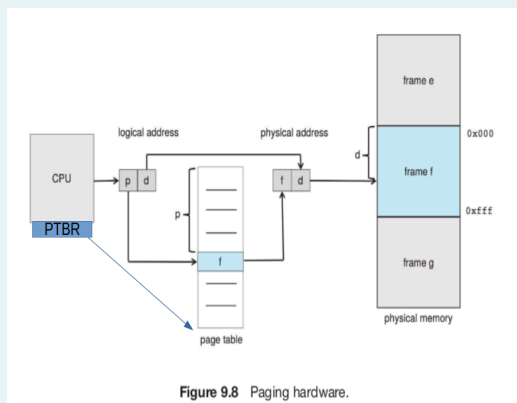
- ☒ a. group descriptor(s) ✓
- ☒ b. inode bitmap in some block group ✓
- ☐ c. inode of `/a/`
- ☒ d. superblock ✓
- ☒ e. inode table in some block group ✓
- ☐ f. link count on `/a/b/` inode
- ☐ g. data blocks of `/a/`
- ☐ h. existing data blocks of `/a/b/`
- ☒ i. block bitmap in some block group ✓
- ☐ j. inode of `/a/b/`
- ☐ k. inode bitmap referring to `/a/b/`
- ☒ l. new data block in some block group ✓

Your answer is partially correct.

You have correctly selected 6.

The correct answers are: superblock, group descriptor(s), inode of `/a/b/`, existing data blocks of `/a/b/`, inode table in some block group, inode bitmap in some block group, block bitmap in some block group, new data block in some block group

Consider the image given below, which explains how paging works.



Mention whether each statement is True or False, with respect to this image.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	The page table is indexed using frame number	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is indexed using page number	✓
<input checked="" type="radio"/>	<input type="radio"/>	The physical address may not be of the same size (in bits) as the logical address	✓
<input type="radio"/>	<input checked="" type="radio"/>	The locating of the page table using PTBR also involves paging translation	✓
<input checked="" type="radio"/>	<input type="radio"/>	The PTBR is present in the CPU as a register	✓
<input type="radio"/>	<input checked="" type="radio"/>	Size of page table is always determined by the size of RAM	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is itself present in Physical memory	✓
<input checked="" type="radio"/>	<input type="radio"/>	Maximum Size of page table is determined by number of bits used for page number	✓

The page table is indexed using frame number: False

The page table is indexed using page number: True

The physical address may not be of the same size (in bits) as the logical address: True

The locating of the page table using PTBR also involves paging translation: False

The PTBR is present in the CPU as a register: True

Size of page table is always determined by the size of RAM: False

The page table is itself present in Physical memory: True

Maximum Size of page table is determined by number of bits used for page number: True

Question **20**

Correct

Mark 0.50 out of 0.50

Which of the following is DONE by allocproc() ?

- ☒ a. setup the trapframe and context pointers appropriately ✓
- ☒ b. allocate PID to the process ✓
- ☐ c. setup the contents of the trapframe of the process properly
- ☐ d. setup kernel memory mappings for the process
- ☒ e. ensure that the process starts in forkret() ✓
- ☒ f. allocate kernel stack for the process ✓
- ☒ g. Select an UNUSED struct proc for use ✓
- ☐ h. ensure that the process starts in trapret()

The correct answers are: Select an UNUSED struct proc for use, allocate PID to the process, allocate kernel stack for the process, setup the trapframe and context pointers appropriately, ensure that the process starts in forkret()

Question **21**

Correct

Mark 0.50 out of 0.50

The "push 0" in vectors.S is

- ☐ a. A placeholder to match the size of struct trapframe
- ☐ b. To be filled in as the return value of the system call
- ☐ c. To indicate that it's a system call and not a hardware interrupt
- ☒ d. Place for the error number value ✓

The correct answer is: Place for the error number value

Question **22**

Partially correct

Mark 0.43 out of 0.50

Mark statements as T/F

All statements are in the context of preventing deadlocks.

True	False			
<input checked="" type="radio"/>	<input type="radio"/>	Circular wait is avoided by enforcing a lock ordering	✓	
<input type="radio"/>	<input checked="" type="radio"/>	If a resource allocation graph contains a cycle then there is a guarantee of a deadlock	✓	
<input checked="" type="radio"/>	<input type="radio"/>	A process holding one resources and waiting for just one more resource can also be involved in a deadlock.	✓	
<input checked="" type="radio"/>	<input type="radio"/>	Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.	✓	
<input type="radio"/>	<input checked="" type="radio"/>	The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order	✗	
<input checked="" type="radio"/>	<input type="radio"/>	Hold and wait means a thread/process holding some locks and waiting for acquiring some.	✓	
<input checked="" type="radio"/>	<input type="radio"/>	Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens	✓	

Circular wait is avoided by enforcing a lock ordering: True

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

A process holding one resources and waiting for just one more resource can also be involved in a deadlock.: True

Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

Hold and wait means a thread/process holding some locks and waiting for acquiring some.: True

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

Question **23**

Correct

Mark 0.50 out of 0.50

Doing a lookup on the pathname /a/b/b/c/d for opening the file "d" requires reading  ✓ no. of inodes. Assume that there are no hard/soft links on the path.

Write the answer as a number.

The correct answer is: 6



Question **24**

Partially correct

Mark 0.31 out of 0.50

It is proposed that when a process does an illegal memory access, xv6 terminate the process by printing the error message "Illegal Memory Access". Select all the changes that need to be done to xv6 for this as True (Note that the changes proposed here may not cover the exhaustive list of all changes required) and the un-necessary/wrong changes as False.

Required	Un-necessary/Wrong			
<input type="radio"/>	<input checked="" type="radio"/>	Change mappages() to set specified permissions on each page table entry		
<input checked="" type="radio"/>	<input type="radio"/>	Change allocuvn() to call mappages() with proper permissions on each page table entry		
<input checked="" type="radio"/>	<input type="radio"/>	Handle the Illegal memory access trap in trap() function, and terminate the currently running process.		
<input checked="" type="radio"/>	<input type="radio"/>	Add code that checks if the illegal memory access trap was due to an actual illegal memory access.		
<input type="radio"/>	<input checked="" type="radio"/>	Ensure that the address 0 is mapped to invalid		
<input type="radio"/>	<input checked="" type="radio"/>	Mark each page as readonly in the page table mappings		
<input checked="" type="radio"/>	<input type="radio"/>	Change exec to treat text/data sections separately and call allocuvn() with proper flags for page table entries		
<input type="radio"/>	<input checked="" type="radio"/>	Change in the Makefile and instruct cc/ld to start the code of each program at some address other than 0		

Change mappages() to set specified permissions on each page table entry: Un-necessary/Wrong  
 Change allocuvn() to call mappages() with proper permissions on each page table entry: Required  
 Handle the Illegal memory access trap in trap() function, and terminate the currently running process.: Required  
 Add code that checks if the illegal memory access trap was due to an actual illegal memory access.: Un-necessary/Wrong  
 Ensure that the address 0 is mapped to invalid: Required  
 Mark each page as readonly in the page table mappings: Un-necessary/Wrong  
 Change exec to treat text/data sections separately and call allocuvn() with proper flags for page table entries: Required  
 Change in the Makefile and instruct cc/ld to start the code of each program at some address other than 0: Required

Question **25**

Correct

Mark 0.50 out of 0.50

## Match pairs

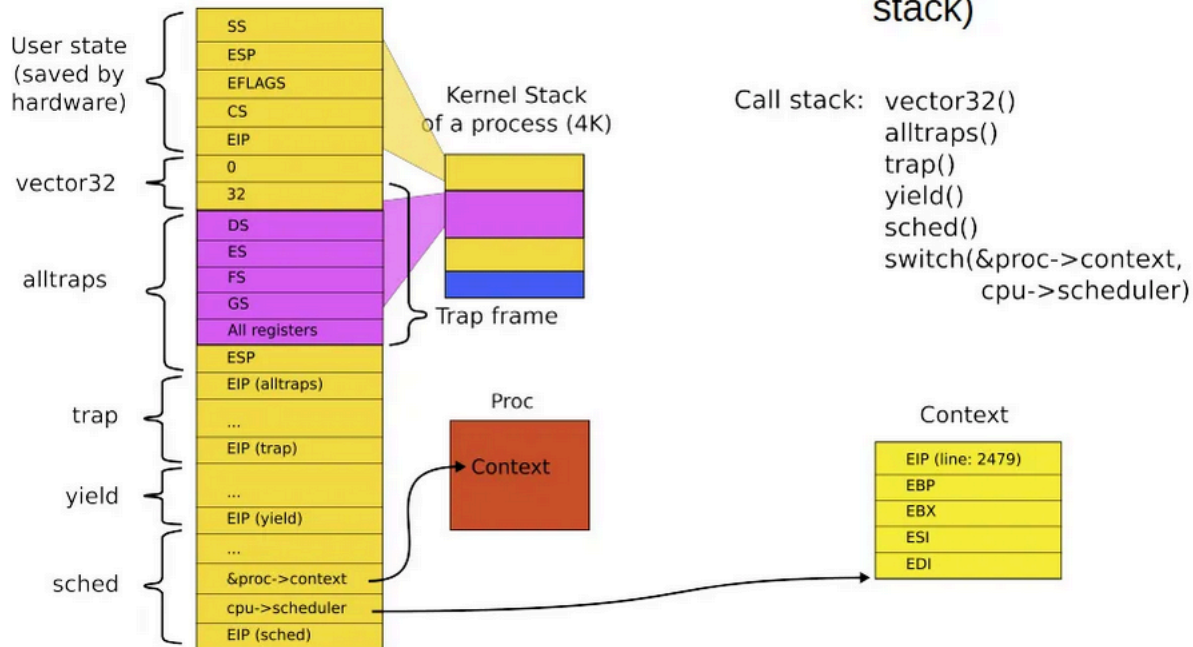
semaphore	<input type="text" value="wait() and signal()"/>	
mutex	<input type="text" value="lock() and unlock()"/>	
peterson	<input type="text" value="per process flag, global turn variable"/>	
spinlock	<input type="text" value="atomic test and set with loop"/>	

Your answer is correct.

The correct answer is: semaphore → wait() and signal(), mutex → lock() and unlock(), peterson → per process flag, global turn variable, spinlock → atomic test and set with loop

Mark statements as True/False, w.r.t. the given diagram

## Stack inside switch() and its two arguments (passed on the stack)



True False

☒ ☐

☐ ☒

The "ESP" (second entry from top) is stack pointer of user-stack of process, while the "ESP" (first entry below pink region) is the trapframe pointer on kernel stack of process.

✓

☐ ☒

☒ ☐

The "context" yellow coloured box, pointed to by `cpu->scheduler` is on the kernel stack of the scheduler.

✓

☒ ☐

☐ ☒

The diagram is correct

✓

☐ ☒

☒ ☐

The diagram is wrong because it shows the user stack and kernel stack together (continuous), but in practice they are separate

✓

diagram shows only kernel stack

☐ ☒

☒ ☐

This is a diagram of `switch()` called from `scheduler()`

✓

No. diagram of `switch()` called from `sched()`

☒ ☐

☐ ☒

The blue shaded part in "kernel stack of a process(4k)" refers to remaining part of stack (not used yet)

✓

The "ESP" (second entry from top) is stack pointer of user-stack of process, while the "ESP" (first entry below pink region) is the trapframe pointer on kernel stack of process.: True

The "context" yellow coloured box, pointed to by `cpu->scheduler` is on the kernel stack of the scheduler.: False

The diagram is correct: True

The diagram is wrong because it shows the user stack and kernel stack together (continuous), but in practice they are separate: False

This is a diagram of `switch()` called from `scheduler()`: False

The blue shaded part in "kernel stack of a process(4k)" refers to remaining part of stack (not used yet): True

Question **27**

Correct

Mark 0.50 out of 0.50

Which of the following is not a task of the code of `swtch()` function

- ☒ a. Save the return value of the old context code ✓
- ☐ b. Save the old context
- ☐ c. Load the new context
- ☐ d. Jump to next context EIP
- ☒ e. Change the kernel stack location ✓
- ☐ f. Switch stacks

The correct answers are: Save the return value of the old context code, Change the kernel stack location

Question **28**

Correct

Mark 0.50 out of 0.50

Mark the statements as True/False, with respect to the use of the variable "chan" in struct proc.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The value of 'chan' is changed only in <code>sleep()</code>	✓
<input type="radio"/>	<input checked="" type="radio"/>	chan is the head pointer to a linked list of processes, waiting for a particular event to occur	✓
<input checked="" type="radio"/>	<input type="radio"/>	When chan is not NULL, the 'state' in struct proc must be SLEEPING	✓
<input type="radio"/>	<input checked="" type="radio"/>	when chan is NULL, the 'state' in proc must be RUNNABLE.	✓
<input checked="" type="radio"/>	<input type="radio"/>	in xv6, the address of an appropriate variable is used as a "condition" for a waiting process.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Changing the state of a process automatically changes the value of 'chan'	✓
<input checked="" type="radio"/>	<input type="radio"/>	'chan' is used only by the <code>sleep()</code> and <code>wakeup1()</code> functions.	✓
<input checked="" type="radio"/>	<input type="radio"/>	chan stores the address of the variable, representing a condition, for which the process is waiting.	✓

The value of 'chan' is changed only in `sleep()`: True

chan is the head pointer to a linked list of processes, waiting for a particular event to occur: False

When chan is not NULL, the 'state' in struct proc must be SLEEPING: True

when chan is NULL, the 'state' in proc must be RUNNABLE.: False

in xv6, the address of an appropriate variable is used as a "condition" for a waiting process.: True

Changing the state of a process automatically changes the value of 'chan': False

'chan' is used only by the `sleep()` and `wakeup1()` functions.: True

chan stores the address of the variable, representing a condition, for which the process is waiting.: True

## Question 29

Correct

Mark 0.50 out of 0.50

Mark statements as True/False w.r.t. the creation of free page list in xv6.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	kmem.use_lock is set to 1 after free page list is created, so that kmem.lock is taken before accessing kmem.freelist.	✓
<input checked="" type="radio"/>	<input type="radio"/>	if(kmem.use_lock) acquire(&kmem.lock); is not done when called from kinit1() because there is no need to take the lock when kinit1() is running because interrupts are disabled and only one processor is running	✓
<input type="radio"/>	<input checked="" type="radio"/>	free page list is a singly circular linked list.	✓ it's singly linked NULL terminated list.
<input checked="" type="radio"/>	<input type="radio"/>	The pointers that link the pages together are in the first 4 bytes of the pages themselves	✓
<input checked="" type="radio"/>	<input type="radio"/>	the kmem.lock is used by kfree() and kalloc() only.	✓
<input type="radio"/>	<input checked="" type="radio"/>	if(kmem.use_lock) acquire(&kmem.lock); this "if" condition is true, when kinit2() runs because multi-processor support has been enabled by now.	✓ No. kinit2() calls kfree() and then initializes use_lock.

kmem.use\_lock is set to 1 after free page list is created, so that kmem.lock is taken before accessing kmem.freelist.: True

if(kmem.use\_lock)

acquire(&amp;kmem.lock);

is not done when called from kinit1() because there is no need to take the lock when kinit1() is running because interrupts are disabled and only one processor is running: True

free page list is a singly circular linked list.: False

The pointers that link the pages together are in the first 4 bytes of the pages themselves: True

the kmem.lock is used by kfree() and kalloc() only.: True

if(kmem.use\_lock)

acquire(&amp;kmem.lock);

this "if" condition is true, when kinit2() runs because multi-processor support has been enabled by now.: False

## Question 30

Correct

Mark 0.50 out of 0.50

Select the statement that most correctly describes what setupkvm() does

- ☒ a. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array ✓
- ☐ b. creates a 1-level page table for the use by the kernel, as specified in kmap[] global array
- ☐ c. creates a 2-level page table for the use of the kernel, as specified in gtdtdesc
- ☐ d. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array and makes kpgdir point to it

The correct answer is: creates a 2-level page table setup with virtual-&gt;physical mappings specified in the kmap[] global array

◀ Quiz-1 (15 Marks)

Jump to...



ESE(60 Marks) ▶