

Started on Friday, 24 February 2023, 2:44 PM

State Finished

Completed on Friday, 24 February 2023, 4:39 PM

Time taken 1 hour 55 mins

Grade **8.78** out of 10.00 (**87.77%**)

Question **1**

Correct

Mark 0.50 out of 0.50

What is meant by formatting a disk/partition?

- a. erasing all data on the disk/partition
- b. storing all the necessary programs on the disk/partition
- c. creating layout of empty directory tree/graph data structure ✓
- d. writing zeroes on all sectors

The correct answer is: creating layout of empty directory tree/graph data structure

Question 2

Partially correct

Mark 0.44 out of 0.50

How does the compiler calculate addresses for the different parts of a C program, when paging is used?

Text	starting with 0	✓
Static variables	Immediately after the text, along with globals	✓
Local variables	An offset with respect to stack pointer (esp)	✓
#include files	No memory allocated, they are handled by linker	✗
#define	No memory allocated, they are handled by pre-processor	✓
typedef	No memory allocated, as they are not variables, but only conceptual definition of a type	✓
Global variables	Immediately after the text	✓
malloced memory	Heap (handled by the malloc-free library, using OS's system calls)	✓

Your answer is partially correct.

You have correctly selected 7.

The correct answer is: Text → starting with 0, Static variables → Immediately after the text, along with globals, Local variables → An offset with respect to stack pointer (esp), #include files → No memory allocated for the file, but if it contains variables, then variables may be allocated memory, #define → No memory allocated, they are handled by pre-processor, typedef → No memory allocated, as they are not variables, but only conceptual definition of a type, Global variables → Immediately after the text, malloced memory → Heap (handled by the malloc-free library, using OS's system calls)

Question 3

Correct

Mark 0.50 out of 0.50

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

```
$ ls ./tmp/asdfksdf >/tmp/ddd 2>&1
```

Program 1

```
int main(int argc, char *argv[]) {  
    int fd, n, i;  
    char buf[128];  
  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    close(1);  
    dup(fd);  
    close(2);  
    dup(fd);  
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);  
}
```

Program 2

```
int main(int argc, char *argv[]) {  
    int fd, n, i;  
    char buf[128];  
  
    close(1);  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    close(2);  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);  
}
```

Select all the correct statements about the programs

Select one or more:

- a. Program 1 makes sure that there is one file offset used for '2' and '1' ✓
- b. Program 1 does 1>&2
- c. Both programs are correct
- d. Both program 1 and 2 are incorrect
- e. Program 2 does 1>&2
- f. Program 2 makes sure that there is one file offset used for '2' and '1'
- g. Program 2 ensures 2>&1 and does not ensure > /tmp/ddd
- h. Program 1 is correct for > /tmp/ddd but not for 2>&1
- i. Program 2 is correct for > /tmp/ddd but not for 2>&1
- j. Program 1 ensures 2>&1 and does not ensure > /tmp/ddd
- k. Only Program 1 is correct ✓
- l. Only Program 2 is correct

Your answer is correct.

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'

Question 4

Partially correct

Mark 0.43 out of 0.50

You must have seen the error message "Segmentation fault, core dumped" very often.

With respect to this error message, mark the statements as True/False.

True	False	
<input type="radio"/> ✗	<input checked="" type="radio"/>	The term "core" refers to the core code of the kernel.
<input type="radio"/> ✗	<input checked="" type="radio"/>	On Linux, the message is printed only because the memory management scheme is segmentation
<input checked="" type="radio"/>	<input type="radio"/> ✗	The process has definitely performed illegal memory access.
<input checked="" type="radio"/>	<input type="radio"/> ✗	On Linux, the process was sent a SIGSEGV signal and the default handler for the signal is "Term", so the process is terminated.
<input type="radio"/> ✗	<input checked="" type="radio"/>	The illegal memory access was detected by the kernel and the process was punished by kernel.
<input checked="" type="radio"/>	<input type="radio"/> ✗	The core file can be analysed later using a debugger, to determine what went wrong.
<input checked="" type="radio"/>	<input type="radio"/> ✗	The image of the process is stored in a file called "core", if the ulimit allows so.

The term "core" refers to the core code of the kernel.: False

On Linux, the message is printed only because the memory management scheme is segmentation: False

The process has definitely performed illegal memory access.: True

On Linux, the process was sent a SIGSEGV signal and the default handler for the signal is "Term", so the process is terminated.: True

The illegal memory access was detected by the kernel and the process was punished by kernel.: False

The core file can be analysed later using a debugger, to determine what went wrong.: True

The image of the process is stored in a file called "core", if the ulimit allows so.: True

Question 5

Correct

Mark 0.50 out of 0.50

Which of the following instructions should be privileged?

Select one or more:

- a. Access memory management unit of the processor ✓
- b. Turn off interrupts. ✓
- c. Read the clock.
- d. Set value of timer. ✓
- e. Switch from user to kernel mode. ✓ This instruction (like INT) is itself privileged - and that is why it not only changes the mode, but also ensures a jump to an ISR (kernel code)
- f. Access a general purpose register
- g. Set value of a memory location
- h. Modify entries in device-status table ✓
- i. Access I/O device. ✓

Your answer is correct.

The correct answers are: Set value of timer., Access memory management unit of the processor, Turn off interrupts., Modify entries in device-status table, Access I/O device., Switch from user to kernel mode.

Question 6

Correct

Mark 0.50 out of 0.50

Map the block allocation scheme with the problem it suffers from

(Match pairs 1-1, match a scheme with the problem that it suffers from relatively the most, compared to others)

Indexed Allocation	Overhead of reading metadata blocks	✓
Continuous allocation	need for compaction	✓
Linked allocation	Too many seeks	✓

Your answer is correct.

The correct answer is: Indexed Allocation → Overhead of reading metadata blocks, Continuous allocation → need for compaction, Linked allocation → Too many seeks

Question 7

Partially correct

Mark 0.43 out of 0.50

Following code claims to implement the command

/bin/ls -l | /usr/bin/head -3 | /usr/bin/tail -1

Fill in the blanks to make the code work.

Note: Do not include space in writing any option. x[1][2] should be written without any space, and so is the case with [1] or [2]. Pay attention to exact syntax and do not write any extra character like ';' or = etc.

```
int main(int argc, char *argv[]) {
    int pid1, pid2;
    int pfd[ 2 ] ✓ ][2];
    pipe( pfd[0] ✓ );
    pid1 = fork() ✓ ;
    if(pid1 != 0) {
        close(pfd[0][0] ✓ );
        close( 1 ✓ );
        dup( pfd[0][1] ✓ );
        execl("/bin/ls", "/bin/ls", " -1 ", NULL);
    }
    pipe( pfd[1] ✓ );
    pid2 ✓ = fork();
    if(pid2 == 0) {
        close( pfd[0][1] ✗ );
        close(0);
        dup( pfd[1][0] ✗ );
        close(pfd[1][0] ✓ );
        close( 1 ✓ );
        dup( pfd[1][1] ✓ );
        execl("/usr/bin/head", "/usr/bin/head", " -3 ", NULL);
    } else {
        close(pfd[1][1] ✓ );
        close( 0 ✓ );
        dup( pfd[1][0] ✓ );
        close(pfd[0][0] ✓ );
    }
}
```

```
execl("/usr/bin/tail", "/usr/bin/tail", " -1 ", NULL);  
}  
}
```

Question 8

Partially correct

Mark 0.45 out of 0.50

Match the elements of C program to their place in memory

Arguments	Stack	✓
Function code	Code	✓
#define MACROS	No memory needed	✗
#include files	No memory needed	✓
Global Static variables	Data	✓
Mallocoed Memory	Heap	✓
Global variables	Data	✓
Local Static variables	Data	✓
Local Variables	Stack	✓
Code of main()	Code	✓

The correct answer is: Arguments → Stack, Function code → Code, #define MACROS → No Memory needed, #include files → No memory needed, Global Static variables → Data, Mallocoed Memory → Heap, Global variables → Data, Local Static variables → Data, Local Variables → Stack, Code of main() → Code

Question 9

Correct

Mark 0.50 out of 0.50

Doing a lookup on the pathname /a/b/b/c/d for opening the file "d" requires reading ✓ no. of inodes. Assume that there are no hard/soft links on the path.

Write the answer as a number.

The correct answer is: 6

Question 10

Partially correct

Mark 0.42 out of 0.50

Mark the statements about device drivers by marking as True or False.

True	False	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	Different devices of the same type (e.g. 2 IDE hard disks) must need different device drivers.
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	It's possible that a particular hardware has multiple device drivers available for it.
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	Device driver is part of OS code
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	Device driver is an intermediary between the hardware controller and OS
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	Device driver is part of hardware
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	Writing a device driver mandatorily demands reading the technical documentation about the hardware.

Different devices of the same type (e.g. 2 IDE hard disks) must need different device drivers.: False

It's possible that a particular hardware has multiple device drivers available for it.: True

Device driver is part of OS code: True

Device driver is an intermediary between the hardware controller and OS: True

Device driver is part of hardware: False

Writing a device driver mandatorily demands reading the technical documentation about the hardware.: True

Question 11

Correct

Mark 0.50 out of 0.50

Map each signal with its meaning

SIGSEGV	Invalid Memory Reference	✓
SIGALRM	Timer Signal from alarm()	✓
SIGPIPE	Broken Pipe	✓
SIGUSR1	User Defined Signal	✓
SIGCHLD	Child Stopped or Terminated	✓

The correct answer is: SIGSEGV → Invalid Memory Reference, SIGALRM → Timer Signal from alarm(), SIGPIPE → Broken Pipe, SIGUSR1 → User Defined Signal, SIGCHLD → Child Stopped or Terminated

Question 12

Partially correct

Mark 0.25 out of 0.50

Select all the blocks that may need to be written back to disk (if updated, of-course), as "Yes", when an operation of deleting a file is carried out on ext2 file system.

An option has to be correct entirely to be marked "Yes"

Block bitmap(s) for all the blocks of the file

 Yes ✓

Data blocks of the file

 Yes ✗

One or multiple data blocks of the parent directory

 Yes ✗

Superblock

 Yes ✓

Possibly one block bitmap corresponding to the parent directory

 Yes ✓

One or more data bitmap blocks for the parent directory

 Yes ✗

Your answer is partially correct.

only one data block of parent directory. multiple blocks not possible. an entry is always contained within one single block

You have correctly selected 3.

The correct answer is: Block bitmap(s) for all the blocks of the file → Yes, Data blocks of the file → No, One or multiple data blocks of the parent directory → No, Superblock → Yes, Possibly one block bitmap corresponding to the parent directory → Yes, One or more data bitmap blocks for the parent directory → No

Question 13

Correct

Mark 0.50 out of 0.50

Select the compiler's view of the process's address space, for each of the following MMU schemes:
(Assume that each scheme,e.g. paging/segmentation/etc is effectively utilised)

Segmentation	many continuous chunks of variable size	✓
Segmentation, then paging	many continuous chunks of variable size	✓
Paging	one continuous chunk	✓
Relocation + Limit	one continuous chunk	✓

Your answer is correct.

The correct answer is: Segmentation → many continuous chunks of variable size, Segmentation, then paging → many continuous chunks of variable size, Paging → one continuous chunk, Relocation + Limit → one continuous chunk

Question 14

Partially correct

Mark 0.30 out of 0.50

Mark the statements about named and un-named pipes as True or False

True	False	
<input checked="" type="radio"/>	<input type="radio"/> X	Named pipe exists as a file
<input type="radio"/>	<input checked="" type="radio"/> X	Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.
<input checked="" type="radio"/>	<input type="radio"/> X	Both types of pipes are an extension of the idea of "message passing".
<input type="radio"/> X	<input checked="" type="radio"/>	A named pipe has a name decided by the kernel.
<input checked="" type="radio"/>	<input type="radio"/> X	Un-named pipes are inherited by a child process from parent.
<input checked="" type="radio"/>	<input type="radio"/> X	Named pipes can exist beyond the life-time of processes using them.
<input checked="" type="radio"/>	<input type="radio"/> X	Both types of pipes provide FIFO communication.
<input type="radio"/> X	<input checked="" type="radio"/>	The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.
<input type="radio"/> X	<input checked="" type="radio"/>	Named pipes can be used for communication between only "related" processes.
<input type="radio"/> X	<input checked="" type="radio"/>	The pipe() system call can be used to create either a named or un-named pipe.

Named pipe exists as a file.: True

Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.: True

Both types of pipes are an extension of the idea of "message passing": True

A named pipe has a name decided by the kernel.: False

Un-named pipes are inherited by a child process from parent.: True

Named pipes can exist beyond the life-time of processes using them.: True

Both types of pipes provide FIFO communication.: True

The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.: False

Named pipes can be used for communication between only "related" processes.: False

The pipe() system call can be used to create either a named or un-named pipe.: False

Question 15

Partially correct

Mark 0.29 out of 0.50

Mark the statements as True/False w.r.t. the basic concepts of memory management.

True	False	
<input type="radio"/> ✗	<input checked="" type="checkbox"/> ✗	The kernel refers to the page table for converting each virtual address to physical address.
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	When a process is executing, each virtual address is converted into physical address by the CPU hardware directly.
<input type="radio"/> ✗	<input checked="" type="checkbox"/> ✗	The compiler interacts with the kernel continuously while compiling a program and obtains the correct set of memory addresses for code/stack/heap/data and then generates the machine code file.
<input type="radio"/> ✗	<input checked="" type="checkbox"/> ✗	The compiler generates the address references for code/data/stack/heap in the executable file as per the memory management schema chosen by the compiler itself, and then the kernel ensures that program is executed with this schema.
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	The kernel ensures that the MMU is setup before scheduling a process and then the CPU/MMU ensures that the address translation takes place.
<input type="radio"/> ✗	<input checked="" type="checkbox"/> ✗	When a process is executing, each virtual address is converted into physical address by the kernel directly.
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	The compiler generates address references for code/data/stack/heap in the executable file, depending on the MM architecture provided by CPU and kernel.

The kernel refers to the page table for converting each virtual address to physical address.: False

When a process is executing, each virtual address is converted into physical address by the CPU hardware directly.: True

The compiler interacts with the kernel continuously while compiling a program and obtains the correct set of memory addresses for code/stack/heap/data and then generates the machine code file.: False

The compiler generates the address references for code/data/stack/heap in the executable file as per the memory management schema chosen by the compiler itself, and then the kernel ensures that program is executed with this schema.: False

The kernel ensures that the MMU is setup before scheduling a process and then the CPU/MMU ensures that the address translation takes place.: True

When a process is executing, each virtual address is converted into physical address by the kernel directly.: False

The compiler generates address references for code/data/stack/heap in the executable file, depending on the MM architecture provided by CPU and kernel.: True

Question 16

Correct

Mark 0.50 out of 0.50

Mark statements True/False w.r.t. change of states of a process. Note that a statement is true only if the claim and argument both are true.

Reference: The process state diagram (and your understanding of how kernel code works). Note - the diagram does not show zombie state!

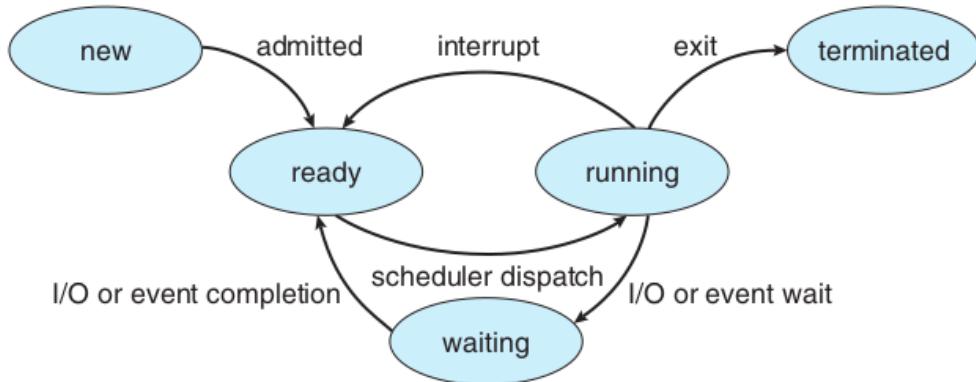


Figure 3.2 Diagram of process state.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	Every forked process has to go through ZOMBIE state, at least for a small duration.
<input type="radio"/>	<input checked="" type="radio"/>	A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first
<input checked="" type="radio"/>	<input type="radio"/>	Only a process in READY state is considered by scheduler
<input checked="" type="radio"/>	<input type="radio"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet
<input type="radio"/>	<input checked="" type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.

Every forked process has to go through ZOMBIE state, at least for a small duration.: True

A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first: False

Only a process in READY state is considered by scheduler: True

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet: True

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

Question 17

Partially correct

Mark 0.45 out of 0.50

Select Yes if the mentioned element should be a part of PCB

Select No otherwise.

Yes	No	
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Memory management information about that process
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Pointer to IDT
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/>	PID of Init
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	PID
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Function pointers to all system calls
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	EIP at the time of context switch
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Process state
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	List of opened files
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Process context
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Pointer to the parent process

Memory management information about that process: Yes

Pointer to IDT: No

PID of Init: No

PID: Yes

Function pointers to all system calls: No

EIP at the time of context switch: Yes

Process state: Yes

List of opened files: Yes

Process context: Yes

Pointer to the parent process: Yes

Question **18**

Correct

Mark 0.50 out of 0.50

Predict the output of the program given here.

Assume that all the path names for the programs are correct. For example "/usr/bin/echo" will actually run echo command.

Assume that there is no mixing of printf output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good

output

should be written as good output

--

```
main() {  
    int i;  
    i = fork();  
    if(i == 0)  
        execl("/usr/bin/echo", "/usr/bin/echo", "hi", 0);  
    else  
        wait(0);  
    fork();  
    execl("/usr/bin/echo", "/usr/bin/echo", "one", 0);  
}
```

Answer: hi one one



The correct answer is: hi one one

Question **19**

Correct

Mark 0.50 out of 0.50

How does the distinction between kernel mode and user mode function as a rudimentary form of protection (security) ?

Select one:

- a. It prohibits one process from accessing other process's memory
- b. It prohibits invocation of kernel code completely, if a user program is running
- c. It prohibits a user mode process from running privileged instructions✓
- d. It disallows hardware interrupts when a process is running

Your answer is correct.

The correct answer is: It prohibits a user mode process from running privileged instructions

Question **20**

Partially correct

Mark 0.33 out of 0.50

Which of the following parts of a C program do not have any corresponding machine code ?

- a. pointer dereference
- b. #directives✓
- c. typedefs✓
- d. global variables
- e. function calls
- f. expressions
- g. local variable declaration

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: #directives, typedefs, global variables

[◀ Quiz-1 Preparation questions](#)

Jump to...

[Quiz - 2 \(17 March 2023\) ►](#)

Started on Friday, 17 March 2023, 2:33 PM

State Finished

Completed on Friday, 17 March 2023, 4:54 PM

Time taken 2 hours 21 mins

Grade 6.37 out of 10.00 (63.73%)

Question 1

Incorrect

Mark 0.00 out of 0.50

The first instruction that runs when you do "make qemu" is

cli

from bootasm.S

Why?

- a. "cli" stands for clear screen and the screen should be cleared before OS boots.
- b. "cli" that is Command Line Interface needs to be enabled first
- c. "cli" clears all registers and makes them zero, so that processor is as good as "new"
- d. "cli" clears the pipeline of the CPU so that it is as good as "fresh" CPU
- e. "cli" enables interrupts, it is required because the kernel supports interrupts.
- f. "cli" disables interrupts. It is required because as of now there are no interrupt handlers available X
- g. It disables interrupts. It is required because the interrupt handlers of kernel are not yet installed.
- h. "cli" enables interrupts, it is required because the kernel must handle interrupts.

Your answer is incorrect.

The correct answer is: It disables interrupts. It is required because the interrupt handlers of kernel are not yet installed.

Question 2

Incorrect

Mark 0.00 out of 0.50

The struct buf has a sleeplock, and not a spinlock, because

- a. struct buf is used as a general purpose cache by kernel and cache operations take lot of time, so better to use sleeplock rather than spinlock
- b. sleeplock is preferable because it is used in interrupt context and spinlock can not be used in interrupt context
- c. It could be a spinlock, but xv6 has chosen sleeplock for purpose of demonstrating how to use a sleeplock.
- d. struct buf is used for disk I/O which takes lot of time, so sleeping/blocking is the only option available. X
- e. struct buf is used for disk I/O which takes lot of time, so sleeping/blocking is preferred to spinning/busy-wait for the desired buf.

Your answer is incorrect.

The correct answer is: struct buf is used for disk I/O which takes lot of time, so sleeping/blocking is preferred to spinning/busy-wait for the desired buf.

Question 3

Partially correct

Mark 0.25 out of 0.50

when is each of the following stacks allocated?

kernel stack of process

during fork() in allocproc() ✓

kernel stack for scheduler, on first processor

on the arrival of a hardware interrupt X

user stack of process

during exec() X

kernel stack for the scheduler, on other processors

in main()->startothers() ✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: kernel stack of process → during fork() in allocproc(), kernel stack for scheduler, on first processor → in entry.S, user stack of process → during fork() in copyuvml(), kernel stack for the scheduler, on other processors → in main()->startothers()

Question 4

Partially correct

Mark 0.56 out of 0.75

Mark statements as True/False w.r.t. ptable.lock

True	False	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	A process can sleep on ptable.lock if it can't acquire it.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	It is taken by one process but released by another process, running on same processor
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	ptable.lock protects the proc[] array and all struct proc in the array
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The swtch() in scheduler() is called without holding the ptable.lock when control jumps to it from sched()
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	the rule of "never block holding a spinlock" does not apply to ptable.lock in xv6
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	One sequence of function calls which takes and releases the ptable.lock is this: iderw->sleep, acquire(ptable.lock)->sched->swtch()->scheduler()->swtch()->yield(), release(ptable.lock)
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	ptable.lock can be held by different processes on different processors at the same time
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	ptable.lock is acquired but never released

A process can sleep on ptable.lock if it can't acquire it.: False

It is taken by one process but released by another process, running on same processor: True

ptable.lock protects the proc[] array and all struct proc in the array: True

The swtch() in scheduler() is called without holding the ptable.lock when control jumps to it from sched(): False

the rule of "never block holding a spinlock" does not apply to ptable.lock in xv6: True

One sequence of function calls which takes and releases the ptable.lock is this:

iderw->sleep, acquire(ptable.lock)->sched->swtch()->scheduler()->swtch()->yield(), release(ptable.lock): True

ptable.lock can be held by different processes on different processors at the same time: False

ptable.lock is acquired but never released: False

Question 5

Correct

Mark 0.50 out of 0.50

Consider the following command and its output:

```
$ ls -lht xv6.img kernel
-rw-rw-r-- 1 abhijit abhijit 4.9M Feb 15 11:09 xv6.img
-rwxrwxr-x 1 abhijit abhijit 209K Feb 15 11:09 kernel*
```

Following code in bootmain()

```
readseg((uchar*)elf, 4096, 0);
```

and following selected lines from Makefile

```
xv6.img: bootblock kernel
    dd if=/dev/zero of=xv6.img count=10000
    dd if=bootblock of=xv6.img conv=notrunc
    dd if=kernel of=xv6.img seek=1 conv=notrunc

kernel: $(OBJS) entry.o entryother initcode kernel.ld
    $(LD) $(LDFLAGS) -T kernel.ld -o kernel entry.o $(OBJS) -b binary initcode entryother
    $(OBJDUMP) -S kernel > kernel.asm
    $(OBJDUMP) -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$$/d' > kernel.sym
```

Also read the code of bootmain() in xv6 kernel.

Select the options that describe the meaning of these lines and their correlation.

- a. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(). ✓
- b. Althought the size of the kernel file is 209 Kb, only 4Kb out of it is the actual kernel code and remaining part is all zeroes.
- c. The bootmain() code does not read the kernel completely in memory
- d. Althought the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes. ✓
- e. readseg() reads first 4k bytes of kernel in memory ✓
- f. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is not read as it is user programs.
- g. The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files ✓
- h. The kernel.ld file contains instructions to the linker to link the kernel properly ✓
- i. The kernel.asm file is the final kernel file

Your answer is correct.

The correct answers are: The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(), readseg() reads first 4k bytes of kernel in memory, The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files, The kernel.ld file contains instructions to the linker to link the kernel properly, Althought the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.

Question 6

Partially correct

Mark 0.43 out of 0.75

code line, MMU setting: Match the line of xv6 code with the MMU setup employed

movw %ax, %gs	real mode	✗
jmp *%eax	protected mode with segmentation and 4 MB pages	✓
inb \$0x64,%al	real mode	✓
ljmp \$(SEG_KCODE<<3), \$start32	protected mode with only segmentation	✗
readseg((uchar*)elf, 4096, 0);	protected mode with only segmentation	✓
orl \$CRO_PE, %eax	real mode	✓
movl \$(V2P_WO(entrypgdir)), %eax	protected mode with segmentation and 4 MB pages	✗

The correct answer is: movw %ax, %gs → protected mode with only segmentation, jmp *%eax → protected mode with segmentation and 4 MB pages, inb \$0x64,%al → real mode, ljmp \$(SEG_KCODE<<3), \$start32 → real mode, readseg((uchar*)elf, 4096, 0); → protected mode with only segmentation, orl \$CRO_PE, %eax → real mode, movl \$(V2P_WO(entrypgdir)), %eax → protected mode with only segmentation

Question 7

Correct

Mark 0.50 out of 0.50

Which of the following is DONE by allocproc() ?

- a. setup kernel memory mappings for the process
- b. Select an UNUSED struct proc for use ✓
- c. allocate kernel stack for the process ✓
- d. ensure that the process starts in trapret()
- e. setup the contents of the trapframe of the process properly
- f. allocate PID to the process ✓
- g. ensure that the process starts in forkret() ✓
- h. setup the trapframe and context pointers appropriately ✓

The correct answers are: Select an UNUSED struct proc for use, allocate PID to the process, allocate kernel stack for the process, setup the trapframe and context pointers appropriately, ensure that the process starts in forkret()

Question 8

Partially correct

Mark 0.25 out of 0.50

Mark statements as True/False w.r.t. the creation of free page list in xv6.

True	False	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	kmem.use_lock is set to 1 after free page list is created, so that kmem.lock is taken before accessing kmem.freelist.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	The pointers that link the pages together are in the first 4 bytes of the pages themselves
<input checked="" type="radio"/>	<input checked="" type="radio"/>	if(kmem.use_lock) acquire(&kmem.lock); is not done when called from kinit1() because there is no need to take the lock when kinit1() is running because interrupts are disabled and only one processor is running
<input checked="" type="radio"/>	<input checked="" type="radio"/>	free page list is a singly circular linked list.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	if(kmem.use_lock) acquire(&kmem.lock); this "if" condition is true, when kinit2() runs because multi-processor support has been enabled by now.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	the kmem.lock is used by kfree() and kalloc() only.

kmem.use_lock is set to 1 after free page list is created, so that kmem.lock is taken before accessing kmem.freelist.: True

The pointers that link the pages together are in the first 4 bytes of the pages themselves: True

if(kmem.use_lock)

acquire(&kmem.lock);

is not done when called from kinit1() because there is no need to take the lock when kinit1() is running because interrupts are disabled and only one processor is running: True

free page list is a singly circular linked list.: False

if(kmem.use_lock)

acquire(&kmem.lock);

this "if" condition is true, when kinit2() runs because multi-processor support has been enabled by now.: False

the kmem.lock is used by kfree() and kalloc() only.: True

Question 9

Partially correct

Mark 0.55 out of 1.00

Given below is code of sleeplock in xv6.

```
// Long-term locks for processes
struct sleeplock {
    uint locked;          // Is the lock held?
    struct spinlock lk;  // spinlock protecting this sleep lock

    // For debugging:
    char *name;           // Name of lock.
    int pid;              // Process holding lock
};

void
acquiresleep(struct sleeplock *lk)
{
    acquire(&lk->lk);
    while (lk->locked) {
        sleep(lk, &lk->lk);
    }
    lk->locked = 1;
    lk->pid = myproc()->pid;
    release(&lk->lk);
}

void
releasesleep(struct sleeplock *lk)
{
    acquire(&lk->lk);
    lk->locked = 0;
    lk->pid = 0;
    wakeup(lk);
    release(&lk->lk);
}
```

Mark the statements as True/False w.r.t. this code.

True	False
<input checked="" type="radio"/> <input type="radio"/>	Wakeup() will wakeup the first process waiting for the lock
<input checked="" type="radio"/> <input type="radio"/>	<pre>acquire(&lk->lk); while (lk->locked) { sleep(lk, &lk->lk); } could also be written as acquire(&lk->lk); if (lk->locked) { sleep(lk, &lk->lk); }</pre>

True	False	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	A process has acquired the sleeplock when it comes out of sleep(): ✗
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	sleep() is called holding a spinlock. This could be avoided by releasing the lock before calling sleep() and acquiring it again after call to sleep(): ✗
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	sleep() is the function which blocks a process. ✓
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	The process which called acquiresleep() and then got blocked, is woken up by the timer interrupt ✓ it's woken up by another process which called releasesleep() and then wakeup()
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	the 'spinlock lk' protects 'locked' variable, but not the 'name' nor the 'pid' ✗
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	the 'spinlock lk' is needed in a sleeplock, because access to the sleeplock for locking/unlocking itself creates a critical section ✓
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	The spinlock lk->lk is held when the process comes out of sleep(): ✓
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	Sleeplock() will ensure that either the process gets the lock or the process gets blocked. ✓
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✗	All processes waiting for the sleeplock will have a race for aquiring lk->lk spinlock, because all are woken up ✓ wakeup() wakes up all processes, and they "thunder" to take the spinlock.

Wakeup() will wakeup the first process waiting for the lock: False

```
acquire(&lk->lk);
while (!lk->locked) {
    sleep(lk, &lk->lk);
}
```

could also be written as

```
acquire(&lk->lk);
if (!lk->locked) {
    sleep(lk, &lk->lk);
}: False
```

A process has acquired the sleeplock when it comes out of sleep(): False

sleep() is called holding a spinlock. This could be avoided by releasing the lock before calling sleep() and acquiring it again after call to sleep(): False

The process which called acquiresleep() and then got blocked, is woken up by the timer interrupt: True

The 'spinlock lk' protects 'locked' variable, but not the 'name' nor the 'pid': False

The 'spinlock lk' is needed in a sleeplock, because access to the sleeplock for locking/unlocking itself creates a critical section: True

The spinlock lk->lk is held when the process comes out of sleep(): True

Sleeplock() will ensure that either the process gets the lock or the process gets blocked.: True

All processes waiting for the sleeplock will have a race for aquiring lk->lk spinlock, because all are woken up: True

Question 10

Partially correct

Mark 0.50 out of 1.00

Mark the statements as True/False w.r.t. swtch()

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	p->context used in scheduler()->swtch() was Generally set when the process was interrupted earlier, and came via sched()->swtch()
<input checked="" type="radio"/>	<input type="radio"/>	swtch() is written in assembly language, because it violates calling convention, by changing the stack itself.
<input type="radio"/>	<input checked="" type="radio"/>	swtch() is written in assembly language because it violates the calling convention by pushing parameters on the stack on its own.
<input type="radio"/>	<input checked="" type="radio"/>	swtch() called from scheduler() changes the stack from the process's kernel stack to the scheduler's kernel stack.
<input checked="" type="radio"/>	<input type="radio"/>	swtch() is called only from sched() or scheduler()
<input type="radio"/>	<input checked="" type="radio"/>	switch stores the old context on new stack, and restores new context from old stack.
<input type="radio"/>	<input checked="" type="radio"/>	sched() is the only place when p->context is set
<input type="radio"/>	<input checked="" type="radio"/>	movl %esp, (%eax) means, *(c->scheduler) = contents of esp When swtch() is called from scheduler()
<input checked="" type="radio"/>	<input type="radio"/>	push in swtch() happens on old stack, while pop happens from new stack
<input checked="" type="radio"/>	<input type="radio"/>	swtch() changes the context from "old" to "new"

p->context used in scheduler()->swtch() was **Generally** set when the process was interrupted earlier, and came via sched()->swtch(): True
 swtch() is written in assembly language, because it violates calling convention, by changing the stack itself.: True

swtch() is written in assembly language because it violates the calling convention by pushing parameters on the stack on its own.: False
 swtch() called from scheduler() changes the stack from the process's kernel stack to the scheduler's kernel stack.: False

swtch() is called only from sched() or scheduler(): True

switch stores the old context on new stack, and restores new context from old stack.: False

sched() is the only place when p->context is set: False

movl %esp, (%eax)

means, *(c->scheduler) = contents of esp

When swtch() is called from scheduler(): False

push in swtch() happens on old stack, while pop happens from new stack: True

swtch() changes the context from "old" to "new": True

Question **11**

Correct

Mark 0.25 out of 0.25

Select the odd one out

- a. Kernel stack of new process to kernel stack of scheduler ✓
- b. Kernel stack of new process to Process stack of new process
- c. Kernel stack of running process to kernel stack of scheduler
- d. Process stack of running process to kernel stack of running process
- e. Kernel stack of scheduler to kernel stack of new process

The correct answer is: Kernel stack of new process to kernel stack of scheduler

Question **12**

Correct

Mark 0.25 out of 0.25

The variable 'end' used as argument to kinit1 has the value

- a. 80102da0
- b. 81000000
- c. 80000000
- d. 80110000
- e. 801154a8 ✓
- f. 8010a48c

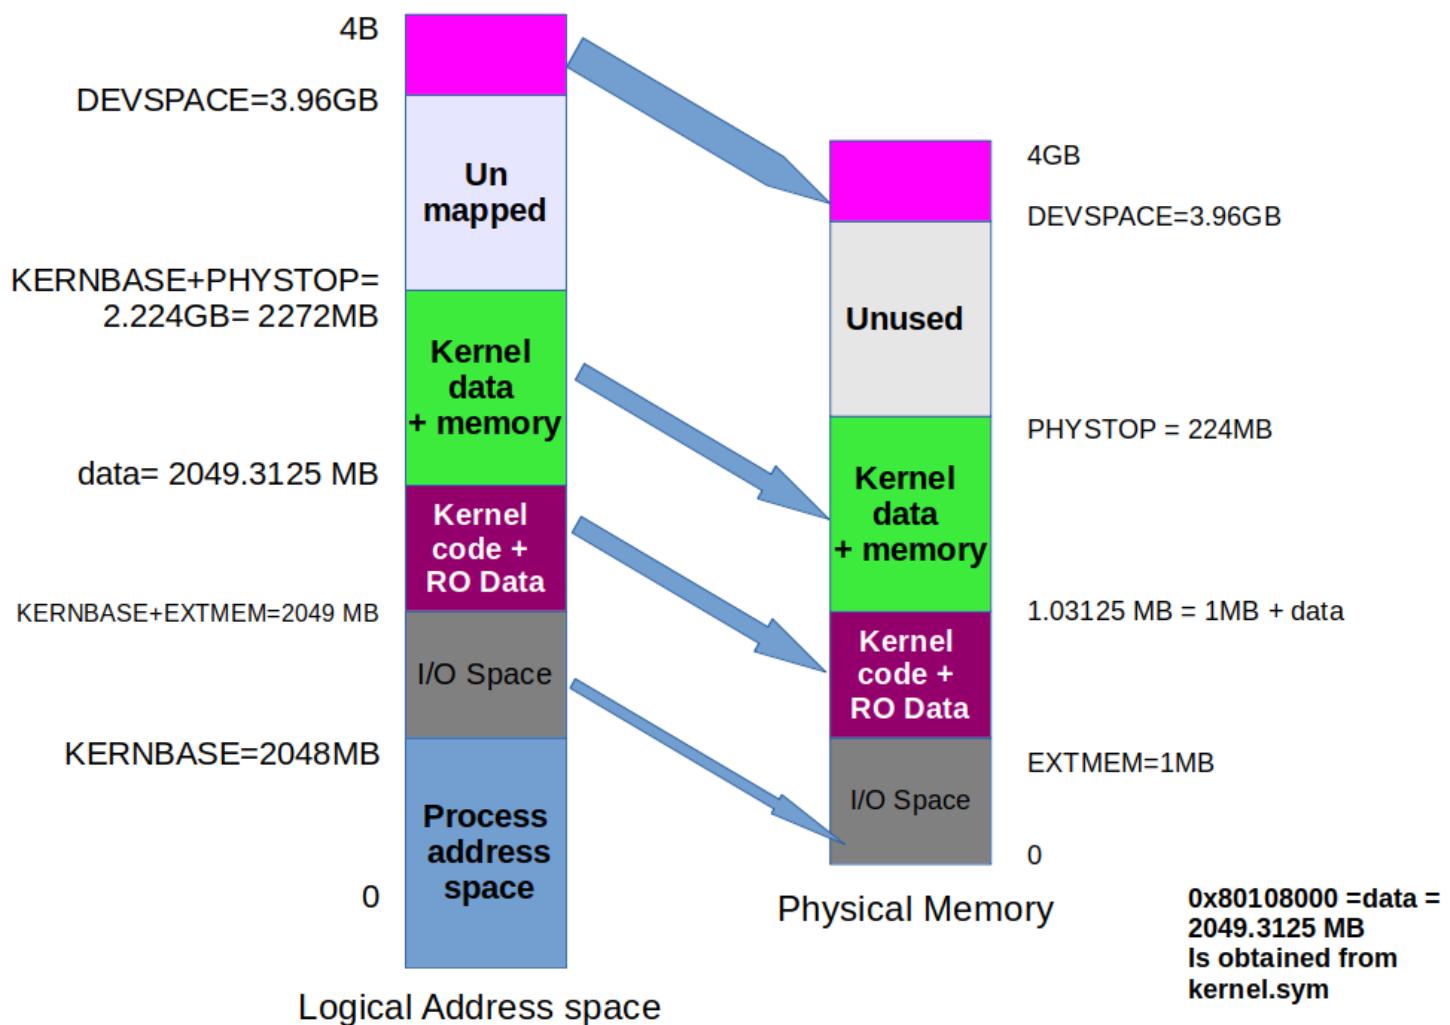
The correct answer is: 801154a8

Question 13

Partially correct

Mark 0.36 out of 0.50

With respect to this diagram, mark statements as True/False.



True False

<input checked="" type="radio"/>	<input type="radio"/>	This diagram only shows the absolutely defined virtual->physical mappings, not the mappings defined at run time by kernel.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/>	"Kernel data + memory" on right side, here refers to the region from which pages are allocated to the kernel and process both.	<input checked="" type="checkbox"/> "Kernel data + memory" on LEFT side, here refers to the virtual addresses of kernel used at run time.
<input checked="" type="radio"/>	<input type="radio"/>	The kernel virtual addresses start from KERNLINK = KERNBASE + EXTMEM	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/>	The process's pages are mapped into physical memory from 1.03125 MB to PHYSSTOP.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/>	The kernel file, after compilation, has maximum virtual address up to "data" as shown in the diagram, which is equal to "end" variable	<input checked="" type="checkbox"/>

True **False**

<input checked="" type="radio"/>	<input checked="" type="radio"/>	When bootloader loads the kernel, then physical memory from EXTMEM upto EXTMEM + data is occupied.	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	PHYSTOP can be changed , but that needs kernel recompilation and re-execution.	✗

This diagram only shows the absolutely defined virtual->physical mappings, not the mappings defined at run time by kernel.: True
"Kernel data + memory" on right side, here refers to the region from which pages are allocated to the kernel and process both.: True

The kernel virtual addresses start from KERNLINK = KERNBASE + EXTMEM: True

The process's pages are mapped into physical memory from 1.03125 MB to PHYSTOP.: True

The kernel file, after compilation, has maximum virtual address up to "data" as shown in the diagram, which is equal to "end" variable: True

When bootloader loads the kernel, then physical memory from EXTMEM upto EXTMEM + data is occupied.: True

PHYSTOP can be changed , but that needs kernel recompilation and re-execution.: True

Question 14

Correct

Mark 0.25 out of 0.25

Which of the following is not a task of the code of swtch() function

- a. Save the return value of the old context code ✓
- b. Load the new context
- c. Change the kernel stack location ✓
- d. Save the old context
- e. Switch stacks
- f. Jump to next context EIP

The correct answers are: Save the return value of the old context code, Change the kernel stack location

Question **15**

Correct

Mark 0.50 out of 0.50

We often use terms like "swtch() changes stack from process's kernel stack to scheduler's stack", or "the values are pushed on stack", or "the stack is initialized to the new page", etc. while discussing xv6 on x86.

Which of the following most accurately describes the meaning of "stack" in such sentences?

- a. The stack variable used in the program being discussed
- b. The region of memory where the kernel remembers all the function calls made
- c. The region of memory allocated by kernel for storing the parameters of functions
- d. The ss:esp pair ✓
- e. The "stack" variable declared in "stack.S" in xv6
- f. the region of memory which is currently used as stack by processor
- g. The stack segment

Your answer is correct.

The correct answer is: The ss:esp pair

Question 16

Partially correct

Mark 0.31 out of 0.50

Mark the statements as True/False, with respect to the use of the variable "chan" in struct proc.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The value of 'chan' is changed only in sleep() ✓
<input type="radio"/>	<input checked="" type="radio"/>	when chan is NULL, the 'state' in proc must be RUNNABLE. ✗
<input checked="" type="radio"/>	<input type="radio"/>	chan stores the address of the variable, representing a condition, for which the process is waiting. ✓
<input checked="" type="radio"/>	<input type="radio"/>	'chan' is used only by the sleep() and wakeup1() functions. ✓
<input type="radio"/>	<input checked="" type="radio"/>	chan is the head pointer to a linked list of processes, waiting for a particular event to occur ✗
<input checked="" type="radio"/>	<input type="radio"/>	When chan is not NULL, the 'state' in struct proc must be SLEEPING ✓
<input checked="" type="radio"/>	<input type="radio"/>	in xv6, the address of an appropriate variable is used as a "condition" for a waiting process. ✓
<input type="radio"/>	<input checked="" type="radio"/>	Changing the state of a process automatically changes the value of 'chan' ✗

The value of 'chan' is changed only in sleep(): True

when chan is NULL, the 'state' in proc must be RUNNABLE.: False

chan stores the address of the variable, representing a condition, for which the process is waiting.: True

'chan' is used only by the sleep() and wakeup1() functions.: True

chan is the head pointer to a linked list of processes, waiting for a particular event to occur: False

When chan is not NULL, the 'state' in struct proc must be SLEEPING: True

in xv6, the address of an appropriate variable is used as a "condition" for a waiting process.: True

Changing the state of a process automatically changes the value of 'chan': False

Question 17

Correct

Mark 0.25 out of 0.25

Match function with its meaning

iderw	Issue a disk read/write for a buffer, block the issuing process	✓
idewait	Wait for disc controller to be ready	✓
idestart	tell disc controller to start I/O for the first buffer on idequeue	✓
ideinit	Initialize the disc controller	✓
ideintr	disk interrupt handler, transfer data from controller to buffer, wake up processes waiting for this buffer, start I/O for next buffer	✓

Your answer is correct.

The correct answer is: iderw → Issue a disk read/write for a buffer, block the issuing process, idewait → Wait for disc controller to be ready, idestart → tell disc controller to start I/O for the first buffer on idequeue, ideinit → Initialize the disc controller, ideintr → disk interrupt handler, transfer data from controller to buffer, wake up processes waiting for this buffer, start I/O for next buffer

Question 18

Correct

Mark 0.25 out of 0.25

Why is there a call to kinit2? Why is it not merged with knit1?

- a. Because there is a limit on the values that the arguments to knit1() can take.
- b. call to seginit() makes it possible to actually use PHYSTOP in argument to kinit2()
- c. When kinit1() is called there is a need for few page frames, but later kinit2() is called to serve need of more page frames
- d. knit2 refers to virtual addresses beyond 4MB, which are not mapped before kalloc() is called✓

The correct answer is: knit2 refers to virtual addresses beyond 4MB, which are not mapped before kalloc() is called

Question **19**

Correct

Mark 0.25 out of 0.25

Which of the following call sequence is impossible in xv6?

- a. Process 1: write() -> sys_write()-> file_write() -- timer interrupt -> trap() -> yield() -> sched() -> switch() (jumps to)-> scheduler() -> swtch() (jumps to)-> Process 2 (return call sequence) sched() -> yield() -> trap-> user-code
- b. Process 1: write() -> sys_write()-> file_write() -> writei() -> bread() -> bget() -> iderw() -> sleep() -> sched() -> switch() (jumps to)-> scheduler() ->switch()(jumps to)-> Process 2 (return call sequence) sched() -> yield() -> trap-> user-code
- c. Process 1: timer interrupt -> trap() -> yield() -> sched() -> switch() -> scheduler()-> Process 2 runs -> write -> sys_write() -> trap()-> ... ✓

Your answer is correct.

The correct answer is: Process 1: timer interrupt -> trap() -> yield() -> sched() -> switch() -> scheduler()-> Process 2 runs -> write -> sys_write() -> trap()-> ...

Question **20**

Partially correct

Mark 0.17 out of 0.50

Mark statements as True/False, w.r.t. the given diagram

Started on Thursday, 2 February 2023, 9:01 PM

State Finished

Completed on Thursday, 2 February 2023, 10:55 PM

Time taken 1 hour 53 mins

Grade 16.46 out of 20.00 (82.28%)

Question 1

Complete

Mark 1.00 out of 1.00

Which of the following are NOT a part of job of a typical compiler?

- a. Check the program for logical errors
- b. Convert high level language code to machine code
- c. Invoke the linker to link the function calls with their code, extern globals with their declaration
- d. Suggest alternative pieces of code that can be written
- e. Check the program for syntactical errors
- f. Process the # directives in a C program

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question 2

Complete

Mark 1.00 out of 1.00

Consider the following code and MAP the file to which each fd points at the end of the code.

```
int main(int argc, char *argv[]) {  
    int fd1, fd2 = 1, fd3 = 1, fd4 = 1;  
  
    fd1 = open("/tmp/1", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);  
    fd2 = open("/tmp/2", O_RDONLY);  
    fd3 = open("/tmp/3", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);  
    close(0);  
    close(1);  
    dup(fd2);  
    dup(fd3);  
    close(fd3);  
    dup2(fd2, fd4);  
    printf("%d %d %d %d\n", fd1, fd2, fd3, fd4);  
    return 0;  
}
```

fd4	/tmp/2
fd2	/tmp/2
1	/tmp/3
fd1	/tmp/1
fd3	closed
0	/tmp/2
2	stderr

The correct answer is: fd4 → /tmp/2, fd2 → /tmp/2, 1 → /tmp/3, fd1 → /tmp/1, fd3 → closed, 0 → /tmp/2, 2 → stderr

Question 3

Complete

Mark 0.75 out of 1.00

Select the sequence of events that are NOT possible, assuming an interruptible kernel code

Select one or more:

- a. P1 running
keyboard hardware interrupt
keyboard interrupt handler running
interrupt handler returns
P1 running
P1 makes system call
system call returns
P1 running
timer interrupt
scheduler
P2 running
- b. P1 running
P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P3 running
Hardware interrupt
Interrupt unblocks P1
Interrupt returns
P3 running
Timer interrupt
Scheduler
P1 running
- c. P1 running
P1 makes system call
system call returns
P1 running
timer interrupt
Scheduler running
P2 running
- d. P1 running
P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again
- e. P1 running
P1 makes system call
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again

- f. P1 running
P1 makes system call
timer interrupt
Scheduler
P2 running
timer interrupt
Scheuler
P1 running
P1's system call return

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question 4

Complete

Mark 0.00 out of 1.00

Select all the correct statements about named pipes and ordinary(unnamed) pipe

Select one or more:

- a. named pipes can be used between multiple processes but ordinary pipes can not be used
- b. named pipe can be used between any processes
- c. both named and unnamed pipes require some kind of agreed protocol to be effectively used among multiple processes
- d. named pipes are more efficient than ordinary pipes
- e. ordinary pipe can only be used between related processes
- f. a named pipe exists as a file on the file system
- g. named pipe exists even if the processes using it do exit()

The correct answers are: ordinary pipe can only be used between related processes, named pipe can be used between any processes, a named pipe exists as a file on the file system, named pipe exists even if the processes using it do exit(), both named and unnamed pipes require some kind of agreed protocol to be effectively used among multiple processes

Question 5

Complete

Mark 4.50 out of 5.00

Following code claims to implement the command

/bin/ls -l | /usr/bin/head -3 | /usr/bin/tail -1

Fill in the blanks to make the code work.

Note: Do not include space in writing any option. x[1][2] should be written without any space, and so is the case with [1] or [2]. Pay attention to exact syntax and do not write any extra character like ';' or = etc.

```
int main(int argc, char *argv[]) {  
    int pid1, pid2;  
    int pfd[ 2 ][2];  
  
    pipe( pfd[0] );  
    pid1 = [ fork() ];  
    if(pid1 != 0) {  
        close(pfd[0][ 0 ]);  
        close( 1 );  
        dup( pfd[0][1] );  
        execl("/bin/ls", "/bin/ls", " -l ", NULL);  
    }  
    pipe( pfd[1] );  
    pid2 = [ fork() ];  
    if(pid2 == 0) {  
        close( pfd[0][1] );  
        close(0);  
        dup( pfd[1][0] );  
        close(pfd[1][ 0 ]);  
        close( 1 );  
        dup( pfd[1][1] );  
        execl("/usr/bin/head", "/usr/bin/head", " -3 ", NULL);  
    } else {  
        close(pfd[1][1]);  
        close( 0 );  
        dup( pfd[1][0] );  
        close(pfd[0][0]);  
    }  
}
```

```

execl("/usr/bin/tail", "/usr/bin/tail", " -1 ", NULL);
}
}

```

Question 6

Complete

Mark 1.00 out of 1.00

Select the compiler's view of the process's address space, for each of the following MMU schemes:

(Assume that each scheme, e.g. paging/segmentation/etc is effectively utilised)

- | | |
|---------------------------|---|
| Paging | one continuous chunk |
| Segmentation, then paging | many continuous chunks of variable size |
| Segmentation | many continuous chunks of variable size |
| Relocation + Limit | one continuous chunk |

The correct answer is: Paging → one continuous chunk, Segmentation, then paging → many continuous chunks of variable size, Segmentation → many continuous chunks of variable size, Relocation + Limit → one continuous chunk

Question 7

Complete

Mark 1.40 out of 2.00

Match the elements of C program to their place in memory

- | | |
|-------------------------|------------------|
| Local Variables | Stack |
| Allocated Memory | Heap |
| Global Static variables | Data |
| Global variables | Stack |
| Code of main() | Code |
| Local Static variables | Data |
| Function code | Data |
| Arguments | Stack |
| #define MACROS | No memory needed |
| #include files | No memory needed |

The correct answer is: Local Variables → Stack, Allocated Memory → Heap, Global Static variables → Data, Global variables → Data, Code of main() → Code, Local Static variables → Data, Function code → Code, Arguments → Stack, #define MACROS → No Memory needed, #include files → No memory needed

Question 8

Complete

Mark 0.63 out of 1.00

Consider the image given below, which explains how paging works.

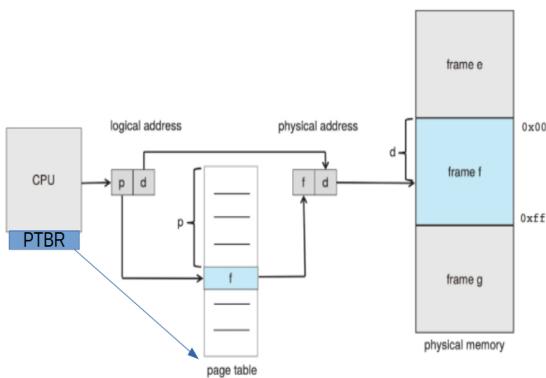


Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

True False

- The physical address may not be of the same size (in bits) as the logical address
- The locating of the page table using PTBR also involves paging translation
- Size of page table is always determined by the size of RAM
- The page table is indexed using frame number
- The PTBR is present in the CPU as a register
- The page table is indexed using page number
- Maximum Size of page table is determined by number of bits used for page number
- The page table is itself present in Physical memory

The physical address may not be of the same size (in bits) as the logical address: True

The locating of the page table using PTBR also involves paging translation: False

Size of page table is always determined by the size of RAM: False

The page table is indexed using frame number: False

The PTBR is present in the CPU as a register: True

The page table is indexed using page number: True

Maximum Size of page table is determined by number of bits used for page number: True

The page table is itself present in Physical memory: True

Question 9

Complete

Mark 0.67 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

(Note: non-interruptible kernel code means, if the kernel code is executing, then interrupts will be disabled).

Note: A possible sequence may have some missing steps in between. An impossible sequence will have n and n+1th steps such that n+1th step can not follow n'th step.

Select one or more:

- a. P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P3 running

Hardware interrupt

Interrupt unblocks P1

Interrupt returns

P3 running

Timer interrupt

Scheduler

P1 running

- b. P1 running

P1 makes system call

system call returns

P1 running

timer interrupt

Scheduler running

P2 running

- c. P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

- d. P1 running

keyboard hardware interrupt

keyboard interrupt handler running

interrupt handler returns

P1 running

P1 makes system call

system call returns

P1 running

timer interrupt

scheduler

P2 running

- e. P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running
timer interrupt
Scheuler
P1 running
P1's system call return

f.

P1 running
P1 makes system call
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again, P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question 10

Complete

Mark 1.00 out of 1.00

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

\$ ls ./tmp/asdfksdf >/tmp/ddd 2>&1

Program 1

```
int main(int argc, char *argv[]) {  
    int fd, n, i;  
    char buf[128];  
  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    close(1);  
    dup(fd);  
    close(2);  
    dup(fd);  
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);  
}
```

Program 2

```
int main(int argc, char *argv[]) {  
    int fd, n, i;  
    char buf[128];  
  
    close(1);  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    close(2);  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);  
}
```

Select all the correct statements about the programs

Select one or more:

- a. Both programs are correct
- b. Program 2 is correct for > /tmp/ddd but not for 2>&1
- c. Program 1 is correct for > /tmp/ddd but not for 2>&1
- d. Program 2 ensures 2>&1 and does not ensure > /tmp/ddd
- e. Program 2 makes sure that there is one file offset used for '2' and '1'
- f. Program 1 ensures 2>&1 and does not ensure > /tmp/ddd
- g. Only Program 2 is correct
- h. Program 1 does 1>&2
- i. Program 1 makes sure that there is one file offset used for '2' and '1'
- j. Program 2 does 1>&2
- k. Both program 1 and 2 are incorrect
- l. Only Program 1 is correct

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'

Question **11**

Complete

Mark 0.71 out of 1.00

Order the events that occur on a timer interrupt:

Jump to a code pointed by IDT

2

Save the context of the currently running process

3

Execute the code of the new process

6

Jump to scheduler code

4

Change to kernel stack of currently running process

1

Select another process for execution

5

Set the context of the new process

7

The correct answer is: Jump to a code pointed by IDT → 2, Save the context of the currently running process → 3, Execute the code of the new process → 7, Jump to scheduler code → 4, Change to kernel stack of currently running process → 1, Select another process for execution → 5, Set the context of the new process → 6

Question **12**

Complete

Mark 0.80 out of 1.00

Select all the correct statements about zombie processes

Select one or more:

- a. A process becomes zombie when its parent finishes
- b. A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it
- c. If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent
- d. Zombie processes are harmless even if OS is up for long time
- e. A process can become zombie if it finishes, but the parent has finished before it
- f. A zombie process occupies space in OS data structures
- g. init() typically keeps calling wait() for zombie processes to get cleaned up
- h. A zombie process remains zombie forever, as there is no way to clean it up

The correct answers are: A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it, A process can become zombie if it finishes, but the parent has finished before it, A zombie process occupies space in OS data structures, If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent, init() typically keeps calling wait() for zombie processes to get cleaned up

Question 13

Complete

Mark 1.00 out of 1.00

Select the state that is not possible after the given state, for a process:

New: RunningReady : WaitingRunning: : None of theseWaiting: Running**Question 14**

Complete

Mark 1.00 out of 1.00

Select the order in which the various stages of a compiler execute.

Pre-processing 1Intermediate code generation 3Linking 4Syntactical Analysis 2Loading does not exist

The correct answer is: Pre-processing → 1, Intermediate code generation → 3, Linking → 4, Syntactical Analysis → 2, Loading → does not exist

Question 15

Complete

Mark 1.00 out of 1.00

A process blocks itself means

- a. The application code calls the scheduler
- b. The kernel code of system call calls scheduler
- c. The kernel code of an interrupt handler, moves the process to a waiting queue and calls scheduler
- d. The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler

The correct answer is: The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler

Jump to...

[Random Quiz 4 : Scheduling, signals, segmentation, paging, compilation, process-state ►](#)

Started on Thursday, 16 February 2023, 9:00 PM

State Finished

Completed on Thursday, 16 February 2023, 10:43 PM

Time taken 1 hour 43 mins

Grade 13.00 out of 15.00 (86.68%)

Question 1

Correct

Mark 1.00 out of 1.00

Which of the following statements is false ?

Select one:

- a. A process scheduling algorithm is preemptive if the CPU can be forcibly removed from a process.
- b. Time sharing systems generally use preemptive CPU scheduling.
- c. Real time systems generally use non preemptive CPU scheduling. ✓
- d. Response time is more predictable in preemptive systems than in non preemptive systems.

Your answer is correct.

The correct answer is: Real time systems generally use non preemptive CPU scheduling.

Question 2

Correct

Mark 1.00 out of 1.00

Order the sequence of events, in scheduling process P1 after process P0

timer interrupt occurs

 2 ✓

Process P1 is running

 6 ✓

context of P0 is saved in P0's PCB

 3 ✓

Control is passed to P1

 5 ✓

context of P1 is loaded from P1's PCB

 4 ✓

Process P0 is running

 1 ✓

Your answer is correct.

The correct answer is: timer interrupt occurs → 2, Process P1 is running → 6, context of P0 is saved in P0's PCB → 3, Control is passed to P1 → 5, context of P1 is loaded from P1's PCB → 4, Process P0 is running → 1

Question 3

Correct

Mark 1.00 out of 1.00

Select the compiler's view of the process's address space, for each of the following MMU schemes:

(Assume that each scheme,e.g. paging/segmentation/etc is effectively utilised)

Paging

 one continuous chunk ✓

Relocation + Limit

 one continuous chunk ✓

Segmentation

 many continuous chunks of variable size ✓

Segmentation, then paging

 many continuous chunks of variable size ✓

Your answer is correct.

The correct answer is: Paging → one continuous chunk, Relocation + Limit → one continuous chunk, Segmentation → many continuous chunks of variable size, Segmentation, then paging → many continuous chunks of variable size

Question 4

Correct

Mark 1.00 out of 1.00

Match the names of PCB structures with kernel

- | | | |
|-------|--------------------|---|
| linux | struct task_struct | ✓ |
| xv6 | struct proc | ✓ |

The correct answer is: linux → struct task_struct, xv6 → struct proc

Question 5

Correct

Mark 1.00 out of 1.00

Select the state that is not possible after the given state, for a process:

- | | | |
|----------|---------------|---|
| New: | Running | ✓ |
| Ready : | Waiting | ✓ |
| Running: | None of these | ✓ |
| Waiting: | Running | ✓ |

Question 6

Partially correct

Mark 0.47 out of 1.00

Select all the correct statements about zombie processes

Select one or more:

- a. A zombie process occupies space in OS data structures
- b. init() typically keeps calling wait() for zombie processes to get cleaned up ✓
- c. If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent ✓
- d. A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it ✓
- e. Zombie processes are harmless even if OS is up for long time ✗
- f. A process can become zombie if it finishes, but the parent has finished before it ✓
- g. A process becomes zombie when its parent finishes
- h. A zombie process remains zombie forever, as there is no way to clean it up

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it, A process can become zombie if it finishes, but the parent has finished before it, A zombie process occupies space in OS data structures, If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent, init() typically keeps calling wait() for zombie processes to get cleaned up

Question 7

Partially correct

Mark 0.86 out of 1.00

Mark True/False

Statements about scheduling and scheduling algorithms

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	xv6 code does not care about Processor Affinity
<input type="radio"/>	<input checked="" type="radio"/>	Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.
<input checked="" type="radio"/>	<input type="radio"/>	Response time will be quite poor on non-interruptible kernels
<input checked="" type="radio"/>	<input type="radio"/>	Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.
<input checked="" type="radio"/>	<input type="radio"/>	Processor Affinity refers to memory accesses of a process being stored on cache of that processor
<input checked="" type="radio"/>	<input type="radio"/>	A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.
<input type="radio"/>	<input checked="" type="radio"/>	On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread It's the negation of this. Time NOT spent in idle thread.

xv6 code does not care about Processor Affinity: True

Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.: True

Response time will be quite poor on non-interruptible kernels: True

Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.: True

Processor Affinity refers to memory accesses of a process being stored on cache of that processor: True

A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.: True

On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread: False

Question 8

Correct

Mark 1.00 out of 1.00

Mark whether the concept is related to scheduling or not.

Yes	No	
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	file-table
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	context-switch
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	timer interrupt
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	ready-queue
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	runnable process

file-table: No

context-switch: Yes

timer interrupt: Yes

ready-queue: Yes

runnable process: Yes

Question 9

Partially correct

Mark 0.50 out of 1.00

Select all the correct statements about signals

Select one or more:

- a. Signals are delivered to a process by another process
- b. The signal handler code runs in user mode of CPU
- c. A signal handler can be invoked asynchronously or synchronously depending on signal type ✓
- d. SIGKILL definitely kills a process because it can't be caught or ignored, and its default action terminates the process ✓
- e. Signal handlers once replaced can't be restored
- f. The signal handler code runs in kernel mode of CPU ✗
- g. Signals are delivered to a process by kernel ✓
- h. SIGKILL definitely kills a process because its code runs in kernel mode of CPU

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: Signals are delivered to a process by kernel, A signal handler can be invoked asynchronously or synchronously depending on signal type, The signal handler code runs in user mode of CPU, SIGKILL definitely kills a process because it can't be caught or ignored, and its default action terminates the process

Question 10

Correct

Mark 1.00 out of 1.00

Map each signal with its meaning

SIGPIPE	Broken Pipe	✓
SIGCHLD	Child Stopped or Terminated	✓
SIGUSR1	User Defined Signal	✓
SIGALRM	Timer Signal from alarm()	✓
SIGSEGV	Invalid Memory Reference	✓

The correct answer is: SIGPIPE → Broken Pipe, SIGCHLD → Child Stopped or Terminated, SIGUSR1 → User Defined Signal, SIGALRM → Timer Signal from alarm(), SIGSEGV → Invalid Memory Reference

Question 11

Partially correct

Mark 0.80 out of 1.00

Mark statements True/False w.r.t. change of states of a process. Note that a statement is true only if the claim and argument both are true.

Reference: The process state diagram (and your understanding of how kernel code works). Note - the diagram does not show zombie state!

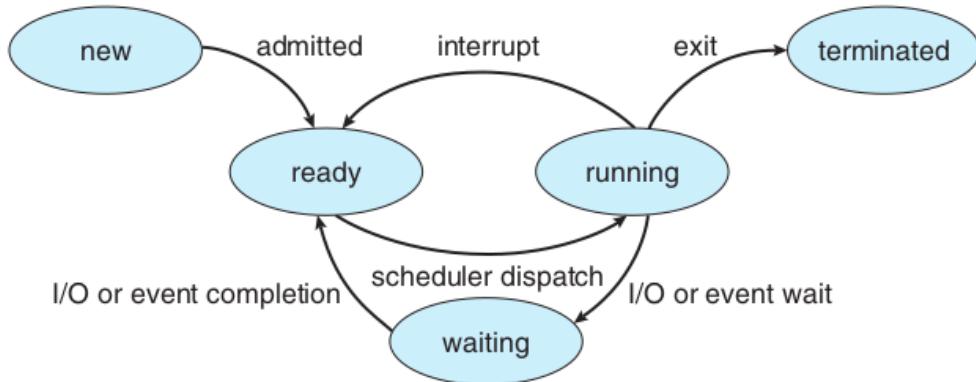


Figure 3.2 Diagram of process state.

True	False	
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Every forked process has to go through ZOMBIE state, at least for a small duration.
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Only a process in READY state is considered by scheduler
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/>	A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.

Every forked process has to go through ZOMBIE state, at least for a small duration.: True

Only a process in READY state is considered by scheduler: True

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet: True

A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first: False

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

Question **12**

Correct

Mark 1.00 out of 1.00

Which of the following parts of a C program do not have any corresponding machine code ?

- a. pointer dereference
- b. function calls
- c. global variables ✓
- d. local variable declaration
- e. #directives ✓
- f. typedefs ✓
- g. expressions

Your answer is correct.

The correct answers are: #directives, typedefs, global variables

Question 13

Partially correct

Mark 1.38 out of 2.00

Select all the correct statements about the state of a process.

- a. A process in ready state is ready to be scheduled ✓
- b. A waiting process starts running after the wait is over
- c. A process changes from running to ready state on a timer interrupt or any I/O wait
- d. A process waiting for I/O completion is typically woken up by the particular interrupt handler code ✓
- e. Processes in the ready queue are in the ready state ✓
- f. A process changes from running to ready state on a timer interrupt
- g. It is not maintained in the data structures by kernel, it is only for conceptual understanding of programmers
- h. A process waiting for any condition is woken up by another process only ✗
- i. A running process may terminate, or go to wait or become ready again ✓
- j. A process can self-terminate only when it's running ✓
- k. Typically, it's represented as a number in the PCB ✓
- l. A process that is running is not on the ready queue ✓
- m. Changing from running state to waiting state results in "giving up the CPU" ✓
- n. A process in ready state is ready to receive interrupts

Your answer is partially correct.

You have correctly selected 8.

The correct answers are: Typically, it's represented as a number in the PCB, A process in ready state is ready to be scheduled, Processes in the ready queue are in the ready state, A process that is running is not on the ready queue, A running process may terminate, or go to wait or become ready again, A process changes from running to ready state on a timer interrupt, Changing from running state to waiting state results in "giving up the CPU", A process can self-terminate only when it's running, A process waiting for I/O completion is typically woken up by the particular interrupt handler code

Question **14**

Correct

Mark 1.00 out of 1.00

Which of the following are NOT a part of job of a typical compiler?

- a. Invoke the linker to link the function calls with their code, extern globals with their declaration
- b. Check the program for logical errors ✓
- c. Convert high level language code to machine code
- d. Check the program for syntactical errors
- e. Suggest alternative pieces of code that can be written ✓
- f. Process the # directives in a C program

Your answer is correct.

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

◀ Random Quiz - 3 (processes, memory management, event driven kernel), compilation-linking-loading, ipc-pipes

Jump to...

Random Quiz - 5: xv6 make, bootloader, interrupt handling, memory management ►

Started on Thursday, 9 March 2023, 6:20 PM

State Finished

Completed on Thursday, 9 March 2023, 7:30 PM

Time taken 1 hour 10 mins

Overdue 14 mins

Grade **5.55** out of 10.00 (**55.53%**)

Question 1

Partially correct

Mark 0.18 out of 1.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only
Mark statements True or False

True	False	
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	xv6 uses physical memory upto 224 MB only
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The stack allocated in entry.S is used as stack for scheduler's context for first processor
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The free page-frame are created out of nearly 222 MB
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	PHYSTOP can be increased to some extent, simply by editing memlayout.h
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The process's address space gets mapped on frames, obtained from ~2MB:224MB range
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The kernel code and data take up less than 2 MB space
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir

xv6 uses physical memory upto 224 MB only: True

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The free page-frame are created out of nearly 222 MB: True

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

The kernel code and data take up less than 2 MB space: True

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

Question **2**

Correct

Mark 1.00 out of 1.00

What's the trapframe in xv6?

- a. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S ✓
- b. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only
- c. The IDT table
- d. A frame of memory that contains all the trap handler's addresses
- e. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- f. A frame of memory that contains all the trap handler code
- g. A frame of memory that contains all the trap handler code's function pointers

Your answer is correct.

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question 3

Partially correct

Mark 0.75 out of 1.00

Select the correct statements about interrupt handling in xv6 code

- a. The trapframe pointer in struct proc, points to a location on process's kernel stack ✓
- b. The CS and EIP are changed only after pushing user code's SS,ESP on stack ✓
- c. On any interrupt/syscall/exception the control first jumps in trapasm.S
- d. All the 256 entries in the IDT are filled in xv6 code ✓
- e. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- f. Before going to alltraps, the kernel stack contains upto 5 entries.
- g. The trapframe pointer in struct proc, points to a location on user stack
- h. The function trap() is the called only in case of hardware interrupt
- i. The CS and EIP are changed immediately (as the first thing) on a hardware interrupt
- j. xv6 uses the 0x64th entry in IDT for system calls
- k. xv6 uses the 64th entry in IDT for system calls ✓
- l. On any interrupt/syscall/exception the control first jumps in vectors.S ✓
- m. The function trap() is the called even if any of the hardware interrupt/system-call/exception occurs ✓

Your answer is partially correct.

You have correctly selected 6.

The correct answers are: All the 256 entries in the IDT are filled in xv6 code, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on process's kernel stack, The function trap() is the called even if any of the hardware interrupt/system-call/exception occurs, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question 4

Partially correct

Mark 0.55 out of 1.00

Consider the following command and its output:

```
$ ls -lht xv6.img kernel
-rw-rw-r-- 1 abhijit abhijit 4.9M Feb 15 11:09 xv6.img
-rwxrwxr-x 1 abhijit abhijit 209K Feb 15 11:09 kernel*
```

Following code in bootmain()

```
readseg((uchar*)elf, 4096, 0);
```

and following selected lines from Makefile

```
xv6.img: bootblock kernel
    dd if=/dev/zero of=xv6.img count=10000
    dd if=bootblock of=xv6.img conv=notrunc
    dd if=kernel of=xv6.img seek=1 conv=notrunc

kernel: $(OBJS) entry.o entryother initcode kernel.ld
    $(LD) $(LDFLAGS) -T kernel.ld -o kernel entry.o $(OBJS) -b binary initcode entryother
    $(OBJDUMP) -S kernel > kernel.asm
    $(OBJDUMP) -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$$/d' > kernel.sym
```

Also read the code of bootmain() in xv6 kernel.

Select the options that describe the meaning of these lines and their correlation.

- a. Althought the size of the kernel file is 209 Kb, only 4Kb out of it is the actual kernel code and remaining part is all zeroes.
- b. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is not X read as it is user programs.
- c. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(). ✓
- d. The kernel.asm file is the final kernel file
- e. Althought the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.
- f. The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files ✓
- g. The kernel.ld file contains instructions to the linker to link the kernel properly ✓
- h. readseg() reads first 4k bytes of kernel in memory ✓
- i. The bootmain() code does not read the kernel completely in memory

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(), readseg() reads first 4k bytes of kernel in memory, The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files, The kernel.ld file contains instructions to the linker to link the kernel properly, Althought the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.

Question 5

Correct

Mark 1.00 out of 1.00

For each function/code-point, select the status of segmentation setup in xv6

bootmain()	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
kvmalloc() in main()	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
bootasm.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	✓
after seginit() in main()	gdt setup with 5 entries (0 to 4) on one processor	✓

Your answer is correct.

The correct answer is: bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor

Question 6

Correct

Mark 1.00 out of 1.00

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?

Select all the appropriate choices

- a. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time
- b. The setting up of the most essential memory management infrastructure needs assembly code ✓
- c. The code for reading ELF file can not be written in assembly
- d. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C ✓

Your answer is correct.

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question 7

Incorrect

Mark 0.00 out of 1.00

xv6.img: bootblock kernel

```
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is INCORRECT?

- a. The kernel is located at block-1 of the xv6.img ✗
- b. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk. ✓
- c. The size of the kernel file is nearly 5 MB ✓
- d. The bootblock is located on block-0 of the xv6.img ✗
- e. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk. ✗
- f. Blocks in xv6.img after kernel may be all zeroes. ✗
- g. xv6.img is the virtual processor used by the qemu emulator
- h. The size of the xv6.img is nearly 5 MB ✗
- i. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file. ✗
- j. The bootblock may be 512 bytes or less (looking at the Makefile instruction) ✗
- k. The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Your answer is incorrect.

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question 8

Partially correct

Mark 0.07 out of 1.00

Select all the correct statements about code of bootmain() in xv6

```
void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry) (void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++) {
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void)) (elf->entry);
    entry();
}
```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- a. The kernel file has only two program headers ✓
- b. The kernel file gets loaded at the Physical address 0x10000 in memory. ✓
- c. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- d. The readseg finally invokes the disk I/O code using assembly instructions
- e. The elf->entry is set by the linker in the kernel file and it's 8010000c
- f. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded
- g. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it. ✓
- h. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- i. The kernel file gets loaded at the Physical address 0x10000 +0x80000000 in memory.

- j. The condition if(ph->memsz > ph->filesz) is never true.
- k. The stosb() is used here, to fill in some space in memory with zeroes ✓

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

Question 9

Correct

Mark 1.00 out of 1.00

In bootasm.S, on the line

```
1jmp    $(SEG_KCODE<<3), $start32
```

The SEG_KCODE << 3, that is shifting of 1 by 3 bits is done because

- a. While indexing the GDT using CS, the value in CS is always divided by 8
- b. The code segment is 16 bit and only upper 13 bits are used for segment number ✓
- c. The value 8 is stored in code segment
- d. The code segment is 16 bit and only lower 13 bits are used for segment number
- e. The ljmp instruction does a divide by 8 on the first argument

Your answer is correct.

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question 10

Incorrect

Mark 0.00 out of 1.00

For each line of code mentioned on the left side, select the location of sp/esp that is in use

`ljmp $(SEG_KCODE<<3), $start32
in bootasm.S`

0x7c00 to 0



`readseg((uchar*)elf, 4096, 0);
in bootmain.c`

Immaterial as the stack is not used here



`jmp *%eax
in entry.S`

0x7c00 to 0



`cli
in bootasm.S`

Choose...

`call bootmain
in bootasm.S`

Immaterial as the stack is not used here



Your answer is incorrect.

The correct answer is: `ljmp $(SEG_KCODE<<3), $start32`

`in bootasm.S → Immaterial as the stack is not used here, readseg((uchar*)elf, 4096, 0);`

`in bootmain.c → 0x7c00 to 0, jmp *%eax`

`in entry.S → The 4KB area in kernel image, loaded in memory, named as 'stack', cli`

`in bootasm.S → Immaterial as the stack is not used here, call bootmain`

`in bootasm.S → 0x7c00 to 0`

[◀ Random Quiz 4 : Scheduling, signals, segmentation, paging, compilation, process-state](#)

Jump to...

[Random Quiz - 6 \(xv6 file system\) ►](#)

Started on Friday, 31 March 2023, 6:18 PM

State Finished

Completed on Friday, 31 March 2023, 7:00 PM

Time taken 41 mins 52 secs

Grade 4.92 out of 15.00 (32.83%)

Question 1

Partially correct

Mark 0.67 out of 1.00

Select all the actions taken by iget()

- a. Returns an inode with given dev+inode-number from cache, if it exists in cache ✓
- b. Returns the inode with reference count incremented ✓
- c. Returns a free-inode , with dev+inode-number set, if not found in cache
- d. Panics if inode does not exist in cache
- e. Returns a valid inode if not found in cache
- f. Returns the inode with inode-cache lock held
- g. Returns the inode locked

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: Returns an inode with given dev+inode-number from cache, if it exists in cache, Returns the inode with reference count incremented, Returns a free-inode , with dev+inode-number set, if not found in cache

Question **2**

Incorrect

Mark 0.00 out of 1.00

Note: for this question you get full marks if you select all and only correct options, you get ZERO if at least one option is wrong or not selected.

Select all the correct statements about log structured file systems.

- a. ext4 is a log structured file system X it's a journaled file system, not log structured
- b. file system recovery recovers all the lost data X
- c. log structured file systems considerably improve the recovery time ✓
- d. xv6 has a log structured file system
- e. ext2 is by default a log structured file system X

Your answer is incorrect.

The correct answers are: xv6 has a log structured file system, log structured file systems considerably improve the recovery time

Question 3

Partially correct

Mark 0.86 out of 2.00

Select T/F w.r.t physical disk handling in xv6 code

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	device files are not supported ✗
<input type="radio"/>	<input checked="" type="radio"/>	The code supports IDE, and not SATA/SCSI ✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	log is kept on the same device as the file system ✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	disk driver handles only one buffer at a time ✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	the superblock does not contain number of free blocks ✓
<input checked="" type="radio"/>	<input type="radio"/>	only direct blocks are supported ✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	only 2 disks are handled by default ✗

device files are not supported: False

The code supports IDE, and not SATA/SCSI: True

log is kept on the same device as the file system: True

disk driver handles only one buffer at a time: True

the superblock does not contain number of free blocks: True

only direct blocks are supported: False

only 2 disks are handled by default: True

Question 4

Partially correct

Mark 1.00 out of 2.00

Marks the statements as True/False w.r.t. "struct buf"

True	False	
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	B_VALID means the buffer is empty and can be reused
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	Lock on a buffer is acquired in bget, and released in brelse
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The reference count (refcnt) in struct buf is = number of processes accessing the buffer
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	A buffer can have both B_VALID and B_DIRTY flags set
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The "next" pointer chain gives the buffers in LRU order
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	A buffer can be both on the MRU/LRU list and also on ideoqueue list.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The buffers are maintained in LRU order, in the function brelse
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	B_DIRTY flag means the buffer contains modified data

B_VALID means the buffer is empty and can be reused: False

Lock on a buffer is acquired in bget, and released in brelse: True

The reference count (refcnt) in struct buf is = number of processes accessing the buffer: True

A buffer can have both B_VALID and B_DIRTY flags set: False

The "next" pointer chain gives the buffers in LRU order: False

A buffer can be both on the MRU/LRU list and also on ideoqueue list.: True

The buffers are maintained in LRU order, in the function brelse: True

B_DIRTY flag means the buffer contains modified data: True

Question 5

Partially correct

Mark 0.40 out of 1.00

Arrange the following in their typical order of use in xv6.

1. iget
2. ilock
3. iunlock
4. iput
5. use inode

Your answer is partially correct.

Grading type: Relative to the next item (including last)

Grade details: 2 / 5 = 40%

Here are the scores for each item in this response:

1. 1 / 1 = 100%
2. 0 / 1 = 0%
3. 1 / 1 = 100%
4. 0 / 1 = 0%
5. 0 / 1 = 0%

The correct order for these items is as follows:

1. iget
2. ilock
3. use inode
4. iunlock
5. iput

Question **6**

Partially correct

Mark 0.50 out of 1.00

Compare XV6 and EXT2 file systems.

Select True/False for each point.

True	False	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✓	In both ext2 and xv6, the superblock gives location of first inode block ✗
<input type="checkbox"/> ✓	<input checked="" type="radio"/> ✗	xv6 contains journal, ext2 does not ✗
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✓	Ext2 contains superblock but xv6 does not. ✓
<input checked="" type="radio"/> ✗	<input checked="" type="radio"/> ✓	xv6 contains inode bitmap, but ext2 does not ✓
<input checked="" type="radio"/> ✗	<input type="radio"/> ✓	Both xv6 and ext2 contain magic number ✗
<input type="checkbox"/> ✓	<input checked="" type="radio"/> ✗	Ext2 contains group descriptors but xv6 does not ✓

In both ext2 and xv6, the superblock gives location of first inode block: False

xv6 contains journal, ext2 does not: True

Ext2 contains superblock but xv6 does not.: False

xv6 contains inode bitmap, but ext2 does not: False

Both xv6 and ext2 contain magic number: False

Ext2 contains group descriptors but xv6 does not: True

Question **7**

Incorrect

Mark 0.00 out of 1.00

Maximum size of a file on xv6 in **bytes** is

(just write a numeric answer)

Answer: 512



The correct answer is: 71680

Question 8

Partially correct

Mark 0.17 out of 1.00

Select all the actions taken by ilock()

- a. Mark the in-memory inode as valid, if needed
- b. Read the inode from disk, if needed ✓
- c. Take the sleeplock on the inode, always
- d. Get the inode from the inode-cache
- e. Copy the on-disk inode into in-memory inode, if neeed ✓
- f. Take the sleeplock on the inode, optionally ✗
- g. Lock all the buffers of the file in memory

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: Read the inode from disk, if needed, Copy the on-disk inode into in-memory inode, if neeed, Take the sleeplock on the inode, always, Mark the in-memory inode as valid, if needed

Question 9

Incorrect

Mark 0.00 out of 1.00

The lines

```
if(ip->type != T_DIR){  
    iunlockput(ip);  
    return 0;  
}
```

in namex() function

mean

- a. The last path component (which is a file, and not a directory) has been resolved, so release the lock (using iunlockput) and return
- b. One of the sub-components on the given path name, was a directory, but it was not supposed to be a directory, hence an error ✗
- c. There was a syntax error in the pathname specified
- d. One of the sub-components on the given path name, was not a directory, hence it's an error
- e. One of the sub-components on the given path name, did not exist, hence it's an error ✗
- f. ilock is held on the inode, and hence it's an error if it is a directory ✗
- g. No directory entry was found for the file to be opened, hence an error

Your answer is incorrect.

The correct answer is: One of the sub-components on the given path name, was not a directory, hence it's an error

Question 10

Partially correct

Mark 0.50 out of 1.00

Suppose an application on xv6 does the following:

```
int main() {
    char arr[128];
    int fd = open("README", O_RDONLY);
    read(fd, arr, 100);
}
```

Assume that the code works.

Which of the following things are true about xv6 kernel code, w.r.t. the above C program.

True	False	
<input checked="" type="radio"/> ✘	<input type="radio"/> ✓	The "memmove(dst, bp->data + off%BSIZE, m);" in readi() will copy the data from the disk to the kernel buffers
<input type="radio"/> ✓	<input checked="" type="radio"/> ✘	The process will be made to sleep only once
<input checked="" type="radio"/> ✘	<input type="radio"/> ✓	The data is transferred from disk to kernel buffers first, and then address of arr is mapped to the kernel buffers
<input type="radio"/> ✓	<input checked="" type="radio"/> ✘	The ONLY function that gets called on return devsw[ip->major].read(ip, dst, n); is consoleread
<input checked="" type="radio"/> ✘	<input type="radio"/> ✓	The loop in readi() will always read a different block using bread()
<input type="radio"/> ✓	<input checked="" type="radio"/> ✘	value of fd will be 3

The "memmove(dst, bp->data + off%BSIZE, m);" in readi() will copy the data from the disk to the kernel buffers: False

The process will be made to sleep only once: True

The data is transferred from disk to kernel buffers first, and then address of arr is mapped to the kernel buffers: False

The ONLY function that gets called on return devsw[ip->major].read(ip, dst, n); is consoleread: True

The loop in readi() will always read a different block using bread(): False

value of fd will be 3: True

Question 11

Partially correct

Mark 0.83 out of 1.00

Map the function in xv6's file system code, to its perceived logical layer.

ideintr	disk driver	✓
filestat()	file descriptor	✓
ialloc	inode	✓
bread	buffer cache	✓
balloc	system call	✗
skipelem	pathname lookup	✓
bmap	pathname lookup	✗
sys_chdir()	system call	✓
stati	inode	✓
commit	logging	✓
namei	pathname lookup	✓
dirlookup	directory	✓

Your answer is partially correct.

You have correctly selected 10.

The correct answer is: ideintr → disk driver, filestat() → file descriptor, ialloc → inode, bread → buffer cache, balloc → block allocation on disk, skipelem → pathname lookup, bmap → inode, sys_chdir() → system call, stati → inode, commit → logging, namei → pathname lookup, dirlookup → directory

Question 12

Incorrect

Mark 0.00 out of 1.00

Match function with its functionality

nameiparent	Write a new entry in a given directory	✗
dirlink	Link a directory with another directory	✗
namex	Search a given name in a given directory	✗
dirlookup	Lookup (search) for a given directory	✗

Your answer is incorrect.

The correct answer is: nameiparent → return in-memory inode for parent directory of a given pathname, dirlink → Write a new entry in a given directory, namex → return in-memory inode for a given pathname, dirlookup → Search a given name in a given directory

Question 13

Incorrect

Mark 0.00 out of 1.00

An inode is read from disk as a part of this function

- a. ilock
- b. sys_read ✗
- c. iget
- d. readi
- e. iread

Your answer is incorrect.

The correct answer is: ilock

◀ Random Quiz - 5: xv6 make, bootloader, interrupt handling, memory management

Jump to...

(Random Quiz - 7) Pre-Endsem Quiz ►

Started on	Wednesday, 19 April 2023, 6:16 PM
State	Finished
Completed on	Wednesday, 19 April 2023, 8:31 PM
Time taken	2 hours 14 mins
Overdue	14 mins 47 secs
Grade	14.08 out of 30.00 (46.92%)

Question **1**

Incorrect

Mark 0.00 out of 1.00

Select all correct statements about file system recovery (without journaling) programs e.g. fsck

Select one or more:

- a. Recovery programs are needed only if the file system has a delayed-write policy. ✓
- b. Recovery is possible due to redundancy in file system data structures
- c. It is possible to lose data as part of recovery ✓
- d. A recovery program, most typically, builds the file system data structure and checks for inconsistencies ✓
- e. Recovery programs recalculate most of the metadata summaries (e.g. free inode count) ✓
- f. They are used to recover deleted files ✗
- g. They may take very long time to execute ✓
- h. Even with a write-through policy, it is possible to need a recovery program.
- i. They can make changes to the on-disk file system

Your answer is incorrect.

The correct answers are: Recovery is possible due to redundancy in file system data structures, A recovery program, most typically, builds the file system data structure and checks for inconsistencies, It is possible to lose data as part of recovery, They may take very long time to execute, They can make changes to the on-disk file system, Recovery programs recalculate most of the metadata summaries (e.g. free inode count), Recovery programs are needed only if the file system has a delayed-write policy., Even with a write-through policy, it is possible to need a recovery program.

Question 2

Partially correct

Mark 0.70 out of 1.00

Mark the statements as True or False, w.r.t. thrashing

True	False	
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Thrashing is particular to demand paging systems, and does not apply to pure paging systems.
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/>	During thrashing the CPU is under-utilised as most time is spent in I/O
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Thrashing occurs when the total size of all process's locality exceeds total memory size.
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Thrashing can be limited if local replacement is used.
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Thrashing occurs because some process is doing lot of disk I/O.
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The working set model is an attempt at approximating the locality of a process.
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Thrashing can occur even if entire memory is not in use.
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/>	mmap() solves the problem of thrashing.
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.

Thrashing is particular to demand paging systems, and does not apply to pure paging systems.: True

During thrashing the CPU is under-utilised as most time is spent in I/O: True

Thrashing occurs when the total size of all process's locality exceeds total memory size.: True

Thrashing can be limited if local replacement is used.: True

Thrashing occurs because some process is doing lot of disk I/O.: False

The working set model is an attempt at approximating the locality of a process.: True

Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.: True

Thrashing can occur even if entire memory is not in use.: False

mmap() solves the problem of thrashing.: False

Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.: False

Question 3

Partially correct

Mark 0.25 out of 1.00

Match each suggested semaphore implementation (discussed in class)

with the problems that it faces

```
struct semaphore {  
    int val;  
    spinlock lk;  
};  
sem_init(semaphore *s, int initval) {  
    s->val = initval;  
    s->sl = 0;  
}  
wait(semaphore *s) {  
    spinlock(&(s->sl));  
    while(s->val <=0) {  
        spinunlock(&(s->sl));  
        spinlock(&(s->sl));  
    }  
    (s->val)--;  
    spinunlock(&(s->sl));  
}
```

too much spinning, bounded wait not guaranteed ✓

```
struct semaphore {  
    int val;  
    spinlock lk;  
};  
sem_init(semaphore *s, int initval) {  
    s->val = initval;  
    s->sl = 0;  
}  
wait(semaphore *s) {  
    spinlock(&(s->sl));  
    while(s->val <=0)  
    ;  
    (s->val)--;  
    spinunlock(&(s->sl));  
}
```

blocks holding a spinlock ✗

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->s1 = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->s1));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->s1));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->s1));
}

signal(semaphore *s) {
    spinlock(*(s->s1));
    (s->val)++;
    x = dequeue(s->s1) and enqueue(readyq, x);
    spinunlock(*(s->s1));
}

```

too much spinning, bounded wait not guaranteed



```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->s1 = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->s1));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->s1));
}

```

deadlock



Your answer is partially correct.

You have correctly selected 1.

The correct answer is:

```

struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        spinunlock(&(s->sl));
        spinlock(&(s->sl));
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

```

→ too much spinning, bounded wait not guaranteed,

```

struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0)
    ;
    (s->val)--;
    spinunlock(&(s->sl));
}

```

→ deadlock,

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->sl));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

signal(semaphore *s) {
    spinlock(*(&(s->sl)));
    (s->val)++;
    x = dequeue(s->sl) and enqueue(readyq, x);
    spinunlock(*(&(s->sl)));
}

```

→ not holding lock after unblock,

```
struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}
```

→ blocks holding a spinlock

Question 4

Partially correct

Mark 0.88 out of 1.00

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

True	False	
<input checked="" type="radio"/>	<input type="radio"/> 	The arguments to system call originally reside on process stack.
<input checked="" type="radio"/>	<input type="radio"/> 	String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer
<input checked="" type="radio"/>	<input type="radio"/> 	The arguments are accessed in the kernel code using esp on the trapframe.
<input type="radio"/> 	<input checked="" type="radio"/>	Integer arguments are stored in eax, ebx, ecx, etc. registers
<input checked="" type="radio"/>	<input type="radio"/> 	Integer arguments are copied from user memory to kernel memory using argint()
<input type="radio"/> 	<input checked="" type="radio"/>	String arguments are first copied to trapframe and then from trapframe to kernel's other variables.
<input checked="" type="radio"/>	<input type="radio"/> 	The functions like argint(), argstr() make the system call arguments available in the kernel.
<input type="radio"/> 	<input checked="" type="radio"/>	The arguments to system call are copied to kernel stack in trapasm.S

The arguments to system call originally reside on process stack.: True

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

The arguments are accessed in the kernel code using esp on the trapframe.: True

Integer arguments are stored in eax, ebx, ecx, etc. registers: False

Integer arguments are copied from user memory to kernel memory using argint(): True

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False

The functions like argint(), argstr() make the system call arguments available in the kernel.: True

The arguments to system call are copied to kernel stack in trapasm.S: False

Question 5

Correct

Mark 1.00 out of 1.00

Map the technique with its feature/problem

static loading	wastage of physical memory	✓
static linking	large executable file	✓
dynamic linking	small executable file	✓
dynamic loading	allocate memory only if needed	✓

The correct answer is: static loading → wastage of physical memory, static linking → large executable file, dynamic linking → small executable file, dynamic loading → allocate memory only if needed

Question 6

Partially correct

Mark 0.10 out of 2.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- a. Demand paging requires additional hardware support, compared to paging.
- b. Demand paging always increases effective memory access time. ✓
- c. Both demand paging and paging support shared memory pages. ✓
- d. TLB hit ratio has zero impact in effective memory access time in demand paging. ✗
- e. Paging requires NO hardware support in CPU
- f. Paging requires some hardware support in CPU ✓
- g. The meaning of valid-invalid bit in page table is different in paging and demand-paging. ✓
- h. Calculations of number of bits for page number and offset are same in paging and demand paging. ✓
- i. With paging, it's possible to have user programs bigger than physical memory. ✗
- j. With demand paging, it's possible to have user programs bigger than physical memory.

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question **7**

Correct

Mark 1.00 out of 1.00

Given that the memory access time is 150 ns, probability of a page fault is 0.9 and page fault handling time is 6 ms,
The effective memory access time in nanoseconds is:

Answer: 5400015 ✓

The correct answer is: 5400015.00

Question **8**

Correct

Mark 1.00 out of 1.00

Assuming a 8- KB page size, what is the page numbers for the address 803160 reference in decimal :
(give answer also in decimal)

Answer: 98 ✓

The correct answer is: 98

Question 9

Partially correct

Mark 1.50 out of 2.00

For Virtual File System to work, which of the following changes are required to be done to an existing OS code (e.g. xv6)?

- a. Each file-system writer needs to provide the set of function pointers for VFS, and these function pointers need to be setup in generic inode of "/" of that file system during mount() ✓
- b. A mount() system call should be provided to mount a partition onto some directory in existing namespace rooted at "/" ✓
- c. The operating system in-memory inode needs to be a generic-inode representing "inode" like data structure across multiple file systems.
- d. The generic inode needs to have a field representing if this inode is a mount point and also to refer/point to the root of the mounted file system's inode. ✓
- e. The filesystem related system calls (e.g. read, write) need to invoke the file system specific functions (e.g. ext2_read, ext2_write, ntfs_read, ntfs_write) using function pointers. ✓
- f. The file system specific function pointers, for file system system-calls, need to be setup in the generic inode during lookup.
- g. Each open() needs to copy the function pointers from the inode of the parent directory into the inode of the child (if not already done), unless it's traversing a mount point. (This may be done as part of lookup() which is called by open()) ✓
- h. The lookup() operation needs to check if it's crossing a mount point and call FS specific operations to read inodes/directories ✓

The correct answers are: A mount() system call should be provided to mount a partition onto some directory in existing namespace rooted at "/", The filesystem related system calls (e.g. read, write) need to invoke the file system specific functions (e.g. ext2_read, ext2_write, ntfs_read, ntfs_write) using function pointers., The file system specific function pointers, for file system system-calls, need to be setup in the generic inode during lookup., The operating system in-memory inode needs to be a generic-inode representing "inode" like data structure across multiple file systems., The generic inode needs to have a field representing if this inode is a mount point and also to refer/point to the root of the mounted file system's inode., The lookup() operation needs to check if it's crossing a mount point and call FS specific operations to read inodes/directories, Each file-system writer needs to provide the set of function pointers for VFS, and these function pointers need to be setup in generic inode of "/" of that file system during mount(), Each open() needs to copy the function pointers from the inode of the parent directory into the inode of the child (if not already done), unless it's traversing a mount point. (This may be done as part of lookup() which is called by open())

Question 10

Partially correct

Mark 0.33 out of 1.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
void  
yield(void)  
{  
...  
release(&ptable.lock);  
}
```

Release the lock held by some another process



```
void  
panic(char *s)  
{  
...  
panicked = 1;
```

If you don't do this, a process may be running on two processors parallelly



```
void  
acquire(struct spinlock *lk)  
{  
...  
getcallerpcs(&lk, lk->pcs);  
}
```

Disable interrupts to avoid another process's pointer being returned



Your answer is partially correct.

You have correctly selected 1.

The correct answer is:

```
void  
yield(void)  
{  
...  
release(&ptable.lock);  
} → Release the lock held by some another process, void  
panic(char *s)  
{  
...  
panicked = 1; → Ensure that no printing happens on other processors, void  
acquire(struct spinlock *lk)  
{  
...  
getcallerpcs(&lk, lk->pcs); → Traverse ebp chain to get sequence of instructions followed in functions calls
```

Question 11

Partially correct

Mark 0.75 out of 1.00

Select all correct statements about journalling (logging) in file systems like ext3

Select one or more:

- a. A different device driver is always needed to access the journal
- b. The purpose of journal is to speed up file system recovery ✓
- c. Most typically a transaction in journal is recorded atomically (full or none) ✓
- d. Journals are often stored circularly
- e. Journals must be maintained on the same device that hosts the file system
- f. the journal contains a summary of all changes made as part of a single transaction ✓
- g. Journal is hosted in the same device that hosts the swap space

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: The purpose of journal is to speed up file system recovery, the journal contains a summary of all changes made as part of a single transaction, Most typically a transaction in journal is recorded atomically (full or none), Journals are often stored circularly

Question **12**

Partially correct

Mark 1.00 out of 2.00

Select all the correct statements about synchronization primitives.

Select one or more:

- a. Thread that is going to block should not be holding any spinlock
- b. Semaphores can be used for synchronization scenarios like ordered execution ✓
- c. Mutexes can be implemented using spinlock ✓
- d. Blocking means one process passing over control to another process
- e. Semaphores are always a good substitute for spinlocks
- f. Blocking means moving the process to a wait queue and spinning
- g. Blocking means moving the process to a wait queue and calling scheduler
- h. All synchronization primitives are implemented essentially with some hardware assistance.
- i. Spinlocks are good for multiprocessor scenarios, for small critical sections
- j. Mutexes can be implemented without any hardware assistance
- k. Spinlocks consume CPU time ✓
- l. Mutexes can be implemented using blocking and wakeup ✓

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: Spinlocks are good for multiprocessor scenarios, for small critical sections, Spinlocks consume CPU time, Semaphores can be used for synchronization scenarios like ordered execution, Mutexes can be implemented using spinlock, Mutexes can be implemented using blocking and wakeup, Thread that is going to block should not be holding any spinlock, Blocking means moving the process to a wait queue and calling scheduler, All synchronization primitives are implemented essentially with some hardware assistance.

Question 13

Partially correct

Mark 0.50 out of 2.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
void  
yield(void)  
{  
...  
release(&ptable.lock);  
}
```

Ensure that no printing happens on other processors



```
struct proc*  
myproc(void) {  
...  
pushcli();  
c = mycpu();  
p = c->proc;  
popcli();  
...  
}
```

Disable interrupts to avoid another process's pointer being returned



```
static inline uint  
xchg(volatile uint *addr, uint newval)  
{  
    uint result;  
  
    // The + in "+m" denotes a read-modify-write  
    // operand.  
    asm volatile("lock; xchgl %0, %1" :  
        "+m" (*addr), "=a" (result) :  
        "1" (newval) :  
        "cc");  
    return result;  
}
```

Atomic compare and swap instruction (to be expanded inline into code)



```
void  
acquire(struct spinlock *lk)  
{  
...  
__sync_synchronize();
```

Avoid a self-deadlock



```
void
acquire(struct spinlock *lk)
{
...
getcallerpcs(&lk, lk->pcs);
```

Disable interrupts to avoid deadlocks



```
void
acquire(struct spinlock *lk)
{
pushcli();
```

Traverse ebp chain to get sequence of instructions followed in functions calls



```
void
sleep(void *chan, struct spinlock *lk)
{
...
if(lk != &phtable.lock){
    acquire(&phtable.lock);
    release(lk);
}
```

Release the lock held by some another process



```
void
panic(char *s)
{
...
panicked = 1;
```

If you don't do this, a process may be running on two processors parallely



Your answer is partially correct.

You have correctly selected 2.

The correct answer is: void

```
yield(void)
{
...
release(&phtable.lock);
}
```

→ Release the lock held by some another process, **struct proc***

```
myproc(void) {
...
pushcli();
c = mycpu();
p = c->proc;
popcli();
...
}
```

→ Disable interrupts to avoid another process's pointer being returned, **static inline uint**

```
xchg(volatile uint *addr, uint newval)
{
    uint result;
```

```
// The + in "+m" denotes a read-modify-write operand.  
asm volatile("lock; xchgl %0, %1" :  
    "+m" (*addr), "=a" (result) :  
    "1" (newval) :  
    "cc");  
return result;  
} → Atomic compare and swap instruction (to be expanded inline into code), void  
acquire(struct spinlock *lk)  
{  
...  
_sync_synchronize();
```

```
→ Tell compiler not to reorder memory access beyond this line, void  
acquire(struct spinlock *lk)  
{  
...  
getcallerpcs(&lk, lk->pcs);
```

```
→ Traverse ebp chain to get sequence of instructions followed in functions calls, void  
acquire(struct spinlock *lk)  
{  
pushcli();  
→ Disable interrupts to avoid deadlocks, void  
sleep(void *chan, struct spinlock *lk)  
{  
...  
if(lk != &ptable.lock){  
    acquire(&ptable.lock);  
    release(lk);  
} → Avoid a self-deadlock, void  
panic(char *s)  
{  
...  
panicked = 1; → Ensure that no printing happens on other processors
```

Question 14

Partially correct

Mark 0.86 out of 1.00

Select T/F for statements about Volume Managers.

Do pay attention to the use of the words physical partition and physical volume.

True	False	
<input checked="" type="radio"/>	<input type="radio"/> X	A logical volume may span across multiple physical partitions
<input type="radio"/> ✓	<input type="radio"/> X	A logical volume may span across multiple physical volumes
<input checked="" type="radio"/>	<input type="radio"/> X	A physical partition should be initialized as a physical volume, before it can be used by volume manager.
<input checked="" type="radio"/>	<input type="radio"/> X	A volume group consists of multiple physical volumes
<input checked="" type="radio"/>	<input type="radio"/> X	The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.
<input checked="" type="radio"/>	<input type="radio"/> X	A logical volume can be extended in size but upto the size of volume group
<input checked="" type="radio"/>	<input type="radio"/> X	The volume manager stores additional metadata on the physical disk partitions

A logical volume may span across multiple physical partitions: True

A logical volume may span across multiple physical volumes: True

A physical partition should be initialized as a physical volume, before it can be used by volume manager.: True

A volume group consists of multiple physical volumes: True

The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.: True

A logical volume can be extended in size but upto the size of volume group: True

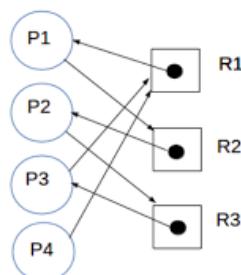
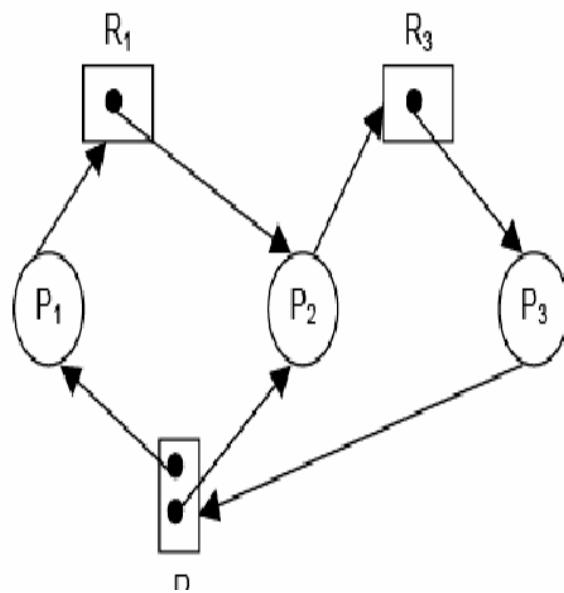
The volume manager stores additional metadata on the physical disk partitions: True

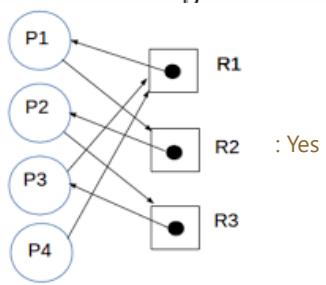
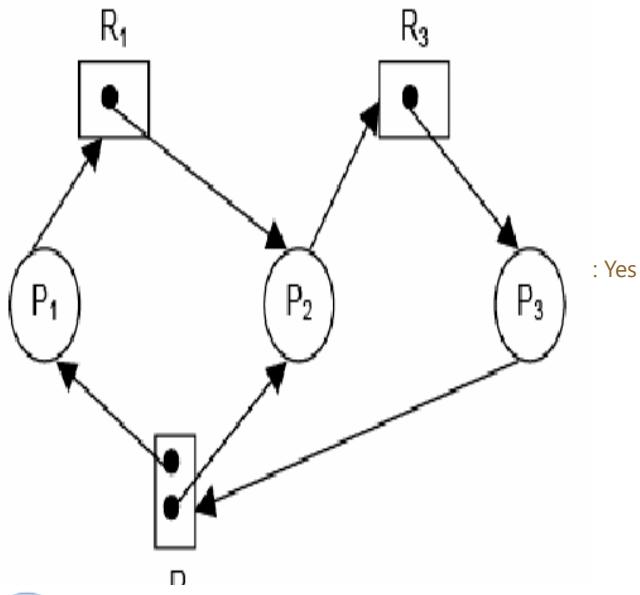
Question 15

Partially correct

Mark 0.50 out of 1.00

For each of the resource allocation diagram shown,
infer whether the graph contains at least one deadlock or not.

Yes **No**



Question **16**

Partially correct

Mark 0.80 out of 1.00

Select all the correct statements w.r.t user and kernel threads

Select one or more:

- a. all three models, that is many-one, one-one, many-many , require a user level thread library ✓
- b. A process may not block in many-one model, if a thread makes a blocking system call
- c. one-one model can be implemented even if there are no kernel threads
- d. one-one model increases kernel's scheduling load ✓
- e. A process blocks in many-one model even if a single thread makes a blocking system call
- f. many-one model can be implemented even if there are no kernel threads ✓
- g. many-one model gives no speedup on multicore processors ✓

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: many-one model can be implemented even if there are no kernel threads, all three models, that is many-one, one-one, many-many , require a user level thread library, one-one model increases kernel's scheduling load, many-one model gives no speedup on multicore processors, A process blocks in many-one model even if a single thread makes a blocking system call

Question 17

Incorrect

Mark 0.00 out of 1.00

Match the code with its functionality

S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statementn2;

Signal(S1);

Execution order P1, P2, P3



P3:

Wait(S1);

Statement S3;

S = 0

P1:

Statement1;

Signal(S)

Execution order P3, P2, P1



P2:

Wait(S)

Statment2;

S = 5

Wait(S)

Critical Section

Execution order P2, then P1



Signal(S)

S = 1

Wait(S)

Critical Section

Counting semaphore



Signal(S);

Your answer is incorrect.

The correct answer is: S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statemetn2;

Signal(S1);

P3:

Wait(S1);

Statement S3; → Execution order P2, P1, P3, S = 0

P1:

Statement1;

Signal(S)

P2:

Wait(S)

Statement2; → Execution order P1, then P2, S = 5

Wait(S)

Critical Section

Signal(S) → Counting semaphore, S = 1

Wait(S)

Critical Section

Signal(S); → Binary Semaphore for mutual exclusion

Question **18**

Not answered

Marked out of 2.00

Write all changes required to xv6 to add a buddy allocator.

Every change should be mentioned in terms of either of the following:

- (a) pseudo-code of new function to be added
- (b) prototype of any new function or new system call to be added
- (c) pseudo-code of changes to an existing function, describing lines to be removed, and lines to be added
- (d) **precise** declaration of new data structures to be added in C, or changes to the existing data structure
- (e) Name and a one-line description of new userland functionality to be added
- (f) Changes to Makefile
- (g) Any other change in a maximum of 20 words per change.

Question 19

Correct

Mark 1.00 out of 1.00

Note: for this question you get full marks if you select all and only correct options, you get ZERO if at least one option is wrong or not selected.

Select all the correct statements about log structured file systems.

- a. a transaction is said to be committed when all operations are written to file system
- b. even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery ✓
- c. file system recovery may end up losing data ✓
- d. file system recovery recovers all the lost data
- e. log may be kept on same block device or another block device ✓

Your answer is correct.

The correct answers are: file system recovery may end up losing data, log may be kept on same block device or another block device, even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery

Question 20

Incorrect

Mark 0.00 out of 1.00

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

11010010

Now, there is a request for a chunk of 45 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: 11001110

✗

The correct answer is: 11011110

Question **21**

Complete

Mark 0.25 out of 3.00

List down all changes required to xv6 code, in order to add the system call chown().

Every change should be mentioned in terms of either of the following:

- (a) pseudo-code of new function to be added
 - (b) prototype of any new function or new system call to be added
 - (c) pseudo-code of changes to an existing function, describing lines to be removed, and lines to be added
 - (d) **precise** declaration of new data structures to be added in C, or changes to the existing data structure
 - (e) Name and a one-line description of new userland functionality to be added
 - (f) Changes to Makefile
 - (g) Any other change in a maximum of 20 words per change.
-
- (a) int chown(char* path, char *owner_name);
 - d) there is no need to change the data structure and no need to add new data structure
 - (e) allow user to change the owner of the file
 - (f) _chown\

Comment:

Question 22

Correct

Mark 1.00 out of 1.00

Calculate the average waiting time using
FCFS scheduling
for the following workload
assuming that they arrive in this order during the first time unit:

Process Burst Time

P1	2
P2	6
P3	2
P4	3

Write only a number in the answer upto two decimal points.

Answer: 5



P2 waits for 2 units

P3 waits for 2+6 units

P4 waits for 2 + 6 +2 units of time

Total waiting = 2 + 2 + 6 + 2 + 6 + 2 = 20 units

Average waiting time = 20/4 = 5

The correct answer is: 5

Question 23

Partially correct

Mark 0.67 out of 1.00

Given that a kernel has 1000 KB of total memory, and holes of sizes (in that order) 300 KB, 200 KB, 100 KB, 250 KB. For each of the requests on the left side, match it with the chunk chosen using the specified algorithm.

Consider each request as first request.

150 KB, first fit	200 KB	✗
150 KB, best fit	300 KB	✗
220 KB, best fit	250 KB	✓
200 KB, first fit	300 KB	✓
100 KB, worst fit	300 KB	✓
50 KB, worst fit	300 KB	✓

The correct answer is: 150 KB, first fit → 300 KB, 150 KB, best fit → 200 KB, 220 KB, best fit → 250 KB, 200 KB, first fit → 300 KB, 100 KB, worst fit → 300 KB, 50 KB, worst fit → 300 KB

Jump to...

[Homework questions: Basics of MM, xv6 booting ►](#)

Started on Wednesday, 19 April 2023, 8:46 PM

State Finished

Completed on Wednesday, 19 April 2023, 9:26 PM

Time taken 39 mins 31 secs

Grade 15.60 out of 30.00 (51.99%)

Question 1

Incorrect

Mark 0.00 out of 2.00

Select all the correct statements about synchronization primitives.

Select one or more:

- a. Blocking means moving the process to a wait queue and calling scheduler
- b. Blocking means moving the process to a wait queue and spinning ✗
- c. Semaphores are always a good substitute for spinlocks ✗
- d. Spinlocks are good for multiprocessor scenarios, for small critical sections
- e. All synchronization primitives are implemented essentially with some hardware assistance.
- f. Blocking means one process passing over control to another process ✗
- g. Semaphores can be used for synchronization scenarios like ordered execution ✓
- h. Mutexes can be implemented without any hardware assistance
- i. Spinlocks consume CPU time
- j. Mutexes can be implemented using spinlock
- k. Thread that is going to block should not be holding any spinlock
- l. Mutexes can be implemented using blocking and wakeup

Your answer is incorrect.

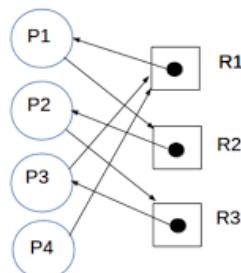
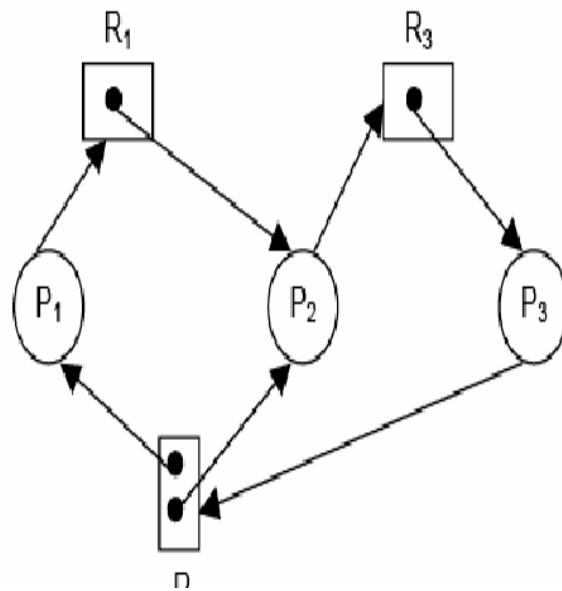
The correct answers are: Spinlocks are good for multiprocessor scenarios, for small critical sections, Spinlocks consume CPU time, Semaphores can be used for synchronization scenarios like ordered execution, Mutexes can be implemented using spinlock, Mutexes can be implemented using blocking and wakeup, Thread that is going to block should not be holding any spinlock, Blocking means moving the process to a wait queue and calling scheduler, All synchronization primitives are implemented essentially with some hardware assistance.

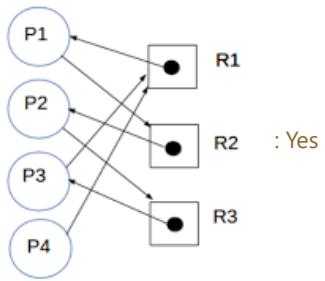
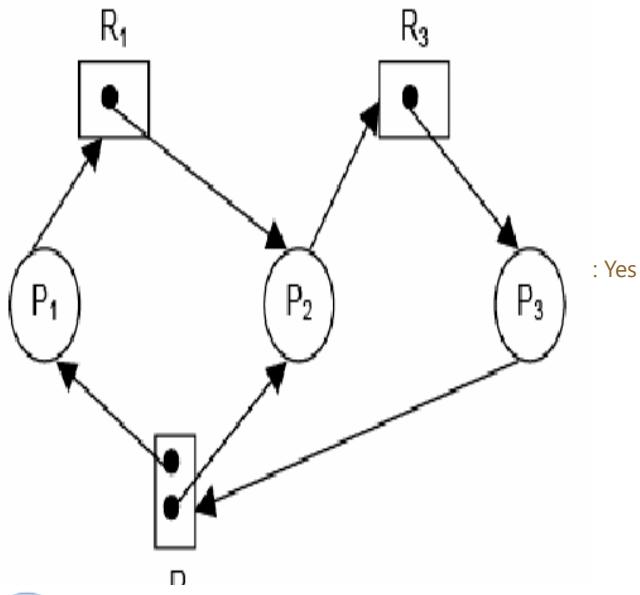
Question 2

Partially correct

Mark 0.50 out of 1.00

For each of the resource allocation diagram shown,
infer whether the graph contains at least one deadlock or not.

Yes **No**



Question 3

Partially correct

Mark 0.25 out of 1.00

Match each suggested semaphore implementation (discussed in class)

with the problems that it faces

```
struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->sl));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

signal(semaphore *s) {
    spinlock(*(s->sl));
    (s->val)++;
    x = dequeue(s->sl) and enqueue(readyq, x);
    spinunlock(*(s->sl));
}
```

deadlock



```
struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}
```

not holding lock after unblock



```
struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        spinunlock(&(s->sl));
        spinlock(&(s->sl));
    }
    (s->val)--;
    spinunlock(&(s->sl));
}
```

too much spinning, bounded wait not guaranteed



```
struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0)
    ;
    (s->val)--;
    spinunlock(&(s->sl));
}
```

blocks holding a spinlock



Your answer is partially correct.

You have correctly selected 1.

The correct answer is:

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->sl));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

signal(semaphore *s) {
    spinlock(*(s->sl));
    (s->val)++;
    x = dequeue(s->sl) and enqueue(readyq, x);
    spinunlock(*(s->sl));
}

```

→ not holding lock after unblock,

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

```

→ blocks holding a spinlock,

```

struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        spinunlock(&(s->sl));
        spinlock(&(s->sl));
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

```

→ too much spinning, bounded wait not guaranteed,

```

struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0)
    ;
    (s->val)--;
    spinunlock(&(s->sl));
}

```

→ deadlock

Question 4

Partially correct

Mark 0.75 out of 1.00

Match the code with its functionality

S = 1
Wait(S)
Critical Section Binary Semaphore for mutual exclusion ✓
Signal(S);

S = 0
P1:
Statement1;
Signal(S) Execution order P1, then P2 ✓
P2:

Wait(S)
Statement2;
S = 5
Wait(S)
Critical Section Execution order P1, P2, P3 ✗
Signal(S)

S1 = 0; S2 = 0;
P2:
Statement1;
Signal(S2);

P1:
Wait(S2); Execution order P2, P1, P3 ✓
Statement2;
Signal(S1);

P3:
Wait(S1);
Statement S3;

Your answer is partially correct.

You have correctly selected 3.

The correct answer is: S = 1

Wait(S)
Critical Section
Signal(S); → Binary Semaphore for mutual exclusion, S = 0

P1:
Statement1;
Signal(S)

P2:
Wait(S)
Statement2; → Execution order P1, then P2, S = 5
Wait(S)
Critical Section
Signal(S) → Counting semaphore, S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statemetn2;

Signal(S1);

P3:

Wait(S1);

Statement S3; → Execution order P2, P1, P3

Question 5

Partially correct

Mark 1.71 out of 2.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- a. With paging, it's possible to have user programs bigger than physical memory.
- b. With demand paging, it's possible to have user programs bigger than physical memory. ✓
- c. Paging requires some hardware support in CPU ✓
- d. Paging requires NO hardware support in CPU
- e. Calculations of number of bits for page number and offset are same in paging and demand paging. ✓
- f. TLB hit ration has zero impact in effective memory access time in demand paging.
- g. Demand paging always increases effective memory access time. ✓
- h. The meaning of valid-invalid bit in page table is different in paging and demand-paging.
- i. Demand paging requires additional hardware support, compared to paging. ✓
- j. Both demand paging and paging support shared memory pages. ✓

Your answer is partially correct.

You have correctly selected 6.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question 6

Partially correct

Mark 0.33 out of 1.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
void  
acquire(struct spinlock *lk)  
{  
...  
    getcallerpcs(&lk, lk->pcs);
```

Disable interrupts to avoid another process's pointer being returned



```
void  
panic(char *s)  
{  
...  
    panicked = 1;
```

If you don't do this, a process may be running on two processors parallelly



```
void  
yield(void)  
{  
...  
    release(&ptable.lock);
```

Release the lock held by some another process



Your answer is partially correct.

You have correctly selected 1.

The correct answer is:

```
void  
acquire(struct spinlock *lk)  
{
```

...
 getcallerpcs(&lk, lk->pcs); → Traverse ebp chain to get sequence of instructions followed in functions calls,

void
panic(char *s)

```
{
```

...
 panicked = 1; → Ensure that no printing happens on other processors,

void
yield(void)

```
{
```

```
...  
release(&ptable.lock);
```

} → Release the lock held by some another process

Question 7

Correct

Mark 1.00 out of 1.00

Given that a kernel has 1000 KB of total memory, and holes of sizes (in that order) 300 KB, 200 KB, 100 KB, 250 KB. For each of the requests on the left side, match it with the chunk chosen using the specified algorithm.

Consider each request as first request.

150 KB, first fit	300 KB	✓
220 KB, best fit	250 KB	✓
100 KB, worst fit	300 KB	✓
200 KB, first fit	300 KB	✓
50 KB, worst fit	300 KB	✓
150 KB, best fit	200 KB	✓

The correct answer is: 150 KB, first fit → 300 KB, 220 KB, best fit → 250 KB, 100 KB, worst fit → 300 KB, 200 KB, first fit → 300 KB, 50 KB, worst fit → 300 KB, 150 KB, best fit → 200 KB

Question 8

Correct

Mark 1.00 out of 1.00

Select T/F for statements about Volume Managers.

Do pay attention to the use of the words physical partition and physical volume.

True	False	
<input checked="" type="radio"/>	<input type="radio"/> X	A logical volume may span across multiple physical partitions
<input checked="" type="radio"/>	<input type="radio"/> X	A volume group consists of multiple physical volumes
<input checked="" type="radio"/>	<input type="radio"/> X	A logical volume may span across multiple physical volumes
<input checked="" type="radio"/>	<input type="radio"/> X	A logical volume can be extended in size but upto the size of volume group
<input checked="" type="radio"/>	<input type="radio"/> X	The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.
<input checked="" type="radio"/>	<input type="radio"/> X	A physical partition should be initialized as a physical volume, before it can be used by volume manager.
<input checked="" type="radio"/>	<input type="radio"/> X	The volume manager stores additional metadata on the physical disk partitions

A logical volume may span across multiple physical partitions: True

A volume group consists of multiple physical volumes: True

A logical volume may span across multiple physical volumes: True

A logical volume can be extended in size but upto the size of volume group: True

The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.: True

A physical partition should be initialized as a physical volume, before it can be used by volume manager.: True

The volume manager stores additional metadata on the physical disk partitions: True

Question **9**

Incorrect

Mark 0.00 out of 1.00

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

11010010

Now, there is a request for a chunk of 45 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: 11001101



The correct answer is: 11011110

Question **10**

Not answered

Marked out of 2.00

Write all changes required to xv6 to add a buddy allocator.

Every change should be mentioned in terms of either of the following:

- (a) pseudo-code of new function to be added
- (b) prototype of any new function or new system call to be added
- (c) pseudo-code of changes to an existing function, describing lines to be removed, and lines to be added
- (d) **precise** declaration of new data structures to be added in C, or changes to the existing data structure
- (e) Name and a one-line description of new userland functionality to be added
- (f) Changes to Makefile
- (g) Any other change in a maximum of 20 words per change.

Question 11

Partially correct

Mark 0.80 out of 1.00

Select all the correct statements w.r.t user and kernel threads

Select one or more:

- a. many-one model can be implemented even if there are no kernel threads ✓
- b. one-one model can be implemented even if there are no kernel threads
- c. all three models, that is many-one, one-one, many-many , require a user level thread library ✓
- d. many-one model gives no speedup on multicore processors ✓
- e. one-one model increases kernel's scheduling load ✓
- f. A process blocks in many-one model even if a single thread makes a blocking system call
- g. A process may not block in many-one model, if a thread makes a blocking system call

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: many-one model can be implemented even if there are no kernel threads, all three models, that is many-one, one-one, many-many , require a user level thread library, one-one model increases kernel's scheduling load, many-one model gives no speedup on multicore processors, A process blocks in many-one model even if a single thread makes a blocking system call

Question 12

Partially correct

Mark 0.50 out of 1.00

Select all correct statements about journalling (logging) in file systems like ext3

Select one or more:

- a. the journal contains a summary of all changes made as part of a single transaction ✓
- b. Journals are often stored circularly
- c. Journals must be maintained on the same device that hosts the file system
- d. The purpose of journal is to speed up file system recovery ✓
- e. Most typically a transaction in journal is recorded atomically (full or none)
- f. Journal is hosted in the same device that hosts the swap space
- g. A different device driver is always needed to access the journal

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: The purpose of journal is to speed up file system recovery, the journal contains a summary of all changes made as part of a single transaction, Most typically a transaction in journal is recorded atomically (full or none), Journals are often stored circularly

Question 13

Correct

Mark 1.00 out of 1.00

Note: for this question you get full marks if you select all and only correct options, you get ZERO if at least one option is wrong or not selected.

Select all the correct statements about log structured file systems.

- a. a transaction is said to be committed when all operations are written to file system
- b. file system recovery may end up losing data ✓
- c. log may be kept on same block device or another block device ✓
- d. even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery ✓
- e. file system recovery recovers all the lost data

Your answer is correct.

The correct answers are: file system recovery may end up losing data, log may be kept on same block device or another block device, even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery

Question **14**

Complete

Mark 0.25 out of 3.00

List down all changes required to xv6 code, in order to add the system call chown().

Every change should be mentioned in terms of either of the following:

- (a) pseudo-code of new function to be added
- (b) prototype of any new function or new system call to be added
- (c) pseudo-code of changes to an existing function, describing lines to be removed, and lines to be added
- (d) **precise** declaration of new data structures to be added in C, or changes to the existing data structure
- (e) Name and a one-line description of new userland functionality to be added
- (f) Changes to Makefile
- (g) Any other change in a maximum of 20 words per change.
 - (a) int chown (char * path , char *ownername);
 - (d)there is no need of new data structures or changes to the existing data structure
 - (e)by using this system call ,user to change the owner if the file
 - (f)_trychown\ at the end of UPROGS

Comment:

Question 15

Partially correct

Mark 1.00 out of 2.00

For Virtual File System to work, which of the following changes are required to be done to an existing OS code (e.g. xv6)?

- a. The operating system in-memory inode needs to be a generic-inode representing "inode" like data structure across multiple file systems.
- b. The generic inode needs to have a field representing if this inode is a mount point and also to refer/point to the root of the mounted file system's inode. ✓
- c. Each file-system writer needs to provide the set of function pointers for VFS, and these function pointers need to be setup in generic inode of "/" of that file system during mount() ✓
- d. A mount() system call should be provided to mount a partition onto some directory in existing namespace rooted at "/" ✓
- e. Each open() needs to copy the function pointers from the inode of the parent directory into the inode of the child (if not already done), unless it's traversing a mount point. (This may be done as part of lookup() which is called by open()) ✓
- f. The file system specific function pointers, for file system system-calls, need to be setup in the generic inode during lookup.
- g. The lookup() operation needs to check if it's crossing a mount point and call FS specific operations to read inodes/directories
- h. The filesystem related system calls (e.g. read, write) need to invoke the file system specific functions (e.g. ext2_read, ext2_write, ntfs_read, ntfs_write) using function pointers.

The correct answers are: A mount() system call should be provided to mount a partition onto some directory in existing namespace rooted at "/", The filesystem related system calls (e.g. read, write) need to invoke the file system specific functions (e.g. ext2_read, ext2_write, ntfs_read, ntfs_write) using function pointers., The file system specific function pointers, for file system system-calls, need to be setup in the generic inode during lookup., The operating system in-memory inode needs to be a generic-inode representing "inode" like data structure across multiple file systems., The generic inode needs to have a field representing if this inode is a mount point and also to refer/point to the root of the mounted file system's inode., The lookup() operation needs to check if it's crossing a mount point and call FS specific operations to read inodes/directories, Each file-system writer needs to provide the set of function pointers for VFS, and these function pointers need to be setup in generic inode of "/" of that file system during mount(), Each open() needs to copy the function pointers from the inode of the parent directory into the inode of the child (if not already done), unless it's traversing a mount point. (This may be done as part of lookup() which is called by open())

Question 16

Partially correct

Mark 0.75 out of 1.00

Map the technique with its feature/problem

dynamic linking	small executable file	✓
static loading	wastage of physical memory	✓
static linking	large executable file	✓
dynamic loading	small executable file	✗

The correct answer is: dynamic linking → small executable file, static loading → wastage of physical memory, static linking → large executable file, dynamic loading → allocate memory only if needed

Question 17

Correct

Mark 1.00 out of 1.00

Mark the statements as True or False, w.r.t. thrashing

True	False	
<input checked="" type="radio"/>	<input type="radio"/> X	Thrashing can be limited if local replacement is used.
<input type="radio"/> X	<input checked="" type="radio"/>	Thrashing occurs because some process is doing lot of disk I/O.
<input checked="" type="radio"/>	<input type="radio"/> X	Thrashing occurs when the total size of all process's locality exceeds total memory size.
<input checked="" type="radio"/>	<input type="radio"/> X	During thrashing the CPU is under-utilised as most time is spent in I/O
<input checked="" type="radio"/>	<input type="radio"/> X	Thrashing is particular to demand paging systems, and does not apply to pure paging systems.
<input type="radio"/> X	<input checked="" type="radio"/>	Thrashing can occur even if entire memory is not in use.
<input checked="" type="radio"/>	<input type="radio"/> X	Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.
<input type="radio"/> X	<input checked="" type="radio"/>	Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.
<input checked="" type="radio"/>	<input type="radio"/> X	The working set model is an attempt at approximating the locality of a process.
<input type="radio"/> X	<input checked="" type="radio"/>	mmap() solves the problem of thrashing.

Thrashing can be limited if local replacement is used.: True

Thrashing occurs because some process is doing lot of disk I/O.: False

Thrashing occurs when the total size of all process's locality exceeds total memory size.: True

During thrashing the CPU is under-utilised as most time is spent in I/O: True

Thrashing is particular to demand paging systems, and does not apply to pure paging systems.: True

Thrashing can occur even if entire memory is not in use.: False

Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.: True

Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.: False

The working set model is an attempt at approximating the locality of a process.: True

mmap() solves the problem of thrashing.: False

Question **18**

Partially correct

Mark 0.75 out of 2.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
struct proc*
myproc(void) {
...
pushcli();
c = mycpu();
p = c->proc;
popcli();
...
}
```

Traverse ebp chain to get sequence of instructions followed in functions calls



```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

// The + in "+m" denotes a read-modify-write
// operand.
asm volatile("lock; xchgl %0, %1" :
    "+m" (*addr), "=a" (result) :
    "1" (newval) :
    "cc");
    return result;
}

void
sleep(void *chan, struct spinlock *lk)
{
...
if(lk != &ptable.lock){
    acquire(&ptable.lock);
    release(lk);
}
}

void
panic(char *s)
{
...
panicked = 1;
```

Atomic compare and swap instruction (to be expanded inline into code)



Release the lock held by some another process



```
void
acquire(struct spinlock *lk)
{
    pushcli();
```

If you don't do this, a process may be running on two processors parallelly



Release the lock held by some another process



```
void
acquire(struct spinlock *lk)
{
...
__sync_synchronize();
```

Tell compiler not to reorder memory access beyond this line



```
void
acquire(struct spinlock *lk)
{
...
getcallerpcs(&lk, lk->pcs);
```

If you don't do this, a process may be running on two processors parallelly



```
void
yield(void)
{
...
release(&ptable.lock);
}
```

Release the lock held by some another process



Your answer is partially correct.

You have correctly selected 3.

The correct answer is: **struct proc***

```
myproc(void) {
...
pushcli();
c = mycpu();
p = c->proc;
popcli();
...
}
```

→ Disable interrupts to avoid another process's pointer being returned, **static inline uint xchg(volatile uint *addr, uint newval)**

```
{
    uint result;

// The + in "+m" denotes a read-modify-write operand.
asm volatile("lock; xchgl %0, %1" :
    "+m" (*addr), "=a" (result) :
    "1" (newval) :
    "cc");
    return result;
} → Atomic compare and swap instruction (to be expanded inline into code), void sleep(void *chan, struct spinlock *lk)
{
```

```
...
if(lk != &ptable.lock){
    acquire(&ptable.lock);
    release(lk);
} → Avoid a self-deadlock, void
panic(char *s)
{
...
panicked = 1; → Ensure that no printing happens on other processors, void
acquire(struct spinlock *lk)
{
    pushcli();
    → Disable interrupts to avoid deadlocks, void
acquire(struct spinlock *lk)
{
...
__sync_synchronize();
```

```
→ Tell compiler not to reorder memory access beyond this line, void
acquire(struct spinlock *lk)
{
...
getcallerpcs(&lk, lk->pcs);
```

```
→ Traverse ebp chain to get sequence of instructions followed in functions calls, void
yield(void)
{
...
release(&ptable.lock);
}
```

→ Release the lock held by some another process

Question **19**

Correct

Mark 1.00 out of 1.00

Calculate the average waiting time using
FCFS scheduling
for the following workload
assuming that they arrive in this order during the first time unit:

Process Burst Time

P1	2
P2	6
P3	2
P4	3

Write only a number in the answer upto two decimal points.

Answer: ✓

P2 waits for 2 units

P3 waits for 2+6 units

P4 waits for 2 + 6 +2 units of time

Total waiting = $2 + 2 + 6 + 2 + 6 + 2 = 20$ units

Average waiting time = $20/4 = 5$

The correct answer is: 5

Question 20

Incorrect

Mark 0.00 out of 1.00

Select all correct statements about file system recovery (without journaling) programs e.g. fsck

Select one or more:

- a. Recovery programs recalculate most of the metadata summaries (e.g. free inode count) ✓
- b. It is possible to lose data as part of recovery ✓
- c. Even with a write-through policy, it is possible to need a recovery program. ✓
- d. They can make changes to the on-disk file system ✓
- e. A recovery program, most typically, builds the file system data structure and checks for inconsistencies ✓
- f. They may take very long time to execute ✓
- g. They are used to recover deleted files ✗
- h. Recovery programs are needed only if the file system has a delayed-write policy.
- i. Recovery is possible due to redundancy in file system data structures

Your answer is incorrect.

The correct answers are: Recovery is possible due to redundancy in file system data structures, A recovery program, most typically, builds the file system data structure and checks for inconsistencies, It is possible to lose data as part of recovery, They may take very long time to execute, They can make changes to the on-disk file system, Recovery programs recalculate most of the metadata summaries (e.g. free inode count), Recovery programs are needed only if the file system has a delayed-write policy., Even with a write-through policy, it is possible to need a recovery program.

Question 21

Correct

Mark 1.00 out of 1.00

Given that the memory access time is 150 ns, probability of a page fault is 0.8 and page fault handling time is 6 ms,
The effective memory access time in nanoseconds is:

Answer: 4800030 ✓

The correct answer is: 4800030.00

Question **22**

Correct

Mark 1.00 out of 1.00

Assuming a 8- KB page size, what is the page numbers for the address 1005699 reference in decimal :

(give answer also in decimal)

Answer: 123



The correct answer is: 123

Question 23

Correct

Mark 1.00 out of 1.00

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

True	False	
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	String arguments are first copied to trapframe and then from trapframe to kernel's other variables.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The arguments to system call are copied to kernel stack in trapasm.S
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The arguments are accessed in the kernel code using esp on the trapframe.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	Integer arguments are copied from user memory to kernel memory using argint()
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The arguments to system call originally reside on process stack.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	Integer arguments are stored in eax, ebx, ecx, etc. registers
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The functions like argint(), argstr() make the system call arguments available in the kernel.

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False

The arguments to system call are copied to kernel stack in trapasm.S: False

The arguments are accessed in the kernel code using esp on the trapframe.: True

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

Integer arguments are copied from user memory to kernel memory using argint(): True

The arguments to system call originally reside on process stack.: True

Integer arguments are stored in eax, ebx, ecx, etc. registers: False

The functions like argint(), argstr() make the system call arguments available in the kernel.: True

[◀ Random Quiz - 6 \(xv6 file system\)](#)

Jump to...

[Homework questions: Basics of MM, xv6 booting ►](#)

Started on Wednesday, 19 April 2023, 8:46 PM

State Finished

Completed on Wednesday, 19 April 2023, 9:26 PM

Time taken 39 mins 31 secs

Grade 15.60 out of 30.00 (51.99%)

Question 1

Incorrect

Mark 0.00 out of 2.00

Select all the correct statements about synchronization primitives.

Select one or more:

- a. Blocking means moving the process to a wait queue and calling scheduler
- b. Blocking means moving the process to a wait queue and spinning ✗
- c. Semaphores are always a good substitute for spinlocks ✗
- d. Spinlocks are good for multiprocessor scenarios, for small critical sections
- e. All synchronization primitives are implemented essentially with some hardware assistance.
- f. Blocking means one process passing over control to another process ✗
- g. Semaphores can be used for synchronization scenarios like ordered execution ✓
- h. Mutexes can be implemented without any hardware assistance
- i. Spinlocks consume CPU time
- j. Mutexes can be implemented using spinlock
- k. Thread that is going to block should not be holding any spinlock
- l. Mutexes can be implemented using blocking and wakeup

Your answer is incorrect.

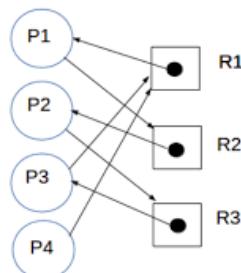
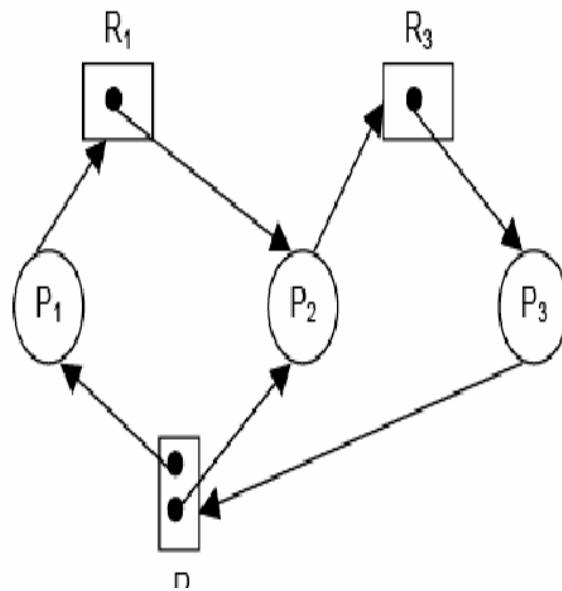
The correct answers are: Spinlocks are good for multiprocessor scenarios, for small critical sections, Spinlocks consume CPU time, Semaphores can be used for synchronization scenarios like ordered execution, Mutexes can be implemented using spinlock, Mutexes can be implemented using blocking and wakeup, Thread that is going to block should not be holding any spinlock, Blocking means moving the process to a wait queue and calling scheduler, All synchronization primitives are implemented essentially with some hardware assistance.

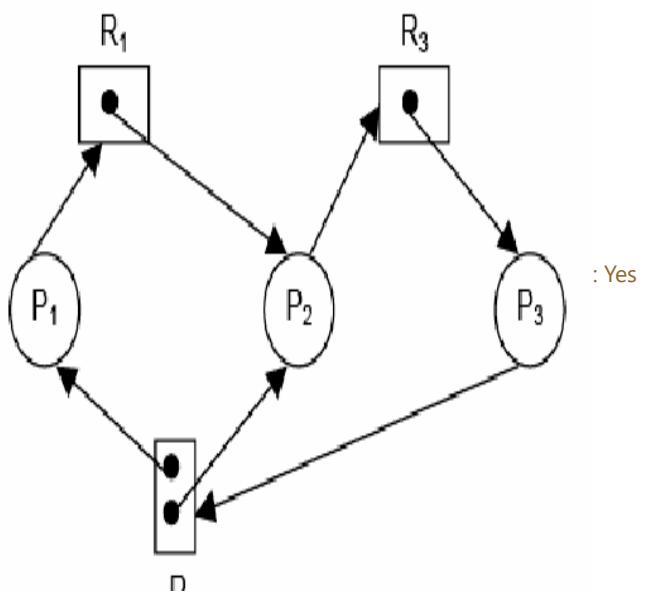
Question 2

Partially correct

Mark 0.50 out of 1.00

For each of the resource allocation diagram shown,
infer whether the graph contains at least one deadlock or not.

Yes **No**



Started on Tuesday, 16 January 2024, 5:08 PM

State Finished

Completed on Tuesday, 16 January 2024, 5:52 PM

Time taken 44 mins 52 secs

Grade 10.20 out of 15.00 (68%)

Question 1

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about the process init on Linuxes/Unixes.

Select one or more:

- a. init can not be killed with SIGKILL ✓
- b. only a process run by 'root' user can exec 'init'
- c. init is created by kernel 'by hand'
- d. any user process can fork and exec init ✗
- e. init typically has a pid=1 ✓
- f. init is created by kernel by forking itself ✗
- g. no process can exec 'init'

Your answer is incorrect.

The correct answers are: init is created by kernel 'by hand', init typically has a pid=1, init can not be killed with SIGKILL, only a process run by 'root' user can exec 'init'

Question 2

Incorrect

Mark 0.00 out of 1.00

Write the possible contents of the file /tmp/xyz after this program.

In the answer if you want to mention any non-text character, then write \0. For example abc\0\0 means abc followed by any two non-text characters

```
int main(int argc, char *argv[]) {  
    int fd1, fd2, n, i;  
    char buf[128];  
  
    fd1 = open("/tmp/xyz", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);  
    write(fd1, "hello", 5);  
    fd2 = open("/tmp/xyz", O_WRONLY, S_IRUSR|S_IWUSR);  
    write(fd2, "bye", 3);  
    close(fd1);  
    close(fd2);  
    return 0;  
}
```

Answer: byelo\0\0



The correct answer is: byelo

Question 3

Partially correct

Mark 0.60 out of 1.00

Select all the correct statements about two modes of CPU operation

Select one or more:

- a. Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode
- b. There is an instruction like 'iret' to return from kernel mode to user mode✓
- c. The two modes are essential for a multitasking system✓
- d. The two modes are essential for a multiprogramming system✓
- e. The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: The two modes are essential for a multiprogramming system, The two modes are essential for a multitasking system, There is an instruction like 'iret' to return from kernel mode to user mode, The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously, Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode

Question 4

Correct

Mark 0.50 out of 0.50

Compare multiprogramming with multitasking

- a. A multitasking system is not necessarily multiprogramming
- b. A multiprogramming system is not necessarily multitasking✓

The correct answer is: A multiprogramming system is not necessarily multitasking

Question 5

Partially correct

Mark 0.25 out of 1.00

Given below is the output of "ps -ef".

Answer the questions based on it.

UID	PID	PPID	C	S	TIME	STIME	TTY	CMD
root	1	0	0	Jan05	?	00:01:08	/sbin/init	splash
root	2	0	0	Jan05	?	00:00:00	[kthreadd]	
root	3	2	0	Jan05	?	00:00:00	[rcu_gp]	
root	4	2	0	Jan05	?	00:00:00	[rcu_par_gp]	
root	9	2	0	Jan05	?	00:00:00	[mm_percpu_wq]	
root	10	2	0	Jan05	?	00:00:00	[rcu_tasks_rude_]	
root	11	2	0	Jan05	?	00:00:00	[rcu_tasks_trace]	
root	12	2	0	Jan05	?	00:00:22	[ksoftirqd/0]	
root	13	2	0	Jan05	?	00:06:29	[rcu_sched]	
root	14	2	0	Jan05	?	00:00:02	[migration/0]	
root	15	2	0	Jan05	?	00:00:00	[idle_inject/0]	
root	16	2	0	Jan05	?	00:00:00	[cpuhp/0]	
root	17	2	0	Jan05	?	00:00:00	[cpuhp/1]	
root	18	2	0	Jan05	?	00:00:00	[idle_inject/1]	
root	19	2	0	Jan05	?	00:00:03	[migration/1]	
root	20	2	0	Jan05	?	00:00:13	[ksoftirqd/1]	
root	22	2	0	Jan05	?	00:00:00	[kworker/1:0H-events_highpri]	
root	23	2	0	Jan05	?	00:00:00	[cpuhp/2]	
root	24	2	0	Jan05	?	00:00:00	[idle_inject/2]	
root	25	2	0	Jan05	?	00:00:01	[migration/2]	
root	26	2	0	Jan05	?	00:00:09	[ksoftirqd/2]	
root	28	2	0	Jan05	?	00:00:00	[kworker/2:0H-kblockd]	
root	29	2	0	Jan05	?	00:00:00	[cpuhp/3]	
root	30	2	0	Jan05	?	00:00:00	[idle_inject/3]	
root	31	2	0	Jan05	?	00:00:02	[migration/3]	
root	32	2	0	Jan05	?	00:00:07	[ksoftirqd/3]	
root	34	2	0	Jan05	?	00:00:00	[kworker/3:0H-events_highpri]	
root	35	2	0	Jan05	?	00:00:00	[cpuhp/4]	
root	36	2	0	Jan05	?	00:00:00	[idle_inject/4]	
root	37	2	0	Jan05	?	00:00:02	[migration/4]	
root	38	2	0	Jan05	?	00:00:06	[ksoftirqd/4]	
root	40	2	0	Jan05	?	00:00:00	[kworker/4:0H-events_highpri]	
root	41	2	0	Jan05	?	00:00:00	[cpuhp/5]	
root	42	2	0	Jan05	?	00:00:00	[idle_inject/5]	
root	43	2	0	Jan05	?	00:00:02	[migration/5]	
root	44	2	0	Jan05	?	00:00:05	[ksoftirqd/5]	
root	46	2	0	Jan05	?	00:00:00	[kworker/5:0H-events_highpri]	
root	47	2	0	Jan05	?	00:00:00	[cpuhp/6]	
root	48	2	0	Jan05	?	00:00:00	[idle_inject/6]	
root	49	2	0	Jan05	?	00:00:02	[migration/6]	
root	50	2	0	Jan05	?	00:00:05	[ksoftirqd/6]	
root	52	2	0	Jan05	?	00:00:00	[kworker/6:0H-events_highpri]	
root	53	2	0	Jan05	?	00:00:00	[cpuhp/7]	
root	54	2	0	Jan05	?	00:00:00	[idle_inject/7]	
root	55	2	0	Jan05	?	00:00:02	[migration/7]	
root	56	2	0	Jan05	?	00:00:06	[ksoftirqd/7]	
root	58	2	0	Jan05	?	00:00:00	[kworker/7:0H-events_highpri]	
root	59	2	0	Jan05	?	00:00:00	[kdevtmpfs]	
root	60	2	0	Jan05	?	00:00:00	[netns]	
root	61	2	0	Jan05	?	00:00:00	[inet_frag_wq]	
root	62	2	0	Jan05	?	00:00:00	[kaudittd]	
root	63	2	0	Jan05	?	00:00:00	[khungtaskd]	
root	64	2	0	Jan05	?	00:00:00	[oom_reaper]	
root	65	2	0	Jan05	?	00:00:00	[writeback]	
root	66	2	0	Jan05	?	00:01:58	[kcompactd0]	
root	67	2	0	Jan05	?	00:00:00	[ksmd]	
root	68	2	0	Jan05	?	00:00:04	[khugepaged]	
root	115	2	0	Jan05	?	00:00:00	[kintegrityd]	
root	116	2	0	Jan05	?	00:00:00	[kblockd]	
root	117	2	0	Jan05	?	00:00:00	[blkcg_punt_bio]	
root	118	2	0	Jan05	?	00:00:00	[tpm_dev_wq]	

root 119 2 0 Jan05 ? 00:00:00 [ata_sff]
root 120 2 0 Jan05 ? 00:00:00 [md]
root 121 2 0 Jan05 ? 00:00:00 [edac-poller]
root 122 2 0 Jan05 ? 00:00:00 [devfreq_wq]
root 123 2 0 Jan05 ? 00:00:00 [watchdogd]
root 129 2 0 Jan05 ? 00:00:00 [irq/25-AMD-Vi]
root 131 2 0 Jan05 ? 00:04:33 [kswapd0]
root 132 2 0 Jan05 ? 00:00:00 [ecryptfs-kthrea]
root 134 2 0 Jan05 ? 00:00:00 [kthrotld]
root 135 2 0 Jan05 ? 00:00:00 [irq/27-pciehp]
root 139 2 0 Jan05 ? 00:00:00 [acpi_thermal_pm]
root 140 2 0 Jan05 ? 00:00:00 [vfio-irqfd-clea]
root 141 2 0 Jan05 ? 00:00:00 [ipv6_addrconf]
root 144 2 0 Jan05 ? 00:00:03 [kworker/6:1H-kblockd]
root 151 2 0 Jan05 ? 00:00:00 [kstrp]
root 154 2 0 Jan05 ? 00:00:00 [zswap-shrink]
root 162 2 0 Jan05 ? 00:00:00 [charger_manager]
root 164 2 0 Jan05 ? 00:00:03 [kworker/4:1H-kblockd]
root 197 2 0 Jan05 ? 00:00:03 [kworker/3:1H-kblockd]
root 213 2 0 Jan05 ? 00:00:03 [kworker/7:1H-kblockd]
root 215 2 0 Jan05 ? 00:00:00 [nvme-wq]
root 216 2 0 Jan05 ? 00:00:00 [nvme-reset-wq]
root 217 2 0 Jan05 ? 00:00:00 [nvme-delete-wq]
root 223 2 0 Jan05 ? 00:00:00 [irq/42-ELAN2513]
root 224 2 0 Jan05 ? 00:04:20 [irq/41-ELAN071B]
root 225 2 0 Jan05 ? 00:00:00 [cryptd]
root 226 2 0 Jan05 ? 00:00:00 [amd_iommu_v2]
root 249 2 0 Jan05 ? 00:00:00 [ttm_swap]
root 250 2 0 Jan05 ? 00:20:29 [gfx]
root 251 2 0 Jan05 ? 00:00:00 [comp_1.0.0]
root 252 2 0 Jan05 ? 00:00:00 [comp_1.1.0]
root 253 2 0 Jan05 ? 00:00:00 [comp_1.2.0]
root 254 2 0 Jan05 ? 00:00:00 [comp_1.3.0]
root 255 2 0 Jan05 ? 00:00:00 [comp_1.0.1]
root 256 2 0 Jan05 ? 00:00:00 [comp_1.1.1]
root 257 2 0 Jan05 ? 00:00:00 [comp_1.2.1]
root 258 2 0 Jan05 ? 00:00:00 [comp_1.3.1]
root 259 2 0 Jan05 ? 00:00:27 [sdma0]
root 260 2 0 Jan05 ? 00:00:00 [vcn_dec]
root 261 2 0 Jan05 ? 00:00:00 [vcn_enc0]
root 262 2 0 Jan05 ? 00:00:00 [vcn_enc1]
root 263 2 0 Jan05 ? 00:00:00 [jpeg_dec]
root 265 2 0 Jan05 ? 00:00:00 [card0-crtc0]
root 266 2 0 Jan05 ? 00:00:00 [card0-crtc1]
root 267 2 0 Jan05 ? 00:00:00 [card0-crtc2]
root 268 2 0 Jan05 ? 00:00:00 [card0-crtc3]
root 271 2 0 Jan05 ? 00:00:03 [kworker/5:1H-kblockd]
root 277 2 0 Jan05 ? 00:00:03 [kworker/1:1H-kblockd]
root 320 2 0 Jan05 ? 00:00:00 [raid5wq]
root 380 2 0 Jan05 ? 00:00:11 [jbd2/nvme0n1p5-]
root 381 2 0 Jan05 ? 00:00:00 [ext4-rsv-conver]
root 432 2 0 Jan05 ? 00:00:03 [kworker/2:1H-kblockd]
root 446 1 0 Jan05 ? 00:01:16 /lib/systemd/systemd-journald
root 470 2 0 Jan05 ? 00:00:00 [rpciod]
root 473 2 0 Jan05 ? 00:00:00 [xpriod]
root 474 2 0 Jan05 ? 00:00:00 bpfILTER_umh
root 503 1 0 Jan05 ? 00:00:07 /lib/systemd/systemd-udevd
root 517 2 0 Jan05 ? 00:00:00 [loop0]
root 554 2 0 Jan05 ? 00:00:00 [loop1]
root 563 2 0 Jan05 ? 00:00:00 [loop2]
root 586 2 0 Jan05 ? 00:00:00 [loop3]
root 588 2 0 Jan05 ? 00:00:00 [loop4]
root 589 2 0 Jan05 ? 00:00:00 [loop5]
root 612 2 0 Jan05 ? 00:00:00 [loop6]
root 613 2 0 Jan05 ? 00:00:00 [loop7]
root 629 2 0 Jan05 ? 00:00:00 [loop8]
root 637 2 0 Jan05 ? 00:00:00 [cfg80211]
root 676 2 0 Jan05 ? 00:05:00 [irq/75-iwlwifi:]
root 678 2 0 Jan05 ? 00:01:07 [irq/76-iwlwifi:]
root 682 2 0 Jan05 ? 00:01:27 [irq/77-iwlwifi:]

root 688 2 0 Jan05 ? 00:00:49 [irq/78-iwlwifi:]
root 695 2 0 Jan05 ? 00:01:39 [irq/79-iwlwifi:]
root 700 2 0 Jan05 ? 00:01:22 [irq/80-iwlwifi:]
root 703 2 0 Jan05 ? 00:01:13 [irq/81-iwlwifi:]
root 704 2 0 Jan05 ? 00:01:38 [irq/82-iwlwifi:]
root 708 2 0 Jan05 ? 00:00:44 [irq/83-iwlwifi:]
root 713 2 0 Jan05 ? 00:00:00 [loop9]
root 715 2 0 Jan05 ? 00:00:00 [irq/84-iwlwifi:]
root 782 2 0 Jan05 ? 00:00:00 [loop10]
root 797 2 0 Jan05 ? 00:00:00 [loop11]
root 811 2 0 Jan05 ? 00:00:00 [loop12]
root 838 2 0 Jan05 ? 00:00:00 [loop13]
root 847 2 0 Jan05 ? 00:00:00 [loop14]
root 879 2 0 Jan05 ? 00:00:00 [loop15]
root 884 2 0 Jan05 ? 00:00:00 [loop16]
root 885 2 0 Jan05 ? 00:00:00 [loop17]
root 945 2 0 Jan05 ? 00:00:00 [loop18]
root 946 2 0 Jan05 ? 00:00:00 [loop19]
root 947 2 0 Jan05 ? 00:00:00 [loop20]
root 1012 2 0 Jan05 ? 00:00:00 [jbd2/nvme0n1p8-]
root 1013 2 0 Jan05 ? 00:00:00 [ext4-rsv-conver]
root 1015 2 0 Jan05 ? 00:01:09 [jbd2/nvme0n1p7-]
root 1016 2 0 Jan05 ? 00:00:00 [ext4-rsv-conver]
_rpc 1062 1 0 Jan05 ? 00:00:00 /sbin/rpcbind -f -w
systemd+ 1063 1 0 Jan05 ? 00:01:24 /lib/systemd/systemd-resolved
systemd+ 1064 1 0 Jan05 ? 00:00:00 /lib/systemd/systemd-timesyncd
root 1144 1 0 Jan05 ? 00:00:46 /usr/sbin/acpid
avahi 1146 1 0 Jan05 ? 00:00:06 avahi-daemon: running [abhijit-laptop.local]
root 1149 1 0 Jan05 ? 00:00:01 /usr/lib/bluetooth/bluetoothd
message+ 1150 1 0 Jan05 ? 00:04:21 /usr/bin/dbus-daemon --system --address=systemd: --nofork --
nopidfile --systemd-activation --syslog-only
root 1152 1 0 Jan05 ? 00:03:12 /usr/sbin/NetworkManager -no-daemon
root 1157 1 0 Jan05 ? 00:01:02 /usr/sbin/iio-sensor-proxy
root 1159 1 0 Jan05 ? 00:00:27 /usr/sbin/irqbalance --foreground
root 1162 1 0 Jan05 ? 00:00:01 /usr/bin/lxcs /var/lib/lxcs
root 1165 1 0 Jan05 ? 00:00:00 /usr/bin/python3 /usr/bin/networkd-dispatcher --run-startup-
triggers
root 1170 1 0 Jan05 ? 00:00:29 /usr/lib/polkit-1/polkitd --no-debug
syslog 1175 1 0 Jan05 ? 00:00:20 /usr/sbin/rsyslogd -n -iNONE
root 1182 1 0 Jan05 ? 00:01:13 /usr/lib/snapd/snapd
root 1187 1 0 Jan05 ? 00:00:12 /usr/lib/accountsservice/accounts-daemon
root 1192 1 0 Jan05 ? 00:00:00 /usr/sbin/cron -f
root 1198 1 0 Jan05 ? 00:00:00 /usr/libexec/switcheroo-control
root 1201 1 0 Jan05 ? 00:00:12 /lib/systemd/systemd-logind
root 1202 1 0 Jan05 ? 00:00:07 /lib/systemd/systemd-machined
root 1203 1 0 Jan05 ? 00:01:28 /usr/lib/udisks2/udisksd
root 1204 1 0 Jan05 ? 00:00:15 /sbin/wpa_supplicant -u -s -0 /run/wpa_supplicant
daemon 1209 1 0 Jan05 ? 00:00:00 /usr/sbin/atd -f
avahi 1216 1146 0 Jan05 ? 00:00:00 avahi-daemon: chroot helper
docker-+ 1279 1 0 Jan05 ? 00:00:22 /usr/bin/docker-registry serve /etc/docker/registry/config.yml
root 1282 1 0 Jan05 ? 00:00:00 /usr/bin/python3 /usr/bin/twistd3 --nodaemon --pidfile= epoptes
jenkins 1285 1 0 Jan05 ? 00:15:01 /usr/bin/java -Djava.awt.headless=true -jar
/usr/share/java/jenkins.war --webroot=/var/cache/jenkins/war --httpPort=8080
root 1296 1 0 Jan05 ? 00:00:18 php-fpm: master process (/etc/php/7.4/fpm/php-fpm.conf)
vnstat 1314 1 0 Jan05 ? 00:00:15 /usr/sbin/vnstatd -n
root 1326 1 0 Jan05 ? 00:00:02 /usr/sbin/ModemManager
root 1327 1 0 Jan05 ? 00:00:31 /usr/bin/anydesk --service
colord 1359 1 0 Jan05 ? 00:00:01 /usr/libexec/colord
root 1374 1 0 Jan05 ? 00:00:00 /usr/sbin/gdm3
root 1420 1 0 Jan05 ? 00:00:14 /usr/sbin/apache2 -k start
www-data 1436 1296 0 Jan05 ? 00:00:00 php-fpm: pool www
www-data 1437 1296 0 Jan05 ? 00:00:00 php-fpm: pool www
mysql 1461 1 0 Jan05 ? 00:38:52 /usr/sbin/mysqld
root 1490 1 0 Jan05 ? 00:00:04 /usr/sbin/libvirtd
root 1491 1 0 Jan05 ? 00:00:00 /usr/bin/python3 /usr/share/unattended-upgrades/unattended-upgrade-
shutdown --wait-for-signal
rtkit 1593 1 0 Jan05 ? 00:00:07 /usr/libexec/rtkit-daemon
libvirt+ 1766 1 0 Jan05 ? 00:00:00 /usr/sbin/dnsmasq --conf-file=/var/lib/libvirt/dnsmasq/default.conf
--leasefile-ro --dhcp-script=/usr/lib/libvirt/libvirt_leaseshelper
root 1767 1766 0 Jan05 ? 00:00:00 /usr/sbin/dnsmasq --conf-file=/var/lib/libvirt/dnsmasq/default.conf

```
--leasefile-ro --dhcp-script=/usr/lib/libvirt/libvirt_leaseshelper
root      1859      1  0 Jan05 ?          00:00:18 /usr/lib/upower/upowerd
root      1995      1  0 Jan05 ?          00:00:00 /opt/saltstack/salt/run/run minion
root      2041    1995  0 Jan05 ?          00:05:18 /opt/saltstack/salt/run/run minion MultiMinionProcessManager
MinionProcessManager
root      2278      1  0 Jan05 ?          00:00:50 /usr/bin/dockerd -H fd:// --
containerd=/run/containerd/containerd.sock
root      2282      1  0 Jan05 ?          00:00:17 /usr/sbin/inetd
whoopsie  2302      1  0 Jan05 ?          00:00:01 /usr/bin/whoopsie -f
kernoops  2330      1  0 Jan05 ?          00:00:19 /usr/sbin/kerneloops --test
kernoops  2341      1  0 Jan05 ?          00:00:19 /usr/sbin/kerneloops
root      2366      2  0 Jan05 ?          00:00:00 [iprt-VBoxWQueue]
lxc-dns+  2370      1  0 Jan05 ?          00:00:00 dnsmasq --conf-file=/dev/null -u lxc-dnsmasq --strict-order --bind-
interfaces --pid-file=/run/lxc/dnsmasq.pid --listen-address 10.0.3.1 --dhcp-range 10.0.3.2,10.0.3.254 --dhcp-lease-
max=253 --dhcp-no-override --except-interface=lo --interface=lxcbr0 --dhcp-
leasefile=/var/lib/misc/dnsmasq.lxcbr0.leases --dhcp-authoritative
root      2404      2  0 Jan05 ?          00:00:00 [iprt-VBoxTscThr]
root      3508      1  0 Jan05 ?          00:00:01 /usr/lib/postfix/sbin/master -w
root      3615    1374  0 Jan05 ?          00:00:01 gdm-session-worker [pam/gdm-password]
abhijit   3629      1  0 Jan05 ?          00:00:17 /lib/systemd/systemd --user
abhijit   3630    3629  0 Jan05 ?          00:00:00 (sd-pam)
abhijit   3636    3629  1 Jan05 ?          04:31:21 /usr/bin/pulseaudio --daemonize=no --log-target=journal
abhijit   3638    3629  0 Jan05 ?          00:20:35 /usr/libexec/tracker-miner-fs
abhijit   3642    3629  0 Jan05 ?          00:01:11 /usr/bin/dbus-daemon --session --address=systemd: --nofork --
nopidfile --systemd-activation --syslog-only
abhijit   3644      1  0 Jan05 ?          00:00:03 /usr/bin/gnome-keyring-daemon --daemonize --login
abhijit   3662    3629  0 Jan05 ?          00:00:01 /usr/libexec/gvfsd
abhijit   3667    3629  0 Jan05 ?          00:00:00 /usr/libexec/gvfsd-fuse /run/user/1000/gvfs -f -o big_writes
abhijit   3673    3629  0 Jan05 ?          00:00:02 /usr/libexec/gvfs-udisks2-volume-monitor
abhijit   3681    3629  0 Jan05 ?          00:00:01 /usr/libexec/gvfs-mtp-volume-monitor
abhijit   3685    3629  0 Jan05 ?          00:00:11 /usr/libexec/gvfs-afc-volume-monitor
abhijit   3690    3629  0 Jan05 ?          00:00:01 /usr/libexec/gvfs-gphoto2-volume-monitor
abhijit   3695    3629  0 Jan05 ?          00:00:01 /usr/libexec/gvfs-goa-volume-monitor
abhijit   3700    3629  0 Jan05 ?          00:00:01 /usr/libexec/goa-daemon
root      3701      2  0 Jan05 ?          00:00:00 [krfcomm]
abhijit   3708    3629  0 Jan05 ?          00:00:04 /usr/libexec/goa-identity-service
abhijit   3724    3615  tty2      00:00:00 /usr/lib/gdm3/gdm-x-session --run-script env
GNOME_SHELL_SESSION_MODE=ubuntu /usr/bin/gnome-session --systemd --session=ubuntu
abhijit   3726    3724  1 Jan05 tty2      02:47:57 /usr/lib/xorg/Xorg vt2 -displayfd 3 -auth
/run/user/1000/gdm/Xauthority -background none -noreset -keeptty -verbose 3
abhijit   3747    3724  0 Jan05 tty2      00:00:00 /usr/libexec/gnome-session-binary --systemd --systemd --
session=ubuntu
abhijit   3816    3747  0 Jan05 ?          00:00:01 /usr/bin/ssh-agent /usr/bin/im-launch env
GNOME_SHELL_SESSION_MODE=ubuntu /usr/bin/gnome-session --systemd --session=ubuntu
abhijit   3845    3629  0 Jan05 ?          00:00:00 /usr/libexec/at-spi-bus-launcher
abhijit   3850    3845  0 Jan05 ?          00:00:07 /usr/bin/dbus-daemon --config-file=/usr/share/defaults/at-
spi2/accessibility.conf --nofork --print-address 3
abhijit   3869    3629  0 Jan05 ?          00:00:00 /usr/libexec/gnome-session-ctl --monitor
abhijit   3876    3629  0 Jan05 ?          00:00:04 /usr/libexec/gnome-session-binary --systemd-service --
session=ubuntu
abhijit   3890    3629  1 Jan05 ?          03:43:47 /usr/bin/gnome-shell
abhijit   3927    3890  0 Jan05 ?          00:31:49 ibus-daemon --panel disable --xim
abhijit   3931    3927  0 Jan05 ?          00:00:00 /usr/libexec/ibus-dconf
abhijit   3932    3927  0 Jan05 ?          00:01:51 /usr/libexec/ibus-extension-gtk3
abhijit   3934    3629  0 Jan05 ?          00:00:04 /usr/libexec/ibus-x11 --kill-daemon
abhijit   3937    3629  0 Jan05 ?          00:00:02 /usr/libexec/ibus-portal
abhijit   3949    3629  0 Jan05 ?          00:00:27 /usr/libexec/at-spi2-registryd --use-gnome-session
abhijit   3953    3629  0 Jan05 ?          00:00:00 /usr/libexec/xdg-permission-store
abhijit   3955    3629  0 Jan05 ?          00:00:01 /usr/libexec/gnome-shell-calendar-server
abhijit   3964    3629  0 Jan05 ?          00:00:00 /usr/libexec/evolution-source-registry
abhijit   3973    3629  0 Jan05 ?          00:00:02 /usr/libexec/evolution-calendar-factory
abhijit   3986    3629  0 Jan05 ?          00:00:01 /usr/libexec/dconf-service
abhijit   3992    3629  0 Jan05 ?          00:00:01 /usr/libexec/evolution-addressbook-factory
abhijit   4007    3629  0 Jan05 ?          00:00:00 /usr/bin/gjs /usr/share/gnome-shell/org.gnome.Shell.Notifications
abhijit   4023    3629  0 Jan05 ?          00:00:00 /usr/libexec/gsd-a11y-settings
abhijit   4025    3629  0 Jan05 ?          00:00:08 /usr/libexec/gsd-color
abhijit   4029    3629  0 Jan05 ?          00:00:00 /usr/libexec/gsd-datetime
abhijit   4032    3629  0 Jan05 ?          00:00:21 /usr/libexec/gsd-housekeeping
abhijit   4033    3629  0 Jan05 ?          00:00:05 /usr/libexec/gsd-keyboard
abhijit   4036    3629  0 Jan05 ?          00:00:12 /usr/libexec/gsd-media-keys
```

abhijit 4037 3629 0 Jan05 ? 00:00:11 /usr/libexec/gsd-power
abhijit 4038 3629 0 Jan05 ? 00:00:00 /usr/libexec/gsd-print-notifications
abhijit 4039 3629 0 Jan05 ? 00:00:01 /usr/libexec/gsd-rfkill
abhijit 4041 3629 0 Jan05 ? 00:00:01 /usr/libexec/gsd-screensaver-proxy
abhijit 4042 3876 0 Jan05 ? 00:01:06 /usr/lib/x86_64-linux-gnu/libexec/kdeconnectd
abhijit 4045 3629 0 Jan05 ? 00:00:35 /usr/libexec/gsd-sharing
abhijit 4047 3629 0 Jan05 ? 00:00:00 /usr/libexec/gsd-smartcard
abhijit 4051 3629 0 Jan05 ? 00:00:00 /usr/libexec/gsd-sound
abhijit 4057 3629 0 Jan05 ? 00:00:00 /usr/libexec/gsd-usb-protection
abhijit 4063 3629 0 Jan05 ? 00:00:05 /usr/libexec/gsd-wacom
abhijit 4071 3629 0 Jan05 ? 00:00:00 /usr/libexec/gsd-wwan
abhijit 4072 3876 0 Jan05 ? 00:01:00 baloo_file
abhijit 4075 3876 0 Jan05 ? 00:00:00 /usr/libexec/gsd-disk-utility-notify
abhijit 4076 3629 0 Jan05 ? 00:00:08 /usr/libexec/gsd-xsettings
abhijit 4078 3876 0 Jan05 ? 00:00:10 /usr/bin/python3 /usr/bin/blueman-applet
abhijit 4082 3876 0 Jan05 ? 00:00:14 /usr/bin/anydesk --tray
abhijit 4108 3876 0 Jan05 ? 00:00:00 /usr/lib/x86_64-linux-gnu/indicator-messages/indicator-messages-service
abhijit 4109 3876 0 Jan05 ? 00:00:06 /usr/libexec/evolution-data-server/evolution-alarm-notify
abhijit 4129 3629 0 Jan05 ? 00:00:50 /snap/snap-store/959/usr/bin/snap-store --gapplication-service
abhijit 4191 3629 0 Jan05 ? 00:00:00 /usr/libexec/gsd-printer
abhijit 4219 3629 0 Jan05 ? 00:00:03 /usr/libexec/xdg-document-portal
abhijit 4265 3927 0 Jan05 ? 00:03:20 /usr/libexec/ibus-engine-simple
abhijit 4291 3629 0 Jan05 ? 00:00:10 /usr/bin/python3 /usr/bin/blueman-tray
abhijit 4301 3629 0 Jan05 ? 00:00:00 /usr/lib/bluetooth/obexd
abhijit 4377 3662 0 Jan05 ? 00:00:02 /usr/libexec/gvfsd-trash --spawner :1.3 /org/gtk/gvfs/exec_spaw/0
abhijit 4395 3629 0 Jan05 ? 00:00:12 /usr/libexec/xdg-desktop-portal
abhijit 4399 3629 0 Jan05 ? 00:01:08 /usr/libexec/xdg-desktop-portal-gtk
abhijit 4480 3629 0 Jan05 ? 00:00:21 /usr/libexec/gvfsd-metadata
abhijit 5687 3662 0 Jan05 ? 00:00:00 /usr/libexec/gvfsd-network --spawner :1.3 /org/gtk/gvfs/exec_spaw/1
abhijit 5701 3662 0 Jan05 ? 00:00:01 /usr/libexec/gvfsd-dnssd --spawner :1.3 /org/gtk/gvfs/exec_spaw/3
root 6228 2 0 Jan05 ? 00:00:00 [kdmflush]
root 6236 2 0 Jan05 ? 00:00:00 [kcryptd_io/253:]
root 6237 2 0 Jan05 ? 00:00:00 [kcryptd/253:0]
root 6238 2 0 Jan05 ? 00:00:23 [dmcrypt_write/2]
root 6260 2 0 Jan05 ? 00:00:24 [jbd2/dm-0-8]
root 6261 2 0 Jan05 ? 00:00:00 [ext4-rsv-conver]
abhijit 6421 3927 0 Jan05 ? 00:00:36 /usr/lib/ibus/ibus-engine-m17n --ibus
abhijit 6434 3876 0 Jan05 ? 00:00:13 update-notifier
abhijit 30565 3629 0 Jan05 ? 00:08:56 /usr/libexec/gnome-terminal-server
abhijit 30576 30565 0 Jan05 pts/0 00:00:00 bash
abhijit 131017 364845 1 Jan09 ? 03:12:27 /usr/lib/virtualbox/VirtualBoxVM --comment ubuntu 18.04 --startvm 45993a5c-3ded-452f-941e-4579d12c1ad9 --no-startvm-errormsgbox
abhijit 159668 30565 0 Jan09 pts/10 00:00:00 bash
abhijit 161637 30565 0 Jan05 pts/1 00:00:01 bash
abhijit 171109 159668 0 Jan09 pts/10 00:00:33 evince.....
abhijit 171114 3629 0 Jan09 ? 00:00:00 /usr/libexec/evinced
abhijit 204055 3629 0 Jan10 ? 00:00:13 /usr/bin/python3 /usr/bin/update-manager --no-update --no-focus-on-map
abhijit 270190 3629 0 Jan05 ? 00:56:34 /usr/lib/x86_64-linux-gnu/libexec/kactivitymanagerd
abhijit 270199 3629 0 Jan05 ? 00:00:24 /usr/bin/kglobalaccel5
abhijit 270207 3629 0 Jan05 ? 00:00:00 kdeinit5: Running...
abhijit 270208 270207 0 Jan05 ? 00:00:19 /usr/lib/x86_64-linux-gnu/libexec/kf5/klauncher --fd=8
abhijit 279971 3629 0 Jan05 ? 01:20:12 telegram-desktop
abhijit 280011 279971 0 Jan05 ? 00:00:00 sh -c /usr/lib/x86_64-linux-gnu/libproxy/0.4.15/pxgsettings
org.gnome.system.proxy org.gnome.system.proxy.http org.gnome.system.proxy.https org.gnome.system.proxy.ftp
org.gnome.system.proxy.socks
abhijit 280015 280011 0 Jan05 ? 00:00:00 /usr/lib/x86_64-linux-gnu/libproxy/0.4.15/pxgsettings
org.gnome.system.proxy org.gnome.system.proxy.http org.gnome.system.proxy.https org.gnome.system.proxy.ftp
org.gnome.system.proxy.socks
abhijit 325756 3629 0 Jan12 ? 00:01:09 /usr/libexec/tracker-store
root 326592 1 0 Jan12 ? 00:00:00 sshd: /usr/sbin/sshd -D [listener] 0 of 10-100 startups
abhijit 347912 3644 0 Jan06 ? 00:00:00 /usr/bin/ssh-agent -D -a /run/user/1000/keyring/.ssh
root 348716 1 0 Jan12 ? 00:00:28 /usr/bin/containerd
abhijit 351306 3629 3 Jan12 ? 03:32:42 /usr/lib/firefox/firefox
abhijit 351429 351306 0 Jan12 ? 00:00:00 /usr/lib/firefox/firefox -contentproc -parentBuildID 20240108143603
-prefsLen 37272 -prefMapSize 247458 -appDir /usr/lib/firefox/browser {83364ade-74ec-4bcc-94f8-9f3779a869f1} 351306 true
socket
abhijit 351458 351306 0 Jan12 ? 00:29:34 /usr/lib/firefox/firefox -contentproc -childID 1 -isForBrowser -
prefsLen 37337 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -

appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {4dbb54ee-e1a2-41e4-b10b-587526532b78} 351306
true tab
abhijit 351495 351306 0 Jan12 ? 00:04:06 /usr/lib/firefox/firefox -contentproc -childID 2 -isForBrowser -
prefsLen 38090 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {02f1a4e7-d831-4a5f-a9fd-59a809bb03e8} 351306
true tab
abhijit 351717 351306 0 Jan12 ? 00:01:14 /usr/lib/firefox/firefox -contentproc -parentBuildID 20240108143603 -
-sandboxingKind 0 -prefsLen 42823 -prefMapSize 247458 -appDir /usr/lib/firefox/browser {a1e46009-cfd1-46c8-ac9c-
19ba8bf3f46a} 351306 true utility
abhijit 351844 351306 0 Jan12 ? 00:07:23 /usr/lib/firefox/firefox -contentproc -parentBuildID 20240108143603 -
-prefsLen 43306 -prefMapSize 247458 -appDir /usr/lib/firefox/browser {91008bef-b6d6-452f-b4a6-e78d887b3569} 351306 true
rdd
abhijit 353500 3662 0 Jan06 ? 00:00:00 /usr/libexec/gvfsd-http --spawner :1.3 /org/gtk/gvfs/exec_spaw/4
abhijit 364809 3890 0 Jan06 ? 00:19:20 /usr/lib/virtualbox/VirtualBox
abhijit 364837 3629 0 Jan06 ? 00:16:49 /usr/lib/virtualbox/VBoxXPComIPCD
abhijit 364845 3629 0 Jan06 ? 00:29:44 /usr/lib/virtualbox/VBoxSVC --auto-shutdown
root 364957 2 0 Jan06 ? 00:00:00 [dio/dm-0]
abhijit 369749 30565 0 Jan06 pts/3 00:00:00 bash
abhijit 369879 30565 0 Jan06 pts/4 00:00:00 bash
abhijit 379620 30565 0 Jan06 pts/6 00:00:00 bash
abhijit 381917 3890 0 Jan06 ? 00:00:46 flameshot
root 386201 2 0 Jan06 ? 00:00:02 [kworker/0:2H-acpi_thermal_pm]
postfix 389291 3508 0 Jan06 ? 00:00:00 qmgr -l -t unix -u
root 486171 2 0 Jan12 ? 00:00:01 [kworker/0:0H-kblockd]
abhijit 527395 3629 0 Jan12 ? 00:00:49 /usr/bin/gedit --gapplication-service
abhijit 551299 351306 0 Jan12 ? 00:01:54 /usr/lib/firefox/firefox -contentproc -childID 438 -isForBrowser -
prefsLen 35109 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {c4ebd70d-bad7-4e7c-9066-39827103c84e} 351306
true tab
abhijit 552002 3890 0 Jan12 ? 00:03:32 /opt/Signal/signal-desktop --no-sandbox
abhijit 552005 552002 0 Jan12 ? 00:00:00 /opt/Signal/signal-desktop --type=zygote --no-zygote-sandbox --no-
sandbox
abhijit 552006 552002 0 Jan12 ? 00:00:00 /opt/Signal/signal-desktop --type=zygote --no-sandbox
abhijit 552037 552005 0 Jan12 ? 00:03:28 /opt/Signal/signal-desktop --type=gpu-process --no-sandbox --
enable-crash-reporter=18887fa1-4d37-46f0-bf9f-b37513c81cf9,no_channel --user-data-dir=/home/abhijit/.config/Signal --
gpu-
preferences=WAAAAAAAAGAAAAAABgAAAAAA4AAAAAAAABAAAAAGAAAAAAAYAAAA
--use-gl=angle --use-angle=swiftshader-webgl --shared-files --field-trial-
handle=0,i,3213315393136307567,15155041764863120512,262144 --disable-
features=HardwareMediaKeyHandling,SpareRendererForSitePerProcess
abhijit 552045 552002 0 Jan12 ? 00:00:04 /opt/Signal/signal-desktop --type=utility --utility-sub-
type=network.mojom.NetworkService --lang=en-GB --service-sandbox-type=none --no-sandbox --enable-crash-
reporter=18887fa1-4d37-46f0-bf9f-b37513c81cf9,no_channel --user-data-dir=/home/abhijit/.config/Signal --shared-
files=v8_context_snapshot_data:100 --field-trial-handle=0,i,3213315393136307567,15155041764863120512,262144 --disable-
features=HardwareMediaKeyHandling,SpareRendererForSitePerProcess
abhijit 552103 552002 1 Jan12 ? 00:57:51 /opt/Signal/signal-desktop --type=renderer --enable-crash-
reporter=18887fa1-4d37-46f0-bf9f-b37513c81cf9,no_channel --user-data-dir=/home/abhijit/.config/Signal --app-
path=/opt/Signal/resources/app.asar --no-sandbox --no-zygote --enable-blink-features=CSSPseudoDir,CSSLogical --disable-
blink-features=Accelerated2dCanvas,AcceleratedSmallCanvases --first-renderer-process --no-sandbox --disable-gpu-
compositing --lang=en-GB --num-raster-threads=4 --enable-main-frame-before-activation --renderer-client-id=4 --time-
ticks-at-unix-epoch=-1704847690836396 --launch-time-ticks=218630770368 --shared-files=v8_context_snapshot_data:100 --
field-trial-handle=0,i,3213315393136307567,15155041764863120512,262144 --disable-
features=HardwareMediaKeyHandling,SpareRendererForSitePerProcess
abhijit 552149 552002 0 Jan12 ? 00:00:12 /opt/Signal/signal-desktop --type=utility --utility-sub-
type=audio.mojom.AudioService --lang=en-GB --service-sandbox-type=none --no-sandbox --enable-crash-reporter=18887fa1-
4d37-46f0-bf9f-b37513c81cf9,no_channel --user-data-dir=/home/abhijit/.config/Signal --shared-
files=v8_context_snapshot_data:100 --field-trial-handle=0,i,3213315393136307567,15155041764863120512,262144 --disable-
features=HardwareMediaKeyHandling,SpareRendererForSitePerProcess
abhijit 554660 351306 0 Jan13 ? 00:04:20 /usr/lib/firefox/firefox -contentproc -childID 442 -isForBrowser -
prefsLen 35109 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {2cffa6f3-6f8c-4a3f-880d-01cc71baf983} 351306
true tab
abhijit 554855 351306 4 Jan13 ? 03:24:06 /usr/lib/firefox/firefox -contentproc -childID 445 -isForBrowser -
prefsLen 35109 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {abcba8e9-5b97-4281-ad64-1356841dcf34} 351306
true tab
abhijit 556349 351306 0 Jan13 ? 00:01:18 /usr/lib/firefox/firefox -contentproc -childID 451 -isForBrowser -
prefsLen 35109 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {6318d453-643c-4b85-9f0f-bfd0a429cd41} 351306
true tab

abhijit 557600 3629 0 Jan13 ? 00:00:08 /usr/lib/speech-dispatcher-modules/sd_espeak-ng /etc/speech-dispatcher/modules/espeak-ng.conf
abhijit 557607 3629 0 Jan13 ? 00:00:08 /usr/lib/speech-dispatcher-modules/sd_dummy /etc/speech-dispatcher/modules/dummy.conf
abhijit 557610 3629 0 Jan13 ? 00:00:08 /usr/lib/speech-dispatcher-modules/sd_generic /etc/speech-dispatcher/modules/mary-generic.conf
abhijit 557613 3629 0 Jan13 ? 00:00:00 /usr/bin/speech-dispatcher --spawn --communication-method unix_socket --socket-path /run/user/1000/speech-dispatcher/speechd.sock
abhijit 571320 351306 0 Jan13 ? 00:01:16 /usr/lib/firefox/firefox -contentproc -childID 486 -isForBrowser -prefsLen 35110 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {71966a59-b1df-4c44-975c-62d7cd41188c} 351306 true tab
abhijit 571410 351306 0 Jan13 ? 00:01:34 /usr/lib/firefox/firefox -contentproc -childID 488 -isForBrowser -prefsLen 35110 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {bfbc15ff-6ca5-4837-b97e-334e0bfc8855} 351306 true tab
abhijit 571582 351306 0 Jan13 ? 00:03:32 /usr/lib/firefox/firefox -contentproc -childID 492 -isForBrowser -prefsLen 35110 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {04e791e5-9cde-4e03-9211-141bdb440139} 351306 true tab
abhijit 591339 4139495 0 Jan13 pts/9 00:00:00 vi mh-pyt.txt
root 594356 2 0 Jan13 ? 00:00:00 [dio/nvme0n1p7]
abhijit 594726 30565 0 Jan13 pts/5 00:00:00 bash
abhijit 791421 351306 1 Jan14 ? 00:40:27 /usr/lib/firefox/firefox -contentproc -childID 839 -isForBrowser -prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {b6f69be3-7ebe-43c9-9d55-c72161180b5e} 351306 true tab
abhijit 794226 351306 0 Jan14 ? 00:02:03 /usr/lib/firefox/firefox -contentproc -childID 848 -isForBrowser -prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {28fddc81-ea39-4496-a1b6-b70d9a8a9c31} 351306 true tab
abhijit 797871 351306 0 Jan14 ? 00:02:41 /usr/lib/firefox/firefox -contentproc -childID 851 -isForBrowser -prefsLen 35226 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {b254a3c7-e1f7-4245-9fcf-820074a4e69f} 351306 true tab
abhijit 799607 30565 0 Jan14 pts/11 00:00:00 bash
abhijit 803503 30565 0 Jan14 pts/12 00:00:00 bash
abhijit 815513 799607 0 Jan15 pts/11 00:00:00 vi timetable-todo2
abhijit 821738 351306 0 Jan15 ? 00:07:30 /usr/lib/firefox/firefox -contentproc -childID 995 -isForBrowser -prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {90cad852-941e-44b3-8fa4-0d44cbcfda7a} 351306 true tab
abhijit 895193 351306 0 Jan15 ? 00:00:04 /usr/lib/firefox/firefox -contentproc -childID 1131 -isForBrowser -prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {218262bd-4f9e-423a-9368-4d2e44f33125} 351306 true tab
root 942398 1 0 10:40 ? 00:00:00 /usr/sbin/cupsd -l
root 942399 1 0 10:40 ? 00:00:00 /usr/sbin/cups-browsed
www-data 942465 1420 0 10:40 ? 00:00:00 /usr/sbin/apache2 -k start
www-data 942466 1420 0 10:40 ? 00:00:00 /usr/sbin/apache2 -k start
www-data 942468 1420 0 10:40 ? 00:00:00 /usr/sbin/apache2 -k start
www-data 942469 1420 0 10:40 ? 00:00:00 /usr/sbin/apache2 -k start
www-data 942470 1420 0 10:40 ? 00:00:00 /usr/sbin/apache2 -k start
www-data 942471 1420 0 10:40 ? 00:00:00 /usr/sbin/apache2 -k start
abhijit 954109 30565 0 11:21 pts/2 00:00:00 bash
abhijit 961628 351306 3 12:36 ? 00:05:34 /usr/lib/firefox/firefox -contentproc -childID 1704 -isForBrowser -prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {cb82b978-ef9d-4a76-b59c-5b13b652a0ec} 351306 true tab
root 961877 2 0 12:36 ? 00:00:02 [kworker/u32:5-events_unbound]
root 962113 2 0 12:39 ? 00:00:08 [kworker/u33:3-hci0]
abhijit 968060 351306 0 13:09 ? 00:00:05 /usr/lib/firefox/firefox -contentproc -childID 1749 -isForBrowser -prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {c0ba3673-966e-48a6-b029-91346e348342} 351306 true tab
root 968299 2 0 13:11 ? 00:00:02 [kworker/u32:4-events_unbound]
root 969560 2 0 13:23 ? 00:00:01 [kworker/5:1-cgroup_destroy]
root 969608 2 0 13:24 ? 00:00:00 [kworker/u32:7-events_unbound]
root 969648 2 0 13:24 ? 00:00:01 [kworker/2:0-inet_frag_wq]
abhijit 969715 379620 0 13:25 pts/6 00:00:00 ssh root@10.1.101.41

root	970435	2 0 13:27 ?	00:00:01 [kworker/3:2-rcu_gp]
root	971921	2 0 13:33 ?	00:00:00 [kworker/1:2-rcu_gp]
root	972048	2 0 13:35 ?	00:00:00 [kworker/7:2-rcu_gp]
root	972127	2 0 13:37 ?	00:00:01 [kworker/0:1-events_long]
root	972207	2 0 13:37 ?	00:00:01 [kworker/4:1-events]
root	972378	2 0 13:39 ?	00:00:00 [kworker/1:0-cgroup_destroy]
root	972435	2 0 13:39 ?	00:00:00 [kworker/6:2-events]
root	972964	2 0 13:41 ?	00:00:00 [kworker/7:0-events]
root	973166	2 0 13:41 ?	00:00:00 [kworker/u32:0-events_unbound]
root	973193	2 0 13:42 ?	00:00:00 [kworker/2:1-events]
root	973282	2 0 13:43 ?	00:00:00 [kworker/4:2-events]
root	973384	2 0 13:44 ?	00:00:00 [kworker/3:0-rcu_gp]
root	973389	2 0 13:44 ?	00:00:01 [kworker/u33:0-hci0]
root	973391	2 0 13:44 ?	00:00:00 [kworker/6:0-rcu_gp]
root	973392	2 0 13:44 ?	00:00:00 [kworker/5:2-events]
root	973655	2 0 13:45 ?	00:00:00 [kworker/1:1-events]
root	973664	2 0 13:46 ?	00:00:00 [kworker/7:1-events]
root	973965	2 0 13:47 ?	00:00:00 [kworker/2:2-events]
root	974126	2 0 13:48 ?	00:00:00 [kworker/u32:1-kcryptd/253:0]
root	974189	2 0 13:48 ?	00:00:00 [kworker/4:0-events]
root	974190	2 0 13:48 ?	00:00:00 [kworker/0:2-events]
root	974369	2 0 13:49 ?	00:00:00 [kworker/u32:2-nvme-wq]
root	974539	2 0 13:49 ?	00:00:00 [kworker/u32:3-events_unbound]
root	974655	2 0 13:50 ?	00:00:00 [kworker/6:1-events]
root	974690	2 0 13:50 ?	00:00:00 [kworker/5:0-events]
root	974740	2 0 13:50 ?	00:00:00 [kworker/7:3-events]
root	974742	2 0 13:50 ?	00:00:00 [kworker/3:1-pm]
root	974743	2 0 13:50 ?	00:00:00 [kworker/u32:6-events_unbound]
root	974744	2 0 13:50 ?	00:00:00 [kworker/u32:8-kcryptd/253:0]
root	974745	2 0 13:50 ?	00:00:00 [kworker/u32:9-events_unbound]
root	974746	2 0 13:50 ?	00:00:00 [kworker/u32:10-events_unbound]
root	974747	2 0 13:50 ?	00:00:00 [kworker/u32:11-events_unbound]
root	974748	2 0 13:50 ?	00:00:00 [kworker/u32:12-kcryptd/253:0]
root	974749	2 0 13:50 ?	00:00:00 [kworker/u32:13-events_unbound]
root	974750	2 0 13:50 ?	00:00:00 [kworker/u32:14-events_unbound]
root	974751	2 0 13:50 ?	00:00:00 [kworker/u32:15-events_unbound]
root	974752	2 0 14:18 ?	00:00:00 [kworker/u32:16-events_unbound]
root	974753	2 0 14:18 ?	00:00:00 [kworker/u32:17-events_unbound]
root	974754	2 0 14:18 ?	00:00:00 [kworker/u32:18-events_unbound]
root	974755	2 0 14:18 ?	00:00:00 [kworker/u32:19-events_unbound]
root	974756	2 0 14:18 ?	00:00:00 [kworker/u32:20-kcryptd/253:0]
root	974757	2 0 14:18 ?	00:00:00 [kworker/u32:21-events_unbound]
root	974758	2 0 14:18 ?	00:00:00 [kworker/u32:22-events_unbound]
root	974759	2 0 14:18 ?	00:00:00 [kworker/u32:23-events_unbound]
root	974760	2 0 14:18 ?	00:00:00 [kworker/u32:24-events_unbound]
root	974761	2 0 14:18 ?	00:00:00 [kworker/u32:25-events_unbound]
root	974762	2 0 14:18 ?	00:00:00 [kworker/u32:26-events_unbound]
root	974763	2 0 14:18 ?	00:00:00 [kworker/u32:27-kcryptd/253:0]
root	974764	2 0 14:18 ?	00:00:00 [kworker/u32:28-kcryptd/253:0]
root	974765	2 0 14:18 ?	00:00:00 [kworker/u32:29-events_unbound]
root	974766	2 0 14:18 ?	00:00:00 [kworker/u32:30+events_unbound]
root	974767	2 0 14:18 ?	00:00:00 [kworker/u32:31-events_unbound]
root	974768	2 0 14:18 ?	00:00:00 [kworker/u32:32-events_unbound]
root	974769	2 0 14:18 ?	00:00:00 [kworker/u32:33-events_unbound]
root	974770	2 0 14:18 ?	00:00:00 [kworker/u32:34-events_unbound]
root	974771	2 0 14:18 ?	00:00:00 [kworker/u32:35-events_unbound]
root	974772	2 0 14:18 ?	00:00:00 [kworker/u32:36-events_unbound]
root	974773	2 0 14:18 ?	00:00:00 [kworker/u32:37-events_unbound]
root	974774	2 0 14:18 ?	00:00:00 [kworker/u32:38-events_unbound]
root	974775	2 0 14:18 ?	00:00:00 [kworker/u32:39-events_unbound]
root	974776	2 0 14:18 ?	00:00:00 [kworker/u32:40-kcryptd/253:0]
root	974778	2 0 14:18 ?	00:00:00 [kworker/u32:42-events_unbound]
root	974779	2 0 14:18 ?	00:00:00 [kworker/3:3-events]
root	974780	2 0 14:18 ?	00:00:00 [kworker/3:4-events]
root	974798	2 0 14:18 ?	00:00:00 [kworker/3:5-events]
root	974995	2 0 14:18 ?	00:00:00 [kworker/u33:2-rb_allocator]
root	975505	2 0 14:20 ?	00:00:00 [kworker/6:3-events]
root	975656	2 0 14:20 ?	00:00:00 [kworker/5:3-events]
root	975657	2 0 14:29 ?	00:00:00 [kworker/1:3-pm]
root	975658	2 0 14:29 ?	00:00:00 [kworker/1:4-events]

```

abhijit 975722 351306 0 14:29 ? 00:00:02 /usr/lib/firefox/firefox -contentproc -childID 1836 -isForBrowser -
prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {d9ea12ea-5b3f-477c-9c08-9a4196ea8d04} 351306
true tab
root 976242 2 0 15:16 ? 00:00:00 [kworker/0:0-events]
root 976243 2 0 15:16 ? 00:00:00 [kworker/0:3-events]
root 976244 2 0 15:16 ? 00:00:00 [kworker/0:4-events]
postfix 976248 3508 0 15:16 ? 00:00:00 pickup -l -t unix -u -c
abhijit 976770 351306 0 15:18 ? 00:00:00 /usr/lib/firefox/firefox -contentproc -childID 1840 -isForBrowser -
prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {cb567d37-1c03-46db-966c-632f4b708354} 351306
true tab
abhijit 976836 351306 0 15:19 ? 00:00:00 /usr/lib/firefox/firefox -contentproc -childID 1841 -isForBrowser -
prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {b62e1986-221e-481e-bda4-62fb37caa67} 351306
true tab
abhijit 977138 351306 0 15:20 ? 00:00:00 /usr/lib/firefox/firefox -contentproc -childID 1842 -isForBrowser -
prefsLen 35227 -prefMapSize 247458 -jsInitLen 229864 -parentBuildID 20240108143603 -greomni /usr/lib/firefox/omni.ja -
appomni /usr/lib/firefox/browser/omni.ja -appDir /usr/lib/firefox/browser {989da731-7bd5-44ce-abd0-200ca4c84b9b} 351306
true tab
abhijit 977184 594726 0 15:21 pts/5 00:00:00 ps -eaf
abhijit 1664880 3629 0 Jan07 ? 00:00:06 /usr/bin/gnome-calendar --gapplication-service
abhijit 1665340 3629 0 Jan07 ? 00:00:00 /usr/bin/gpg-agent --supervised
abhijit 3872409 3629 0 Jan07 ? 00:00:05 /usr/bin/seahorse --gapplication-service
root 3873244 1 0 Jan07 ? 00:00:55 /sbin/mount.ntfs /dev/nvme0n1p3 /media/abhijit/windows -o
rw,nodev,nosuid,windows_names,uid=1000,gid=1000,uhelper=udisks2
abhijit 3884359 30565 0 Jan07 pts/7 00:00:00 bash
abhijit 4108623 30565 0 Jan08 pts/8 00:00:00 bash
abhijit 4136834 3890 0 Jan08 ? 00:00:00 /usr/lib/libreoffice/program/oosplash --calc
abhijit 4136869 4136834 0 Jan08 ? 00:32:11 /usr/lib/libreoffice/program/soffice.bin --calc
abhijit 4139495 30565 0 Jan08 pts/9 00:00:00 bash

```

The PID of the grand-parent of the process with PID 4108623 is :

3629



The two processes which were created by kernel have PIDs (in increasing order)

0

and

1



The process that created most of the "graphical" processes is having PID

2



Question 6

Partially correct

Mark 0.67 out of 1.00

Order the following events in boot process (from 1 onwards)

Shell	5	✗
Boot loader	2	✓
Init	4	✓
BIOS	1	✓
Login interface	6	✗
OS	3	✓

Your answer is partially correct.

You have correctly selected 4.

The correct answer is: Shell → 6, Boot loader → 2, Init → 4, BIOS → 1, Login interface → 5, OS → 3

Question 7

Correct

Mark 0.50 out of 0.50

Is the terminal a part of the kernel on GNU/Linux systems?

- a. yes
 b. no ✓ wrong

The correct answer is: no

Question 8

Correct

Mark 1.00 out of 1.00

Consider the following programs

exec1.c

```
#include <unistd.h>
#include <stdio.h>
int main() {
    execl("./exec2", "./exec2", NULL);
}
```

exec2.c

```
#include <unistd.h>
#include <stdio.h>
int main() {
    execl("/bin/ls", "/bin/ls", NULL);
    printf("hello\n");
}
```

Compiled as

```
cc  exec1.c -o exec1
cc  exec2.c -o exec2
```

And run as

```
$ ./exec1
```

Explain the output of the above command (./exec1)

Assume that /bin/ls , i.e. the 'ls' program exists.

Select one:

- a. Execution fails as the call to execl() in exec1 fails
- b. Execution fails as the call to execl() in exec2 fails
- c. Execution fails as one exec can't invoke another exec
- d. Program prints hello
- e. "ls" runs on current directory✓

Your answer is correct.

The correct answer is: "ls" runs on current directory

Question 9

Correct

Mark 0.50 out of 0.50

When you turn your computer ON, on BIOS based systems, you are often shown an option like "Press F9 for boot options". What does this mean?

- a. The choice of booting slowly or fast
- b. The BIOS allows us to choose the boot device, the device from which the boot loader will be loaded✓
- c. The choice of which OS to boot from
- d. The choice of the boot loader (e.g. GRUB or Windows-Loader)

The correct answer is: The BIOS allows us to choose the boot device, the device from which the boot loader will be loaded

Question 10

Correct

Mark 1.00 out of 1.00

How does the distinction between kernel mode and user mode function as a rudimentary form of protection (security) ?

Select one:

- a. It prohibits invocation of kernel code completely, if a user program is running
- b. It prohibits one process from accessing other process's memory
- c. It disallows hardware interrupts when a process is running
- d. It prohibits a user mode process from running privileged instructions✓

Your answer is correct.

The correct answer is: It prohibits a user mode process from running privileged instructions

Question 11

Correct

Mark 1.00 out of 1.00

Select all the correct statements about bootloader.

Every wrong selection will deduct marks proportional to $1/n$ where n is total wrong choices in the question.

You will get minimum a zero.

- a. Bootloader must be one sector in length
- b. The bootloader loads the BIOS
- c. Bootloaders allow selection of OS to boot from✓
- d. Modern Bootloaders often allow configuring the way an OS boots✓
- e. LILO is a bootloader✓

Your answer is correct.

The correct answers are: LILO is a bootloader, Modern Bootloaders often allow configuring the way an OS boots, Bootloaders allow selection of OS to boot from

Question 12

Correct

Mark 1.00 out of 1.00

Predict the output of the program given here.

Assume that all the path names for the programs are correct. For example "/usr/bin/echo" will actually run echo command.

Assume that there is no mixing of printf output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good
output
should be written as good output

--

```
main() {  
    int i;  
    i = fork();  
    if(i == 0)  
        execl("/usr/bin/echo", "/usr/bin/echo", "hi", 0);  
    else  
        wait(0);  
    fork();  
    execl("/usr/bin/echo", "/usr/bin/echo", "one", 0);  
}
```

Answer: hi one one ✓

The correct answer is: hi one one

Question 13

Correct

Mark 1.00 out of 1.00

Select all statements that correctly explain the use/purpose of system calls.

Select one or more:

- a. Provide an environment for process creation ✓
- b. Handle exceptions like division by zero
- c. Allow I/O device access to user processes ✓
- d. Provide services for accessing files ✓
- e. Handle ALL types of interrupts
- f. Run each instruction of an application program
- g. Switch from user mode to kernel mode ✓

Your answer is correct.

The correct answers are: Switch from user mode to kernel mode, Provide services for accessing files, Allow I/O device access to user processes, Provide an environment for process creation

Question 14

Partially correct

Mark 0.60 out of 1.00

Select all the correct statements about two modes of CPU operation

Select one or more:

- a. Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode
- b. There is an instruction like 'iret' to return from kernel mode to user mode✓
- c. The two modes are essential for a multiprogramming system✓
- d. The two modes are essential for a multitasking system✓
- e. The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: The two modes are essential for a multiprogramming system, The two modes are essential for a multitasking system, There is an instruction like 'iret' to return from kernel mode to user mode, The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously, Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode

Question 15

Partially correct

Mark 0.25 out of 0.50

Select all the correct statements about bootloader.

Every wrong selection will deduct marks proportional to $1/n$ where n is total wrong choices in the question.

You will get minimum a zero.

- a. Bootloaders allow selection of OS to boot from✓
- b. Bootloader must be one sector in length✗
- c. LILO is a bootloader✓
- d. The bootloader loads the BIOS
- e. Modern Bootloaders often allow configuring the way an OS boots✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: LILO is a bootloader, Modern Bootloaders often allow configuring the way an OS boots, Bootloaders allow selection of OS to boot from

Question 16

Partially correct

Mark 0.33 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

(Note: non-interruptible kernel code means, if the kernel code is executing, then interrupts will be disabled).

Note: A possible sequence may have some missing steps in between. An impossible sequence will have n and n+1th steps such that n+1th step can not follow n'th step.

Select one or more:

- a. P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P3 running

Hardware interrupt

Interrupt unblocks P1

Interrupt returns

P3 running

Timer interrupt

Scheduler

P1 running

- b. P1 running

Keyboard hardware interrupt

Keyboard interrupt handler running

interrupt handler returns

P1 running

P1 makes system call

System call returns

P1 running

timer interrupt

Scheduler

P2 running

- c. P1 running ✓

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

- d.

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

- e. P1 running

P1 makes system call

System call returns

P1 running

timer interrupt

Scheduler running

P2 running

- f. P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return

Your answer is partially correct.

You have correctly selected 1.

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again, P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheduler

P1 running

P1's system call return,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question 17

Correct

Mark 1.00 out of 1.00

What will this program do?

```
int main() {  
    fork();  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("hello");  
}
```

- a. run ls twice and print hello twice
- b. run ls twice and print hello twice, but output will appear in some random order
- c. run ls twice✓
- d. one process will run ls, another will print hello
- e. run ls once

Your answer is correct.

The correct answer is: run ls twice

[◀ Surprise Quiz - 1 \(pre-requisites\)](#)

Jump to...

[Surprise Quiz - 3 \(processes, memory management, event driven kernel\), compilation-linking-loading](#) ▶

Started on Wednesday, 7 February 2024, 6:09 PM

State Finished

Completed on Wednesday, 7 February 2024, 7:10 PM

Time taken 1 hour

Grade 18.14 out of 20.00 (90.68%)

Question 1

Correct

Mark 1.00 out of 1.00

Consider the following code and MAP the file to which each fd points at the end of the code. Assume that files/folders exist when needed with proper permissions and open() calls work.

```
int main(int argc, char *argv[]) {
    int fd1, fd2 = 1, fd3 = 1, fd4 = 1;

    fd1 = open("/tmp/1", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    fd2 = open("/tmp/2", O_RDONLY);
    fd3 = open("/tmp/3", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    close(0);
    close(1);
    dup(fd2);
    dup(fd3);
    close(fd3);
    dup2(fd2, fd4);
    printf("%d %d %d %d\n", fd1, fd2, fd3, fd4);
    return 0;
}
```

fd2	/tmp/2	✓
fd1	/tmp/1	✓
0	/tmp/2	✓
fd3	closed	✓
2	stderr	✓
1	/tmp/3	✓
fd4	/tmp/2	✓

The correct answer is: fd2 → /tmp/2, fd1 → /tmp/1, 0 → /tmp/2, fd3 → closed, 2 → stderr, 1 → /tmp/3, fd4 → /tmp/2

Question 2

Partially correct

Mark 1.43 out of 2.00

Order the events that occur on a timer interrupt:

Save the context of the currently running process	3	✓
Jump to scheduler code	4	✓
Set the context of the new process	6	✓
Jump to a code pointed by IDT	1	✗
Change to kernel stack of currently running process	2	✗
Select another process for execution	5	✓
Execute the code of the new process	7	✓

The correct answer is: Save the context of the currently running process → 3, Jump to scheduler code → 4, Set the context of the new process → 6, Jump to a code pointed by IDT → 2, Change to kernel stack of currently running process → 1, Select another process for execution → 5, Execute the code of the new process → 7

Question 3

Partially correct

Mark 0.75 out of 1.00

Select the sequence of events that are NOT possible, assuming an interruptible kernel code

Select one or more:

- a. P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P3 running

Hardware interrupt

Interrupt unblocks P1

Interrupt returns

P3 running

Timer interrupt

Scheduler

P1 running

- b. P1 running

keyboard hardware interrupt

keyboard interrupt handler running

interrupt handler returns

P1 running

P1 makes system call

system call returns

P1 running

timer interrupt

scheduler

P2 running

- c. P1 running ✖

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return

- d. P1 running ✓

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

- e. ✓

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

- f. P1 running

P1 makes system call

system call returns

P1 running
timer interrupt
Scheduler running
P2 running

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question 4

Partially correct

Mark 1.60 out of 2.00

Match the elements of C program to their place in memory

Allocated Memory	Heap	✓
Global variables	Data	✓
Arguments	Stack	✓
#include files	No Memory needed	✗
Local Static variables	Data	✓
Global Static variables	Data	✓
Function code	Code	✓
Local Variables	Stack	✓
#define MACROS	No Memory needed	✓
Code of main()	Main_Code	✗

The correct answer is: Allocated Memory → Heap, Global variables → Data, Arguments → Stack, #include files → No memory needed, Local Static variables → Data, Global Static variables → Data, Function code → Code, Local Variables → Stack, #define MACROS → No Memory needed, Code of main() → Code

Question 5

Correct

Mark 1.00 out of 1.00

Select the order in which the various stages of a compiler execute.

Intermediate code generation	3	✓
Syntactical Analysis	2	✓
Pre-processing	1	✓
Linking	4	✓
Loading	does not exist	✓

The correct answer is: Intermediate code generation → 3, Syntactical Analysis → 2, Pre-processing → 1, Linking → 4, Loading → does not exist

Question 6

Correct

Mark 2.00 out of 2.00

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

```
$ ls . /tmp/asdfksdf >/tmp/ddd 2>&1
```

Program 1

```
int main(int argc, char *argv[]) {  
    int fd, n, i;  
    char buf[128];  
  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    close(1);  
    dup(fd);  
    close(2);  
    dup(fd);  
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);  
}
```

Program 2

```
int main(int argc, char *argv[]) {  
    int fd, n, i;  
    char buf[128];  
  
    close(1);  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    close(2);  
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);  
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);  
}
```

Select all the correct statements about the programs

Select one or more:

- a. Program 1 ensures 2>&1 and does not ensure >/tmp/ddd
- b. Program 1 is correct for > /tmp/ddd but not for 2>&1
- c. Only Program 1 is correct✓
- d. Program 1 does 1>&2
- e. Program 2 does 1>&2
- f. Both program 1 and 2 are incorrect
- g. Program 2 makes sure that there is one file offset used for '2' and '1'
- h. Program 2 is correct for > /tmp/ddd but not for 2>&1
- i. Program 1 makes sure that there is one file offset used for '2' and '1'✓
- j. Only Program 2 is correct
- k. Program 2 ensures 2>&1 and does not ensure > /tmp/ddd
- l. Both programs are correct

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'

Question 7

Correct

Mark 1.00 out of 1.00

Select all the correct statements about zombie processes

Select one or more:

- a. A process can become zombie if it finishes, but the parent has finished before it✓
- b. init() typically keeps calling wait() for zombie processes to get cleaned up✓
- c. A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it✓
- d. Zombie processes are harmless even if OS is up for long time
- e. A zombie process remains zombie forever, as there is no way to clean it up
- f. If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent✓
- g. A process becomes zombie when its parent finishes
- h. A zombie process occupies space in OS data structures✓

The correct answers are: A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it, A process can become zombie if it finishes, but the parent has finished before it, A zombie process occupies space in OS data structures, If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent, init() typically keeps calling wait() for zombie processes to get cleaned up

Question 8

Correct

Mark 1.00 out of 1.00

Consider the image given below, which explains how paging works.

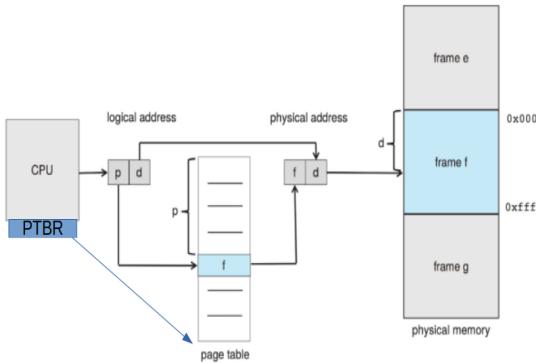


Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The page table is itself present in Physical memory
<input checked="" type="radio"/>	<input type="radio"/>	Maximum Size of page table is determined by number of bits used for page number
<input checked="" type="radio"/>	<input type="radio"/>	The page table is indexed using page number
<input type="radio"/>	<input checked="" type="radio"/>	The page table is indexed using frame number
<input checked="" type="radio"/>	<input type="radio"/>	The PTBR is present in the CPU as a register
<input checked="" type="radio"/>	<input type="radio"/>	The physical address may not be of the same size (in bits) as the logical address
<input type="radio"/>	<input checked="" type="radio"/>	Size of page table is always determined by the size of RAM
<input type="radio"/>	<input checked="" type="radio"/>	The locating of the page table using PTBR also involves paging translation

The page table is itself present in Physical memory: True

Maximum Size of page table is determined by number of bits used for page number: True

The page table is indexed using page number: True

The page table is indexed using frame number: False

The PTBR is present in the CPU as a register: True

The physical address may not be of the same size (in bits) as the logical address: True

Size of page table is always determined by the size of RAM: False

The locating of the page table using PTBR also involves paging translation: False

Question 9

Correct

Mark 1.00 out of 1.00

Select all the correct statements about MMU and its functionality (on a non-demand paged system)

Select one or more:

- a. The operating system interacts with MMU for every single address translation
- b. Illegal memory access is detected by operating system
- c. Logical to physical address translations in MMU are done in hardware, automatically ✓
- d. MMU is a separate chip outside the processor
- e. MMU is inside the processor ✓
- f. Illegal memory access is detected in hardware by MMU and a trap is raised ✓
- g. The Operating system sets up relevant CPU registers to enable proper MMU translations ✓
- h. Logical to physical address translations in MMU are done with specific machine instructions

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

Question 10

Partially correct

Mark 0.86 out of 1.00

Mark the statements as True/False w.r.t. the basic concepts of memory management.

True	False	
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	When a process is executing, each virtual address is converted into physical address by the kernel directly.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The compiler generates address references for code/data/stack/heap in the executable file, depending on the MM architecture provided by CPU and kernel.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The compiler generates the address references for code/data/stack/heap in the executable file as per the memory management schema chosen by the compiler itself, and then the kernel ensures that program is executed with this schema.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The compiler interacts with the kernel continuously while compiling a program and obtains the correct set of memory addresses for code/stack/heap/data and then generates the machine code file.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The kernel ensures that the MMU is setup before scheduling a process and then the CPU/MMU ensures that the address translation takes place.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	When a process is executing, each virtual address is converted into physical address by the CPU hardware directly.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The kernel refers to the page table for converting each virtual address to physical address.

When a process is executing, each virtual address is converted into physical address by the kernel directly.: False

The compiler generates address references for code/data/stack/heap in the executable file, depending on the MM architecture provided by CPU and kernel.: True

The compiler generates the address references for code/data/stack/heap in the executable file as per the memory management schema chosen by the compiler itself, and then the kernel ensures that program is executed with this schema.: False

The compiler interacts with the kernel continuously while compiling a program and obtains the correct set of memory addresses for code/stack/heap/data and then generates the machine code file.: False

The kernel ensures that the MMU is setup before scheduling a process and then the CPU/MMU ensures that the address translation takes place.: True

When a process is executing, each virtual address is converted into physical address by the CPU hardware directly.: True

The kernel refers to the page table for converting each virtual address to physical address.: False

Question 11

Partially correct

Mark 0.50 out of 1.00

Select the correct statements about paging (not demand paging) mechanism

Select one or more:

- a. User process can update its own PTBR
- b. An invalid entry on a page means, it was an illegal memory reference
- c. The PTBR is loaded by the OS ✓
- d. Page table is accessed by the OS as part of execution of an instruction
- e. User process can update its own page table entries
- f. Page table is accessed by the MMU as part of execution of an instruction ✓
- g. OS creates the page table for every process ✓
- h. An invalid entry on a page means, either it was illegal memory reference or the page was not present in memory. ✗

The correct answers are: OS creates the page table for every process, The PTBR is loaded by the OS, Page table is accessed by the MMU as part of execution of an instruction, An invalid entry on a page means, it was an illegal memory reference

Question 12

Correct

Mark 1.00 out of 1.00

Select the compiler's view of the process's address space, for each of the following MMU schemes:

(Assume that each scheme, e.g. paging/segmentation/etc is effectively utilised)

Segmentation	many continuous chunks of variable size	✓
Segmentation, then paging	many continuous chunks of variable size	✓
Relocation + Limit	one continuous chunk	✓
Paging	one continuous chunk	✓

The correct answer is: Segmentation → many continuous chunks of variable size, Segmentation, then paging → many continuous chunks of variable size, Relocation + Limit → one continuous chunk, Paging → one continuous chunk

Question 13

Correct

Mark 1.00 out of 1.00

Select the state that is not possible after the given state, for a process:

New:	Running	✓
Ready :	Waiting	✓
Running:	None of these	✓
Waiting:	Running	✓

Question 14

Correct

Mark 1.00 out of 1.00

A process blocks itself means

- a. The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler✓
- b. The application code calls the scheduler
- c. The kernel code of an interrupt handler, moves the process to a waiting queue and calls scheduler
- d. The kernel code of system call calls scheduler

The correct answer is: The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler

Question 15

Correct

Mark 1.00 out of 1.00

which of the following is not a difference between real mode and protected mode

- a. in real mode the segment is multiplied by 16, in protected mode segment is used as index in GDT
- b. in real mode the addressable memory is less than in protected mode
- c. in real mode general purpose registers are 16 bit, in protected mode they are 32 bit
- d. processor starts in real mode
- e. in real mode the addressable memory is more than in protected mode✓

The correct answer is: in real mode the addressable memory is more than in protected mode

Question 16

Correct

Mark 1.00 out of 1.00

Predict the output of the program given here.

Assume that there is no mixing of printf output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good
output
should be written as good output

--

```
int main() {  
    int pid;  
    printf("hi\n");  
    pid = fork();  
    if(pid == 0) {  
        exit(0);  
    }  
    printf("bye\n");  
    fork();  
    printf("ok\n");  
}
```

Answer: hi bye ok ok



The correct answer is: hi bye ok ok

Question 17

Correct

Mark 1.00 out of 1.00

Which of the following are NOT a part of job of a typical compiler?

- a. Invoke the linker to link the function calls with their code, extern globals with their declaration
- b. Check the program for syntactical errors
- c. Convert high level language code to machine code
- d. Suggest alternative pieces of code that can be written ✓
- e. Check the program for logical errors ✓
- f. Process the # directives in a C program

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

[◀ Surprise Quiz - 2](#)

Jump to...

[Questions for test on kalloc/kfree/kvmalloc, etc. ►](#)

Started on Saturday, 20 February 2021, 2:51 PM

State Finished

Completed on Saturday, 20 February 2021, 3:55 PM

Time taken 1 hour 3 mins

Grade 7.30 out of 20.00 (37%)

Question 1

Partially correct

Mark 0.80 out of 1.00

Select all the correct statements about the state of a process.

- a. A process can self-terminate only when it's running ✓
- b. Typically, it's represented as a number in the PCB ✓
- c. A process that is running is not on the ready queue ✓
- d. Processes in the ready queue are in the ready state ✓
- e. It is not maintained in the data structures by kernel, it is only for conceptual understanding of programmers
- f. Changing from running state to waiting state results in "giving up the CPU" ✓
- g. A process in ready state is ready to receive interrupts
- h. A waiting process starts running after the wait is over ✗
- i. A process changes from running to ready state on a timer interrupt ✓
- j. A process in ready state is ready to be scheduled ✓
- k. A running process may terminate, or go to wait or become ready again ✓
- l. A process waiting for I/O completion is typically woken up by the particular interrupt handler code ✓
- m. A process waiting for any condition is woken up by another process only
- n. A process changes from running to ready state on a timer interrupt or any I/O wait

Your answer is partially correct.

You have selected too many options.

The correct answers are: Typically, it's represented as a number in the PCB, A process in ready state is ready to be scheduled, Processes in the ready queue are in the ready state, A process that is running is not on the ready queue, A running process may terminate, or go to wait or become ready again, A process changes from running to ready state on a timer interrupt, Changing from running state to waiting state results in "giving up the CPU", A process can self-terminate only when it's running, A process waiting for I/O completion is typically woken up by the particular interrupt handler code

Question 2

Incorrect

Mark 0.00 out of 1.00

For each line of code mentioned on the left side, select the location of sp/esp that is in use

`jmp *%eax`
in entry.S

0x7c00 to 0x10000



`ljmp $(SEG_KCODE<<3), $start32`
in bootasm.S

0x10000 to 0x7c00



`call bootmain`
in bootasm.S

0x7c00 to 0x10000



`cli`
in bootasm.S

0x7c00 to 0



`readseg((uchar*)elf, 4096, 0);`
in bootmain.c

The 4KB area in kernel image, loaded in memory, named as 'stack'



Your answer is incorrect.

The correct answer is: `jmp *%eax`

`in entry.S` → The 4KB area in kernel image, loaded in memory, named as 'stack', `ljmp $(SEG_KCODE<<3), $start32`

`in bootasm.S` → Immaterail as the stack is not used here, `call bootmain`

`in bootasm.S` → 0x7c00 to 0, `cli`

`in bootasm.S` → Immaterail as the stack is not used here, `readseg((uchar*)elf, 4096, 0);`

`in bootmain.c` → 0x7c00 to 0

Question 3

Correct

Mark 0.25 out of 0.25

Order the following events in boot process (from 1 onwards)

Boot loader	2	✓
Shell	6	✓
BIOS	1	✓
OS	3	✓
Init	4	✓
Login interface	5	✓

Your answer is correct.

The correct answer is: Boot loader → 2, Shell → 6, BIOS → 1, OS → 3, Init → 4, Login interface → 5

Question 4

Partially correct

Mark 0.30 out of 0.50

Consider the following command and its output:

```
$ ls -lht xv6.img kernel
-rw-rw-r-- 1 abhijit abhijit 4.9M Feb 15 11:09 xv6.img
-rwxrwxr-x 1 abhijit abhijit 209K Feb 15 11:09 kernel*
```

Following code in bootmain()

```
readseg((uchar*)elf, 4096, 0);
```

and following selected lines from Makefile

```
xv6.img: bootblock kernel
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

```
kernel: $(OBJS) entry.o entryother initcode kernel.ld
$(LD) $(LDFLAGS) -T kernel.ld -o kernel entry.o $(OBJS) -b binary initcode entryother
$(OBJDUMP) -S kernel > kernel.asm
$(OBJDUMP) -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$$/d' > kernel.sym
```

Also read the code of bootmain() in xv6 kernel.

Select the options that describe the meaning of these lines and their correlation.

- a. Although the size of the kernel file is 209 Kb, only 4Kb out of it is the actual kernel code and remaining part is all zeroes.
- b. The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files ✓
- c. The kernel.ld file contains instructions to the linker to link the kernel properly ✓
- d. The bootmain() code does not read the kernel completely in memory
- e. readseg() reads first 4k bytes of kernel in memory
- f. Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.
- g. The kernel.asm file is the final kernel file
- h. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is not read as it is user programs.
- i. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(). ✓

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(), readseg() reads first 4k bytes of kernel in memory, The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files, The kernel.ld file contains instructions to the linker to link the kernel properly, Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.

Question 5

Partially correct

Mark 0.50 out of 1.00

```
int f() {  
    int count;  
    for (count = 0; count < 2; count++) {  
        if (fork() == 0)  
            printf("Operating-System\\n");  
    }  
    printf("TYCOMP\\n");  
}
```

The number of times "Operating-System" is printed, is:

Answer:

The correct answer is: 7.00

Question 6

Partially correct

Mark 0.40 out of 0.50

Select Yes/True if the mentioned element must be a part of PCB

Select No/False otherwise.

Yes**No**

<input checked="" type="radio"/>	<input checked="" type="radio"/>	PID	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Process context	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	List of opened files	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Process state	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Parent's PID	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Pointer to IDT	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Function pointers to all system calls	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Memory management information about that process	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Pointer to the parent process	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	EIP at the time of context switch	✓

PID: Yes

Process context: Yes

List of opened files: Yes

Process state: Yes

Parent's PID: No

Pointer to IDT: No

Function pointers to all system calls: No

Memory management information about that process: Yes

Pointer to the parent process: Yes

EIP at the time of context switch: Yes

Question 7

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about code of bootmain() in xv6

```

void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void))(elf->entry);
    entry();
}

```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- a. The kernel file gets loaded at the Physical address 0x10000 +0x80000000 in memory. ✗
- b. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- c. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded ✓
- d. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it. ✓
- e. The kernel file has only two program headers ✓
- f. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- g. The readseg finally invokes the disk I/O code using assembly instructions ✓
- h. The elf->entry is set by the linker in the kernel file and it's 8010000c ✓
- i. The kernel file gets loaded at the Physical address 0x10000 in memory. ✓
- j. The condition if(ph->memsz > ph->filesz) is never true. ✗
- k. The stosb() is used here, to fill in some space in memory with zeroes ✓

Your answer is incorrect.

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

Question 8

Partially correct

Mark 0.13 out of 0.25

Which of the following are NOT a part of job of a typical compiler?

- a. Check the program for logical errors ✓
- b. Convert high level language code to machine code
- c. Process the # directives in a C program
- d. Invoke the linker to link the function calls with their code, extern globals with their declaration
- e. Check the program for syntactical errors
- f. Suggest alternative pieces of code that can be written

Your answer is partially correct.

You have correctly selected 1.

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question 9

Correct

Mark 0.25 out of 0.25

Rank the following storage systems from slowest (first) to fastest(last)

Cache	6	✓
Hard Disk	3	✓
RAM	5	✓
Optical Disks	2	✓
Non volatile memory	4	✓
Registers	7	✓
Magnetic Tapes	1	✓

Your answer is correct.

The correct answer is: Cache → 6, Hard Disk → 3, RAM → 5, Optical Disks → 2, Non volatile memory → 4, Registers → 7, Magnetic Tapes → 1

Question 10

Partially correct

Mark 0.21 out of 0.50

Which of the following parts of a C program do not have any corresponding machine code ?

- a. local variable declaration
- b. global variables
- c. function calls ✗
- d. #directives ✓
- e. expressions
- f. pointer dereference
- g. typedefs ✓

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: #directives, typedefs, global variables

Question 11

Correct

Mark 0.25 out of 0.25

Match a system call with it's description

pipe	create an unnamed FIFO storage with 2 ends - one for reading and another for writing	✓
dup	create a copy of the specified file descriptor into smallest available file descriptor	✓
dup2	create a copy of the specified file descriptor into another specified file descriptor	✓
exec	execute a binary file overlaying the image of current process	✓
fork	create an identical child process	✓

Your answer is correct.

The correct answer is: pipe → create an unnamed FIFO storage with 2 ends - one for reading and another for writing, dup → create a copy of the specified file descriptor into smallest available file descriptor, dup2 → create a copy of the specified file descriptor into another specified file descriptor, exec → execute a binary file overlaying the image of current process, fork → create an identical child process

Question 12

Correct

Mark 0.25 out of 0.25

Match the register with the segment used with it.

eip	cs	✓
edi	es	✓
esi	ds	✓
ebp	ss	✓
esp	ss	✓

Your answer is correct.

The correct answer is: eip → cs, edi → es, esi → ds, ebp → ss, esp → ss

Question 13

Correct

Mark 0.25 out of 0.25

What's the trapframe in xv6?

- a. A frame of memory that contains all the trap handler code
- b. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- c. The IDT table
- d. A frame of memory that contains all the trap handler code's function pointers
- e. A frame of memory that contains all the trap handler's addresses
- f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S ✓
- g. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only

Your answer is correct.

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question 14

Incorrect

Mark 0.00 out of 0.50

Select all the correct statements about linking and loading.

Select one or more:

- a. Continuous memory management schemes can support dynamic linking and dynamic loading. ✗
- b. Loader is last stage of the linker program ✗
- c. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently) ✓
- d. Dynamic linking and loading is not possible without demand paging or demand segmentation. ✓
- e. Dynamic linking essentially results in relocatable code. ✓
- f. Continuous memory management schemes can support static linking and static loading. (may be inefficiently) ✓
- g. Loader is part of the operating system ✓
- h. Static linking leads to non-relocatable code ✗
- i. Dynamic linking is possible with continuous memory management, but variable sized partitions only. ✗

Your answer is incorrect.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

Question 15

Incorrect

Mark 0.00 out of 0.25

In bootasm.S, on the line

```
ljmp    $(SEG_KCODE<<3), $start32
```

The SEG_KCODE << 3, that is shifting of 1 by 3 bits is done because

- a. The value 8 is stored in code segment
- b. The code segment is 16 bit and only upper 13 bits are used for segment number
- c. The code segment is 16 bit and only lower 13 bits are used for segment number ✗
- d. While indexing the GDT using CS, the value in CS is always divided by 8
- e. The ljmp instruction does a divide by 8 on the first argument

Your answer is incorrect.

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question 16

Partially correct

Mark 0.07 out of 0.50

Order the events that occur on a timer interrupt:

Change to kernel stack

1	✗
---	---

Jump to a code pointed by IDT

2	✗
---	---

Jump to scheduler code

5	✗
---	---

Set the context of the new process

4	✗
---	---

Save the context of the currently running process

3	✓
---	---

Execute the code of the new process

6	✗
---	---

Select another process for execution

7	✗
---	---

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: Change to kernel stack → 2, Jump to a code pointed by IDT → 1, Jump to scheduler code → 4, Set the context of the new process → 6, Save the context of the currently running process → 3, Execute the code of the new process → 7, Select another process for execution → 5

Question 17

Incorrect

Mark 0.00 out of 1.00

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

```
$ ls . /tmp/asdfksdf >/tmp/ddd 2>&1
```

Program 1

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    dup(fd);
    close(2);
    dup(fd);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Program 2

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    close(1);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(2);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Select all the correct statements about the programs

Select one or more:

- a. Both programs are correct ✗
- b. Program 2 makes sure that there is one file offset used for '2' and '1' ✗
- c. Only Program 2 is correct ✗
- d. Program 2 does 1>&2 ✗
- e. Program 2 ensures 2>&1 and does not ensure >/tmp/ddd ✗
- f. Program 1 makes sure that there is one file offset used for '2' and '1' ✓
- g. Program 1 is correct for >/tmp/ddd but not for 2>&1 ✗
- h. Program 1 does 1>&2 ✗
- i. Both program 1 and 2 are incorrect ✗
- j. Program 2 is correct for >/tmp/ddd but not for 2>&1 ✗
- k. Only Program 1 is correct ✓
- l. Program 1 ensures 2>&1 and does not ensure >/tmp/ddd ✗

Your answer is incorrect.

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'

Question 18

Correct

Mark 0.25 out of 0.25

Select the option which best describes what the CPU does during its powered ON lifetime

- a. Ask the user what is to be done, and execute that task
- b. Ask the OS what is to be done, and execute that task
- c. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask the User or the OS what is to be done next, repeat
- d. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per ✓ the instruction itself, repeat
- e. Fetch instruction specified by OS, Decode and execute it, repeat
- f. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask OS what is to be done next, repeat

The correct answer is: Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, repeat

Question 19

Partially correct

Mark 0.86 out of 1.00

Consider the following code and MAP the file to which each fd points at the end of the code.

```
int main(int argc, char *argv[]) {
    int fd1, fd2 = 1, fd3 = 1, fd4 = 1;

    fd1 = open("/tmp/1", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    fd2 = open("/tmp/2", O_RDONLY);
    fd3 = open("/tmp/3", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    close(0);
    close(1);
    dup(fd2);
    dup(fd3);
    close(fd3);
    dup2(fd2, fd4);
    printf("%d %d %d %d\n", fd1, fd2, fd3, fd4);
    return 0;
}
```

1	closed	✗
fd4	/tmp/2	✓
fd2	/tmp/2	✓
fd1	/tmp/1	✓
2	stderr	✓
0	/tmp/2	✓
fd3	closed	✓

Your answer is partially correct.

You have correctly selected 6.

The correct answer is: 1 → /tmp/3, fd4 → /tmp/2, fd2 → /tmp/2, fd1 → /tmp/1, 2 → stderr, 0 → /tmp/2, fd3 → closed

Question 20

Incorrect

Mark 0.00 out of 2.00

Following code claims to implement the command

```
/bin/ls -l | /usr/bin/head -3 | /usr/bin/tail -1
```

Fill in the blanks to make the code work.

Note: Do not include space in writing any option. x[1][2] should be written without any space, and so is the case with [1] or [2]. Pay attention to exact syntax and do not write any extra character like ';' or = etc.

```
int main(int argc, char *argv[]) {
```

```
    int pid1, pid2;
```

```
    int pfd[
```

```
    x ] [2];
```

```
    pipe(
```

```
    x );
```

```
    pid1 =
```

```
    x ;
```

```
    if(pid1 != 0) {
```

```
        close(pfd[0]
```

```
    x );
```

```
        close(
```

```
    x );
```

```
        dup(
```

```
    x );
```

```
        execl("/bin/ls", "/bin/ls", "
```

```
    x ", NULL);
```

```
    }
```

```
    pipe(
```

```
    x );
```

```
    x = fork();
```

```
    if(pid2 == 0) {
```

```
        close(
```

```
        x ;
```

```
        close(0);
```

```
        dup(
```

```
        x );
```

```
        close(pfd[1]
```

```
✗ );
close(
  
✗ );
dup(
  
✗ );
execl("/usr/bin/head", "/usr/bin/head", "  
  
✗ ", NULL);
} else {
close(pfd
  
✗ );
close(
  
✗ );
dup(
  
✗ );
close(pfd
  
✗ );
execl("/usr/bin/tail", "/usr/bin/tail", "  
  
✗ ", NULL);
}  
}
```

Question 21

Partially correct

Mark 0.11 out of 1.00

Select all the correct statements about calling convention on x86 32-bit.

- a. Return address is one location above the ebp ✓
- b. Parameters may be passed in registers or on stack ✓
- c. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function ✓
- d. The ebp pointers saved on the stack constitute a chain of activation records ✓
- e. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables ✗
- f. Parameters may be passed in registers or on stack ✓
- g. The return value is either stored on the stack or returned in the eax register ✗
- h. Parameters are pushed on the stack in left-right order
- i. during execution of a function, ebp is pointing to the old ebp
- j. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function ✗
- k. Compiler may allocate more memory on stack than needed ✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

Question 22

Correct

Mark 1.00 out of 1.00

Match the program with its output (ignore newlines in the output. Just focus on the count of the number of 'hi')

main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi ✓

main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi hi ✓

main() { int i = NULL; fork(); printf("hi\n"); }

hi hi ✓

main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi ✓

Your answer is correct.

The correct answer is: main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi, main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi hi, main() { int i = NULL; fork(); printf("hi\n"); } → hi hi, main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi

Question 23

Incorrect

Mark 0.00 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?

Select all the appropriate choices

- a. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time ✗
- b. The setting up of the most essential memory management infrastructure needs assembly code ✓
- c. The code for reading ELF file can not be written in assembly ✗
- d. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C ✓

Your answer is incorrect.

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question 24

Incorrect

Mark 0.00 out of 0.50

```
xv6.img: bootblock kernel
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- a. The size of the kernel file is nearly 5 MB ✓
- b. The kernel is located at block-1 of the xv6.img ✗
- c. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk. ✗
- d. The size of xv6.img is exactly = (size of bootblock) + (size of kernel) ✗
- e. The bootblock is located on block-0 of the xv6.img ✗
- f. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk. ✓
- g. The bootblock may be 512 bytes or less (looking at the Makefile instruction) ✗
- h. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file. ✗
- i. The size of the xv6.img is nearly 5 MB ✗
- j. xv6.img is the virtual processor used by the qemu emulator ✓
- k. Blocks in xv6.img after kernel may be all zeroes. ✗

Your answer is incorrect.

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question 25

Incorrect

Mark 0.00 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

Select one or more:

a. P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return

b. P1 running

P1 makes sytem call and blocks

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P1 running again



c. P1 running

P1 makes system call

system call returns

P1 running

timer interrupt

Scheduler running

P2 running

d. P1 running

P1 makes sytem call and blocks

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P3 running

Hardware interrupt

Interrupt unblocks P1

Interrupt returns

P3 running

Timer interrupt

Scheduler

P1 running



e.

P1 running

P1 makes sytem call

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P1 running again

f. P1 running

keyboard hardware interrupt

keyboard interrupt handler running

interrupt handler returns

P1 running

P1 makes sytem call

system call returns



P1 running
timer interrupt
scheduler
P2 running

Your answer is incorrect.

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again, P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question 26

Correct

Mark 0.25 out of 0.25

Which of the following are the files related to bootloader in xv6?

- a. bootasm.s and entry.S
- b. bootasm.S and bootmain.c ✓
- c. bootasm.S, bootmain.c and bootblock.c
- d. bootmain.c and bootblock.S

Your answer is correct.

The correct answer is: bootasm.S and bootmain.c

Question 27

Correct

Mark 0.25 out of 0.25

Match the following parts of a C program to the layout of the process in memory

Instructions	Text section	✓
Local Variables	Stack Section	✓
Dynamically allocated memory	Heap Section	✓
Global and static data	Data section	✓

Your answer is correct.

The correct answer is:

Instructions → Text section, Local Variables → Stack Section,
Dynamically allocated memory → Heap Section,
Global and static data → Data section

Question 28

Incorrect

Mark 0.00 out of 0.50

What will this program do?

```
int main() {  
    fork();  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("hello");  
}
```

- a. one process will run ls, another will print hello
- b. run ls once ✗
- c. run ls twice
- d. run ls twice and print hello twice
- e. run ls twice and print hello twice, but output will appear in some random order

Your answer is incorrect.

The correct answer is: run ls twice

Question 29

Correct

Mark 0.25 out of 0.25

What is the OS Kernel?

- a. The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run ✓ correct
- b. The set of tools like compiler, linker, loader, terminal, shell, etc.
- c. Only the system programs like compiler, linker, loader, etc.
- d. Everything that I see on my screen

The correct answer is: The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run

Question 30

Correct

Mark 0.50 out of 0.50

Which of the following is/are not saved during context switch?

- a. Program Counter
- b. General Purpose Registers
- c. Bus ✓
- d. Stack Pointer
- e. MMU related registers/information
- f. Cache ✓
- g. TLB ✓

Your answer is correct.

The correct answers are: TLB, Cache, Bus

Question 31

Partially correct

Mark 0.10 out of 0.25

Select the order in which the various stages of a compiler execute.

Linking	3	✗
Syntactical Analysis	2	✓
Pre-processing	1	✓
Intermediate code generation	does not exist	✗
Loading	4	✗

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Linking → 4, Syntactical Analysis → 2, Pre-processing → 1, Intermediate code generation → 3, Loading → does not exist

Question 32

Partially correct

Mark 0.08 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

context of P0 is saved in P0's PCB	2	✗
context of P1 is loaded from P1's PCB	3	✗
Process P1 is running	5	✗
timer interrupt occurs	6	✗
Process P0 is running	1	✓
Control is passed to P1	4	✗

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: context of P0 is saved in P0's PCB → 3, context of P1 is loaded from P1's PCB → 4, Process P1 is running → 6, timer interrupt occurs → 2, Process P0 is running → 1, Control is passed to P1 → 5

Question 33

Not answered

Marked out of 1.00

Select the correct statements about interrupt handling in xv6 code

- a. On any interrupt/syscall/exception the control first jumps in vectors.S
- b. The trapframe pointer in struct proc, points to a location on user stack
- c. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- d. xv6 uses the 64th entry in IDT for system calls
- e. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- f. The trapframe pointer in struct proc, points to a location on kernel stack
- g. The function trap() is called only in case of hardware interrupt
- h. The CS and EIP are changed only immediately on a hardware interrupt
- i. All the 256 entries in the IDT are filled

- j. On any interrupt/syscall/exception the control first jumps in trapasm.S
- k. The function trap() is called irrespective of hardware interrupt/system-call/exception
- l. xv6 uses the 0x64th entry in IDT for system calls
- m. Before going to alltraps, the kernel stack contains upto 5 entries.

Your answer is incorrect.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in trapasm.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

[◀ \(Assignment\) Change free list management in xv6](#)

Jump to...

Started on Thursday, 18 March 2021, 2:46 PM

State Finished

Completed on Thursday, 18 March 2021, 3:50 PM

Time taken 1 hour 4 mins

Grade 10.36 out of 20.00 (52%)

Question 1

Partially correct

Mark 0.57 out of 1.00

Mark True, the actions done as part of code of swtch() in swtch.S, in xv6

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Restore new callee saved registers from kernel stack of new context	
<input checked="" type="radio"/>	<input type="radio"/>	Save old callee saved registers on kernel stack of old context	
	<input checked="" type="radio"/>	Save old callee saved registers on user stack of old context	
	<input checked="" type="radio"/>	Switch from old process context to new process context	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Switch from one stack (old) to another(new)	
	<input checked="" type="radio"/>	Restore new callee saved registers from user stack of new context	
	<input checked="" type="radio"/>	Jump to code in new context	

Restore new callee saved registers from kernel stack of new context: True

Save old callee saved registers on kernel stack of old context: True

Save old callee saved registers on user stack of old context: False

Switch from old process context to new process context: False

Switch from one stack (old) to another(new): True

Restore new callee saved registers from user stack of new context: False

Jump to code in new context: False

Question 2

Partially correct

Mark 0.17 out of 0.50

For each function/code-point, select the status of segmentation setup in xv6

bootmain()	gdt setup with 3 entries, right from first line of code of bootloader	✗
kvmalloc() in main()	gdt setup with 5 entries (0 to 4) on one processor	✗
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	✓
after seginit() in main()	gdt setup with 5 entries (0 to 4) on all processors	✗
bootasm.S	gdt setup with 3 entries, right from first line of code of bootloader	✗
entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S

Question 3

Partially correct

Mark 0.38 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- a. The meaning of valid-invalid bit in page table is different in paging and demand-paging. ✓
- b. Demand paging requires additional hardware support, compared to paging. ✓
- c. Paging requires some hardware support in CPU
- d. With paging, it's possible to have user programs bigger than physical memory. ✗
- e. Both demand paging and paging support shared memory pages. ✓
- f. Demand paging always increases effective memory access time.
- g. With demand paging, it's possible to have user programs bigger than physical memory. ✓
- h. Calculations of number of bits for page number and offset are same in paging and demand paging. ✓
- i. TLB hit ration has zero impact in effective memory access time in demand paging.
- j. Paging requires NO hardware support in CPU

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question 4

Partially correct

Mark 0.44 out of 0.50

Suppose a processor supports base(relocation register) + limit scheme of MMU.

Assuming this, mark the statements as True/False

True	False	
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS may terminate the process while handling the interrupt of memory violation
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The hardware detects any memory access beyond the limit value and raises an interrupt
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/>	The hardware may terminate the process while handling the interrupt of memory violation
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS sets up the relocation and limit registers when the process is scheduled
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The process sets up its own relocation and limit registers when the process is scheduled
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS detects any memory access beyond the limit value and raises an interrupt
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The compiler generates machine code assuming appropriately sized segments for code, data and stack.

The OS may terminate the process while handling the interrupt of memory violation: True

The hardware detects any memory access beyond the limit value and raises an interrupt: True

The hardware may terminate the process while handling the interrupt of memory violation: False

The OS sets up the relocation and limit registers when the process is scheduled: True

The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;: True

The process sets up its own relocation and limit registers when the process is scheduled: False

The OS detects any memory access beyond the limit value and raises an interrupt: False

The compiler generates machine code assuming appropriately sized segments for code, data and stack.: False

Question 5

Correct

Mark 0.50 out of 0.50

Consider the following list of free chunks, in continuous memory management:

10k, 25k, 12k, 7k, 9k, 13k

Suppose there is a request for chunk of size 9k, then the free chunk selected under each of the following schemes will be

Best fit:

9k



First fit:

10k



Worst fit:

25k

**Question 6**

Partially correct

Mark 0.50 out of 1.00

Select all the correct statements about MMU and its functionality

Select one or more:

- a. MMU is a separate chip outside the processor
- b. MMU is inside the processor ✓
- c. Logical to physical address translations in MMU are done with specific machine instructions
- d. The operating system interacts with MMU for every single address translation ✗
- e. Illegal memory access is detected in hardware by MMU and a trap is raised ✓
- f. The Operating system sets up relevant CPU registers to enable proper MMU translations
- g. Logical to physical address translations in MMU are done in hardware, automatically ✓
- h. Illegal memory access is detected by operating system

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

Question 7

Incorrect

Mark 0.00 out of 0.50

Assuming a 8- KB page size, what is the page numbers for the address 874815 reference in decimal :
(give answer also in decimal)

Answer: 2186



The correct answer is: 107

Question 8

Incorrect

Mark 0.00 out of 0.25

Select the compiler's view of the process's address space, for each of the following MMU schemes:
(Assume that each scheme,e.g. paging/segmentation/etc is effectively utilised)

Segmentation, then paging	Many continuous chunks each of page size	
Relocation + Limit	Many continuous chunks of same size	
Segmentation	one continuous chunk	
Paging	many continuous chunks of variable size	

Your answer is incorrect.

The correct answer is: Segmentation, then paging → many continuous chunks of variable size, Relocation + Limit → one continuous chunk, Segmentation → many continuous chunks of variable size, Paging → one continuous chunk

Question 9

Incorrect

Mark 0.00 out of 0.50

Suppose the memory access time is 180ns and TLB hit ratio is 0.3, then effective memory access time is (in nanoseconds);

Answer: 192



The correct answer is: 306.00

Question 10

Correct

Mark 0.50 out of 0.50

In xv6, The struct context is given as

```
struct context {
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};
```

Select all the reasons that explain why only these 5 registers are included in the struct context.

- a. The segment registers are same across all contexts, hence they need not be saved ✓
- b. esp is not saved in context, because context{} is on stack and it's address is always argument to swtch() ✓
- c. xv6 tries to minimize the size of context to save memory space
- d. esp is not saved in context, because it's not part of the context
- e. eax, ecx, edx are caller save, hence no need to save ✓

Your answer is correct.

The correct answers are: The segment registers are same across all contexts, hence they need not be saved, eax, ecx, edx are caller save, hence no need to save, esp is not saved in context, because context{} is on stack and it's address is always argument to swtch()

Question 11

Partially correct

Mark 0.83 out of 1.50

Arrange the following events in order, in page fault handling:

Disk interrupt wakes up the process

7	✓
---	---

The reference bit is found to be invalid by MMU

1	✓
---	---

OS makes available an empty frame

6	✗
---	---

Restart the instruction that caused the page fault

9	✓
---	---

A hardware interrupt is issued

3	✗
---	---

OS schedules a disk read for the page (from backing store)

5	✓
---	---

Process is kept in wait state

4	✗
---	---

Page tables are updated for the process

8	✓
---	---

Operating system decides that the page was not in memory

2	✗
---	---

Your answer is partially correct.

You have correctly selected 5.

The correct answer is: Disk interrupt wakes up the process → 7, The reference bit is found to be invalid by MMU → 1, OS makes available an empty frame → 4, Restart the instruction that caused the page fault → 9, A hardware interrupt is issued → 2, OS schedules a disk read for the page (from backing store) → 5, Process is kept in wait state → 6, Page tables are updated for the process → 8, Operating system decides that the page was not in memory → 3

Question 12

Incorrect

Mark 0.00 out of 0.50

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

00001010

Now, there is a request for a chunk of 70 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: 11101010



The correct answer is: 11111010

Question 13

Incorrect

Mark 0.00 out of 0.25

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- a. end, 4MB
- b. P2V(end), P2V(PHYSTOP)
- c. end, P2V(4MB + PHYSTOP)
- d. P2V(end), PHYSTOP ✗
- e. end, (4MB + PHYSTOP)
- f. end, PHYSTOP
- g. end, P2V(PHYSTOP)

Your answer is incorrect.

The correct answer is: end, P2V(PHYSTOP)

Question 14

Partially correct

Mark 0.33 out of 0.50

Match the pair

Hashed page table	Linear search on collision done by OS (e.g. SPARC Solaris) typically	✓
Inverted Page table	Linear/Parallel search using frame number in page table	✗
Hierarchical Paging	More memory access time per hierarchy	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Hashed page table → Linear search on collision done by OS (e.g. SPARC Solaris) typically, Inverted Page table → Linear/Parallel search using page number in page table, Hierarchical Paging → More memory access time per hierarchy

Question 15

Partially correct

Mark 0.29 out of 0.50

After virtual memory is implemented

(select T/F for each of the following) One Program's size can be larger than physical memory size

True	False	
<input checked="" type="radio"/>	<input type="radio"/> ✗	Code need not be completely in memory
<input checked="" type="radio"/>	<input type="radio"/> ✗	Cumulative size of all programs can be larger than physical memory size
<input type="radio"/> ✗	<input checked="" type="radio"/>	Virtual access to memory is granted
<input checked="" type="radio"/>	<input type="radio"/> ✗	Logical address space could be larger than physical address space
<input type="radio"/> ✗	<input checked="" type="radio"/>	Virtual addresses are available
<input checked="" type="radio"/>	<input checked="" type="radio"/> ✗	Relatively less I/O may be possible during process execution
<input checked="" type="radio"/>	<input type="radio"/> ✗	One Program's size can be larger than physical memory size

Code need not be completely in memory: True

Cumulative size of all programs can be larger than physical memory size: True

Virtual access to memory is granted: False

Logical address space could be larger than physical address space: True

Virtual addresses are available: False

Relatively less I/O may be possible during process execution: True

One Program's size can be larger than physical memory size: True

Question 16

Partially correct

Mark 0.64 out of 1.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only
Mark statements True or False**True False**

<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The stack allocated in entry.S is used as stack for scheduler's context for first processor	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The free page-frame are created out of nearly 222 MB	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The kernel code and data take up less than 2 MB space	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	PHYSTOP can be increased to some extent, simply by editing memlayout.h	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	xv6 uses physical memory upto 224 MB only	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The process's address space gets mapped on frames, obtained from ~2MB:224MB range	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context	✗

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

The free page-frame are created out of nearly 222 MB: True

The kernel code and data take up less than 2 MB space: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

xv6 uses physical memory upto 224 MB only: True

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

Question 17

Incorrect

Mark 0.00 out of 1.50

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using LRU replacement is:

Answer: ✖

#6# 6,4# 6,4,2 # 0,4,2#0,1,2#6,1,2#6,9,2#0,9,2#0,5,2

The correct answer is: 9

Question 18

Partially correct

Mark 0.31 out of 0.50

Consider the image given below, which explains how paging works.

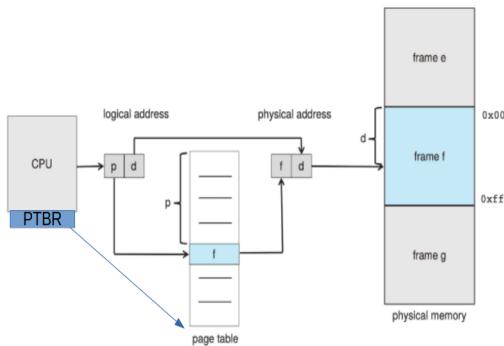


Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The PTBR is present in the CPU as a register
<input type="radio"/>	<input checked="" type="radio"/>	The page table is indexed using frame number
<input checked="" type="radio"/>	<input type="radio"/>	The page table is indexed using page number
<input type="radio"/>	<input checked="" type="radio"/>	The locating of the page table using PTBR also involves paging translation
<input type="radio"/>	<input checked="" type="radio"/>	Size of page table is always determined by the size of RAM
<input checked="" type="radio"/>	<input type="radio"/>	The page table is itself present in Physical memory
<input checked="" type="radio"/>	<input type="radio"/>	Maximum Size of page table is determined by number of bits used for page number
<input checked="" type="radio"/>	<input type="radio"/>	The physical address may not be of the same size (in bits) as the logical address

The PTBR is present in the CPU as a register: True

The page table is indexed using frame number: False

The page table is indexed using page number: True

The locating of the page table using PTBR also involves paging translation: False

Size of page table is always determined by the size of RAM: False

The page table is itself present in Physical memory: True

Maximum Size of page table is determined by number of bits used for page number: True

The physical address may not be of the same size (in bits) as the logical address: True

Question 19

Correct

Mark 2.00 out of 2.00

Given below is shared memory code with two processes sharing a memory segment.

The first process sends a user input string to second process. The second capitalizes the string. Then the first process prints the capitalized version.

Fill in the blanks to complete the code.

// First process

```
#define SHMSZ 27

int main()
{
    char c;
    int shmid;
    key_t key;
    char *shm, *s, string[128];
    key = 5679;
    if ((shmid =
        shmget
        ✓ (key, SHMSZ, IPC_CREAT | 0666)) < 0) {
        perror("shmget");
        exit(1);
    }
    if ((shm =
        shmat
        ✓ (shmid, NULL, 0)) == (char *) -1) {
        perror("shmat");
        exit(1);
    }
    s = shm;
    *s = '$';
    scanf("%s", string);
    strcpy(s + 1, string);
    *s =
        @
        ✓ ';' //note the quotes
    while(*s != '
        $
        ')
        sleep(1);
        printf("%s\n", s + 1);
        exit(0);
}
```

//Second process

```
#define SHMSZ 27

int main()
{
    int shmid;
    key_t key;
    char *shm, *s;
    int i;
    char string[128];
    key =
        5679
```

```

✓ ;
if ((shmid = shmget(key, SHMSZ, 0666)) < 0) {
    perror("shmget");
    exit(1);
}
if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
    perror("shmat");
    exit(1);
}
s =

✓ ;
while(*s != '@')
    sleep(1);
for(i = 0; i < strlen(s + 1); i++)
    s[i + 1] = toupper(s[i + 1]);
*s = '$';
exit(0);
}

```

Question 20

Partially correct

Mark 0.25 out of 0.50

Map the functionality/use with function/variable in xv6 code.

return a free page, if available; 0, otherwise

Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed

Array listing the kernel memory mappings, to be used by setupkvm()

Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

Setup kernel part of a page table, and switch to that page table

 kinit1()

 mappages()

 kmap[]

 kvmalloc()

 walkpgdir()

 setupkvm()

Your answer is partially correct.

You have correctly selected 3.

The correct answer is: return a free page, if available; 0, otherwise → kalloc(), Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → mappages(), Array listing the kernel memory mappings, to be used by setupkvm() → kmap[], Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → setupkvm(), Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → walkpgdir(), Setup kernel part of a page table, and switch to that page table → kvmalloc()

Question 21

Partially correct

Mark 1.53 out of 2.50

Order events in xv6 timer interrupt code

(Transition from process P1 to P2's code.)

P2 is selected and marked RUNNING

12 ✓

Change of stack from user stack to kernel stack of P1

3 ✓

Timer interrupt occurs

2 ✓

alltraps() will call iret

17 ✗

change to context of P2, P2's kernel stack in use now

13 ✓

P2's trap() will return to alltraps

16 ✗

jump in vector.S

4 ✓

P2 will return from sched() in yield()

14 ✗

yield() is called

8 ✓

trap() is called

7 ✓

Process P2 is executing

18 ✗

P1 is marked as RUNNABLE

9 ✓

P2's yield() will return in trap()

15 ✗

Process P1 is executing

1 ✓

sched() is called,

11 ✗

change to context of the scheduler, scheduler's stack in use now

10 ✗

jump to alltraps

5 ✓

Trapframe is built on kernel stack of P1

6 ✓

Your answer is partially correct.

You have correctly selected 11.

The correct answer is: P2 is selected and marked RUNNING → 12, Change of stack from user stack to kernel stack of P1 → 3, Timer interrupt occurs → 2, alltraps() will call iret → 18, change to context of P2, P2's kernel stack in use now → 13, P2's trap() will return to alltraps → 17, jump in vector.S → 4, P2 will return from sched() in yield() → 15, yield() is called → 8, trap() is called → 7, Process P2 is executing → 14, P1 is marked as RUNNABLE → 9, P2's yield() will return in trap() → 16, Process P1 is executing → 1, sched() is called, → 10, change to context of the scheduler, scheduler's stack in use now → 11, jump to alltraps → 5, Trapframe is built on kernel stack of P1 → 6

Question 22

Incorrect

Mark 0.00 out of 1.00

Given that the memory access time is 200 ns, probability of a page fault is 0.7 and page fault handling time is 8 ms,
The effective memory access time in nanoseconds is:

Answer: ✖

The correct answer is: 5600060.00

Question 23

Correct

Mark 0.25 out of 0.25

Select the state that is not possible after the given state, for a process:

- New: Running ✓
- Ready : Waiting ✓
- Running: None of these ✓
- Waiting: Running ✓

Question 24

Partially correct

Mark 0.63 out of 1.00

Select the correct statements about sched() and scheduler() in xv6 code

- a. scheduler() switches to the selected process's context ✓
- b. When either sched() or scheduler() is called, it does not return immediately to caller ✓
- c. After call to swtch() in sched(), the control moves to code in scheduler()
- d. Each call to sched() or scheduler() involves change of one stack inside swtch() ✓
- e. After call to swtch() in scheduler(), the control moves to code in sched()
- f. When either sched() or scheduler() is called, it results in a context switch ✓
- g. sched() switches to the scheduler's context ✓
- h. sched() and scheduler() are co-routines

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: sched() and scheduler() are co-routines, When either sched() or scheduler() is called, it does not return immediately to caller, When either sched() or scheduler() is called, it results in a context switch, sched() switches to the scheduler's context, scheduler() switches to the selected process's context, After call to swtch() in scheduler(), the control moves to code in sched(), After call to swtch() in sched(), the control moves to code in scheduler(), Each call to sched() or scheduler() involves change of one stack inside swtch()

Question 25

Correct

Mark 0.25 out of 0.25

The data structure used in kalloc() and kfree() in xv6 is

- a. Doubly linked circular list
- b. Singly linked circular list
- c. Double linked NULL terminated list
- d. Singly linked NULL terminated list



Your answer is correct.

The correct answer is: Singly linked NULL terminated list

[◀ \(Assignment\) lseek system call in xv6](#)

Jump to...

Question **14**

Not yet answered

Marked out of 3.00

List down all changes required to xv6 code, in order to add the system call chown().

Time left 1:53:37

Every change should be mentioned in terms of either of the following:

- (a) pseudo-code of new function to be added
- (b) prototype of any new function or new system call to be added
- (c) pseudo-code of changes to an existing function, describing lines to be removed, and lines to be added
- (d) **precise** declaration of new data structures to be added in C, or changes to the existing data structure
- (e) Name and a one-line description of new userland functionality to be added
- (f) Changes to Makefile
- (g) Any other change in a maximum of 20 words per change.

Paragraph

(a) int chown (char * path , char *ownername);
(d)there is no need of new data structures or changes to the existing data structure
(e)by using this system call ,user to change the owner if the file
(f)_trychown\ at the end of UPROGS

◀ ▶

Path: p

◀ Random Quiz - 6 (xv6 file system)

Jump to...

Homework questions: Basics of MM, xv6 booting ►

Question **18**

Not yet answered

Marked out of 2.00

Time left 1:34:55

Write all changes required to xv6 to add a buddy allocator.

Every change should be mentioned in terms of either of the following:

- (a) pseudo-code of new function to be added
- (b) prototype of any new function or new system call to be added
- (c) pseudo-code of changes to an existing function, describing lines to be removed, and lines to be added
- (d) **precise** declaration of new data structures to be added in C, or changes to the existing data structure
- (e) Name and a one-line description of new userland functionality to be added
- (f) Changes to Makefile
- (g) Any other change in a maximum of 20 words per change.

Paragraph

Path: p

[◀ Random Quiz - 6 \(xv6 file system\)](#)

Jump to...

[Homework questions: Basics of MM, xv6 booting ▶](#)

Question 17

Not yet answered

Marked out of 1.00

Match the code with its functionality

S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statement2;

Execution order P1, P2, P3

Time left 1:35:43

Signal(S1);

P3:

Wait(S1);

Statement S3;

S = 0

P1:

Statement1;

Signal(S)

Execution order P3, P2, P1

P2:

Wait(S)

Statement2;

S = 5

Wait(S)

Critical Section

Execution order P2, then P1

Signal(S)

S = 1

Wait(S)

Critical Section

Counting semaphore

Signal(S);

[◀ Random Quiz - 6 \(xv6 file system\)](#)

Jump to...

[Homework questions: Basics of MM, xv6 booting ►](#)

Time left 1:48:32

Question 13

Not yet answered

Marked out of 2.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
void  
yield(void)  
{  
...  
release(&ptable.lock);  
}
```

Ensure that no printing happens on other processors

```
struct proc*  
myproc(void) {  
...  
pushcli();  
c = mycpu();  
p = c->proc;  
popcli();  
...  
}
```

Disable interrupts to avoid another process's pointer being returned

```
static inline uint  
xchg(volatile uint *addr, uint newval)  
{  
    uint result;  
  
    // The + in "+m" denotes a read-modify-write operand.  
    asm volatile("lock; xchgl %0, %1" :  
        "+m" (*addr), "=a" (result) :  
        "1" (newval) :  
        "cc");  
    return result;  
}  
  
void  
acquire(struct spinlock *lk)  
{  
...  
    __sync_synchronize();  
}
```

Atomic compare and swap instruction (to be expanded inline into code)

Avoid a self-deadlock

```
void  
acquire(struct spinlock *lk)  
{  
...  
getcallerpcs(&lk, lk->pcs);
```

Disable interrupts to avoid deadlocks

```
void  
acquire(struct spinlock *lk)  
{  
pushcli();
```

Traverse ebp chain to get sequence of instructions followed in functions calls

```
void  
sleep(void *chan, struct spinlock *lk)  
{  
...  
if(lk != &phtable.lock){  
    acquire(&phtable.lock);  
    release(lk);  
}
```

Release the lock held by some another process

```
void  
panic(char *s)  
{  
...  
panicked = 1;
```

If you don't do this, a process may be running on two processors parallely

◀ Random Quiz - 6 (xv6 file system)

Jump to...

Homework questions: Basics of MM, xv6 booting ►

Question 12

Not yet answered

Marked out of 2.00

Select all the correct statements about synchronization primitives.

Select one or more:

Time left 2:04:57

- a. Thread that is going to block should not be holding any spinlock
- b. Semaphores can be used for synchronization scenarios like ordered execution
- c. Mutexes can be implemented using spinlock
- d. Blocking means one process passing over control to another process
- e. Semaphores are always a good substitute for spinlocks
- f. Blocking means moving the process to a wait queue and spinning
- g. Blocking means moving the process to a wait queue and calling scheduler
- h. All synchronization primitives are implemented essentially with some hardware assistance.
- i. Spinlocks are good for multiprocessor scenarios, for small critical sections
- j. Mutexes can be implemented without any hardware assistance
- k. Spinlocks consume CPU time
- l. Mutexes can be implemented using blocking and wakeup

[◀ Random Quiz - 6 \(xv6 file system\)](#)

Jump to...

[Homework questions: Basics of MM, xv6 booting ►](#)

Question 10

Not yet answered

Marked out of 1.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
void  
yield(void)  
{  
...  
release(&ptable.lock);  
}
```

Time left 2:12:36

Release the lock held by some another process

```
void  
panic(char *s)  
{  
...  
panicked = 1;
```

If you don't do this, a process may be running on two processors parallely

```
void  
acquire(struct spinlock *lk)  
{  
...  
getcallerpcs(&lk, lk->pcs);
```

Disable interrupts to avoid another process's pointer being returned

[◀ Random Quiz - 6 \(xv6 file system\)](#)

Jump to...

Homework questions: Basics of MM, xv6 booting ►

Time left 2:50:14

Question 3

Not yet answered

Marked out of 1.00

Match each suggested semaphore implementation (discussed in class)

with the problems that it faces

```
struct semaphore {  
    int val;  
    spinlock lk;  
};  
sem_init(semaphore *s, int initval) {  
    s->val = initval;  
    s->sl = 0;  
}  
wait(semaphore *s) {  
    spinlock(&(s->sl));  
    while(s->val <=0) {  
        spinunlock(&(s->sl));  
        spinlock(&(s->sl));  
    }  
    (s->val)--;  
    spinunlock(&(s->sl));  
}
```

Choose...

```
struct semaphore {  
    int val;  
    spinlock lk;  
};  
sem_init(semaphore *s, int initval) {  
    s->val = initval;  
    s->sl = 0;  
}  
wait(semaphore *s) {  
    spinlock(&(s->sl));  
    while(s->val <=0)  
    ;  
    (s->val)--;  
    spinunlock(&(s->sl));  
}
```

Choose...

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->sl));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

signal(semaphore *s) {
    spinlock(*(s->sl));
    (s->val)++;
    x = dequeue(s->sl) and enqueue(readyq, x);
    spinunlock(*(s->sl));
}

```

Choose...

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

```

Choose...

[◀ Random Quiz - 6 \(xv6 file system\)](#)

Jump to...

Homework questions: Basics of MM, xv6 booting ►

Time left 2:50:14

Question 3

Not yet answered

Marked out of 1.00

Match each suggested semaphore implementation (discussed in class)

with the problems that it faces

```
struct semaphore {  
    int val;  
    spinlock lk;  
};  
sem_init(semaphore *s, int initval) {  
    s->val = initval;  
    s->sl = 0;  
}  
wait(semaphore *s) {  
    spinlock(&(s->sl));  
    while(s->val <=0) {  
        spinunlock(&(s->sl));  
        spinlock(&(s->sl));  
    }  
    (s->val)--;  
    spinunlock(&(s->sl));  
}
```

Choose...

```
struct semaphore {  
    int val;  
    spinlock lk;  
};  
sem_init(semaphore *s, int initval) {  
    s->val = initval;  
    s->sl = 0;  
}  
wait(semaphore *s) {  
    spinlock(&(s->sl));  
    while(s->val <=0)  
    ;  
    (s->val)--;  
    spinunlock(&(s->sl));  
}
```

Choose...

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->sl));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

signal(semaphore *s) {
    spinlock(*(s->sl));
    (s->val)++;
    x = dequeue(s->sl) and enqueue(readyq, x);
    spinunlock(*(s->sl));
}

```

Choose...

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

```

Choose...

[◀ Random Quiz - 6 \(xv6 file system\)](#)

Jump to...

Homework questions: Basics of MM, xv6 booting ►

Started on Thursday, 18 March 2021, 2:46 PM

State Finished

Completed on Thursday, 18 March 2021, 3:50 PM

Time taken 1 hour 4 mins

Grade 10.36 out of 20.00 (52%)

Question 1

Partially correct

Mark 0.57 out of 1.00

Mark True, the actions done as part of code of swtch() in swtch.S, in xv6

True

False

<input checked="" type="radio"/>	<input checked="" type="radio"/>	Restore new callee saved registers from kernel stack of new context	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Save old callee saved registers on kernel stack of old context	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Save old callee saved registers on user stack of old context	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Switch from old process context to new process context	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Switch from one stack (old) to another(new)	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Restore new callee saved registers from user stack of new context	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Jump to code in new context	✗

Restore new callee saved registers from kernel stack of new context: True

Save old callee saved registers on kernel stack of old context: True

Save old callee saved registers on user stack of old context: False

Switch from old process context to new process context: False

Switch from one stack (old) to another(new): True

Restore new callee saved registers from user stack of new context: False

Jump to code in new context: False

Question 2

Partially correct

Mark 0.17 out of 0.50

For each function/code-point, select the status of segmentation setup in xv6

bootmain()	gdt setup with 3 entries, right from first line of code of bootloader	✗
kvmalloc() in main()	gdt setup with 5 entries (0 to 4) on one processor	✗
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	✓
after seginit() in main()	gdt setup with 5 entries (0 to 4) on all processors	✗
bootasm.S	gdt setup with 3 entries, right from first line of code of bootloader	✗
entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S

Question 3

Partially correct

Mark 0.38 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- a. The meaning of valid-invalid bit in page table is different in paging and demand-paging. ✓
- b. Demand paging requires additional hardware support, compared to paging. ✓
- c. Paging requires some hardware support in CPU
- d. With paging, it's possible to have user programs bigger than physical memory. ✗
- e. Both demand paging and paging support shared memory pages. ✓
- f. Demand paging always increases effective memory access time.
- g. With demand paging, it's possible to have user programs bigger than physical memory. ✓
- h. Calculations of number of bits for page number and offset are same in paging and demand paging. ✓
- i. TLB hit ration has zero impact in effective memory access time in demand paging.
- j. Paging requires NO hardware support in CPU

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question 4

Partially correct

Mark 0.44 out of 0.50

Suppose a processor supports base(relocation register) + limit scheme of MMU.

Assuming this, mark the statements as True/False

True	False	
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS may terminate the process while handling the interrupt of memory violation
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The hardware detects any memory access beyond the limit value and raises an interrupt
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/>	The hardware may terminate the process while handling the interrupt of memory violation
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS sets up the relocation and limit registers when the process is scheduled
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The process sets up its own relocation and limit registers when the process is scheduled
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS detects any memory access beyond the limit value and raises an interrupt
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The compiler generates machine code assuming appropriately sized segments for code, data and stack.

The OS may terminate the process while handling the interrupt of memory violation: True

The hardware detects any memory access beyond the limit value and raises an interrupt: True

The hardware may terminate the process while handling the interrupt of memory violation: False

The OS sets up the relocation and limit registers when the process is scheduled: True

The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;: True

The process sets up its own relocation and limit registers when the process is scheduled: False

The OS detects any memory access beyond the limit value and raises an interrupt: False

The compiler generates machine code assuming appropriately sized segments for code, data and stack.: False

Question 5

Correct

Mark 0.50 out of 0.50

Consider the following list of free chunks, in continuous memory management:

10k, 25k, 12k, 7k, 9k, 13k

Suppose there is a request for chunk of size 9k, then the free chunk selected under each of the following schemes will be

Best fit:

9k



First fit:

10k



Worst fit:

25k

**Question 6**

Partially correct

Mark 0.50 out of 1.00

Select all the correct statements about MMU and its functionality

Select one or more:

- a. MMU is a separate chip outside the processor
- b. MMU is inside the processor ✓
- c. Logical to physical address translations in MMU are done with specific machine instructions
- d. The operating system interacts with MMU for every single address translation ✗
- e. Illegal memory access is detected in hardware by MMU and a trap is raised ✓
- f. The Operating system sets up relevant CPU registers to enable proper MMU translations
- g. Logical to physical address translations in MMU are done in hardware, automatically ✓
- h. Illegal memory access is detected by operating system

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

Question 7

Incorrect

Mark 0.00 out of 0.50

Assuming a 8- KB page size, what is the page numbers for the address 874815 reference in decimal :
(give answer also in decimal)

Answer: 2186



The correct answer is: 107

Question 8

Incorrect

Mark 0.00 out of 0.25

Select the compiler's view of the process's address space, for each of the following MMU schemes:
(Assume that each scheme,e.g. paging/segmentation/etc is effectively utilised)

Segmentation, then paging	Many continuous chunks each of page size	
Relocation + Limit	Many continuous chunks of same size	
Segmentation	one continuous chunk	
Paging	many continuous chunks of variable size	

Your answer is incorrect.

The correct answer is: Segmentation, then paging → many continuous chunks of variable size, Relocation + Limit → one continuous chunk, Segmentation → many continuous chunks of variable size, Paging → one continuous chunk

Question 9

Incorrect

Mark 0.00 out of 0.50

Suppose the memory access time is 180ns and TLB hit ratio is 0.3, then effective memory access time is (in nanoseconds);

Answer: 192



The correct answer is: 306.00

Question 10

Correct

Mark 0.50 out of 0.50

In xv6, The struct context is given as

```
struct context {
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};
```

Select all the reasons that explain why only these 5 registers are included in the struct context.

- a. The segment registers are same across all contexts, hence they need not be saved ✓
- b. esp is not saved in context, because context{} is on stack and it's address is always argument to swtch() ✓
- c. xv6 tries to minimize the size of context to save memory space
- d. esp is not saved in context, because it's not part of the context
- e. eax, ecx, edx are caller save, hence no need to save ✓

Your answer is correct.

The correct answers are: The segment registers are same across all contexts, hence they need not be saved, eax, ecx, edx are caller save, hence no need to save, esp is not saved in context, because context{} is on stack and it's address is always argument to swtch()

Question 11

Partially correct

Mark 0.83 out of 1.50

Arrange the following events in order, in page fault handling:

Disk interrupt wakes up the process

7	✓
---	---

The reference bit is found to be invalid by MMU

1	✓
---	---

OS makes available an empty frame

6	✗
---	---

Restart the instruction that caused the page fault

9	✓
---	---

A hardware interrupt is issued

3	✗
---	---

OS schedules a disk read for the page (from backing store)

5	✓
---	---

Process is kept in wait state

4	✗
---	---

Page tables are updated for the process

8	✓
---	---

Operating system decides that the page was not in memory

2	✗
---	---

Your answer is partially correct.

You have correctly selected 5.

The correct answer is: Disk interrupt wakes up the process → 7, The reference bit is found to be invalid by MMU → 1, OS makes available an empty frame → 4, Restart the instruction that caused the page fault → 9, A hardware interrupt is issued → 2, OS schedules a disk read for the page (from backing store) → 5, Process is kept in wait state → 6, Page tables are updated for the process → 8, Operating system decides that the page was not in memory → 3

Question 12

Incorrect

Mark 0.00 out of 0.50

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

00001010

Now, there is a request for a chunk of 70 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: 11101010



The correct answer is: 11111010

Question 13

Incorrect

Mark 0.00 out of 0.25

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- a. end, 4MB
- b. P2V(end), P2V(PHYSTOP)
- c. end, P2V(4MB + PHYSTOP)
- d. P2V(end), PHYSTOP ✗
- e. end, (4MB + PHYSTOP)
- f. end, PHYSTOP
- g. end, P2V(PHYSTOP)

Your answer is incorrect.

The correct answer is: end, P2V(PHYSTOP)

Question 14

Partially correct

Mark 0.33 out of 0.50

Match the pair

Hashed page table	Linear search on collision done by OS (e.g. SPARC Solaris) typically	✓
Inverted Page table	Linear/Parallel search using frame number in page table	✗
Hierarchical Paging	More memory access time per hierarchy	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Hashed page table → Linear search on collision done by OS (e.g. SPARC Solaris) typically, Inverted Page table → Linear/Parallel search using page number in page table, Hierarchical Paging → More memory access time per hierarchy

Question 15

Partially correct

Mark 0.29 out of 0.50

After virtual memory is implemented

(select T/F for each of the following) One Program's size can be larger than physical memory size

True	False	
<input checked="" type="radio"/>	<input type="radio"/> ✗	Code need not be completely in memory
<input checked="" type="radio"/>	<input type="radio"/> ✗	Cumulative size of all programs can be larger than physical memory size
<input type="radio"/> ✗	<input checked="" type="radio"/>	Virtual access to memory is granted
<input checked="" type="radio"/>	<input type="radio"/> ✗	Logical address space could be larger than physical address space
<input type="radio"/> ✗	<input checked="" type="radio"/>	Virtual addresses are available
<input checked="" type="radio"/>	<input checked="" type="radio"/> ✗	Relatively less I/O may be possible during process execution
<input checked="" type="radio"/>	<input type="radio"/> ✗	One Program's size can be larger than physical memory size

Code need not be completely in memory: True

Cumulative size of all programs can be larger than physical memory size: True

Virtual access to memory is granted: False

Logical address space could be larger than physical address space: True

Virtual addresses are available: False

Relatively less I/O may be possible during process execution: True

One Program's size can be larger than physical memory size: True

Question 16

Partially correct

Mark 0.64 out of 1.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only
Mark statements True or False**True False**

<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The stack allocated in entry.S is used as stack for scheduler's context for first processor	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The free page-frame are created out of nearly 222 MB	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The kernel code and data take up less than 2 MB space	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	PHYSTOP can be increased to some extent, simply by editing memlayout.h	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	xv6 uses physical memory upto 224 MB only	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The process's address space gets mapped on frames, obtained from ~2MB:224MB range	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context	✗

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

The free page-frame are created out of nearly 222 MB: True

The kernel code and data take up less than 2 MB space: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

xv6 uses physical memory upto 224 MB only: True

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

Question 17

Incorrect

Mark 0.00 out of 1.50

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using LRU replacement is:

Answer: ✖

#6# 6,4# 6,4,2 # 0,4,2#0,1,2#6,1,2#6,9,2#0,9,2#0,5,2

The correct answer is: 9

Question 18

Partially correct

Mark 0.31 out of 0.50

Consider the image given below, which explains how paging works.

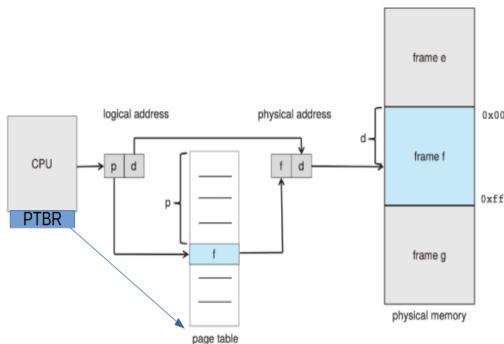


Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The PTBR is present in the CPU as a register
<input type="radio"/>	<input checked="" type="radio"/>	The page table is indexed using frame number
<input checked="" type="radio"/>	<input type="radio"/>	The page table is indexed using page number
<input type="radio"/>	<input checked="" type="radio"/>	The locating of the page table using PTBR also involves paging translation
<input type="radio"/>	<input checked="" type="radio"/>	Size of page table is always determined by the size of RAM
<input checked="" type="radio"/>	<input type="radio"/>	The page table is itself present in Physical memory
<input checked="" type="radio"/>	<input type="radio"/>	Maximum Size of page table is determined by number of bits used for page number
<input checked="" type="radio"/>	<input type="radio"/>	The physical address may not be of the same size (in bits) as the logical address

The PTBR is present in the CPU as a register: True

The page table is indexed using frame number: False

The page table is indexed using page number: True

The locating of the page table using PTBR also involves paging translation: False

Size of page table is always determined by the size of RAM: False

The page table is itself present in Physical memory: True

Maximum Size of page table is determined by number of bits used for page number: True

The physical address may not be of the same size (in bits) as the logical address: True

Question 19

Correct

Mark 2.00 out of 2.00

Given below is shared memory code with two processes sharing a memory segment.

The first process sends a user input string to second process. The second capitalizes the string. Then the first process prints the capitalized version.

Fill in the blanks to complete the code.

// First process

```
#define SHMSZ 27

int main()
{
    char c;
    int shmid;
    key_t key;
    char *shm, *s, string[128];
    key = 5679;
    if ((shmid =
        shmget
        ✓ (key, SHMSZ, IPC_CREAT | 0666)) < 0) {
        perror("shmget");
        exit(1);
    }
    if ((shm =
        shmat
        ✓ (shmid, NULL, 0)) == (char *) -1) {
        perror("shmat");
        exit(1);
    }
    s = shm;
    *s = '$';
    scanf("%s", string);
    strcpy(s + 1, string);
    *s =
        @
        ✓ ';' //note the quotes
    while(*s != '
        $
        ')
        sleep(1);
        printf("%s\n", s + 1);
        exit(0);
}
```

//Second process

```
#define SHMSZ 27

int main()
{
    int shmid;
    key_t key;
    char *shm, *s;
    int i;
    char string[128];
    key =
        5679
```

```

✓ ;
if ((shmid = shmget(key, SHMSZ, 0666)) < 0) {
    perror("shmget");
    exit(1);
}
if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
    perror("shmat");
    exit(1);
}
s =

✓ ;
while(*s != '@')
    sleep(1);
for(i = 0; i < strlen(s + 1); i++)
    s[i + 1] = toupper(s[i + 1]);
*s = '$';
exit(0);
}

```

Question 20

Partially correct

Mark 0.25 out of 0.50

Map the functionality/use with function/variable in xv6 code.

return a free page, if available; 0, otherwise

Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed

Array listing the kernel memory mappings, to be used by setupkvm()

Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

Setup kernel part of a page table, and switch to that page table

 kinit1()

 mappages()

 kmap[]

 kvmalloc()

 walkpgdir()

 setupkvm()

Your answer is partially correct.

You have correctly selected 3.

The correct answer is: return a free page, if available; 0, otherwise → kalloc(), Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → mappages(), Array listing the kernel memory mappings, to be used by setupkvm() → kmap[], Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → setupkvm(), Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → walkpgdir(), Setup kernel part of a page table, and switch to that page table → kvmalloc()

Question 21

Partially correct

Mark 1.53 out of 2.50

Order events in xv6 timer interrupt code

(Transition from process P1 to P2's code.)

P2 is selected and marked RUNNING

12 ✓

Change of stack from user stack to kernel stack of P1

3 ✓

Timer interrupt occurs

2 ✓

alltraps() will call iret

17 ✗

change to context of P2, P2's kernel stack in use now

13 ✓

P2's trap() will return to alltraps

16 ✗

jump in vector.S

4 ✓

P2 will return from sched() in yield()

14 ✗

yield() is called

8 ✓

trap() is called

7 ✓

Process P2 is executing

18 ✗

P1 is marked as RUNNABLE

9 ✓

P2's yield() will return in trap()

15 ✗

Process P1 is executing

1 ✓

sched() is called,

11 ✗

change to context of the scheduler, scheduler's stack in use now

10 ✗

jump to alltraps

5 ✓

Trapframe is built on kernel stack of P1

6 ✓

Your answer is partially correct.

You have correctly selected 11.

The correct answer is: P2 is selected and marked RUNNING → 12, Change of stack from user stack to kernel stack of P1 → 3, Timer interrupt occurs → 2, alltraps() will call iret → 18, change to context of P2, P2's kernel stack in use now → 13, P2's trap() will return to alltraps → 17, jump in vector.S → 4, P2 will return from sched() in yield() → 15, yield() is called → 8, trap() is called → 7, Process P2 is executing → 14, P1 is marked as RUNNABLE → 9, P2's yield() will return in trap() → 16, Process P1 is executing → 1, sched() is called, → 10, change to context of the scheduler, scheduler's stack in use now → 11, jump to alltraps → 5, Trapframe is built on kernel stack of P1 → 6

Question 22

Incorrect

Mark 0.00 out of 1.00

Given that the memory access time is 200 ns, probability of a page fault is 0.7 and page fault handling time is 8 ms,
The effective memory access time in nanoseconds is:

Answer: ✖

The correct answer is: 5600060.00

Question 23

Correct

Mark 0.25 out of 0.25

Select the state that is not possible after the given state, for a process:

- New: Running ✓
- Ready : Waiting ✓
- Running: None of these ✓
- Waiting: Running ✓

Question 24

Partially correct

Mark 0.63 out of 1.00

Select the correct statements about sched() and scheduler() in xv6 code

- a. scheduler() switches to the selected process's context ✓
- b. When either sched() or scheduler() is called, it does not return immediately to caller ✓
- c. After call to swtch() in sched(), the control moves to code in scheduler()
- d. Each call to sched() or scheduler() involves change of one stack inside swtch() ✓
- e. After call to swtch() in scheduler(), the control moves to code in sched()
- f. When either sched() or scheduler() is called, it results in a context switch ✓
- g. sched() switches to the scheduler's context ✓
- h. sched() and scheduler() are co-routines

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: sched() and scheduler() are co-routines, When either sched() or scheduler() is called, it does not return immediately to caller, When either sched() or scheduler() is called, it results in a context switch, sched() switches to the scheduler's context, scheduler() switches to the selected process's context, After call to swtch() in scheduler(), the control moves to code in sched(), After call to swtch() in sched(), the control moves to code in scheduler(), Each call to sched() or scheduler() involves change of one stack inside swtch()

Question 25

Correct

Mark 0.25 out of 0.25

The data structure used in kalloc() and kfree() in xv6 is

- a. Doubly linked circular list
- b. Singly linked circular list
- c. Double linked NULL terminated list
- d. Singly linked NULL terminated list



Your answer is correct.

The correct answer is: Singly linked NULL terminated list

[◀ \(Assignment\) lseek system call in xv6](#)

Jump to...

Started on Saturday, 20 February 2021, 2:51 PM

State Finished

Completed on Saturday, 20 February 2021, 3:55 PM

Time taken 1 hour 3 mins

Grade 7.30 out of 20.00 (37%)

Question 1

Partially correct

Mark 0.80 out of 1.00

Select all the correct statements about the state of a process.

- a. A process can self-terminate only when it's running ✓
- b. Typically, it's represented as a number in the PCB ✓
- c. A process that is running is not on the ready queue ✓
- d. Processes in the ready queue are in the ready state ✓
- e. It is not maintained in the data structures by kernel, it is only for conceptual understanding of programmers
- f. Changing from running state to waiting state results in "giving up the CPU" ✓
- g. A process in ready state is ready to receive interrupts
- h. A waiting process starts running after the wait is over ✗
- i. A process changes from running to ready state on a timer interrupt ✓
- j. A process in ready state is ready to be scheduled ✓
- k. A running process may terminate, or go to wait or become ready again ✓
- l. A process waiting for I/O completion is typically woken up by the particular interrupt handler code ✓
- m. A process waiting for any condition is woken up by another process only
- n. A process changes from running to ready state on a timer interrupt or any I/O wait

Your answer is partially correct.

You have selected too many options.

The correct answers are: Typically, it's represented as a number in the PCB, A process in ready state is ready to be scheduled, Processes in the ready queue are in the ready state, A process that is running is not on the ready queue, A running process may terminate, or go to wait or become ready again, A process changes from running to ready state on a timer interrupt, Changing from running state to waiting state results in "giving up the CPU", A process can self-terminate only when it's running, A process waiting for I/O completion is typically woken up by the particular interrupt handler code

Question 2

Incorrect

Mark 0.00 out of 1.00

For each line of code mentioned on the left side, select the location of sp/esp that is in use

`jmp *%eax`
in entry.S

0x7c00 to 0x10000



`ljmp $(SEG_KCODE<<3), $start32`
in bootasm.S

0x10000 to 0x7c00



`call bootmain`
in bootasm.S

0x7c00 to 0x10000



`cli`
in bootasm.S

0x7c00 to 0



`readseg((uchar*)elf, 4096, 0);`
in bootmain.c

The 4KB area in kernel image, loaded in memory, named as 'stack'



Your answer is incorrect.

The correct answer is: `jmp *%eax`

`in entry.S` → The 4KB area in kernel image, loaded in memory, named as 'stack', `ljmp $(SEG_KCODE<<3), $start32`

`in bootasm.S` → Immaterail as the stack is not used here, `call bootmain`

`in bootasm.S` → 0x7c00 to 0, `cli`

`in bootasm.S` → Immaterail as the stack is not used here, `readseg((uchar*)elf, 4096, 0);`

`in bootmain.c` → 0x7c00 to 0

Question 3

Correct

Mark 0.25 out of 0.25

Order the following events in boot process (from 1 onwards)

Boot loader	2	✓
Shell	6	✓
BIOS	1	✓
OS	3	✓
Init	4	✓
Login interface	5	✓

Your answer is correct.

The correct answer is: Boot loader → 2, Shell → 6, BIOS → 1, OS → 3, Init → 4, Login interface → 5

Question 4

Partially correct

Mark 0.30 out of 0.50

Consider the following command and its output:

```
$ ls -lht xv6.img kernel
-rw-rw-r-- 1 abhijit abhijit 4.9M Feb 15 11:09 xv6.img
-rwxrwxr-x 1 abhijit abhijit 209K Feb 15 11:09 kernel*
```

Following code in bootmain()

```
readseg((uchar*)elf, 4096, 0);
```

and following selected lines from Makefile

```
xv6.img: bootblock kernel
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

```
kernel: $(OBJS) entry.o entryother initcode kernel.ld
$(LD) $(LDFLAGS) -T kernel.ld -o kernel entry.o $(OBJS) -b binary initcode entryother
$(OBJDUMP) -S kernel > kernel.asm
$(OBJDUMP) -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$$/d' > kernel.sym
```

Also read the code of bootmain() in xv6 kernel.

Select the options that describe the meaning of these lines and their correlation.

- a. Although the size of the kernel file is 209 Kb, only 4Kb out of it is the actual kernel code and remaining part is all zeroes.
- b. The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files ✓
- c. The kernel.ld file contains instructions to the linker to link the kernel properly ✓
- d. The bootmain() code does not read the kernel completely in memory
- e. readseg() reads first 4k bytes of kernel in memory
- f. Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.
- g. The kernel.asm file is the final kernel file
- h. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is not read as it is user programs.
- i. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(). ✓

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(), readseg() reads first 4k bytes of kernel in memory, The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files, The kernel.ld file contains instructions to the linker to link the kernel properly, Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.

Question 5

Partially correct

Mark 0.50 out of 1.00

```
int f() {  
    int count;  
    for (count = 0; count < 2; count++) {  
        if (fork() == 0)  
            printf("Operating-System\n");  
    }  
    printf("TYCOMP\n");  
}
```

The number of times "Operating-System" is printed, is:

Answer:

The correct answer is: 7.00

Question 6

Partially correct

Mark 0.40 out of 0.50

Select Yes/True if the mentioned element must be a part of PCB

Select No/False otherwise.

Yes**No**

<input checked="" type="radio"/>	<input checked="" type="radio"/>	PID	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Process context	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	List of opened files	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Process state	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Parent's PID	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Pointer to IDT	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Function pointers to all system calls	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Memory management information about that process	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Pointer to the parent process	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	EIP at the time of context switch	✓

PID: Yes

Process context: Yes

List of opened files: Yes

Process state: Yes

Parent's PID: No

Pointer to IDT: No

Function pointers to all system calls: No

Memory management information about that process: Yes

Pointer to the parent process: Yes

EIP at the time of context switch: Yes

Question 7

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about code of bootmain() in xv6

```

void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void))(elf->entry);
    entry();
}

```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- a. The kernel file gets loaded at the Physical address 0x10000 +0x80000000 in memory. ✗
- b. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- c. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded ✓
- d. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it. ✓
- e. The kernel file has only two program headers ✓
- f. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- g. The readseg finally invokes the disk I/O code using assembly instructions ✓
- h. The elf->entry is set by the linker in the kernel file and it's 8010000c ✓
- i. The kernel file gets loaded at the Physical address 0x10000 in memory. ✓
- j. The condition if(ph->memsz > ph->filesz) is never true. ✗
- k. The stosb() is used here, to fill in some space in memory with zeroes ✓

Your answer is incorrect.

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

Question 8

Partially correct

Mark 0.13 out of 0.25

Which of the following are NOT a part of job of a typical compiler?

- a. Check the program for logical errors ✓
- b. Convert high level language code to machine code
- c. Process the # directives in a C program
- d. Invoke the linker to link the function calls with their code, extern globals with their declaration
- e. Check the program for syntactical errors
- f. Suggest alternative pieces of code that can be written

Your answer is partially correct.

You have correctly selected 1.

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question 9

Correct

Mark 0.25 out of 0.25

Rank the following storage systems from slowest (first) to fastest(last)

Cache	6	✓
Hard Disk	3	✓
RAM	5	✓
Optical Disks	2	✓
Non volatile memory	4	✓
Registers	7	✓
Magnetic Tapes	1	✓

Your answer is correct.

The correct answer is: Cache → 6, Hard Disk → 3, RAM → 5, Optical Disks → 2, Non volatile memory → 4, Registers → 7, Magnetic Tapes → 1

Question 10

Partially correct

Mark 0.21 out of 0.50

Which of the following parts of a C program do not have any corresponding machine code ?

- a. local variable declaration
- b. global variables
- c. function calls ✗
- d. #directives ✓
- e. expressions
- f. pointer dereference
- g. typedefs ✓

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: #directives, typedefs, global variables

Question 11

Correct

Mark 0.25 out of 0.25

Match a system call with it's description

pipe	create an unnamed FIFO storage with 2 ends - one for reading and another for writing	✓
dup	create a copy of the specified file descriptor into smallest available file descriptor	✓
dup2	create a copy of the specified file descriptor into another specified file descriptor	✓
exec	execute a binary file overlaying the image of current process	✓
fork	create an identical child process	✓

Your answer is correct.

The correct answer is: pipe → create an unnamed FIFO storage with 2 ends - one for reading and another for writing, dup → create a copy of the specified file descriptor into smallest available file descriptor, dup2 → create a copy of the specified file descriptor into another specified file descriptor, exec → execute a binary file overlaying the image of current process, fork → create an identical child process

Question 12

Correct

Mark 0.25 out of 0.25

Match the register with the segment used with it.

eip	cs	✓
edi	es	✓
esi	ds	✓
ebp	ss	✓
esp	ss	✓

Your answer is correct.

The correct answer is: eip → cs, edi → es, esi → ds, ebp → ss, esp → ss

Question 13

Correct

Mark 0.25 out of 0.25

What's the trapframe in xv6?

- a. A frame of memory that contains all the trap handler code
- b. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- c. The IDT table
- d. A frame of memory that contains all the trap handler code's function pointers
- e. A frame of memory that contains all the trap handler's addresses
- f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S ✓
- g. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only

Your answer is correct.

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question 14

Incorrect

Mark 0.00 out of 0.50

Select all the correct statements about linking and loading.

Select one or more:

- a. Continuous memory management schemes can support dynamic linking and dynamic loading. ✗
- b. Loader is last stage of the linker program ✗
- c. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently) ✓
- d. Dynamic linking and loading is not possible without demand paging or demand segmentation. ✓
- e. Dynamic linking essentially results in relocatable code. ✓
- f. Continuous memory management schemes can support static linking and static loading. (may be inefficiently) ✓
- g. Loader is part of the operating system ✓
- h. Static linking leads to non-relocatable code ✗
- i. Dynamic linking is possible with continuous memory management, but variable sized partitions only. ✗

Your answer is incorrect.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

Question 15

Incorrect

Mark 0.00 out of 0.25

In bootasm.S, on the line

```
ljmp    $(SEG_KCODE<<3), $start32
```

The SEG_KCODE << 3, that is shifting of 1 by 3 bits is done because

- a. The value 8 is stored in code segment
- b. The code segment is 16 bit and only upper 13 bits are used for segment number
- c. The code segment is 16 bit and only lower 13 bits are used for segment number ✗
- d. While indexing the GDT using CS, the value in CS is always divided by 8
- e. The ljmp instruction does a divide by 8 on the first argument

Your answer is incorrect.

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question 16

Partially correct

Mark 0.07 out of 0.50

Order the events that occur on a timer interrupt:

Change to kernel stack

1	✗
---	---

Jump to a code pointed by IDT

2	✗
---	---

Jump to scheduler code

5	✗
---	---

Set the context of the new process

4	✗
---	---

Save the context of the currently running process

3	✓
---	---

Execute the code of the new process

6	✗
---	---

Select another process for execution

7	✗
---	---

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: Change to kernel stack → 2, Jump to a code pointed by IDT → 1, Jump to scheduler code → 4, Set the context of the new process → 6, Save the context of the currently running process → 3, Execute the code of the new process → 7, Select another process for execution → 5

Question 17

Incorrect

Mark 0.00 out of 1.00

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

```
$ ls . /tmp/asdfksdf >/tmp/ddd 2>&1
```

Program 1

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    dup(fd);
    close(2);
    dup(fd);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Program 2

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    close(1);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(2);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Select all the correct statements about the programs

Select one or more:

- a. Both programs are correct ✗
- b. Program 2 makes sure that there is one file offset used for '2' and '1' ✗
- c. Only Program 2 is correct ✗
- d. Program 2 does 1>&2 ✗
- e. Program 2 ensures 2>&1 and does not ensure >/tmp/ddd ✗
- f. Program 1 makes sure that there is one file offset used for '2' and '1' ✓
- g. Program 1 is correct for >/tmp/ddd but not for 2>&1 ✗
- h. Program 1 does 1>&2 ✗
- i. Both program 1 and 2 are incorrect ✗
- j. Program 2 is correct for >/tmp/ddd but not for 2>&1 ✗
- k. Only Program 1 is correct ✓
- l. Program 1 ensures 2>&1 and does not ensure >/tmp/ddd ✗

Your answer is incorrect.

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'

Question 18

Correct

Mark 0.25 out of 0.25

Select the option which best describes what the CPU does during its powered ON lifetime

- a. Ask the user what is to be done, and execute that task
- b. Ask the OS what is to be done, and execute that task
- c. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask the User or the OS what is to be done next, repeat
- d. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per ✓ the instruction itself, repeat
- e. Fetch instruction specified by OS, Decode and execute it, repeat
- f. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask OS what is to be done next, repeat

The correct answer is: Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, repeat

Question 19

Partially correct

Mark 0.86 out of 1.00

Consider the following code and MAP the file to which each fd points at the end of the code.

```
int main(int argc, char *argv[]) {
    int fd1, fd2 = 1, fd3 = 1, fd4 = 1;

    fd1 = open("/tmp/1", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    fd2 = open("/tmp/2", O_RDONLY);
    fd3 = open("/tmp/3", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    close(0);
    close(1);
    dup(fd2);
    dup(fd3);
    close(fd3);
    dup2(fd2, fd4);
    printf("%d %d %d %d\n", fd1, fd2, fd3, fd4);
    return 0;
}
```

1	closed	✗
fd4	/tmp/2	✓
fd2	/tmp/2	✓
fd1	/tmp/1	✓
2	stderr	✓
0	/tmp/2	✓
fd3	closed	✓

Your answer is partially correct.

You have correctly selected 6.

The correct answer is: 1 → /tmp/3, fd4 → /tmp/2, fd2 → /tmp/2, fd1 → /tmp/1, 2 → stderr, 0 → /tmp/2, fd3 → closed

Question 20

Incorrect

Mark 0.00 out of 2.00

Following code claims to implement the command

```
/bin/ls -l | /usr/bin/head -3 | /usr/bin/tail -1
```

Fill in the blanks to make the code work.

Note: Do not include space in writing any option. x[1][2] should be written without any space, and so is the case with [1] or [2]. Pay attention to exact syntax and do not write any extra character like ';' or = etc.

```
int main(int argc, char *argv[]) {
```

```
    int pid1, pid2;
```

```
    int pfd[
```

```
    x ] [2];
```

```
    pipe(
```

```
    x );
```

```
    pid1 =
```

```
    x ;
```

```
    if(pid1 != 0) {
```

```
        close(pfd[0]
```

```
    x );
```

```
        close(
```

```
    x );
```

```
        dup(
```

```
    x );
```

```
        execl("/bin/ls", "/bin/ls", "
```

```
    x ", NULL);
```

```
    }
```

```
    pipe(
```

```
    x );
```

```
    x = fork();
```

```
    if(pid2 == 0) {
```

```
        close(
```

```
        x ;
```

```
        close(0);
```

```
        dup(
```

```
        x );
```

```
        close(pfd[1]
```

```
✗ );
close(
  
✗ );
dup(
  
✗ );
execl("/usr/bin/head", "/usr/bin/head", "  
  
", NULL);
} else {
close(pfd
  
✗ );
close(
  
✗ );
dup(
  
✗ );
close(pfd
  
✗ );
execl("/usr/bin/tail", "/usr/bin/tail", "  
  
", NULL);
}  
}
```

Question 21

Partially correct

Mark 0.11 out of 1.00

Select all the correct statements about calling convention on x86 32-bit.

- a. Return address is one location above the ebp ✓
- b. Parameters may be passed in registers or on stack ✓
- c. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function ✓
- d. The ebp pointers saved on the stack constitute a chain of activation records ✓
- e. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables ✗
- f. Parameters may be passed in registers or on stack ✓
- g. The return value is either stored on the stack or returned in the eax register ✗
- h. Parameters are pushed on the stack in left-right order
- i. during execution of a function, ebp is pointing to the old ebp
- j. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function ✗
- k. Compiler may allocate more memory on stack than needed ✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

Question 22

Correct

Mark 1.00 out of 1.00

Match the program with its output (ignore newlines in the output. Just focus on the count of the number of 'hi')

main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi ✓

main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi hi ✓

main() { int i = NULL; fork(); printf("hi\n"); }

hi hi ✓

main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi ✓

Your answer is correct.

The correct answer is: main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi, main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi hi, main() { int i = NULL; fork(); printf("hi\n"); } → hi hi, main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi

Question 23

Incorrect

Mark 0.00 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?

Select all the appropriate choices

- a. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time ✗
- b. The setting up of the most essential memory management infrastructure needs assembly code ✓
- c. The code for reading ELF file can not be written in assembly ✗
- d. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C ✓

Your answer is incorrect.

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question 24

Incorrect

Mark 0.00 out of 0.50

```
xv6.img: bootblock kernel
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- a. The size of the kernel file is nearly 5 MB ✓
- b. The kernel is located at block-1 of the xv6.img ✗
- c. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk. ✗
- d. The size of xv6.img is exactly = (size of bootblock) + (size of kernel) ✗
- e. The bootblock is located on block-0 of the xv6.img ✗
- f. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk. ✓
- g. The bootblock may be 512 bytes or less (looking at the Makefile instruction) ✗
- h. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file. ✗
- i. The size of the xv6.img is nearly 5 MB ✗
- j. xv6.img is the virtual processor used by the qemu emulator ✓
- k. Blocks in xv6.img after kernel may be all zeroes. ✗

Your answer is incorrect.

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question 25

Incorrect

Mark 0.00 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

Select one or more:

a. P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return

b. P1 running

P1 makes sytem call and blocks

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P1 running again



c. P1 running

P1 makes system call

system call returns

P1 running

timer interrupt

Scheduler running

P2 running

d. P1 running

P1 makes sytem call and blocks

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P3 running

Hardware interrupt

Interrupt unblocks P1

Interrupt returns

P3 running

Timer interrupt

Scheduler

P1 running



e.

P1 running

P1 makes sytem call

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P1 running again

f. P1 running

keyboard hardware interrupt

keyboard interrupt handler running

interrupt handler returns

P1 running

P1 makes sytem call

system call returns



P1 running
timer interrupt
scheduler
P2 running

Your answer is incorrect.

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again, P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question 26

Correct

Mark 0.25 out of 0.25

Which of the following are the files related to bootloader in xv6?

- a. bootasm.s and entry.S
- b. bootasm.S and bootmain.c ✓
- c. bootasm.S, bootmain.c and bootblock.c
- d. bootmain.c and bootblock.S

Your answer is correct.

The correct answer is: bootasm.S and bootmain.c

Question 27

Correct

Mark 0.25 out of 0.25

Match the following parts of a C program to the layout of the process in memory

Instructions	Text section	✓
Local Variables	Stack Section	✓
Dynamically allocated memory	Heap Section	✓
Global and static data	Data section	✓

Your answer is correct.

The correct answer is:

Instructions → Text section, Local Variables → Stack Section,
Dynamically allocated memory → Heap Section,
Global and static data → Data section

Question 28

Incorrect

Mark 0.00 out of 0.50

What will this program do?

```
int main() {  
    fork();  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("hello");  
}
```

- a. one process will run ls, another will print hello
- b. run ls once ✗
- c. run ls twice
- d. run ls twice and print hello twice
- e. run ls twice and print hello twice, but output will appear in some random order

Your answer is incorrect.

The correct answer is: run ls twice

Question 29

Correct

Mark 0.25 out of 0.25

What is the OS Kernel?

- a. The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run ✓ correct
- b. The set of tools like compiler, linker, loader, terminal, shell, etc.
- c. Only the system programs like compiler, linker, loader, etc.
- d. Everything that I see on my screen

The correct answer is: The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run

Question 30

Correct

Mark 0.50 out of 0.50

Which of the following is/are not saved during context switch?

- a. Program Counter
- b. General Purpose Registers
- c. Bus ✓
- d. Stack Pointer
- e. MMU related registers/information
- f. Cache ✓
- g. TLB ✓

Your answer is correct.

The correct answers are: TLB, Cache, Bus

Question 31

Partially correct

Mark 0.10 out of 0.25

Select the order in which the various stages of a compiler execute.

Linking	3	
Syntactical Analysis	2	
Pre-processing	1	
Intermediate code generation	does not exist	
Loading	4	

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Linking → 4, Syntactical Analysis → 2, Pre-processing → 1, Intermediate code generation → 3, Loading → does not exist

Question 32

Partially correct

Mark 0.08 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

context of P0 is saved in P0's PCB	2	
context of P1 is loaded from P1's PCB	3	
Process P1 is running	5	
timer interrupt occurs	6	
Process P0 is running	1	
Control is passed to P1	4	

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: context of P0 is saved in P0's PCB → 3, context of P1 is loaded from P1's PCB → 4, Process P1 is running → 6, timer interrupt occurs → 2, Process P0 is running → 1, Control is passed to P1 → 5

Question 33

Not answered

Marked out of 1.00

Select the correct statements about interrupt handling in xv6 code

- a. On any interrupt/syscall/exception the control first jumps in vectors.S
- b. The trapframe pointer in struct proc, points to a location on user stack
- c. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- d. xv6 uses the 64th entry in IDT for system calls
- e. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- f. The trapframe pointer in struct proc, points to a location on kernel stack
- g. The function trap() is called only in case of hardware interrupt
- h. The CS and EIP are changed only immediately on a hardware interrupt
- i. All the 256 entries in the IDT are filled

- j. On any interrupt/syscall/exception the control first jumps in trapasm.S
- k. The function trap() is called irrespective of hardware interrupt/system-call/exception
- l. xv6 uses the 0x64th entry in IDT for system calls
- m. Before going to alltraps, the kernel stack contains upto 5 entries.

Your answer is incorrect.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in trapasm.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

[◀ \(Assignment\) Change free list management in xv6](#)

Jump to...



Started on Tuesday, 22 March 2022, 2:06:46 PM

State Finished

Completed on Tuesday, 22 March 2022, 5:29:11 PM

Time taken 3 hours 22 mins

Grade 23.12 out of 40.00 (58%)

Question 1

Partially correct

Mark 0.50 out of 1.00

Mark whether the given sequence of events is possible or not-possible. Also, select the reason for your answer.

For each sequence it's a not-possible sequence if some important event is not mentioned in the sequence.

Assume that the kernel code is non-interruptible and uniprocessor system.

```
Process P1, user code executing
Timer interrupt
Context changes to kernel context
Generic interrupt handler runs
Generic interrupt handler calls Scheduler
Scheduler selects P2 for execution
After scheduler, Process P2 user code executing
```

This sequence of events is: possible ✗

Because

Process P2 has to return from interrupt context before it's user code executes ✓

Question 2

Correct

Mark 1.00 out of 1.00

The data structure used in kalloc() and kfree() in xv6 is

- a. Double linked NULL terminated list
- b. Singly linked circular list
- c. Singly linked NULL terminated list
- d. Doubly linked circular list



Your answer is correct.

The correct answer is: Singly linked NULL terminated list

Question 3

Incorrect

Mark 0.00 out of 1.00

Given that a kernel has 1000 KB of total memory, and holes of sizes (in that order) 300 KB, 200 KB, 100 KB, 250 KB. For each of the requests on the left side, match it with the chunk chosen using the specified algorithm.

Consider each request as first request.

220 KB, best fit	100 KB	▼	✗
50 KB, worst fit	100 KB	▼	✗
100 KB, worst fit	100 KB	▼	✗
150 KB, best fit	100 KB	▼	✗
200 KB, first fit	100 KB	▼	✗
150 KB, first fit	100 KB	▼	✗

The correct answer is: 220 KB, best fit → 250 KB, 50 KB, worst fit → 300 KB, 100 KB, worst fit → 300 KB, 150 KB, best fit → 200 KB, 200 KB, first fit → 300 KB, 150 KB, first fit → 300 KB

Question 4

Incorrect

Mark 0.00 out of 1.00

Select the most common causes of use of IPC by processes

- a. Sharing of information of common interest
- b. Breaking up a large task into small tasks and speeding up computation, on multiple core machines ✓
- c. More modular code
- d. More security checks
- e. Get the kernel performance statistics ✗

The correct answers are: Sharing of information of common interest, Breaking up a large task into small tasks and speeding up computation, on multiple core machines, More modular code

Question 5

Partially correct

Mark 0.33 out of 1.00

Select all correct statements w.r.t. Major and Minor page faults on Linux

- a. Major page faults are likely to occur in more numbers at the beginning of the process
- b. Minor page fault may occur because of a page fault during fork(), on code of an already running process
- c. Minor page faults are an improvement of the page buffering techniques
- d. Minor page fault may occur because the page was a shared memory page
- e. Minor page fault may occur because the page was freed, but still tagged and available in the free page list ✓
- f. Thrashing is possible only due to major page faults ✓

The correct answers are: Minor page fault may occur because the page was a shared memory page, Minor page fault may occur because of a page fault during fork(), on code of an already running process, Minor page fault may occur because the page was freed, but still tagged and available in the free page list, Major page faults are likely to occur in more numbers at the beginning of the process, Thrashing is possible only due to major page faults, Minor page faults are an improvement of the page buffering techniques

Question 6

Correct

Mark 1.00 out of 1.00

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

10011010

Now, there is a request for a chunk of 50 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: ✓

The correct answer is: 11111010

Question 7

Partially correct

Mark 0.67 out of 1.00

W.r.t. xv6 code, match the state of a process with a code that sets the state

SLEEPING	sleep(), called by any process blocking itself	❖ ✓
EMBRYO	fork()->allocproc() before setting up the UVM	❖ ✓
UNUSED	exit(), called by an interrupt handler	❖ ✗
ZOMBIE	exit(), called by process itself	❖ ✓
RUNNING	fork()->allocproc() before setting up the UVM	❖ ✗
RUNNABLE	wakeup(), called by an interrupt handler	❖ ✓

The correct answer is: SLEEPING → sleep(), called by any process blocking itself, EMBRYO → fork()->allocproc() before setting up the UVM, UNUSED → wait(), called by parent process, ZOMBIE → exit(), called by process itself, RUNNING → scheduler(), RUNNABLE → wakeup(), called by an interrupt handler

Question 8

Partially correct

Mark 0.67 out of 1.00

For each function/code-point, select the status of segmentation setup in xv6

entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	❖ ✓
bootmain()	gdt setup with 3 entries, right from first line of code of bootloader	❖ ✗
after seginit() in main()	gdt setup with 5 entries (0 to 4) on one processor	❖ ✓
kvmalloc() in main()	gdt setup with 3 entries, at start32 symbol of bootasm.S	❖ ✓
bootasm.S	gdt setup with 3 entries, right from first line of code of bootloader	❖ ✗
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	❖ ✓

Your answer is partially correct.

You have correctly selected 4.

The correct answer is: entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors

Question 9

Correct

Mark 2.00 out of 2.00

For the reference string

3 4 3 5 2

using FIFO replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.

Select the correct statement.

Select one:

- a. Exhibit Balady's anomaly between 3 and 4 frames
- b. Do not exhibit Balady's anomaly
- c. Exhibit Balady's anomaly between 2 and 3 frames



Your answer is correct.

The correct answer is: Do not exhibit Balady's anomaly

Question 10

Partially correct

Mark 0.60 out of 1.00

Choice of the global or local replacement strategy is a subjective choice for kernel programmers. There are advantages and disadvantages on either side. Out of the following statements, that advocate either global or local replacement strategy, select those statements that have a logically CONSISTENT argument. (That is any statement that is logically correct about either global or local replacement)

Consistent Inconsistent

<input checked="" type="radio"/>	<input checked="" type="radio"/>	Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Global replacement may give highly variable per process completion time because number of page faults become unpredictable.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.	<input checked="" type="checkbox"/>

Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).: Consistent

Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.: Consistent

Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.: Consistent

Global replacement may give highly variable per process completion time because number of page faults become un-predictable.: Consistent

Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.: Consistent

Question 11

Partially correct

Mark 0.60 out of 1.00

Mark the statements as True or False, w.r.t. thrashing

True**False**

<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	mmap() solves the problem of thrashing.	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	Thrashing can be limited if local replacement is used.	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	During thrashing the CPU is under-utilised as most time is spent in I/O	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	Thrashing is particular to demand paging systems, and does not apply to pure paging systems.	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	Thrashing occurs when the total size of all processes's locality exceeds total memory size.	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	Thrashing can occur even if entire memory is not in use.	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	The working set model is an attempt at approximating the locality of a process.	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input checked="" type="radio"/>	<input checked="" type="radio"/> <input checked="" type="radio"/>	Thrashing occurs because some process is doing lot of disk I/O.	<input checked="" type="checkbox"/>	

mmap() solves the problem of thrashing.: False

Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.: True

Thrashing can be limited if local replacement is used.: True

During thrashing the CPU is under-utilised as most time is spent in I/O: True

Thrashing is particular to demand paging systems, and does not apply to pure paging systems.: True

Thrashing occurs when the total size of all processes's locality exceeds total memory size.: True

Thrashing can occur even if entire memory is not in use.: False

Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.: False

The working set model is an attempt at approximating the locality of a process.: True

Thrashing occurs because some process is doing lot of disk I/O.: False

Question 12

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about signals

Select one or more:

- a. Signals are delivered to a process by kernel ✓
- b. A signal handler can be invoked asynchronously or synchronously depending on signal type
- c. Signals are delivered to a process by another process
- d. The signal handler code runs in user mode of CPU
- e. SIGKILL definitely kills a process because it can't be caught or ignored, and its default action terminates the process ✓
- f. Signal handlers once replaced can't be restored
- g. The signal handler code runs in kernel mode of CPU ✗
- h. SIGKILL definitely kills a process because its code runs in kernel mode of CPU ✗

Your answer is incorrect.

The correct answers are: Signals are delivered to a process by kernel, A signal handler can be invoked asynchronously or synchronously depending on signal type, The signal handler code runs in user mode of CPU, SIGKILL definitely kills a process because it can't be caught or ignored, and its default action terminates the process

Question 13

Incorrect

Mark 0.00 out of 1.00

Select the correct statements about interrupt handling in xv6 code

- a. On any interrupt/syscall/exception the control first jumps in vectors.S ✓
- b. Before going to alltraps, the kernel stack contains upto 5 entries. ✗
- c. xv6 uses the 0x64th entry in IDT for system calls ✗
- d. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt ✓
- e. The CS and EIP are changed only immediately on a hardware interrupt ✗
- f. xv6 uses the 64th entry in IDT for system calls ✗
- g. All the 256 entries in the IDT are filled ✗
- h. The function trap() is the called only in case of hardware interrupt ✗
- i. The trapframe pointer in struct proc, points to a location on kernel stack ✗
- j. The CS and EIP are changed only after pushing user code's SS,ESP on stack ✗
- k. The function trap() is the called irrespective of hardware interrupt/system-call/exception ✗
- l. On any interrupt/syscall/exception the control first jumps in trapasm.S ✗
- m. The trapframe pointer in struct proc, points to a location on user stack ✗

Your answer is incorrect.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is the called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question 14

Correct

Mark 1.00 out of 1.00

For the reference string

3 4 3 5 2

using LRU replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.

Select the most correct statement.

Select one:

- a. Exhibit Balady's anomaly between 2 and 3 frames
- b. Exhibit Balady's anomaly between 3 and 4 frames
- c. LRU will never exhibit Balady's anomaly
- d. This example does not exhibit Balady's anomaly



Your answer is correct.

The correct answer is: LRU will never exhibit Balady's anomaly

Question 15

Partially correct

Mark 0.10 out of 1.00

Select all the correct statements about process states.

Note that in this question you lose marks for every incorrect choice that you make, proportional to actual number of incorrect choices.

- a. Process state is changed only by interrupt handlers
- b. The scheduler can change state of a process from RUNNABLE to RUNNING
- c. A process becomes ZOMBIE when it calls exit() ✓
- d. Process state is stored in the processor
- e. A process becomes ZOMBIE when another process bites into its memory
- f. The scheduler can change state of a process from RUNNABLE to RUNNING and vice-versa ✗
- g. Process state can be implemented as just a number
- h. Process state is stored in the PCB ✓
- i. Process state is implemented as a string ✗

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: Process state is stored in the PCB, Process state can be implemented as just a number, The scheduler can change state of a process from RUNNABLE to RUNNING, A process becomes ZOMBIE when it calls exit()

Question 16

Correct

Mark 1.00 out of 1.00

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- a. end, P2V(PHYSTOP) ✓
- b. end, (4MB + PHYSTOP)
- c. end, P2V(4MB + PHYSTOP)
- d. P2V(end), PHYSTOP
- e. end, 4MB
- f. P2V(end), P2V(PHYSTOP)
- g. end, PHYSTOP

Your answer is correct.

The correct answer is: end, P2V(PHYSTOP)

Question 17

Incorrect

Mark 0.00 out of 1.00

If one thread opens a file with read privileges then

Select one:

- a. any other thread cannot read from that file
- b. other threads in the another process can also read from that file ✗
- c. none of these
- d. other threads in the same process can also read from that file

Your answer is incorrect.

The correct answer is: other threads in the same process can also read from that file

Question 18

Correct

Mark 1.00 out of 1.00

Map the functionality/use with function/variable in xv6 code.

Setup kernel part of a page table, and switch to that page table

kvmalloc()



mappages()



setupkvm()



walkpgdir()



kmap[]



kalloc()



Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

Array listing the kernel memory mappings, to be used by setupkvm()

return a free page, if available; 0, otherwise

Your answer is correct.

The correct answer is: Setup kernel part of a page table, and switch to that page table → kvmalloc(), Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → mappages(), Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → setupkvm(), Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → walkpgdir(), Array listing the kernel memory mappings, to be used by setupkvm() → kmap[], return a free page, if available; 0, otherwise → kalloc()

Question 19

Partially correct

Mark 0.63 out of 1.00

Consider a demand-paging system with the following time-measured utilizations:

CPU utilization : 20%

Paging disk: 97.7%

Other I/O devices: 5%

For each of the following, indicate whether it will (or is likely to) improve CPU utilization (even if by a small amount). Explain your answers.

a. Install a faster CPU : Yes

b. Install a bigger paging disk. : Yes

c. Increase the degree of multiprogramming. : Yes

d. Decrease the degree of multiprogramming. : Yes

e. Install more main memory.: Yes

f. Install a faster hard disk or multiple controllers with multiple hard disks. : Yes

g. Add prepaging to the page-fetch algorithms. :

Yes

h. Increase the page size. : Yes

Question 20

Not answered

Marked out of 1.00

Given below is a sequence of reference bits on pages before the second chance algorithm runs. Before the algorithm runs, the counter is at the page marked (x). Write the sequence of reference bits after the second chance algorithm has executed once. In the answer write PRECISELY one space BETWEEN each number and do not mention (x).

0 0 1(x) 1 0 1 1

Answer:



The correct answer is: 0 0 0 0 0 1 1

Question 21

Correct

Mark 1.00 out of 1.00

Consider a computer system with a 32-bit logical address and 4- KB page size. The system supports up to 512 MB of physical memory. How many entries are there in each of the following?

Write answer as a decimal number.

A conventional, single-level page table:

1048576



An inverted page table:

131072



Question 22

Partially correct

Mark 0.38 out of 1.00

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

True**False**

<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	String arguments are first copied to trapframe and then from trapframe to kernel's other variables.	✓	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The arguments to system call originally reside on process stack.	✗	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer	✗	
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The arguments to system call are copied to kernel stack in trapasm.S	✓	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The arguments are accessed in the kernel code using esp on the trapframe.	✗	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	Integer arguments are copied from user memory to kernel memory using argint()	✓	
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The functions like argint(), argstr() make the system call arguments available in the kernel.	✗	
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	Integer arguments are stored in eax, ebx, ecx, etc. registers	✗	

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False

The arguments to system call originally reside on process stack.: True

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

The arguments to system call are copied to kernel stack in trapasm.S: False

The arguments are accessed in the kernel code using esp on the trapframe.: True

Integer arguments are copied from user memory to kernel memory using argint(): True

The functions like argint(), argstr() make the system call arguments available in the kernel.: True

Integer arguments are stored in eax, ebx, ecx, etc. registers: False

Question 23

Partially correct

Mark 0.60 out of 1.00

Select all the correct statements w.r.t user and kernel threads

Select one or more:

- a. one-one model can be implemented even if there are no kernel threads
- b. one-one model increases kernel's scheduling load
- c. A process blocks in many-one model even if a single thread makes a blocking system call ✓
- d. A process may not block in many-one model, if a thread makes a blocking system call
- e. many-one model can be implemented even if there are no kernel threads
- f. all three models, that is many-one, one-one, many-many , require a user level thread library ✓
- g. many-one model gives no speedup on multicore processors ✓

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: many-one model can be implemented even if there are no kernel threads, all three models, that is many-one, one-one, many-many , require a user level thread library, one-one model increases kernel's scheduling load, many-one model gives no speedup on multicore processors, A process blocks in many-one model even if a single thread makes a blocking system call

Question 24

Correct

Mark 1.00 out of 1.00

After virtual memory is implemented

(select T/F for each of the following) One Program's size can be larger than physical memory size

True**False**

<input checked="" type="radio"/>	<input type="radio"/> X	Cumulative size of all programs can be larger than physical memory size	<input checked="" type="checkbox"/>	
<input type="radio"/> X	<input checked="" type="radio"/>	Virtual addresses become available to executing process	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/>	<input type="radio"/> X	Code need not be completely in memory	<input checked="" type="checkbox"/>	
<input type="radio"/> X	<input checked="" type="radio"/>	Virtual access to memory is granted to all processes	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/>	<input type="radio"/> X	Relatively less I/O may be possible during process execution	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/>	<input type="radio"/> X	Logical address space could be larger than physical address space	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/>	<input type="radio"/> X	One Program's size can be larger than physical memory size	<input checked="" type="checkbox"/>	

Cumulative size of all programs can be larger than physical memory size: True

Virtual addresses become available to executing process: False

Code need not be completely in memory: True

Virtual access to memory is granted to all processes: False

Relatively less I/O may be possible during process execution: True

Logical address space could be larger than physical address space: True

One Program's size can be larger than physical memory size: True

Question 25

Partially correct

Mark 0.50 out of 1.00

Select the correct points of comparison between POSIX and System V shared memory.

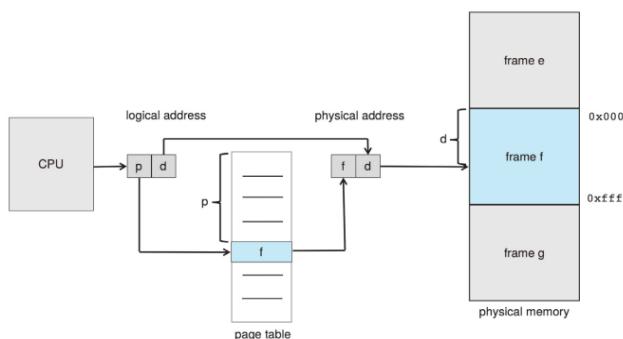
- a. POSIX shared memory is newer than System V shared memory
- b. POSIX allows giving name to shared memory, System V does not
- c. System V is more prevalent than POSIX even today
- d. POSIX shared memory is "thread safe", System V is not

The correct answers are: POSIX shared memory is newer than System V shared memory, POSIX shared memory is "thread safe", System V is not, POSIX allows giving name to shared memory, System V does not, System V is more prevalent than POSIX even today

Question 26

Partially correct

Mark 0.50 out of 1.00

**Figure 9.8** Paging hardware.

Mark the statements as True or False, w.r.t. the above diagram (note that the diagram does not cover all details of what actually happens!)

True False

<input checked="" type="radio"/>	<input type="radio"/>	The logical address issued by CPU is the same one generated by compiler	✓	
<input checked="" type="radio"/>	<input type="radio"/>	The page table is in physical memory and must be continuous	✓	
<input checked="" type="radio"/>	<input type="radio"/>	The combining of f and d is done by MMU	✗	
<input type="radio"/>	<input checked="" type="radio"/>	Using the offset d in the physical page-frame is done by MMU	✓	
<input checked="" type="radio"/>	<input type="radio"/>	The split of logical address into p and d is done by MMU	✗	
<input type="radio"/>	<input checked="" type="radio"/>	There are total 3 memory references in this diagram	✗	

The logical address issued by CPU is the same one generated by compiler: True

The page table is in physical memory and must be continuous: True

The combining of f and d is done by MMU: True

Using the offset d in the physical page-frame is done by MMU: False

The split of logical address into p and d is done by MMU: True

There are total 3 memory references in this diagram: False

Question 27

Partially correct

Mark 0.60 out of 1.00

Mark the statements about named and un-named pipes as True or False

True**False**

<input checked="" type="radio"/> ✗	<input type="checkbox"/> ✗	The pipe() system call can be used to create either a named or un-named pipe.	<input checked="" type="checkbox"/> ✗	
<input checked="" type="radio"/> ✗	<input type="checkbox"/> ✗	Named pipes can be used for communication between only "related" processes.	<input checked="" type="checkbox"/> ✗	
<input checked="" type="checkbox"/> ✗	<input checked="" type="checkbox"/> ✗	Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.	<input checked="" type="checkbox"/> ✗	
<input checked="" type="radio"/> ✗	<input type="checkbox"/> ✗	A named pipe has a name decided by the kernel.	<input checked="" type="checkbox"/> ✗	
<input checked="" type="checkbox"/> ✗	<input checked="" type="checkbox"/> ✗	Both types of pipes are an extension of the idea of "message passing".	<input checked="" type="checkbox"/> ✗	
<input checked="" type="checkbox"/> ✗	<input checked="" type="checkbox"/> ✗	Named pipe exists as a file	<input checked="" type="checkbox"/> ✗	
<input checked="" type="checkbox"/> ✗	<input checked="" type="checkbox"/> ✗	Both types of pipes provide FIFO communication.	<input checked="" type="checkbox"/> ✗	
<input checked="" type="checkbox"/> ✗	<input checked="" type="checkbox"/> ✗	Un-named pipes are inherited by a child process from parent.	<input checked="" type="checkbox"/> ✗	
<input checked="" type="checkbox"/> ✗	<input checked="" type="checkbox"/> ✗	Named pipes can exist beyond the life-time of processes using them.	<input checked="" type="checkbox"/> ✗	
<input checked="" type="radio"/> ✗	<input type="checkbox"/> ✗	The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.	<input checked="" type="checkbox"/> ✗	

The pipe() system call can be used to create either a named or un-named pipe.: False

Named pipes can be used for communication between only "related" processes.: False

Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.: True

A named pipe has a name decided by the kernel.: False

Both types of pipes are an extension of the idea of "message passing": True

Named pipe exists as a file: True

Both types of pipes provide FIFO communication.: True

Un-named pipes are inherited by a child process from parent.: True

Named pipes can exist beyond the life-time of processes using them.: True

The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.: False

Question 28

Partially correct

Mark 0.20 out of 2.00

Order the following events, in the creation of init() process in xv6:

1. userinit() is called
2. trapframe and context pointers are set to proper location
3. page table mappings of 'initcode' are replaced by makpings of 'init'
4. initcode process runs
5. code is set to start in forkret() when process gets scheduled
6. initcode calls exec system call
7. initcode is selected by scheduler for execution
8. Stack is allocated for "/init" process
9. Arguments on setup on process stack for /init
10. sys_exec runs
11. the header of "/init" ELF file is ready by kernel
12. trap() runs
13. kernel stack is allocated for initcode process
14. function pointer from syscalls[] array is invoked
15. kernel memory mappings are created for initcode
16. initcode process is set to be runnable
17. empty struct proc is obtained for initcode
18. memory mappings are created for "/init" process
19. name of process "/init" is copied in struct proc
20. values are set in the trapframe of initcode

Your answer is partially correct.

Grading type: Relative to the next item (including last)

Grade details: 2 / 20 = 10%

Here are the scores for each item in this response:

1. 0 / 1 = 0%
2. 0 / 1 = 0%
3. 0 / 1 = 0%
4. 0 / 1 = 0%
5. 0 / 1 = 0%
6. 0 / 1 = 0%
7. 0 / 1 = 0%
8. 1 / 1 = 100%
9. 0 / 1 = 0%
10. 1 / 1 = 100%
11. 0 / 1 = 0%
12. 0 / 1 = 0%
13. 0 / 1 = 0%

- 14. $0 / 1 = 0\%$
- 15. $0 / 1 = 0\%$
- 16. $0 / 1 = 0\%$
- 17. $0 / 1 = 0\%$
- 18. $0 / 1 = 0\%$
- 19. $0 / 1 = 0\%$
- 20. $0 / 1 = 0\%$

The correct order for these items is as follows:

1. userinit() is called
2. empty struct proc is obtained for initcode
3. kernel stack is allocated for initcode process
4. trapframe and context pointers are set to proper location
5. code is set to start in forkret() when process gets scheduled
6. kernel memory mappings are created for initcode
7. values are set in the trapframe of initcode
8. initcode process is set to be runnable
9. initcode is selected by scheduler for execution
10. initcode process runs
11. initcode calls exec system call
12. trap() runs
13. function pointer from syscalls[] array is invoked
14. sys_exec runs
15. the header of "/init" ELF file is ready by kernel
16. memory mappings are created for "/init" process
17. Stack is allocated for "/init" process
18. Arguments on setup on process stack for /init
19. name of process "/init" is copied in struct proc
20. page table mappings of 'initcode' are replaced by makpings of 'init'

Question 29

Correct

Mark 2.00 out of 2.00

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using FIFO replacement is:

Answer: 10



#6# 6,4# 6,4,2 #0,4,2# 0,1,2 #0,1,6 #9,1,6# 9,2,6# 9,2,0 #5,2,0

The correct answer is: 10

Question 30

Correct

Mark 2.00 out of 2.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only Mark statements True or False

True False

<input checked="" type="radio"/>	<input type="radio"/> X	xv6 uses physical memory upto 224 MB only	✓	
<input checked="" type="radio"/>	<input type="radio"/> X	The process's address space gets mapped on frames, obtained from ~2MB:224MB range	✓	
<input checked="" type="radio"/>	<input type="radio"/> X	The stack allocated in entry.S is used as stack for scheduler's context for first processor	✓	
<input type="radio"/> X	<input checked="" type="radio"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context	✓	
<input checked="" type="radio"/>	<input type="radio"/> X	The free page-frame are created out of nearly 222 MB	✓	
<input type="radio"/> X	<input checked="" type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process	✓	
<input checked="" type="radio"/>	<input type="radio"/> X	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context	✓	
<input type="radio"/> X	<input checked="" type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context	✓	
<input checked="" type="radio"/>	<input type="radio"/> X	PHYSTOP can be increased to some extent, simply by editing memlayout.h	✓	
<input checked="" type="radio"/>	<input type="radio"/> X	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir	✓	
<input checked="" type="radio"/>	<input type="radio"/> X	The kernel code and data take up less than 2 MB space	✓	

xv6 uses physical memory upto 224 MB only: True

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

The free page-frame are created out of nearly 222 MB: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

The kernel code and data take up less than 2 MB space: True

Question 31

Partially correct

Mark 0.60 out of 1.00

Mark statements True/False w.r.t. change of states of a process. Note that a statement is true only if the claim and argument both are true.

Reference: The process state diagram (and your understanding of how kernel code works). Note - the diagram does not show zombie state!

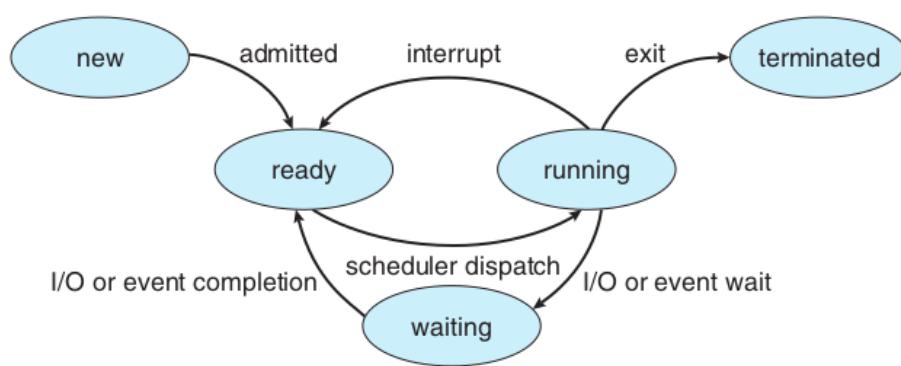


Figure 3.2 Diagram of process state.

True False

<input checked="" type="radio"/>	<input type="radio"/>	Only a process in READY state is considered by scheduler	<input checked="" type="checkbox"/>	
<input type="radio"/>	<input checked="" type="radio"/>	A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first	<input type="checkbox"/>	
<input checked="" type="radio"/>	<input type="radio"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/>	<input type="radio"/>	Every forked process has to go through ZOMBIE state, at least for a small duration.	<input checked="" type="checkbox"/>	
<input type="radio"/>	<input checked="" type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.	<input type="checkbox"/>	

Only a process in READY state is considered by scheduler: True

A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first: False

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet: True

Every forked process has to go through ZOMBIE state, at least for a small duration.: True

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

Question 32

Partially correct

Mark 0.89 out of 2.00

Match the description of a memory management function with the name of the function that provides it, in xv6

Copy the code pages of a process	copyuvvm()	✗
Setup and load the user page table for initcode process	setupkvm()	✗
Load contents from ELF into existing pages	loaduvvm()	✓
Create a copy of the page table of a process	clearpteul()	✗
Switch to kernel page table	switchuvvm()	✗
Mark the page as in-accessible	clearpteul()	✓
Switch to user page table	switchuvvm()	✓
setup the kernel part in the page table	setupkvm()	✓
Load contents from ELF into pages after allocating the pages first	loaduvvm()	✗

The correct answer is: Copy the code pages of a process → No such function, Setup and load the user page table for initcode process → inituvvm(), Load contents from ELF into existing pages → loaduvvm(), Create a copy of the page table of a process → copyuvvm(), Switch to kernel page table → switchkvm(), Mark the page as in-accessible → clearpteul(), Switch to user page table → switchuvvm(), setup the kernel part in the page table → setupkvm(), Load contents from ELF into pages after allocating the pages first → No such function

Question 33

Correct

Mark 1.00 out of 1.00

Select all the correct statements about MMU and its functionality (on a non-demand paged system)

Select one or more:

- a. The operating system interacts with MMU for every single address translation
- b. Illegal memory access is detected by operating system
- c. MMU is a separate chip outside the processor
- d. The Operating system sets up relevant CPU registers to enable proper MMU translations ✓
- e. MMU is inside the processor ✓
- f. Illegal memory access is detected in hardware by MMU and a trap is raised ✓
- g. Logical to physical address translations in MMU are done with specific machine instructions
- h. Logical to physical address translations in MMU are done in hardware, automatically ✓

Your answer is correct.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

Question 34

Partially correct

Mark 0.67 out of 1.00

Mark the statements as True or False, w.r.t. mmap()

True**False**

<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/> <input type="checkbox"/>	MMap_FIXED guarantees that the mapping is always done at the specified address	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	mmap() results in changes to page table of a process.	<input checked="" type="checkbox"/>	
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	mmap() can be implemented on both demand paged and non-demand paged systems.	<input checked="" type="checkbox"/>	
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/> <input type="checkbox"/>	mmap() results in changes to buffer-cache of the kernel.	<input checked="" type="checkbox"/>	
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/> <input type="checkbox"/>	on failure mmap() returns NULL	<input checked="" type="checkbox"/>	
<input type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/> <input type="checkbox"/>	MAP_SHARED leads to a mapping that is copy-on-write	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	MAP_PRIVATE leads to a mapping that is copy-on-write	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	mmap() is a system call	<input checked="" type="checkbox"/>	
<input checked="" type="radio"/> <input type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	on failure mmap() returns (void *)-1	<input checked="" type="checkbox"/>	

MAP_FIXED guarantees that the mapping is always done at the specified address.: False

mmap() results in changes to page table of a process.: True

mmap() can be implemented on both demand paged and non-demand paged systems.: True

mmap() results in changes to buffer-cache of the kernel.: False

on failure mmap() returns NULL: False

MAP_SHARED leads to a mapping that is copy-on-write: False

MAP_PRIVATE leads to a mapping that is copy-on-write: True

mmap() is a system call: True

on failure mmap() returns (void *)-1: True

Question 35

Partially correct

Mark 0.10 out of 1.00

Select all the correct statements about linking and loading.

Select one or more:

- a. Dynamic linking is possible with continuous memory management, but variable sized partitions only. ✗
- b. Static linking leads to non-relocatable code
- c. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently)
- d. Continuous memory management schemes can support dynamic linking and dynamic loading. ✗
- e. Dynamic linking essentially results in relocatable code. ✓
- f. Loader is last stage of the linker program
- g. Continuous memory management schemes can support static linking and static loading. (may be inefficiently) ✓
- h. Dynamic linking and loading is not possible without demand paging or demand segmentation.
- i. Loader is part of the operating system ✓

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

[◀ \(Optional Assignment\) Iseek system call in xv6](#)

Jump to...



[Feedback on Quiz-2 ►](#)

Started on Monday, 7 March 2022, 7:00:46 PM

State Finished

Completed on Monday, 7 March 2022, 8:20:00 PM

Time taken 1 hour 19 mins

Grade 11.42 out of 15.00 (76%)

Question 1

Complete

Mark 1.00 out of 1.00

Map the virtual address to physical address in xv6

KERNBASE	0	▼
KERNLINK	0x100000	▼
80108000	0x108000	▼
0xFE000000	0xFE000000	▼

The correct answer is: KERNBASE → 0, KERNLINK → 0x100000, 80108000 → 0x108000, 0xFE000000 → 0xFE000000

Question 2

Complete

Mark 1.00 out of 1.00

Does exec() code around clearptau() lead to wastage of one page frame?

- a. no
- b. yes

The correct answer is: yes

Question 3

Complete

Mark 1.00 out of 1.00

What does seginit() do?

- a. Nothing significant, just repetition of earlier GDT setup but with kernel page table allocated now
- b. Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 0
- c. Nothing significant, just repetition of earlier GDT setup but with 2-level paging setup done
- d. Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 3
- e. Nothing significant, just repetition of earlier GDT setup but with free frames list created now

The correct answer is: Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 3

Question 4

Complete

Mark 1.50 out of 1.50

Arrange the following in the correct order of execution (w.r.t. 'init')

initcode() returns from trapret()	7	◆
initcode() calls exec("/init", ...)	8	◆
'initcode' process is marked RUNNABLE	3	◆
userinit() is called	1	◆
scheduler() schedules initcode() process	5	◆
'initcode' struct proc is created	2	◆
mpmain() calls scheduler()	4	◆
initcode() returns in forkret()	6	◆

The correct answer is: initcode() returns from trapret() → 7, initcode() calls exec("/init", ...) → 8, 'initcode' process is marked RUNNABLE → 3, userinit() is called → 1, scheduler() schedules initcode() process → 5, 'initcode' struct proc is created → 2, mpmain() calls scheduler() → 4, initcode() returns in forkret() → 6

Question 5

Complete

Mark 1.00 out of 1.00

Why is there a call to kinit2? Why is it not merged with knit1?

- a. knit2 refers to virtual addresses beyond 4MB, which are not mapped before kvalloc() is called
- b. call to seginit() makes it possible to actually use PHYSTOP in argument to kinit2()
- c. When kinit1() is called there is a need for few page frames, but later knit2() is called to serve need of more page frames
- d. Because there is a limit on the values that the arguments to knit1() can take.

The correct answer is: knit2 refers to virtual addresses beyond 4MB, which are not mapped before kvalloc() is called

Question 6

Complete

Mark 0.67 out of 1.00

Which of the following is done by mappages()?

- a. create page table mappings to the range given by "pa" and "pa + size"
- b. allocate page table if required
- c. allocate page frame if required
- d. allocate page directory if required
- e. create page table mappings for the range given by "va" and "va + size"

The correct answers are: create page table mappings for the range given by "va" and "va + size", allocate page table if required, create page table mappings to the range given by "pa" and "pa + size"

Question 7

Complete

Mark 0.08 out of 1.00

Select all the correct statements about initcode

- a. code of initcode is loaded in memory by the kernel during userinit()
- b. The data and stack of initcode is mapped to one single page in userinit()
- c. code of 'initcode' is loaded along with the kernel during booting
- d. initcode essentially calls exec("/init",...)
- e. the size of 'initcode' is 2c
- f. code of initcode is loaded at virtual address 0
- g. initcode is the 'init' process

The correct answers are: code of 'initcode' is loaded along with the kernel during booting, the size of 'initcode' is 2c, The data and stack of initcode is mapped to one single page in userinit(), initcode essentially calls exec("/init",...)

Question 8

Complete

Mark 1.00 out of 1.00

The variable 'end' used as argument to kinit1 has the value

- a. 8010a48c
- b. 80110000
- c. 80102da0
- d. 81000000
- e. 80000000
- f. 801154a8

The correct answer is: 801154a8

Question 9

Complete

Mark 1.00 out of 1.00

What does userinit() do ?

- a. initializes the users
- b. sets up the 'init' process to start execution in forkret()
- c. sets up the 'initcode' process to start execution in forkret ()
- d. sets up the 'initcode' process to start execution in trapret()
- e. initializes the process 'init' and starts executing it
- f. sets up the 'initcode' process to start execution in forkret()

The correct answer is: sets up the 'initcode' process to start execution in forkret()

Question 10

Complete

Mark 1.50 out of 1.50

Which of the following is DONE by allocproc() ?

- a. allocate PID to the process
- b. setup the trapframe and context pointers appropriately
- c. setup the contents of the trapframe of the process properly
- d. ensure that the process starts in trapret()
- e. setup kernel memory mappings for the process
- f. ensure that the process starts in forkret()
- g. allocate kernel stack for the process
- h. Select an UNUSED struct proc for use

The correct answers are: Select an UNUSED struct proc for use, allocate PID to the process, allocate kernel stack for the process, setup the trapframe and context pointers appropriately, ensure that the process starts in forkret()

Question 11

Complete

Mark 0.67 out of 2.00

exec() does this: curproc->tf->eip = elf.entry, but userinit() does this: p->tf->eip = 0; Select all the statements from below, that collectively explain this

- a. elf.entry is anyways 0, so both statements mean the same
- b. the initcode is created using objcopy, which discards all relocation information and symbols (like entry)
- c. the 'entry' in initcode is anyways 0
- d. exec() loads from ELF file and the address of first instruction to be executed is given by 'entry'
- e. the code of 'initcode' is loaded at physical address 0
- f. In userinit() the function inituvm() has mapped the code of 'initcode' to be starting at virtual address 0

The correct answers are: exec() loads from ELF file and the address of first instruction to be executed is given by 'entry', In userinit() the function inituvm() has mapped the code of 'initcode' to be starting at virtual address 0, the initcode is created using objcopy, which discards all relocation information and symbols (like entry)

Question 12

Complete

Mark 0.00 out of 1.00

The approximate number of page frames created by kinit1 is

- a. 10
- b. 2000
- c. 3000
- d. 1000
- e. 4
- f. 16
- g. 4000

The correct answer is: 3000

Question 13

Complete

Mark 1.00 out of 1.00

Select the statement that most correctly describes what setupkvm() does

- a. creates a 1-level page table for the use by the kernel, as specified in kmap[] global array
- b. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array
- c. creates a 2-level page table for the use of the kernel, as specified in gdtdesc
- d. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global arrray and makes kpgdir point to it

The correct answer is: creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array

[◀ Questions for test on kalloc/kfree/kvmalloc, etc.](#)

Jump to...



(Optional Assignment) Slab allocator in xv6 ►

Started on Saturday, 26 February 2022, 5:20:40 PM

State Finished

Completed on Saturday, 26 February 2022, 6:35:00 PM

Time taken 1 hour 14 mins

Grade 9.05 out of 15.00 (60%)

Question 1

Complete

Mark 0.75 out of 1.00

which of the following, do you think, are valid concerns for making the kernel pageable?

- a. No data structure of kernel should be pageable
- b. The kernel's own page tables should not be pageable
- c. No part of kernel code should be pageable.
- d. The disk driver and disk interrupt handler should not be pageable
- e. The page fault handler should not be pageable
- f. The kernel must have some dedicated frames for it's own work

The correct answers are: The kernel's own page tables should not be pageable, The page fault handler should not be pageable, The kernel must have some dedicated frames for it's own work, The disk driver and disk interrupt handler should not be pageable

Question 2

Complete

Mark 0.67 out of 1.00

Shared memory is possible with which of the following memory management schemes ?

Select one or more:

- a. demand paging
- b. continuous memory management
- c. segmentation
- d. paging

The correct answers are: paging, segmentation, demand paging

Question 3

Complete

Mark 0.29 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- a. With demand paging, it's possible to have user programs bigger than physical memory.
- b. Both demand paging and paging support shared memory pages.
- c. Demand paging requires additional hardware support, compared to paging.
- d. Paging requires NO hardware support in CPU
- e. The meaning of valid-invalid bit in page table is different in paging and demand-paging.
- f. Paging requires some hardware support in CPU
- g. TLB hit ration has zero impact in effective memory access time in demand paging.
- h. With paging, it's possible to have user programs bigger than physical memory.
- i. Calculations of number of bits for page number and offset are same in paging and demand paging.
- j. Demand paging always increases effective memory access time.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question 4

Complete

Mark 0.00 out of 1.00

Select all the correct statements, w.r.t. Copy on Write

- a. use of COW during fork() is useless if exec() is called by the child
- b. Vfork() assumes that there will be no write, but rather exec()
- c. use of COW during fork() is useless if child called exit()
- d. Fork() used COW technique to improve performance of new process creation.
- e. If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated
- f. COW helps us save memory

The correct answers are: Fork() used COW technique to improve performance of new process creation., If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated, COW helps us save memory, Vfork() assumes that there will be no write, but rather exec()

Question 5

Complete

Mark 0.14 out of 1.00

Suppose two processes share a library between them. The library consists of 5 pages, and these 5 pages are mapped to frames 9, 15, 23, 4, 7 respectively. Process P1 has got 6 pages, first 3 of which consist of process's own code/data and 3 correspond to library's pages 0, 2, 4. Process P2 has got 7 pages, first 3 of which consist of process's own code/data and remaining 4 correspond to library's pages 0, 1, 3, 4. Fill in the blanks for page table entries of P1 and P2.

Page table of P1, Page 5	23	▼
Page table of P2, Page 1	5	▼
Page table of P2, Page 0	6	▼
Page table of P1, Page 3	9	▼
Page table of P1, Page 4	9	▼
Page table of P2, Page 3	7	▼
Page table of P2, Page 4	9	▼

The correct answer is: Page table of P1, Page 5 → 7, Page table of P2, Page 1 → 15, Page table of P2, Page 0 → 9, Page table of P1, Page 3 → 9, Page table of P1, Page 4 → 23, Page table of P2, Page 3 → 4, Page table of P2, Page 4 → 7

Question 6

Complete

Mark 1.00 out of 1.00

Page sizes are a power of 2 because

Select one:

- a. MMU only understands numbers that are power of 2
- b. Power of 2 calculations are highly efficient
- c. Certain bits are reserved for offset in logical address. Hence page size = $2^{(\text{no.of offset bits})}$
- d. operating system calculations happen using power of 2
- e. Certain bits are reserved for offset in logical address. Hence page size = $2^{(32 - \text{no.of offset bits})}$

The correct answer is: Certain bits are reserved for offset in logical address. Hence page size = $2^{(\text{no.of offset bits})}$

Question 7

Complete

Mark 0.60 out of 1.00

Given below is the "maps" file for a particular instance of "vim.basic" process.

Mark the given statements as True or False, w.r.t. the contents of the map file.

```

55a43501b000-55a435049000 r--p 00000000 103:05 917529          /usr/bin/vim.basic
55a435049000-55a435248000 r-xp 0002e000 103:05 917529          /usr/bin/vim.basic
55a435248000-55a4352b6000 r--p 0022d000 103:05 917529          /usr/bin/vim.basic
55a4352b7000-55a4352c5000 r--p 0029b000 103:05 917529          /usr/bin/vim.basic
55a4352c5000-55a4352e2000 rw-p 002a9000 103:05 917529          /usr/bin/vim.basic
55a4352e2000-55a4352f0000 rw-p 00000000 00:00 0                  [heap]
55a436bc9000-55a436e5b000 rw-p 00000000 00:00 0
7f275b0a3000-7f275b0a6000 r--p 00000000 103:05 917901          /usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so
7f275b0a6000-7f275b0ad000 r-xp 00003000 103:05 917901          /usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so
7f275b0ad000-7f275b0af000 r--p 0000a000 103:05 917901          /usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so
7f275b0af000-7f275b0b0000 r--p 0000b000 103:05 917901          /usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so
7f275b0b0000-7f275b0b1000 rw-p 0000c000 103:05 917901          /usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so
7f275b0b1000-7f275b0b7000 rw-p 00000000 00:00 0
7f275b0b7000-7f275b8f5000 r--p 00000000 103:05 925247          /usr/lib/locale/locale-archive
7f275b8f5000-7f275b8fa000 rw-p 00000000 00:00 0
7f275b8fa000-7f275b8fc000 r--p 00000000 103:05 924216          /usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4
7f275b8fc000-7f275b901000 r-xp 00002000 103:05 924216          /usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4
7f275b901000-7f275b904000 r--p 00007000 103:05 924216          /usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4
7f275b904000-7f275b905000 ---p 0000a000 103:05 924216          /usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4
7f275b905000-7f275b906000 r--p 0000a000 103:05 924216          /usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4
7f275b906000-7f275b907000 rw-p 0000b000 103:05 924216          /usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4
7f275b907000-7f275b90a000 r--p 00000000 103:05 924627          /usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8
7f275b90a000-7f275b921000 r-xp 00003000 103:05 924627          /usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8
7f275b921000-7f275b932000 r--p 0001a000 103:05 924627          /usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8
7f275b932000-7f275b933000 ---p 0002b000 103:05 924627          /usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8
7f275b933000-7f275b934000 r--p 0002b000 103:05 924627          /usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8
7f275b934000-7f275b935000 rw-p 0002c000 103:05 924627          /usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8
7f275b935000-7f275b937000 rw-p 00000000 00:00 0
7f275b937000-7f275b938000 r--p 00000000 103:05 917914          /usr/lib/x86_64-linux-
gnu/libutil-2.31.so
7f275b938000-7f275b939000 r-xp 00001000 103:05 917914          /usr/lib/x86_64-linux-
gnu/libutil-2.31.so
7f275b939000-7f275b93a000 r--p 00002000 103:05 917914          /usr/lib/x86_64-linux-
gnu/libutil-2.31.so
7f275b93a000-7f275b93b000 r--p 00002000 103:05 917914          /usr/lib/x86_64-linux-
gnu/libutil-2.31.so
7f275b93b000-7f275b93c000 rw-p 00003000 103:05 917914          /usr/lib/x86_64-linux-
```

```

gnu/libutil-2.31.so
7f275b93c000-7f275b93e000 r--p 00000000 103:05 915906          /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b93e000-7f275b94f000 r-xp 00002000 103:05 915906          /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b94f000-7f275b955000 r--p 00013000 103:05 915906          /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b955000-7f275b956000 ---p 00019000 103:05 915906          /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b956000-7f275b957000 r--p 00019000 103:05 915906          /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b957000-7f275b958000 rw-p 0001a000 103:05 915906          /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b958000-7f275b95c000 r--p 00000000 103:05 923645          /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b95c000-7f275b978000 r-xp 00004000 103:05 923645          /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b978000-7f275b982000 r--p 00020000 103:05 923645          /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b982000-7f275b983000 ---p 0002a000 103:05 923645          /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b983000-7f275b985000 r--p 0002a000 103:05 923645          /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b985000-7f275b986000 rw-p 0002c000 103:05 923645          /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b986000-7f275b988000 r--p 00000000 103:05 924057          /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b988000-7f275b98d000 r-xp 00002000 103:05 924057          /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b98d000-7f275b98f000 r--p 00007000 103:05 924057          /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b98f000-7f275b990000 r--p 00008000 103:05 924057          /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b990000-7f275b991000 rw-p 00009000 103:05 924057          /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b991000-7f275b995000 r--p 00000000 103:05 921934          /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b995000-7f275b9a3000 r-xp 00004000 103:05 921934          /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9a3000-7f275b9a9000 r--p 00012000 103:05 921934          /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9a9000-7f275b9aa000 r--p 00017000 103:05 921934          /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9aa000-7f275b9ab000 rw-p 00018000 103:05 921934          /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9ab000-7f275b9ad000 rw-p 00000000 00:00 0             /usr/lib/x86_64-linux-
7f275b9ad000-7f275b9af000 r--p 00000000 103:05 924631          /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9af000-7f275b9b4000 r-xp 00002000 103:05 924631          /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b4000-7f275b9b5000 r--p 00007000 103:05 924631          /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b5000-7f275b9b6000 ---p 00008000 103:05 924631          /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b6000-7f275b9b7000 r--p 00008000 103:05 924631          /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b7000-7f275b9b8000 rw-p 00009000 103:05 924631          /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b8000-7f275b9ba000 r--p 00000000 103:05 924277          /usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0
7f275b9ba000-7f275ba1e000 r-xp 00002000 103:05 924277          /usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0

```

7f275ba1e000-7f275ba46000 r--p 00066000 103:05 924277	/usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0	
7f275ba46000-7f275ba47000 r--p 0008d000 103:05 924277	/usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0	
7f275ba47000-7f275ba48000 rw-p 0008e000 103:05 924277	/usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0	
7f275ba48000-7f275ba4d000 r--p 00000000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275ba4d000-7f275bbe5000 r-xp 00025000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bbe5000-7f275bc2f000 r--p 0019d000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bc2f000-7f275bc30000 ---p 001e7000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bc30000-7f275bc33000 r--p 001e7000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bc33000-7f275bc36000 rw-p 001ea000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bc36000-7f275bc3a000 rw-p 00000000 00:00 0	
7f275bc3a000-7f275bc41000 r--p 00000000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc41000-7f275bc52000 r-xp 00007000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc52000-7f275bc57000 r--p 00018000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc57000-7f275bc58000 r--p 0001c000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc58000-7f275bc59000 rw-p 0001d000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc59000-7f275bc5d000 rw-p 00000000 00:00 0	
7f275bc5d000-7f275bcce000 r--p 00000000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275bcce000-7f275bf29000 r-xp 00071000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275bf29000-7f275c142000 r--p 002cc000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275c142000-7f275c143000 ---p 004e5000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275c143000-7f275c149000 r--p 004e5000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275c149000-7f275c190000 rw-p 004eb000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275c190000-7f275c1b3000 rw-p 00000000 00:00 0	
7f275c1b3000-7f275c1b4000 r--p 00000000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b4000-7f275c1b6000 r-xp 00001000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b6000-7f275c1b7000 r--p 00003000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b7000-7f275c1b8000 r--p 00003000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b8000-7f275c1b9000 rw-p 00004000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b9000-7f275c1bb000 rw-p 00000000 00:00 0	
7f275c1bb000-7f275c1c0000 r-xp 00000000 103:05 923815	/usr/lib/x86_64-linux-
gnu/libgpm.so.2	
7f275c1c0000-7f275c3bf000 ---p 00005000 103:05 923815	/usr/lib/x86_64-linux-
gnu/libgpm.so.2	
7f275c3bf000-7f275c3c0000 r--p 00004000 103:05 923815	/usr/lib/x86_64-linux-
gnu/libgpm.so.2	
7f275c3c0000-7f275c3c1000 rw-p 00005000 103:05 923815	/usr/lib/x86_64-linux-
gnu/libgpm.so.2	

7f275c3c1000-7f275c3c3000 r--p 00000000 103:05 923315	/usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253	
7f275c3c3000-7f275c3c8000 r-xp 00002000 103:05 923315	/usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253	
7f275c3c8000-7f275c3ca000 r--p 00007000 103:05 923315	/usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253	
7f275c3ca000-7f275c3cb000 r--p 00008000 103:05 923315	/usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253	
7f275c3cb000-7f275c3cc000 rw-p 00009000 103:05 923315	/usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253	
7f275c3cc000-7f275c3cf000 r--p 00000000 103:05 923446	/usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5	
7f275c3cf000-7f275c3d9000 r-xp 00003000 103:05 923446	/usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5	
7f275c3d9000-7f275c3dd000 r--p 0000d000 103:05 923446	/usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5	
7f275c3dd000-7f275c3de000 r--p 00010000 103:05 923446	/usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5	
7f275c3de000-7f275c3df000 rw-p 00011000 103:05 923446	/usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5	
7f275c3df000-7f275c3e5000 r--p 00000000 103:05 924431	/usr/lib/x86_64-linux-
gnu/libselinux.so.1	
7f275c3e5000-7f275c3fe000 r-xp 00006000 103:05 924431	/usr/lib/x86_64-linux-
gnu/libselinux.so.1	
7f275c3fe000-7f275c405000 r--p 0001f000 103:05 924431	/usr/lib/x86_64-linux-
gnu/libselinux.so.1	
7f275c405000-7f275c406000 ---p 00026000 103:05 924431	/usr/lib/x86_64-linux-
gnu/libselinux.so.1	
7f275c406000-7f275c407000 r--p 00026000 103:05 924431	/usr/lib/x86_64-linux-
gnu/libselinux.so.1	
7f275c407000-7f275c408000 rw-p 00027000 103:05 924431	/usr/lib/x86_64-linux-
gnu/libselinux.so.1	
7f275c408000-7f275c40a000 rw-p 00000000 00:00 0	
7f275c40a000-7f275c418000 r--p 00000000 103:05 924540	/usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2	
7f275c418000-7f275c427000 r-xp 0000e000 103:05 924540	/usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2	
7f275c427000-7f275c435000 r--p 0001d000 103:05 924540	/usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2	
7f275c435000-7f275c439000 r--p 0002a000 103:05 924540	/usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2	
7f275c439000-7f275c43a000 rw-p 0002e000 103:05 924540	/usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2	
7f275c43a000-7f275c449000 r--p 00000000 103:05 917895	/usr/lib/x86_64-linux-
gnu/libm-2.31.so	
7f275c449000-7f275c4f0000 r-xp 0000f000 103:05 917895	/usr/lib/x86_64-linux-
gnu/libm-2.31.so	
7f275c4f0000-7f275c587000 r--p 000b6000 103:05 917895	/usr/lib/x86_64-linux-
gnu/libm-2.31.so	
7f275c587000-7f275c588000 r--p 0014c000 103:05 917895	/usr/lib/x86_64-linux-
gnu/libm-2.31.so	
7f275c588000-7f275c589000 rw-p 0014d000 103:05 917895	/usr/lib/x86_64-linux-
gnu/libm-2.31.so	
7f275c589000-7f275c58b000 rw-p 00000000 00:00 0	
7f275c5ae000-7f275c5af000 r--p 00000000 103:05 917889	/usr/lib/x86_64-linux-gnu/ld-
2.31.so	
7f275c5af000-7f275c5d2000 r-xp 00001000 103:05 917889	/usr/lib/x86_64-linux-gnu/ld-
2.31.so	
7f275c5d2000-7f275c5da000 r--p 00024000 103:05 917889	/usr/lib/x86_64-linux-gnu/ld-
2.31.so	
7f275c5db000-7f275c5dc000 r--p 0002c000 103:05 917889	/usr/lib/x86_64-linux-gnu/ld-
2.31.so	

```
7f275c5dc000-7f275c5dd000 rw-p 0002d000 103:05 917889          /usr/lib/x86_64-linux-gnu/ld-
2.31.so
7f275c5dd000-7f275c5de000 rw-p 00000000 00:00 0
7ffd22d2f000-7ffd22d50000 rw-p 00000000 00:00 0          [stack]
7ffd22db0000-7ffd22db4000 r--p 00000000 00:00 0          [vvar]
7ffd22db4000-7ffd22db6000 r-xp 00000000 00:00 0          [vdso]
fffffffff600000-ffffffffffff601000 --xp 00000000 00:00 0          [vsySCALL]
```

True False

This is a virtual memory map (not physical memory map)

The 5th entry 55a4352c5000-55a4352e2000 **may** correspond to "data" of the vim.basic

vim.basic uses the math library

The size of the stack is one page

The size of the heap is one page

This is a virtual memory map (not physical memory map): True

The 5th entry 55a4352c5000-55a4352e2000 **may** correspond to "data" of the vim.basic: True

vim.basic uses the math library: True

The size of the stack is one page: False

The size of the heap is one page: False

Question 8

Complete

Mark 0.50 out of 0.50

Map the technique with its feature/problem

static linking	large executable file
dynamic loading	allocate memory only if needed
static loading	wastage of physical memory
dynamic linking	small executable file

The correct answer is: static linking → large executable file, dynamic loading → allocate memory only if needed, static loading → wastage of physical memory, dynamic linking → small executable file

Question 9

Complete

Mark 0.36 out of 0.50

Map the parts of a C code to the memory regions they are related to

function arguments	stack	◆
functions	stack	◆
local variables	stack	◆
static variables	data	◆
global un-initialized variables	bss	◆
global initialized variables	data	◆
malloced memory	stack	◆

The correct answer is: function arguments → stack, functions → code, local variables → stack, static variables → data, global un-initialized variables → bss, global initialized variables → data, malloced memory → heap

Question 10

Complete

Mark 0.75 out of 1.00

W.r.t the figure given below, mark the given statements as True or False.

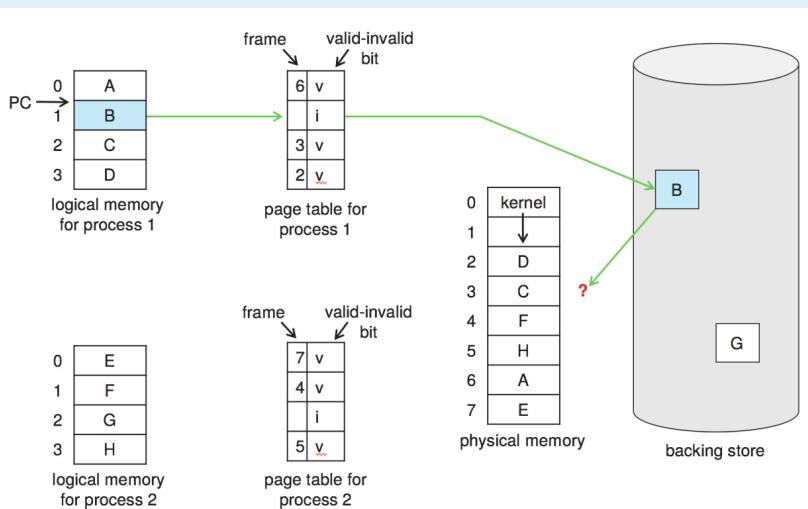


Figure 10.9 Need for page replacement.

True False

- | | | |
|----------------------------------|----------------------------------|--|
| <input type="radio"/> | <input checked="" type="radio"/> | Global replacement means chose any of the frame from 2 to 7 |
| <input checked="" type="radio"/> | <input type="radio"/> | Page 1 of process 1 needs a replacement |
| <input checked="" type="radio"/> | <input type="radio"/> | Handling this scenario demands two disk I/Os |
| <input checked="" type="radio"/> | <input type="radio"/> | Local replacement means chose any of the frames 2, 3, 6 |
| <input type="radio"/> | <input checked="" type="radio"/> | The kernel's pages can not used for replacement if kernel is not pageable. |
| <input checked="" type="radio"/> | <input type="radio"/> | Global replacement means chose any of the frame from 0 to 7 |
| <input checked="" type="radio"/> | <input type="radio"/> | Kernel occupies two page frames |
| <input type="radio"/> | <input checked="" type="radio"/> | Local replacement means chose any of the frame from 2 to 7 |

Global replacement means chose any of the frame from 2 to 7: True

Page 1 of process 1 needs a replacement: True

Handling this scenario demands two disk I/Os: True

Local replacement means chose any of the frames 2, 3, 6: True

The kernel's pages can not used for replacement if kernel is not pageable.: True

Global replacement means chose any of the frame from 0 to 7: False

Kernel occupies two page frames: True

Local replacement means chose any of the frame from 2 to 7: False

Question 11

Complete

Mark 0.75 out of 1.00

Given below is the output of the command "ps -eo min_flt,maj_flt,cmd" on a Linux Desktop system. Select the statements that are consistent with the output

```
626729 482768 /usr/lib/firefox/firefox -contentproc -parentBuildID 20220202182137 -prefsLen 9256 -prefMapSize 264738 -appDir /usr/lib/firefox/browser 6094 true rdd
2167 687 /usr/sbin/apache2 -k start
1265185 222 /usr/bin/gnome-shell
102648 111 /usr/sbin/mysqld
9813 0 bash
15497 370 /usr/bin/gedit --gapplication-service
```

- a. All of the processes here exhibit some good locality of reference
- b. The bash shell is mostly busy doing work within a particular locality
- c. Apache web-server has not been doing much work
- d. Firefox has likely been running for a large amount of time

The correct answers are: Firefox has likely been running for a large amount of time, Apache web-server has not been doing much work, The bash shell is mostly busy doing work within a particular locality, All of the processes here exhibit some good locality of reference

Question 12

Complete

Mark 0.00 out of 1.00

Calculate the EAT in NANO-seconds (upto 2 decimal points) w.r.t. a page fault, given

Memory access time = 105 ns

Average page fault service time = 13 ms

Page fault rate = 0.6

Answer:

The correct answer is: 7800042.00

Question 13

Complete

Mark 0.50 out of 1.00

For the reference string

3 4 3 5 2

the number of page faults (including initial ones) using

FIFO replacement and 2 page frames is :

4

FIFO replacement and 3 page frames is :

3

Question 14

Complete

Mark 1.00 out of 1.00

Assuming a 8- KB page size, what is the page numbers for the address 874815 reference in decimal :

(give answer also in decimal)

Answer: 107

The correct answer is: 107

Question 15

Complete

Mark 0.75 out of 1.00

Order the following events, related to page fault handling, in correct order

1. MMU detects that a page table entry is marked "invalid"
2. Page fault interrupt is generated
3. Page fault handler in kernel starts executing
4. Page fault handler detects that it's a page fault and not illegal memory access
5. Page faulting process is made to wait in a queue
6. Other processes scheduled by scheduler
7. Empty frame is found
8. Disk read is issued
9. Disk Interrupt occurs
10. Disk interrupt handler runs
11. Page table of page faulted process is updated
12. Page faulted process is moved to ready-queue

The correct order for these items is as follows:

1. MMU detects that a page table entry is marked "invalid"
2. Page fault interrupt is generated
3. Page fault handler in kernel starts executing
4. Page fault handler detects that it's a page fault and not illegal memory access
5. Empty frame is found
6. Disk read is issued
7. Page faulting process is made to wait in a queue
8. Other processes scheduled by scheduler
9. Disk Interrupt occurs
10. Disk interrupt handler runs
11. Page table of page faulted process is updated
12. Page faulted process is moved to ready-queue

Question 16

Complete

Mark 1.00 out of 1.00

Given six memory partitions of 300 KB , 600 KB , 350 KB , 200 KB , 750 KB , and 125 KB (in order), how would the first-fit, best-fit, and worst-fit algorithms place processes of size 115 KB and 500 KB (in order)?

- worst fit 500 KB 635 KB ▲
- best fit 500 KB 600 KB ▲
- worst fit 115 KB 750 KB ▲
- best fit 115 KB 125 KB ▲
- first fit 115 KB 300 KB ▲
- first fit 500 KB 600 KB ▲

The correct answer is: worst fit 500 KB → 635 KB, best fit 500 KB → 600 KB, worst fit 115 KB → 750 KB, best fit 115 KB → 125 KB, first fit 115 KB → 300 KB, first fit 500 KB → 600 KB

[◀ \(Code\) mmap related programs](#)

Jump to...

[Points from Mid-term feedback ►](#)

Operating Systems Even-sem-21-22

[Dashboard](#) / My courses / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [14 February - 20 February](#)
 / [Topic-wise Quiz-3 \(processes, trap handling, sched...](#)

Topic-wise Quiz-3 (processes, trap handling, scheduler)

Topic-wise Quiz-3 (processes, trap handling, scheduler)

Attempts allowed: 1

This quiz closed on Monday, 21 February 2022, 8:00:00 PM

To continue with this quiz attempt you must open your webcam, and it will be used to monitor you during the quiz.

Time limit: 55 mins

Summary of your previous attempts

State	Grade / 10.00	Review
Finished Submitted Monday, 21 February 2022, 7:59:27 PM	8.90	Review

Your final grade for this quiz is 8.90/10.00.

No more attempts are allowed

[Back to the course](#)

[◀ Description of some possible course mini projects](#)

Jump to...

[\(Code\) mmap related programs ▶](#)

You are logged in as [111903018 Ankit Nehul Div1](#) ([Log out](#))
[OS-even-sem-21-22](#)

Search/Create/Approve Courses

[Search Courses](#)

[New Course Request\(for Faculty only\).](#)

[Approve Requests\(For Admins only\).](#)

Learn Moodle

[COEP Moodle videos for teachers](#)

[Moodle Video Tutorials](#)

[Student Tutorials](#)

[Teacher Tutorials](#)

[Change Password](#)

[Preferences](#)

[English \(en\)](#)

[English \(en\)](#)

[मराठी \(mr\)](#)

[हिन्दी \(hi\)](#)

[Data retention summary](#)[Get the mobile app](#)[!\[\]\(0e47951b6066b873243d98ff863ad40a_img.jpg\) Give feedback about this software](#)

Started on	Saturday, 12 February 2022, 10:01:48 AM
State	Finished
Completed on	Saturday, 12 February 2022, 11:56:09 AM
Time taken	1 hour 54 mins
Grade	5.82 out of 10.00 (58%)

Question 1

Complete

Mark 0.75 out of 1.00

Select the correct statements about interrupt handling in xv6 code

- a. The trapframe pointer in struct proc, points to a location on user stack
- b. On any interrupt/syscall/exception the control first jumps in vectors.S
- c. On any interrupt/syscall/exception the control first jumps in trapasm.S
- d. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- e. xv6 uses the 0x64th entry in IDT for system calls
- f. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- g. The trapframe pointer in struct proc, points to a location on kernel stack
- h. All the 256 entries in the IDT are filled

- i. xv6 uses the 64th entry in IDT for system calls
- j. The CS and EIP are changed only immediately on a hardware interrupt
- k. The function trap() is called only in case of hardware interrupt
- l. Before going to alltraps, the kernel stack contains upto 5 entries.
- m. The function trap() is called irrespective of hardware interrupt/system-call/exception

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question 2

Complete

Mark 0.00 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

(Note: non-interruptible kernel code means, if the kernel code is executing, then interrupts will be disabled).

Note: A possible sequence may have some missing steps in between. An impossible sequence will have n and n+1th steps such that n+1th step can not follow n'th step.

Select one or more:

a. P1 running

P1 makes system call
system call returns
P1 running
timer interrupt
Scheduler running
P2 running

b. P1 running

P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P3 running
Hardware interrupt
Interrupt unblocks P1
Interrupt returns
P3 running
Timer interrupt
Scheduler
P1 running

c. P1 running

P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again

d. P1 running

keyboard hardware interrupt
keyboard interrupt handler running
interrupt handler returns
P1 running
P1 makes system call
system call returns
P1 running
timer interrupt
scheduler
P2 running

e.

P1 running
P1 makes system call
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again

f. P1 running

P1 makes system call
 timer interrupt
 Scheduler
 P2 running
 timer interrupt
 Scheuler
 P1 running
 P1's system call return

The correct answers are: P1 running

P1 makes system call and blocks
 Scheduler
 P2 running
 P2 makes system call and blocks
 Scheduler
 P1 running again, P1 running
 P1 makes system call
 timer interrupt
 Scheduler
 P2 running
 timer interrupt
 Scheuler
 P1 running
 P1's system call return,
 P1 running
 P1 makes system call
 Scheduler
 P2 running
 P2 makes system call and blocks
 Scheduler
 P1 running again

Question 3

Complete

Mark 0.00 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?

Select all the appropriate choices

- a. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time
- b. The code for reading ELF file can not be written in assembly
- c. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C
- d. The setting up of the most essential memory management infrastructure needs assembly code

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question 4

Complete

Mark 0.13 out of 0.50

Select all the correct statements about zombie processes

Select one or more:

- a. A zombie process occupies space in OS data structures
- b. A zombie process remains zombie forever, as there is no way to clean it up
- c. If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent
- d. A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it
- e. A process becomes zombie when its parent finishes
- f. Zombie processes are harmless even if OS is up for long time
- g. A process can become zombie if it finishes, but the parent has finished before it
- h. init() typically keeps calling wait() for zombie processes to get cleaned up

The correct answers are: A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it, A process can become zombie if it finishes, but the parent has finished before it, A zombie process occupies space in OS data structures, If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent, init() typically keeps calling wait() for zombie processes to get cleaned up

Question 5

Complete

Mark 0.50 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

Process P1 is running

6	◆
---	---

timer interrupt occurs

2	◆
---	---

context of P1 is loaded from P1's PCB

4	◆
---	---

context of P0 is saved in P0's PCB

3	◆
---	---

Process P0 is running

1	◆
---	---

Control is passed to P1

5	◆
---	---

The correct answer is: Process P1 is running → 6, timer interrupt occurs → 2, context of P1 is loaded from P1's PCB → 4, context of P0 is saved in P0's PCB → 3, Process P0 is running → 1, Control is passed to P1 → 5

Question 6

Complete

Mark 0.50 out of 0.50

Suppose a program does a `scanf()` call.

Essentially the `scanf` does a `read()` system call.

This call will obviously "block" waiting for the user input.

In terms of OS data structures and execution of code, what does it mean?

Select one:

- a. `read()` will return and process will be taken to a wait queue
- b. OS code for `read()` will move the PCB of this process to a wait queue and return from the system call
- c. OS code for `read()` will move PCB of current process to a wait queue and call scheduler
- d. OS code for `read()` will call scheduler
- e. `read()` returns and process calls scheduler()

The correct answer is: OS code for `read()` will move PCB of current process to a wait queue and call scheduler

Question 7

Complete

Mark 0.50 out of 0.50

Consider the following programs

exec1.c

```
#include <unistd.h>
#include <stdio.h>
int main() {
    exec("./exec2", "./exec2", NULL);
}
```

exec2.c

```
#include <unistd.h>
#include <stdio.h>
int main() {
    exec("/bin/ls", "/bin/ls", NULL);
    printf("hello\n");
}
```

Compiled as

```
cc  exec1.c -o exec1
cc  exec2.c -o exec2
```

And run as

```
$ ./exec1
```

Explain the output of the above command (./exec1)

Assume that /bin/ls , i.e. the 'ls' program exists.

Select one:

- a. Program prints hello
- b. Execution fails as one exec can't invoke another exec
- c. "ls" runs on current directory
- d. Execution fails as the call to execl() in exec2 fails
- e. Execution fails as the call to execl() in exec1 fails

The correct answer is: "ls" runs on current directory

Question 8

Complete

Mark 0.25 out of 0.50

The bootmain() function has this code

```
elf = (struct elfhdr*)0x10000; // scratch space
readseg((uchar*)elf, 4096, 0);
```

Mark the statements as True or False with respect to this code.

In these statements 0x1000 is referred to as ADDRESS

True**False**

This line effectively loads the ELF header and the program headers at ADDRESS



If the value of ADDRESS is changed, then the program will not work



If the value of ADDRESS is changed to a higher number (upto a limit), the program could still work



If the value of ADDRESS is changed to a lower number (upto a limit), the program could still work



If the value ADDRESS is changed to a 0 the program could still work



This line loads the kernel code at ADDRESS

This line effectively loads the ELF header and the program headers at ADDRESS: False

If the value of ADDRESS is changed, then the program will not work: False

If the value of ADDRESS is changed to a higher number (upto a limit), the program could still work: True

If the value of ADDRESS is changed to a lower number (upto a limit), the program could still work: True

The value ADDRESS is changed to a 0 the program could still work: False

This line loads the kernel code at ADDRESS: False

Question 9

Complete

Mark 0.80 out of 1.00

Mark the statements, w.r.t. the scheduler of xv6 as True or False

True**False**

The work of selecting and scheduling a process is done only in `scheduler()` and not in `sched()`

The variable `c->scheduler` on first processor uses the stack allocated entry.S

`sched()` and `scheduler()` are co-routines

the control returns to `switchkvm();` after `swtch(&(c->scheduler), p->context);` in `scheduler()`

When a process is scheduled for execution, it resumes execution in `sched()` after the call to `swtch()`

`swtch` is a function that saves old context, loads new context, and returns to last EIP in the new context

`sched()` calls `scheduler()` and `scheduler()` calls `sched()`

`swtch` is a function that does not return to the caller

The function `scheduler()` executes using the kernel-only stack

the control returns to `mycpu()->intena = intena; ()`; after `swtch(&p->context, mycpu()->scheduler);` in `sched()`

The work of selecting and scheduling a process is done only in `scheduler()` and not in `sched()`: True

The variable `c->scheduler` on first processor uses the stack allocated entry.S: True

`sched()` and `scheduler()` are co-routines: True

the control returns to `switchkvm();` after `swtch(&(c->scheduler), p->context);` in `scheduler()`: False

When a process is scheduled for execution, it resumes execution in `sched()` after the call to `swtch()`

: True

`swtch` is a function that saves old context, loads new context, and returns to last EIP in the new context: True

`sched()` calls `scheduler()` and `scheduler()` calls `sched()`: False

`swtch` is a function that does not return to the caller: True

The function `scheduler()` executes using the kernel-only stack: True

the control returns to `mycpu()->intena = intena; ()`; after `swtch(&p->context, mycpu()->scheduler);` in `sched()`:

False

Question 10

Complete

Mark 0.36 out of 0.50

Order the events that occur on a timer interrupt:

- | | | |
|---|---|---|
| Change to kernel stack of currently running process | 1 | ◆ |
| Select another process for execution | 5 | ◆ |
| Execute the code of the new process | 7 | ◆ |
| Jump to a code pointed by IDT | 3 | ◆ |
| Jump to scheduler code | 4 | ◆ |
| Save the context of the currently running process | 2 | ◆ |
| Set the context of the new process | 6 | ◆ |

The correct answer is: Change to kernel stack of currently running process → 1, Select another process for execution → 5, Execute the code of the new process → 7, Jump to a code pointed by IDT → 2, Jump to scheduler code → 4, Save the context of the currently running process → 3, Set the context of the new process → 6

Question 11

Complete

Mark 0.20 out of 0.50

For each line of code mentioned on the left side, select the location of sp/esp that is in use

`jmp *%eax`
in entry.S

The 4KB area in kernel image, loaded in memory, named as 'stack' ◆

`cli`
in bootasm.S

Immaterial as the stack is not used here ◆

`readseg((uchar*)elf, 4096, 0);`
in bootmain.c

0x10000 to 0x7c00 ◆

`call bootmain`
in bootasm.S

0x7c00 to 0x10000 ◆

`ljmp $(SEG_KCODE<<3), $start32`
in bootasm.S

0x7c00 to 0 ◆

The correct answer is: `jmp *%eax`
in entry.S → The 4KB area in kernel image, loaded in memory, named as 'stack', `cli`
`in bootasm.S` → Immortal as the stack is not used here, `readseg((uchar*)elf, 4096, 0);`
in bootmain.c → 0x7c00 to 0, `call bootmain`
`in bootasm.S` → 0x7c00 to 0, `ljmp $(SEG_KCODE<<3), $start32`
`in bootasm.S` → Immortal as the stack is not used here

Question 12

Complete

Mark 0.50 out of 0.50

What's the trapframe in xv6?

- a. A frame of memory that contains all the trap handler code's function pointers
- b. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S
- c. A frame of memory that contains all the trap handler code
- d. The IDT table
- e. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only
- f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- g. A frame of memory that contains all the trap handler's addresses

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question 13

Complete

Mark 0.50 out of 0.50

In bootasm.S, on the line

```
ljmp      $(SEG_KCODE<<3), $start32
```

The SEG_KCODE << 3, that is shifting of 1 by 3 bits is done because

- a. The ljmp instruction does a divide by 8 on the first argument
- b. The code segment is 16 bit and only lower 13 bits are used for segment number
- c. The code segment is 16 bit and only upper 13 bits are used for segment number
- d. While indexing the GDT using CS, the value in CS is always divided by 8
- e. The value 8 is stored in code segment

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question 14

Complete

Mark 0.43 out of 0.50

Select all the correct statements about code of bootmain() in xv6

```

void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry) (void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++) {
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void))(elf->entry);
    entry();
}

```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- a. The elf->entry is set by the linker in the kernel file and it's 8010000c
- b. The kernel file gets loaded at the Physical address 0x10000 in memory.
- c. The elf->entry is set by the linker in the kernel file and it's 0x80000000
- d. The readseg finally invokes the disk I/O code using assembly instructions
- e. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it.
- f. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded
- g. The stosb() is used here, to fill in some space in memory with zeroes
- h. The kernel file gets loaded at the Physical address 0x10000 +0x80000000 in memory.
- i. The elf->entry is set by the linker in the kernel file and it's 0x80000000
- j. The kernel file has only two program headers
- k. The condition if(ph->memsz > ph->filesz) is never true.

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

Question 15

Complete

Mark 0.40 out of 0.50

Select Yes if the mentioned element should be a part of PCB

Select No otherwise.

Yes**No**

<input type="radio"/>	<input checked="" type="radio"/> PID of Init	
<input type="radio"/>	<input checked="" type="radio"/> Pointer to IDT	
<input checked="" type="radio"/>	<input type="radio"/> EIP at the time of context switch	
<input type="radio"/>	<input checked="" type="radio"/> Pointer to the parent process	
<input checked="" type="radio"/>	<input type="radio"/> PID	
<input checked="" type="radio"/>	<input type="radio"/> Process state	
<input checked="" type="radio"/>	<input type="radio"/> Function pointers to all system calls	
<input checked="" type="radio"/>	<input type="radio"/> Process context	
<input type="radio"/>	<input type="radio"/> List of opened files	
<input checked="" type="radio"/>	<input type="radio"/> Memory management information about that process	

PID of Init: No

Pointer to IDT: No

EIP at the time of context switch: Yes

Pointer to the parent process: Yes

PID: Yes

Process state: Yes

Function pointers to all system calls: No

Process context: Yes

List of opened files: Yes

Memory management information about that process: Yes

Question 16

Not answered

Marked out of 1.00

Which parts of the xv6 code in bootasm.S bootmain.c , entry.S and in the codepath related to scheduler() and trap handling() can also be written in some other way, and still ensure that xv6 works properly?

Writing code is not necessary. You only need to comment on which part of the code could be changed to something else or written in another fashion.

Maximum two points to be written.

◀ Extra Reading on Linkers: A writeup by Ian Taylor (keep changing url string from 38 to 39, and so on)

Jump to...



(Code) IPC - Shm, Messages ►

Started on Wednesday, 9 February 2022, 7:01:59 PM

State Finished

Completed on Wednesday, 9 February 2022, 8:01:50 PM

Time taken 59 mins 51 secs

Grade 10.00 out of 11.00 (91%)

Question 1

Complete

Mark 1.00 out of 1.00

The ljmp instruction in general does

- a. change the CS and EIP to 32 bit mode, and jumps to new value of EIP
- b. change the CS and EIP to 32 bit mode, and jumps to next line of code
- c. change the CS and EIP to 32 bit mode
- d. change the CS and EIP to 32 bit mode, and jumps to kernel code

The correct answer is: change the CS and EIP to 32 bit mode, and jumps to new value of EIP

Question 2

Complete

Mark 1.00 out of 1.00

The kernel is loaded at Physical Address

- a. 0x0010000
- b. 0x80000000
- c. 0x00100000
- d. 0x80100000

The correct answer is: 0x00100000

Question 3

Complete

Mark 1.00 out of 1.00

The number of GDT entries setup during boot process of xv6 is

- a. 0
- b. 3
- c. 256
- d. 2
- e. 255
- f. 4

The correct answer is: 3

Question 4

Complete

Mark 1.00 out of 1.00

The variable \$stack in entry.S is

- a. located at less than 0x7c00
- b. a memory region allocated as a part of entry.S
- c. located at the value given by %esp as setup by bootmain()
- d. located at 0x7c00
- e. located at 0

The correct answer is: a memory region allocated as a part of entry.S

Question 5

Complete

Mark 1.00 out of 1.00

which of the following is not a difference between real mode and protected mode

- a. in real mode the segment is multiplied by 16, in protected mode segment is used as index in GDT
- b. in real mode general purpose registers are 16 bit, in protected mode they are 32 bit
- c. in real mode the addressable memory is more than in protected mode
- d. processor starts in real mode
- e. in real mode the addressable memory is less than in protected mode

The correct answer is: in real mode the addressable memory is more than in protected mode

Question 6

Complete

Mark 1.00 out of 1.00

The kernel ELF file contains how many Program headers?

- a. 2
- b. 4
- c. 9
- d. 3
- e. 10

The correct answer is: 3

Question 7

Complete

Mark 1.00 out of 1.00

ELF Magic number is

- a. 0xELF
- b. 0xFFFFFFFF
- c. 0
- d. 0x0x464CELF
- e. 0xELFELFELF
- f. 0x464C457FL
- g. 0x464C457FU

The correct answer is: 0x464C457FU

Question 8

Complete

Mark 1.00 out of 1.00

x86 provides which of the following type of memory management options?

- a. segmentation and two level paging
- b. segmentation and one level paging
- c. segmentation only
- d. segmentation or one or two level paging
- e. segmentation and one or two level paging
- f. segmentation or paging

The correct answer is: segmentation and one or two level paging

Question 9

Complete

Mark 1.00 out of 1.00

Why is the code of entry() in Assembly and not in C?

- a. There is no particular reason, it could also be in C
- b. Because the kernel code must begin in assembly
- c. Because it needs to setup paging
- d. Because the symbol entry() is inside the ELF file

The correct answer is: Because it needs to setup paging

Question 10

Not answered

Marked out of 0.50

code line, MMU setting: Match the line of xv6 code with the MMU setup employed

Answer:

The correct answer is: inb \$0x64,%al

Question 11

Not answered

Marked out of 0.50

Match the pairs of which action is taken by whom

Answer:

The correct answer is: kernel

Question 12

Complete

Mark 1.00 out of 1.00

The right side of line of code "entry = (void(*)(void))(elf->entry)" means

- a. Get the "entry" in ELF structure and convert it into a function pointer accepting no arguments and returning nothing
- b. Get the "entry" in ELF structure and convert it into a function void pointer
- c. Convert the "entry" in ELF structure into void
- d. Get the "entry" in ELF structure and convert it into a void pointer

The correct answer is: Get the "entry" in ELF structure and convert it into a function pointer accepting no arguments and returning nothing

[◀ Homework questions: Basics of MM, xv6 booting](#)

Jump to...

[\(Code\)](#) [Files](#), [redirection](#), [dup](#), [\(IPC\)pipe](#) ►

[Dashboard](#) / My courses / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [24 January - 30 January](#)
 / [Topic-wise Quiz: 1 \(system calls, x86, calling convention\)](#)

Started on Monday, 24 January 2022, 7:02:26 PM

State Finished

Completed on Monday, 24 January 2022, 8:57:38 PM

Time taken 1 hour 55 mins

Grade **10.40** out of 20.00 (52%)

Question 1

Complete

Mark 1.00 out of 1.00

```
int value = 5;
int main()
{
    pid_t pid;
    pid = fork();
    if (pid == 0) { /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("%d", value); /* LINE A */
    }
    return 0;
}
```

What's the value printed here at LINE A?

Answer:

The correct answer is: 5

Question 2

Complete

Mark 1.00 out of 1.00

Why should a program exist in memory before it starts executing ?

- a. Because the hard disk is a slow medium
- b. Because the memory is volatile
- c. Because the variables of the program are stored in memory
- d. Because the processor can run instructions and access data only from memory

The correct answer is: Because the processor can run instructions and access data only from memory

Question 3

Complete

Mark 2.00 out of 2.00

Match the program with its output (ignore newlines in the output. Just focus on the count of the number of 'hi')

`main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }`

hi

hi hi

hi

hi hi

`main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }`

`main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }`

`main() { int i = NULL; fork(); printf("hi\n"); }`

The correct answer is: `main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi, main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi hi, main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi, main() { int i = NULL; fork(); printf("hi\n"); } → hi hi`

Question 4

Complete

Mark 1.33 out of 2.00

Which of the following instructions should be privileged?

Select one or more:

- a. Switch from user to kernel mode.
- b. Turn off interrupts.
- c. Set value of timer.
- d. Read the clock.
- e. Set value of a memory location
- f. Access I/O device.
- g. Access memory management unit of the processor
- h. Modify entries in device-status table
- i. Access a general purpose register

The correct answers are: Set value of timer., Access memory management unit of the processor, Turn off interrupts., Modify entries in device-status table, Access I/O device., Switch from user to kernel mode.

Question 5

Complete

Mark 0.00 out of 0.50

Is the command "cat README > done &" possible on xv6? (Note the & in the end)

- a. yes
- b. no

The correct answer is: yes

Question 6

Complete

Mark 1.00 out of 1.00

Match the register with the segment used with it.

esp	ss	▼
ebp	ss	▼
eip	cs	▼
esi	ds	▼
edi	es	▼

The correct answer is: esp → ss, ebp → ss, eip → cs, esi → ds, edi → es

Question 7

Complete

Mark 0.00 out of 2.00

```
xv6.img: bootblock kernel
```

```
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- a. Blocks in xv6.img after kernel may be all zeroes.
- b. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file.
- c. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk.
- d. The size of the kernel file is nearly 5 MB
- e. The size of the xv6.img is nearly 5 MB
- f. The bootblock is located on block-0 of the xv6.img
- g. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk.
- h. The kernel is located at block-1 of the xv6.img
- i. xv6.img is the virtual processor used by the qemu emulator
- j. The bootblock may be 512 bytes or less (looking at the Makefile instruction)
- k. The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question 8

Complete

Mark 0.50 out of 0.50

Is the terminal a part of the kernel on GNU/Linux systems?

- a. yes
- b. no

The correct answer is: no

Question 9

Complete

Mark 1.00 out of 2.00

Which of the following are NOT a part of job of a typical compiler?

- a. Process the # directives in a C program
- b. Convert high level language code to machine code
- c. Invoke the linker to link the function calls with their code, extern globals with their declaration
- d. Suggest alternative pieces of code that can be written
- e. Check the program for syntactical errors
- f. Check the program for logical errors

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question 10

Complete

Mark 0.00 out of 0.50

Compare multiprogramming with multitasking

- a. A multiprogramming system is not necessarily multitasking
- b. A multitasking system is not necessarily multiprogramming

The correct answer is: A multiprogramming system is not necessarily multitasking

Question 11

Complete

Mark 0.40 out of 1.00

Select all the correct statements about two modes of CPU operation

Select one or more:

- a. The two modes are essential for a multiprogramming system
- b. Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode
- c. The two modes are essential for a multitasking system
- d. The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously
- e. There is an instruction like 'iret' to return from kernel mode to user mode

The correct answers are: The two modes are essential for a multiprogramming system, The two modes are essential for a multitasking system, There is an instruction like 'iret' to return from kernel mode to user mode, The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously, Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode

Question 12

Complete

Mark 1.00 out of 1.00

Rank the following storage systems from slowest (first) to fastest(last)

You can drag and drop the items below/above each other.

Magnetic tapes

Optical disk

Hard-disk drives

Nonvolatile memory

Main memory

Cache

Registers

Question 13

Complete

Mark 0.83 out of 2.00

Select all statements that correctly explain the use/purpose of system calls.

Select one or more:

- a. Handle exceptions like division by zero
- b. Switch from user mode to kernel mode
- c. Allow I/O device access to user processes
- d. Run each instruction of an application program
- e. Provide services for accessing files
- f. Provide an environment for process creation
- g. Handle ALL types of interrupts

The correct answers are: Switch from user mode to kernel mode, Provide services for accessing files, Allow I/O device access to user processes, Provide an environment for process creation

Question 14

Complete

Mark 0.33 out of 0.50

Order the following events in boot process (from 1 onwards)

Login interface	6	◆
Shell	4	◆
Boot loader	2	◆
OS	3	◆
Init	4	◆
BIOS	1	◆

The correct answer is: Login interface → 5, Shell → 6, Boot loader → 2, OS → 3, Init → 4, BIOS → 1

Question 15

Complete

Mark 0.00 out of 2.00

Select all the correct statements about calling convention on x86 32-bit.

- a. The return value is either stored on the stack or returned in the eax register
- b. during execution of a function, ebp is pointing to the old ebp
- c. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function
- d. Parameters may be passed in registers or on stack
- e. Parameters may be passed in registers or on stack
- f. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables
- g. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function
- h. Return address is one location above the ebp
- i. Compiler may allocate more memory on stack than needed
- j. Parameters are pushed on the stack in left-right order
- k. The ebp pointers saved on the stack constitute a chain of activation records

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

Question 16

Complete

Mark 0.00 out of 1.00

How does the distinction between kernel mode and user mode function as a rudimentary form of protection (security) ?

Select one:

- a. It prohibits invocation of kernel code completely, if a user program is running
- b. It prohibits one process from accessing other process's memory
- c. It prohibits a user mode process from running privileged instructions
- d. It disallows hardware interrupts when a process is running

The correct answer is: It prohibits a user mode process from running privileged instructions

[◀ \(Task\) Compulsory xv6 task](#)

Jump to...



[\(Optional Assignment\) Shell Programming\(Conformance tests\) ▶](#)

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [7 February - 13 February](#)

/ [Quiz-1: 10 AM](#)

Started on Saturday, 12 February 2022, 10:00:21 AM

State Finished

Completed on Saturday, 12 February 2022, 11:25:53 AM

Time taken 1 hour 25 mins

Grade 4.94 out of 10.00 (49%)

Question 1

Complete

Mark 0.00 out of 0.50

Select all the correct statements about code of bootmain() in xv6

```

void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stobs(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void))(elf->entry);
    entry();
}

```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- a. The elf->entry is set by the linker in the kernel file and it's 0x80000000
- b. The condition if(ph->memsz > ph->filesz) is never true.
- c. The readseg finally invokes the disk I/O code using assembly instructions
- d. The elf->entry is set by the linker in the kernel file and it's 0x80000000
- e. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded
- f. The kernel file gets loaded at the Physical address 0x10000 +0x80000000 in memory.
- g. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it.
- h. The elf->entry is set by the linker in the kernel file and it's 8010000c
- i. The kernel file gets loaded at the Physical address 0x10000 in memory.
- j. The stobs() is used here, to fill in some space in memory with zeroes
- k. The kernel file has only two program headers

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

Question 2

Complete

Mark 0.50 out of 0.50

What's the trapframe in xv6?

- a. A frame of memory that contains all the trap handler's addresses
- b. A frame of memory that contains all the trap handler code's function pointers
- c. The IDT table
- d. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S
- e. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only
- f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- g. A frame of memory that contains all the trap handler code

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question 3

Complete

Mark 0.21 out of 0.50

Order the events that occur on a timer interrupt:

- | | |
|---|---|
| Jump to a code pointed by IDT | 3 |
| Jump to scheduler code | 2 |
| Select another process for execution | 5 |
| Change to kernel stack of currently running process | 4 |
| Set the context of the new process | 6 |
| Save the context of the currently running process | 1 |
| Execute the code of the new process | 7 |

The correct answer is: Jump to a code pointed by IDT → 2, Jump to scheduler code → 4, Select another process for execution → 5, Change to kernel stack of currently running process → 1, Set the context of the new process → 6, Save the context of the currently running process → 3, Execute the code of the new process → 7

Question 4

Complete

Mark 0.50 out of 0.50

Suppose a program does a scanf() call.

Essentially the scanf does a read() system call.

This call will obviously "block" waiting for the user input.

In terms of OS data structures and execution of code, what does it mean?

Select one:

- a. OS code for read() will move PCB of current process to a wait queue and call scheduler
- b. OS code for read() will call scheduler
- c. OS code for read() will move the PCB of this process to a wait queue and return from the system call
- d. read() will return and process will be taken to a wait queue
- e. read() returns and process calls scheduler()

The correct answer is: OS code for read() will move PCB of current process to a wait queue and call scheduler

Question 5

Complete

Mark 0.50 out of 0.50

In bootasm.S, on the line

```
ljmp    $(SEG_KCODE<<3), $start32
```

The SEG_KCODE << 3, that is shifting of 1 by 3 bits is done because

- a. The value 8 is stored in code segment
- b. The code segment is 16 bit and only upper 13 bits are used for segment number
- c. While indexing the GDT using CS, the value in CS is always divided by 8
- d. The ljmp instruction does a divide by 8 on the first argument
- e. The code segment is 16 bit and only lower 13 bits are used for segment number

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question 6

Complete

Mark 0.40 out of 0.50

Select Yes if the mentioned element should be a part of PCB

Select No otherwise.

Yes**No**

- Memory management information about that process

- Process context

- Function pointers to all system calls

- PID

- EIP at the time of context switch

- List of opened files

- PID of Init

- Process state

- Pointer to the parent process

- Pointer to IDT

Memory management information about that process: Yes

Process context: Yes

Function pointers to all system calls: No

PID: Yes

EIP at the time of context switch: Yes

List of opened files: Yes

PID of Init: No

Process state: Yes

Pointer to the parent process: Yes

Pointer to IDT: No

Question 7

Complete

Mark 0.40 out of 1.00

Mark the statements, w.r.t. the scheduler of xv6 as True or False

True False

The work of selecting and scheduling a process is done only in `scheduler()` and not in `sched()`

`swtch` is a function that saves old context, loads new context, and returns to last EIP in the new context

The variable `c->scheduler` on first processor uses the stack allocated entry.S

`sched()` and `scheduler()` are co-routines

The function `scheduler()` executes using the kernel-only stack

When a process is scheduled for execution, it resumes execution in `sched()` after the call to `swtch()`

`swtch` is a function that does not return to the caller

the control returns to `mycpu()->intena = intena; ()`; after `swtch(&p->context, mycpu()->scheduler);` in `sched()`

the control returns to `switch kvm();` after `swtch(&(c->scheduler), p->context);` in `scheduler()`

`sched()` calls `scheduler()` and `scheduler()` calls `sched()`

The work of selecting and scheduling a process is done only in `scheduler()` and not in `sched()`: True

`swtch` is a function that saves old context, loads new context, and returns to last EIP in the new context: True

The variable `c->scheduler` on first processor uses the stack allocated entry.S: True

`sched()` and `scheduler()` are co-routines: True

The function `scheduler()` executes using the kernel-only stack: True

When a process is scheduled for execution, it resumes execution in `sched()` after the call to `swtch()`

: True

`swtch` is a function that does not return to the caller: True

the control returns to `mycpu()->intena = intena; ()`; after `swtch(&p->context, mycpu()->scheduler);` in `sched()`:

False

the control returns to `switch kvm();` after `swtch(&(c->scheduler), p->context);` in `scheduler()`: False

`sched()` calls `scheduler()` and `scheduler()` calls `sched()`: False

Question 8

Complete

Mark 0.00 out of 0.50

Consider the following programs

exec1.c

```
#include <unistd.h>
#include <stdio.h>
int main() {
    exec("./exec2", "./exec2", NULL);
}
```

exec2.c

```
#include <unistd.h>
#include <stdio.h>
int main() {
    exec("/bin/ls", "/bin/ls", NULL);
    printf("hello\n");
}
```

Compiled as

```
cc  exec1.c -o exec1
cc  exec2.c -o exec2
```

And run as

```
$ ./exec1
```

Explain the output of the above command (./exec1)

Assume that /bin/ls , i.e. the 'ls' program exists.

Select one:

- a. "ls" runs on current directory
- b. Execution fails as the call to execl() in exec1 fails
- c. Execution fails as one exec can't invoke another exec
- d. Program prints hello
- e. Execution fails as the call to execl() in exec2 fails

The correct answer is: "ls" runs on current directory

Question 9

Complete

Mark 0.00 out of 0.50

For each line of code mentioned on the left side, select the location of sp/esp that is in use

cli**in bootasm.S**

0x7c00 to 0

readseg((uchar*)elf, 4096, 0);**in bootmain.c**

Immaterial as the stack is not used here

ljmp \$(SEG_KCODE<<3), \$start32**in bootasm.S**

0x10000 to 0x7c00

call bootmain**in bootasm.S**

The 4KB area in kernel image, loaded in memory, named as 'stack'

jmp *%eax**in entry.S**

0x7c00 to 0x10000

The correct answer is: **cli**

in bootasm.S → Immateral as the stack is not used here, **readseg((uchar*)elf, 4096, 0);**

in bootmain.c → 0x7c00 to 0, **ljmp \$(SEG_KCODE<<3), \$start32**

in bootasm.S → Immateral as the stack is not used here, **call bootmain**

in bootasm.S → 0x7c00 to 0, **jmp *%eax**

in entry.S → The 4KB area in kernel image, loaded in memory, named as 'stack'

Question 10

Complete

Mark 0.63 out of 1.00

Select the correct statements about interrupt handling in xv6 code

- a. xv6 uses the 64th entry in IDT for system calls
- b. xv6 uses the 0x64th entry in IDT for system calls
- c. On any interrupt/syscall/exception the control first jumps in vectors.S
- d. The function trap() is called irrespective of hardware interrupt/system-call/exception
- e. Before going to altraps, the kernel stack contains upto 5 entries.
- f. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- g. The trapframe pointer in struct proc, points to a location on kernel stack
- h. The function trap() is called only in case of hardware interrupt
- i. The trapframe pointer in struct proc, points to a location on user stack
- j. On any interrupt/syscall/exception the control first jumps in trapasm.S
- k. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- l. The CS and EIP are changed only immediately on a hardware interrupt
- m. All the 256 entries in the IDT are filled

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to altraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question 11

Complete

Mark 0.50 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

timer interrupt occurs

2

Process P0 is running

1

context of P1 is loaded from P1's PCB

4

Process P1 is running

6

Control is passed to P1

5

context of P0 is saved in P0's PCB

3

The correct answer is: timer interrupt occurs → 2, Process P0 is running → 1, context of P1 is loaded from P1's PCB → 4, Process P1 is running → 6, Control is passed to P1 → 5, context of P0 is saved in P0's PCB → 3

Question 12

Complete

Mark 0.00 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

(Note: non-interruptible kernel code means, if the kernel code is executing, then interrupts will be disabled).

Note: A possible sequence may have some missing steps in between. An impossible sequence will have n and n+1th steps such that n+1th step can not follow n'th step.

Select one or more:

a. P1 running

P1 makes system call
timer interrupt
Scheduler
P2 running
timer interrupt
Scheduler
P1 running
P1's system call return

b. P1 running

P1 makes system call
system call returns
P1 running
timer interrupt
Scheduler running
P2 running

c. P1 running

P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P3 running
Hardware interrupt
Interrupt unblocks P1
Interrupt returns
P3 running
Timer interrupt
Scheduler
P1 running

d. P1 running

keyboard hardware interrupt
keyboard interrupt handler running
interrupt handler returns
P1 running
P1 makes system call
system call returns
P1 running
timer interrupt
scheduler
P2 running

e.

P1 running
P1 makes system call
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again

f. P1 running

P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again, P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question 13

Complete

Mark 0.50 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?

Select all the appropriate choices

- a. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C
- b. The setting up of the most essential memory management infrastructure needs assembly code
- c. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time
- d. The code for reading ELF file can not be written in assembly

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question 14

Complete

Mark 0.17 out of 0.50

The bootmain() function has this code

```
elf = (struct elfhdr*)0x10000; // scratch space  
readseg((uchar*)elf, 4096, 0);
```

Mark the statements as True or False with respect to this code.

In these statements 0x1000 is referred to as ADDRESS

True False

- The value ADDRESS is changed to a 0 the program could still work
- If the value of ADDRESS is changed to a higher number (upto a limit), the program could still work
- This line loads the kernel code at ADDRESS
- If the value of ADDRESS is changed to a lower number (upto a limit), the program could still work
- This line effectively loads the ELF header and the program headers at ADDRESS
- If the value of ADDRESS is changed, then the program will not work

The value ADDRESS is changed to a 0 the program could still work: False

If the value of ADDRESS is changed to a higher number (upto a limit), the program could still work: True

This line loads the kernel code at ADDRESS: False

If the value of ADDRESS is changed to a lower number (upto a limit), the program could still work: True

This line effectively loads the ELF header and the program headers at ADDRESS: False

If the value of ADDRESS is changed, then the program will not work: False

Question 15

Complete

Mark 0.50 out of 1.00

Which parts of the xv6 code in bootasm.S bootmain.c , entry.S and in the codepath related to scheduler() and trap handling() can also be written in some other way, and still ensure that xv6 works properly?

Writing code is not necessary. You only need to comment on which part of the code could be changed to something else or written in another fashion.

Maximum two points to be written.

We can use a scheduling algorithm. We can use the kernel stack in scheduler function in entry.S and bootmain.c .

Question 16

Complete

Mark 0.13 out of 0.50

Select all the correct statements about zombie processes

Select one or more:

- a. If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent
- b. Zombie processes are harmless even if OS is up for long time
- c. A zombie process occupies space in OS data structures
- d. A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it
- e. A process can become zombie if it finishes, but the parent has finished before it
- f. init() typically keeps calling wait() for zombie processes to get cleaned up
- g. A process becomes zombie when its parent finishes
- h. A zombie process remains zombie forever, as there is no way to clean it up

The correct answers are: A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it, A process can become zombie if it finishes, but the parent has finished before it, A zombie process occupies space in OS data structures, If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent, init() typically keeps calling wait() for zombie processes to get cleaned up

◀ Extra Reading on Linkers: A writeup by Ian Taylor (keep changing url string from 38 to 39, and so on)

Jump to...

Started on Saturday, 20 February 2021, 2:51 PM

State Finished

Completed on Saturday, 20 February 2021, 3:55 PM

Time taken 1 hour 3 mins

Grade 7.30 out of 20.00 (37%)

Question 1

Partially correct

Mark 0.80 out of 1.00

Select all the correct statements about the state of a process.

- a. A process can self-terminate only when it's running ✓
- b. Typically, it's represented as a number in the PCB ✓
- c. A process that is running is not on the ready queue ✓
- d. Processes in the ready queue are in the ready state ✓
- e. It is not maintained in the data structures by kernel, it is only for conceptual understanding of programmers
- f. Changing from running state to waiting state results in "giving up the CPU" ✓
- g. A process in ready state is ready to receive interrupts
- h. A waiting process starts running after the wait is over ✗
- i. A process changes from running to ready state on a timer interrupt ✓
- j. A process in ready state is ready to be scheduled ✓
- k. A running process may terminate, or go to wait or become ready again ✓
- l. A process waiting for I/O completion is typically woken up by the particular interrupt handler code ✓
- m. A process waiting for any condition is woken up by another process only
- n. A process changes from running to ready state on a timer interrupt or any I/O wait

Your answer is partially correct.

You have selected too many options.

The correct answers are: Typically, it's represented as a number in the PCB, A process in ready state is ready to be scheduled, Processes in the ready queue are in the ready state, A process that is running is not on the ready queue, A running process may terminate, or go to wait or become ready again, A process changes from running to ready state on a timer interrupt, Changing from running state to waiting state results in "giving up the CPU", A process can self-terminate only when it's running, A process waiting for I/O completion is typically woken up by the particular interrupt handler code

Question 2

Incorrect

Mark 0.00 out of 1.00

For each line of code mentioned on the left side, select the location of sp/esp that is in use

`jmp *%eax`
in entry.S

0x7c00 to 0x10000



`ljmp $(SEG_KCODE<<3), $start32`
in bootasm.S

0x10000 to 0x7c00



`call bootmain`
in bootasm.S

0x7c00 to 0x10000



`cli`
in bootasm.S

0x7c00 to 0



`readseg((uchar*)elf, 4096, 0);`
in bootmain.c

The 4KB area in kernel image, loaded in memory, named as 'stack'



Your answer is incorrect.

The correct answer is: `jmp *%eax`

`in entry.S` → The 4KB area in kernel image, loaded in memory, named as 'stack', `ljmp $(SEG_KCODE<<3), $start32`

`in bootasm.S` → Immaterail as the stack is not used here, `call bootmain`

`in bootasm.S` → 0x7c00 to 0, `cli`

`in bootasm.S` → Immaterail as the stack is not used here, `readseg((uchar*)elf, 4096, 0);`

`in bootmain.c` → 0x7c00 to 0

Question 3

Correct

Mark 0.25 out of 0.25

Order the following events in boot process (from 1 onwards)

Boot loader	2	✓
Shell	6	✓
BIOS	1	✓
OS	3	✓
Init	4	✓
Login interface	5	✓

Your answer is correct.

The correct answer is: Boot loader → 2, Shell → 6, BIOS → 1, OS → 3, Init → 4, Login interface → 5

Question 4

Partially correct

Mark 0.30 out of 0.50

Consider the following command and its output:

```
$ ls -lht xv6.img kernel
-rw-rw-r-- 1 abhijit abhijit 4.9M Feb 15 11:09 xv6.img
-rwxrwxr-x 1 abhijit abhijit 209K Feb 15 11:09 kernel*
```

Following code in bootmain()

```
readseg((uchar*)elf, 4096, 0);
```

and following selected lines from Makefile

```
xv6.img: bootblock kernel
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

```
kernel: $(OBJS) entry.o entryother initcode kernel.ld
$(LD) $(LDFLAGS) -T kernel.ld -o kernel entry.o $(OBJS) -b binary initcode entryother
$(OBJDUMP) -S kernel > kernel.asm
$(OBJDUMP) -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$$/d' > kernel.sym
```

Also read the code of bootmain() in xv6 kernel.

Select the options that describe the meaning of these lines and their correlation.

- a. Although the size of the kernel file is 209 Kb, only 4Kb out of it is the actual kernel code and remaining part is all zeroes.
- b. The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files ✓
- c. The kernel.ld file contains instructions to the linker to link the kernel properly ✓
- d. The bootmain() code does not read the kernel completely in memory
- e. readseg() reads first 4k bytes of kernel in memory
- f. Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.
- g. The kernel.asm file is the final kernel file
- h. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is not read as it is user programs.
- i. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(). ✓

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(), readseg() reads first 4k bytes of kernel in memory, The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files, The kernel.ld file contains instructions to the linker to link the kernel properly, Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.

Question 5

Partially correct

Mark 0.50 out of 1.00

```
int f() {  
    int count;  
    for (count = 0; count < 2; count++) {  
        if (fork() == 0)  
            printf("Operating-System\n");  
    }  
    printf("TYCOMP\n");  
}
```

The number of times "Operating-System" is printed, is:

Answer:

The correct answer is: 7.00

Question 6

Partially correct

Mark 0.40 out of 0.50

Select Yes/True if the mentioned element must be a part of PCB

Select No/False otherwise.

Yes**No**

<input checked="" type="radio"/>	<input checked="" type="radio"/>	PID	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Process context	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	List of opened files	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Process state	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Parent's PID	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Pointer to IDT	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Function pointers to all system calls	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Memory management information about that process	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Pointer to the parent process	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	EIP at the time of context switch	✓

PID: Yes

Process context: Yes

List of opened files: Yes

Process state: Yes

Parent's PID: No

Pointer to IDT: No

Function pointers to all system calls: No

Memory management information about that process: Yes

Pointer to the parent process: Yes

EIP at the time of context switch: Yes

Question 7

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about code of bootmain() in xv6

```

void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void))(elf->entry);
    entry();
}

```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- a. The kernel file gets loaded at the Physical address 0x10000 +0x80000000 in memory. ✗
- b. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- c. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded ✓
- d. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it. ✓
- e. The kernel file has only two program headers ✓
- f. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- g. The readseg finally invokes the disk I/O code using assembly instructions ✓
- h. The elf->entry is set by the linker in the kernel file and it's 8010000c ✓
- i. The kernel file gets loaded at the Physical address 0x10000 in memory. ✓
- j. The condition if(ph->memsz > ph->filesz) is never true. ✗
- k. The stosb() is used here, to fill in some space in memory with zeroes ✓

Your answer is incorrect.

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

Question 8

Partially correct

Mark 0.13 out of 0.25

Which of the following are NOT a part of job of a typical compiler?

- a. Check the program for logical errors ✓
- b. Convert high level language code to machine code
- c. Process the # directives in a C program
- d. Invoke the linker to link the function calls with their code, extern globals with their declaration
- e. Check the program for syntactical errors
- f. Suggest alternative pieces of code that can be written

Your answer is partially correct.

You have correctly selected 1.

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question 9

Correct

Mark 0.25 out of 0.25

Rank the following storage systems from slowest (first) to fastest(last)

Cache	6	✓
Hard Disk	3	✓
RAM	5	✓
Optical Disks	2	✓
Non volatile memory	4	✓
Registers	7	✓
Magnetic Tapes	1	✓

Your answer is correct.

The correct answer is: Cache → 6, Hard Disk → 3, RAM → 5, Optical Disks → 2, Non volatile memory → 4, Registers → 7, Magnetic Tapes → 1

Question 10

Partially correct

Mark 0.21 out of 0.50

Which of the following parts of a C program do not have any corresponding machine code ?

- a. local variable declaration
- b. global variables
- c. function calls ✗
- d. #directives ✓
- e. expressions
- f. pointer dereference
- g. typedefs ✓

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: #directives, typedefs, global variables

Question 11

Correct

Mark 0.25 out of 0.25

Match a system call with it's description

pipe	create an unnamed FIFO storage with 2 ends - one for reading and another for writing	✓
dup	create a copy of the specified file descriptor into smallest available file descriptor	✓
dup2	create a copy of the specified file descriptor into another specified file descriptor	✓
exec	execute a binary file overlaying the image of current process	✓
fork	create an identical child process	✓

Your answer is correct.

The correct answer is: pipe → create an unnamed FIFO storage with 2 ends - one for reading and another for writing, dup → create a copy of the specified file descriptor into smallest available file descriptor, dup2 → create a copy of the specified file descriptor into another specified file descriptor, exec → execute a binary file overlaying the image of current process, fork → create an identical child process

Question 12

Correct

Mark 0.25 out of 0.25

Match the register with the segment used with it.

eip	cs	✓
edi	es	✓
esi	ds	✓
ebp	ss	✓
esp	ss	✓

Your answer is correct.

The correct answer is: eip → cs, edi → es, esi → ds, ebp → ss, esp → ss

Question 13

Correct

Mark 0.25 out of 0.25

What's the trapframe in xv6?

- a. A frame of memory that contains all the trap handler code
- b. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- c. The IDT table
- d. A frame of memory that contains all the trap handler code's function pointers
- e. A frame of memory that contains all the trap handler's addresses
- f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S ✓
- g. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only

Your answer is correct.

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question 14

Incorrect

Mark 0.00 out of 0.50

Select all the correct statements about linking and loading.

Select one or more:

- a. Continuous memory management schemes can support dynamic linking and dynamic loading. ✗
- b. Loader is last stage of the linker program ✗
- c. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently) ✓
- d. Dynamic linking and loading is not possible without demand paging or demand segmentation. ✓
- e. Dynamic linking essentially results in relocatable code. ✓
- f. Continuous memory management schemes can support static linking and static loading. (may be inefficiently) ✓
- g. Loader is part of the operating system ✓
- h. Static linking leads to non-relocatable code ✗
- i. Dynamic linking is possible with continuous memory management, but variable sized partitions only. ✗

Your answer is incorrect.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

Question 15

Incorrect

Mark 0.00 out of 0.25

In bootasm.S, on the line

```
ljmp    $(SEG_KCODE<<3), $start32
```

The SEG_KCODE << 3, that is shifting of 1 by 3 bits is done because

- a. The value 8 is stored in code segment
- b. The code segment is 16 bit and only upper 13 bits are used for segment number
- c. The code segment is 16 bit and only lower 13 bits are used for segment number ✗
- d. While indexing the GDT using CS, the value in CS is always divided by 8
- e. The ljmp instruction does a divide by 8 on the first argument

Your answer is incorrect.

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question 16

Partially correct

Mark 0.07 out of 0.50

Order the events that occur on a timer interrupt:

Change to kernel stack

1	✗
---	---

Jump to a code pointed by IDT

2	✗
---	---

Jump to scheduler code

5	✗
---	---

Set the context of the new process

4	✗
---	---

Save the context of the currently running process

3	✓
---	---

Execute the code of the new process

6	✗
---	---

Select another process for execution

7	✗
---	---

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: Change to kernel stack → 2, Jump to a code pointed by IDT → 1, Jump to scheduler code → 4, Set the context of the new process → 6, Save the context of the currently running process → 3, Execute the code of the new process → 7, Select another process for execution → 5

Question 17

Incorrect

Mark 0.00 out of 1.00

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

```
$ ls . /tmp/asdfksdf >/tmp/ddd 2>&1
```

Program 1

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    dup(fd);
    close(2);
    dup(fd);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Program 2

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    close(1);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(2);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Select all the correct statements about the programs

Select one or more:

- a. Both programs are correct ✗
- b. Program 2 makes sure that there is one file offset used for '2' and '1' ✗
- c. Only Program 2 is correct ✗
- d. Program 2 does 1>&2 ✗
- e. Program 2 ensures 2>&1 and does not ensure >/tmp/ddd ✗
- f. Program 1 makes sure that there is one file offset used for '2' and '1' ✓
- g. Program 1 is correct for >/tmp/ddd but not for 2>&1 ✗
- h. Program 1 does 1>&2 ✗
- i. Both program 1 and 2 are incorrect ✗
- j. Program 2 is correct for >/tmp/ddd but not for 2>&1 ✗
- k. Only Program 1 is correct ✓
- l. Program 1 ensures 2>&1 and does not ensure >/tmp/ddd ✗

Your answer is incorrect.

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'

Question 18

Correct

Mark 0.25 out of 0.25

Select the option which best describes what the CPU does during its powered ON lifetime

- a. Ask the user what is to be done, and execute that task
- b. Ask the OS what is to be done, and execute that task
- c. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask the User or the OS what is to be done next, repeat
- d. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per ✓ the instruction itself, repeat
- e. Fetch instruction specified by OS, Decode and execute it, repeat
- f. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask OS what is to be done next, repeat

The correct answer is: Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, repeat

Question 19

Partially correct

Mark 0.86 out of 1.00

Consider the following code and MAP the file to which each fd points at the end of the code.

```
int main(int argc, char *argv[]) {
    int fd1, fd2 = 1, fd3 = 1, fd4 = 1;

    fd1 = open("/tmp/1", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    fd2 = open("/tmp/2", O_RDONLY);
    fd3 = open("/tmp/3", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    close(0);
    close(1);
    dup(fd2);
    dup(fd3);
    close(fd3);
    dup2(fd2, fd4);
    printf("%d %d %d %d\n", fd1, fd2, fd3, fd4);
    return 0;
}
```

1	closed	✗
fd4	/tmp/2	✓
fd2	/tmp/2	✓
fd1	/tmp/1	✓
2	stderr	✓
0	/tmp/2	✓
fd3	closed	✓

Your answer is partially correct.

You have correctly selected 6.

The correct answer is: 1 → /tmp/3, fd4 → /tmp/2, fd2 → /tmp/2, fd1 → /tmp/1, 2 → stderr, 0 → /tmp/2, fd3 → closed

Question 20

Incorrect

Mark 0.00 out of 2.00

Following code claims to implement the command

```
/bin/ls -l | /usr/bin/head -3 | /usr/bin/tail -1
```

Fill in the blanks to make the code work.

Note: Do not include space in writing any option. x[1][2] should be written without any space, and so is the case with [1] or [2]. Pay attention to exact syntax and do not write any extra character like ';' or = etc.

```
int main(int argc, char *argv[]) {
```

```
    int pid1, pid2;
```

```
    int pfd[
```

```
    x ] [2];
```

```
    pipe(
```

```
    x );
```

```
    pid1 =
```

```
    x ;
```

```
    if(pid1 != 0) {
```

```
        close(pfd[0]
```

```
    x );
```

```
        close(
```

```
    x );
```

```
        dup(
```

```
    x );
```

```
        execl("/bin/ls", "/bin/ls", "
```

```
    x ", NULL);
```

```
    }
```

```
    pipe(
```

```
    x );
```

```
    x = fork();
```

```
    if(pid2 == 0) {
```

```
        close(
```

```
    x ;
```

```
        close(0);
```

```
        dup(
```

```
    x );
```

```
        close(pfd[1]
```

```
✗ );
close(
  
✗ );
dup(
  
✗ );
execl("/usr/bin/head", "/usr/bin/head", "  
  
✗ ", NULL);
} else {
close(pfd
  
✗ );
close(
  
✗ );
dup(
  
✗ );
close(pfd
  
✗ );
execl("/usr/bin/tail", "/usr/bin/tail", "  
  
✗ ", NULL);
}  
}
```

Question 21

Partially correct

Mark 0.11 out of 1.00

Select all the correct statements about calling convention on x86 32-bit.

- a. Return address is one location above the ebp ✓
- b. Parameters may be passed in registers or on stack ✓
- c. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function ✓
- d. The ebp pointers saved on the stack constitute a chain of activation records ✓
- e. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables ✗
- f. Parameters may be passed in registers or on stack ✓
- g. The return value is either stored on the stack or returned in the eax register ✗
- h. Parameters are pushed on the stack in left-right order
- i. during execution of a function, ebp is pointing to the old ebp
- j. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function ✗
- k. Compiler may allocate more memory on stack than needed ✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

Question 22

Correct

Mark 1.00 out of 1.00

Match the program with its output (ignore newlines in the output. Just focus on the count of the number of 'hi')

main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi ✓

main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi hi ✓

main() { int i = NULL; fork(); printf("hi\n"); }

hi hi ✓

main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi ✓

Your answer is correct.

The correct answer is: main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi, main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi hi, main() { int i = NULL; fork(); printf("hi\n"); } → hi hi, main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi

Question 23

Incorrect

Mark 0.00 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?

Select all the appropriate choices

- a. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time ✗
- b. The setting up of the most essential memory management infrastructure needs assembly code ✓
- c. The code for reading ELF file can not be written in assembly ✗
- d. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C ✓

Your answer is incorrect.

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question 24

Incorrect

Mark 0.00 out of 0.50

```
xv6.img: bootblock kernel
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- a. The size of the kernel file is nearly 5 MB ✓
- b. The kernel is located at block-1 of the xv6.img ✗
- c. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk. ✗
- d. The size of xv6.img is exactly = (size of bootblock) + (size of kernel) ✗
- e. The bootblock is located on block-0 of the xv6.img ✗
- f. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk. ✓
- g. The bootblock may be 512 bytes or less (looking at the Makefile instruction) ✗
- h. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file. ✗
- i. The size of the xv6.img is nearly 5 MB ✗
- j. xv6.img is the virtual processor used by the qemu emulator ✓
- k. Blocks in xv6.img after kernel may be all zeroes. ✗

Your answer is incorrect.

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question 25

Incorrect

Mark 0.00 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

Select one or more:

a. P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return

b. P1 running

P1 makes sytem call and blocks

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P1 running again



c. P1 running

P1 makes system call

system call returns

P1 running

timer interrupt

Scheduler running

P2 running

d. P1 running

P1 makes sytem call and blocks

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P3 running

Hardware interrupt

Interrupt unblocks P1

Interrupt returns

P3 running

Timer interrupt

Scheduler

P1 running



e.

P1 running

P1 makes sytem call

Scheduler

P2 running

P2 makes sytem call and blocks

Scheduler

P1 running again

f. P1 running

keyboard hardware interrupt

keyboard interrupt handler running

interrupt handler returns

P1 running

P1 makes sytem call

system call returns



P1 running
timer interrupt
scheduler
P2 running

Your answer is incorrect.

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again, P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question 26

Correct

Mark 0.25 out of 0.25

Which of the following are the files related to bootloader in xv6?

- a. bootasm.s and entry.S
- b. bootasm.S and bootmain.c ✓
- c. bootasm.S, bootmain.c and bootblock.c
- d. bootmain.c and bootblock.S

Your answer is correct.

The correct answer is: bootasm.S and bootmain.c

Question 27

Correct

Mark 0.25 out of 0.25

Match the following parts of a C program to the layout of the process in memory

Instructions	Text section	✓
Local Variables	Stack Section	✓
Dynamically allocated memory	Heap Section	✓
Global and static data	Data section	✓

Your answer is correct.

The correct answer is:

Instructions → Text section, Local Variables → Stack Section,
 Dynamically allocated memory → Heap Section,
 Global and static data → Data section

Question 28

Incorrect

Mark 0.00 out of 0.50

What will this program do?

```
int main() {
    fork();
    execl("/bin/ls", "/bin/ls", NULL);
    printf("hello");
}
```

- a. one process will run ls, another will print hello
- b. run ls once ✗
- c. run ls twice
- d. run ls twice and print hello twice
- e. run ls twice and print hello twice, but output will appear in some random order

Your answer is incorrect.

The correct answer is: run ls twice

Question 29

Correct

Mark 0.25 out of 0.25

What is the OS Kernel?

- a. The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run ✓ correct
- b. The set of tools like compiler, linker, loader, terminal, shell, etc.
- c. Only the system programs like compiler, linker, loader, etc.
- d. Everything that I see on my screen

The correct answer is: The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run

Question 30

Correct

Mark 0.50 out of 0.50

Which of the following is/are not saved during context switch?

- a. Program Counter
- b. General Purpose Registers
- c. Bus ✓
- d. Stack Pointer
- e. MMU related registers/information
- f. Cache ✓
- g. TLB ✓

Your answer is correct.

The correct answers are: TLB, Cache, Bus

Question 31

Partially correct

Mark 0.10 out of 0.25

Select the order in which the various stages of a compiler execute.

Linking	3	
Syntactical Analysis	2	
Pre-processing	1	
Intermediate code generation	does not exist	
Loading	4	

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Linking → 4, Syntactical Analysis → 2, Pre-processing → 1, Intermediate code generation → 3, Loading → does not exist

Question 32

Partially correct

Mark 0.08 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

context of P0 is saved in P0's PCB	2	
context of P1 is loaded from P1's PCB	3	
Process P1 is running	5	
timer interrupt occurs	6	
Process P0 is running	1	
Control is passed to P1	4	

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: context of P0 is saved in P0's PCB → 3, context of P1 is loaded from P1's PCB → 4, Process P1 is running → 6, timer interrupt occurs → 2, Process P0 is running → 1, Control is passed to P1 → 5

Question 33

Not answered

Marked out of 1.00

Select the correct statements about interrupt handling in xv6 code

- a. On any interrupt/syscall/exception the control first jumps in vectors.S
- b. The trapframe pointer in struct proc, points to a location on user stack
- c. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- d. xv6 uses the 64th entry in IDT for system calls
- e. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- f. The trapframe pointer in struct proc, points to a location on kernel stack
- g. The function trap() is called only in case of hardware interrupt
- h. The CS and EIP are changed only immediately on a hardware interrupt
- i. All the 256 entries in the IDT are filled

- j. On any interrupt/syscall/exception the control first jumps in trapasm.S
- k. The function trap() is called irrespective of hardware interrupt/system-call/exception
- l. xv6 uses the 0x64th entry in IDT for system calls
- m. Before going to alltraps, the kernel stack contains upto 5 entries.

Your answer is incorrect.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in trapasm.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

[◀ \(Assignment\) Change free list management in xv6](#)

Jump to...

Started on Tuesday, 22 March 2022, 1:59:34 PM

State Finished

Completed on Tuesday, 22 March 2022, 4:14:53 PM

Time taken 2 hours 15 mins

Grade 24.57 out of 40.00 (61%)

Question 1

Partially correct

Mark 0.10 out of 1.00

Select all the correct statements w.r.t user and kernel threads

Select one or more:

a. many-one model can be implemented even if there are no kernel threads ✓

b. many-one model gives no speedup on multicore processors

c. all three models, that is many-one, one-one, many-many , require a user level thread library ✓

d. A process blocks in many-one model even if a single thread makes a blocking system call

e. A process may not block in many-one model, if a thread makes a blocking system call ✗

f. one-one model increases kernel's scheduling load ✓

g. one-one model can be implemented even if there are no kernel threads

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: many-one model can be implemented even if there are no kernel threads, all three models, that is many-one, one-one, many-many , require a user level thread library, one-one model increases kernel's scheduling load, many-one model gives no speedup on multicore processors, A process blocks in many-one model even if a single thread makes a blocking system call

Question 2

Partially correct

Mark 0.50 out of 1.00

Select all correct statements w.r.t. Major and Minor page faults on Linux

- a. Thrashing is possible only due to major page faults
- b. Minor page fault may occur because the page was freed, but still tagged and available in the free page list
- c. Minor page fault may occur because of a page fault during fork(), on code of an already running process ✓
- d. Major page faults are likely to occur in more numbers at the beginning of the process ✓
- e. Minor page fault may occur because the page was a shared memory page ✓
- f. Minor page faults are an improvement of the page buffering techniques

The correct answers are: Minor page fault may occur because the page was a shared memory page, Minor page fault may occur because of a page fault during fork(), on code of an already running process, Minor page fault may occur because the page was freed, but still tagged and available in the free page list, Major page faults are likely to occur in more numbers at the beginning of the process, Thrashing is possible only due to major page faults, Minor page faults are an improvement of the page buffering techniques

Question 3

Correct

Mark 2.00 out of 2.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only
Mark statements True or False

True	False	
<input checked="" type="radio"/>	<input type="radio"/> ✗	The stack allocated in entry.S is used as stack for scheduler's context for first processor
<input type="radio"/> ✗	<input checked="" type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context
<input checked="" type="radio"/>	<input type="radio"/> ✗	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir
<input checked="" type="radio"/>	<input type="radio"/> ✗	The kernel code and data take up less than 2 MB space
<input checked="" type="radio"/>	<input type="radio"/> ✗	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context
<input checked="" type="radio"/>	<input type="radio"/> ✗	xv6 uses physical memory upto 224 MB only
<input checked="" type="radio"/>	<input type="radio"/> ✗	The free page-frame are created out of nearly 222 MB
<input type="radio"/> ✗	<input checked="" type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process
<input checked="" type="radio"/>	<input type="radio"/> ✗	PHYSTOP can be increased to some extent, simply by editing memlayout.h
<input checked="" type="radio"/>	<input type="radio"/> ✗	The process's address space gets mapped on frames, obtained from ~2MB:224MB range
<input type="radio"/> ✗	<input checked="" type="radio"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

The kernel code and data take up less than 2 MB space: True

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

xv6 uses physical memory upto 224 MB only: True

The free page-frame are created out of nearly 222 MB: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

Question 4

Correct

Mark 1.00 out of 1.00

Given that a kernel has 1000 KB of total memory, and holes of sizes (in that order) 300 KB, 200 KB, 100 KB, 250 KB. For each of the requests on the left side, match it with the chunk chosen using the specified algorithm.

Consider each request as first request.

50 KB, worst fit	300 KB	✓
100 KB, worst fit	300 KB	✓
200 KB, first fit	300 KB	✓
150 KB, best fit	200 KB	✓
220 KB, best fit	250 KB	✓
150 KB, first fit	300 KB	✓

The correct answer is: 50 KB, worst fit → 300 KB, 100 KB, worst fit → 300 KB, 200 KB, first fit → 300 KB, 150 KB, best fit → 200 KB, 220 KB, best fit → 250 KB, 150 KB, first fit → 300 KB

Question 5

Partially correct

Mark 0.60 out of 1.00

Choice of the global or local replacement strategy is a subjective choice for kernel programmers. There are advantages and disadvantages on either side. Out of the following statements, that advocate either global or local replacement strategy, select those statements that have a logically **CONSISTENT** argument. (That is any statement that is logically correct about either global or local replacement)

Consistent **Inconsistent**

<input checked="" type="radio"/>	<input type="radio"/> X	Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> X	Global replacement may give highly variable per process completion time because number of page faults become un-predictable.	X
<input checked="" type="radio"/>	<input type="radio"/> X	Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> X	Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> X	Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.	X

Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.: Consistent

Global replacement may give highly variable per process completion time because number of page faults become un-predictable.: Consistent

Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).: Consistent

Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.: Consistent

Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.: Consistent

Question 6

Correct

Mark 1.00 out of 1.00

Map the functionality/use with function/variable in xv6 code.

Setup kernel part of a page table, and switch to that page table

kvmalloc()

Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

✓ setupkvm()

return a free page, if available; 0, otherwise

✓ kalloc()

Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed

mappages()

Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

✓ walkpgdir()

Array listing the kernel memory mappings, to be used by setupkvm()

✓ kmap[]

Your answer is correct.

The correct answer is: Setup kernel part of a page table, and switch to that page table → kvmalloc(), Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → setupkvm(), return a free page, if available; 0, otherwise → kalloc(), Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → mappages(), Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → walkpgdir(), Array listing the kernel memory mappings, to be used by setupkvm() → kmap[]

Question 7

Partially correct

Mark 0.20 out of 1.00

Mark statements True/False w.r.t. change of states of a process. Note that a statement is true only if the claim and argument both are true.

Reference: The process state diagram (and your understanding of how kernel code works). Note - the diagram does not show zombie state!

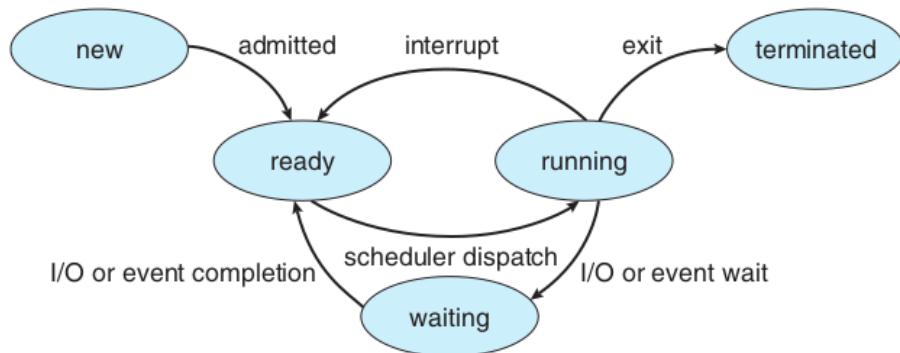


Figure 3.2 Diagram of process state.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	Every forked process has to go through ZOMBIE state, at least for a small duration.
<input type="radio"/>	<input checked="" type="radio"/>	A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first
<input checked="" type="radio"/>	<input type="radio"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet
<input checked="" type="radio"/>	<input type="radio"/>	Only a process in READY state is considered by scheduler
<input type="radio"/>	<input checked="" type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.

Every forked process has to go through ZOMBIE state, at least for a small duration.: True

A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first: False

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet: True

Only a process in READY state is considered by scheduler: True

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

Question 8

Correct

Mark 2.00 out of 2.00

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using FIFO replacement is:

Answer:

#6# 6,4# 6,4,2 #0,4,2# 0,1,2 #0,1,6 #9,1,6# 9,2,6# 9,2,0 #5,2,0

The correct answer is: 10

Question 9

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about linking and loading.

Select one or more:

- a. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently)
- b. Dynamic linking essentially results in relocatable code.
- c. Continuous memory management schemes can support static linking and static loading. (may be inefficiently)
- d. Loader is last stage of the linker program
- e. Dynamic linking and loading is not possible without demand paging or demand segmentation.
- f. Static linking leads to non-relocatable code
- g. Continuous memory management schemes can support dynamic linking and dynamic loading.
- h. Dynamic linking is possible with continuous memory management, but variable sized partitions only.
- i. Loader is part of the operating system

Your answer is incorrect.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

Question 10

Correct

Mark 1.00 out of 1.00

Consider a computer system with a 32-bit logical address and 4- KB page size. The system supports up to 512 MB of physical memory. How many entries are there in each of the following?

Write answer as a decimal number.

A conventional, single-level page table:

1048576



An inverted page table:

131072

**Question 11**

Incorrect

Mark 0.00 out of 1.00

W.r.t. xv6 code, match the state of a process with a code that sets the state

RUNNING	Choose...
ZOMBIE	Choose...
EMBRYO	Choose...
SLEEPING	Choose...
UNUSED	Choose...
RUNNABLE	scheduler()



The correct answer is: RUNNING → scheduler(), ZOMBIE → exit(), called by process itself, EMBRYO → fork()->allocproc() before setting up the UVM, SLEEPING → sleep(), called by any process blocking itself, UNUSED → wait(), called by parent process, RUNNABLE → wakeup(), called by an interrupt handler

Question 12

Not answered

Marked out of 1.00

Select the correct statements about interrupt handling in xv6 code

- a. The trapframe pointer in struct proc, points to a location on user stack
- b. The CS and EIP are changed only immediately on a hardware interrupt
- c. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- d. The trapframe pointer in struct proc, points to a location on kernel stack
- e. On any interrupt/syscall/exception the control first jumps in vectors.S
- f. The function trap() is called irrespective of hardware interrupt/system-call/exception
- g. All the 256 entries in the IDT are filled
- h. The function trap() is called only in case of hardware interrupt
- i. xv6 uses the 64th entry in IDT for system calls
- j. xv6 uses the 0x64th entry in IDT for system calls
- k. Before going to alptraps, the kernel stack contains upto 5 entries.
- l. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- m. On any interrupt/syscall/exception the control first jumps in trapasm.S

Your answer is incorrect.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alptraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question 13

Correct

Mark 1.00 out of 1.00

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- a. end, (4MB + PHYSTOP)
- b. P2V(end), PHYSTOP
- c. end, P2V(4MB + PHYSTOP)
- d. end, PHYSTOP
- e. end, 4MB
- f. end, P2V(PHYSTOP) ✓
- g. P2V(end), P2V(PHYSTOP)

Your answer is correct.

The correct answer is: end, P2V(PHYSTOP)

Question 14

Correct

Mark 1.00 out of 1.00

Select all the correct statements about MMU and its functionality (on a non-demand paged system)

Select one or more:

- a. Illegal memory access is detected in hardware by MMU and a trap is raised ✓
- b. Illegal memory access is detected by operating system
- c. MMU is a separate chip outside the processor
- d. The operating system interacts with MMU for every single address translation
- e. Logical to physical address translations in MMU are done in hardware, automatically ✓
- f. Logical to physical address translations in MMU are done with specific machine instructions
- g. MMU is inside the processor ✓
- h. The Operating system sets up relevant CPU registers to enable proper MMU translations ✓

Your answer is correct.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

Question 15

Incorrect

Mark 0.00 out of 2.00

Order the following events, in the creation of init() process in xv6:

1. ✘ initcode is selected by scheduler for execution
2. ✘ kernel stack is allocated for initcode process
3. ✘ values are set in the trapframe of initcode
4. ✘ sys_exec runs
5. ✘ initcode process is set to be runnable
6. ✘ code is set to start in forkret() when process gets scheduled
7. ✘ Arguments are setup on process stack for /init
8. ✘ trapframe and context pointers are set to proper location
9. ✘ trap() runs
10. ✘ userinit() is called
11. ✘ the header of "/init" ELF file is ready by kernel
12. ✘ empty struct proc is obtained for initcode
13. ✘ Stack is allocated for "/init" process
14. ✘ function pointer from syscalls[] array is invoked
15. ✘ memory mappings are created for "/init" process
16. ✘ page table mappings of 'initcode' are replaced by mappings of 'init'
17. ✘ initcode calls exec system call
18. ✘ initcode process runs
19. ✘ name of process "/init" is copied in struct proc

Your answer is incorrect.

Grading type: Relative to the next item (including last)

Grade details: 0 / 19 = 0%

Here are the scores for each item in this response:

1. 0 / 1 = 0%
2. 0 / 1 = 0%
3. 0 / 1 = 0%
4. 0 / 1 = 0%
5. 0 / 1 = 0%
6. 0 / 1 = 0%
7. 0 / 1 = 0%
8. 0 / 1 = 0%
9. 0 / 1 = 0%
10. 0 / 1 = 0%
11. 0 / 1 = 0%
12. 0 / 1 = 0%
13. 0 / 1 = 0%
14. 0 / 1 = 0%
15. 0 / 1 = 0%

- 16. 0 / 1 = 0%
- 17. 0 / 1 = 0%
- 18. 0 / 1 = 0%
- 19. 0 / 1 = 0%

The correct order for these items is as follows:

1. userinit() is called
2. empty struct proc is obtained for initcode
3. kernel stack is allocated for initcode process
4. trapframe and context pointers are set to proper location
5. code is set to start in forkret() when process gets scheduled
6. kernel memory mappings are created for initcode
7. values are set in the trapframe of initcode
8. initcode process is set to be runnable
9. initcode is selected by scheduler for execution
10. initcode process runs
11. initcode calls exec system call
12. trap() runs
13. function pointer from syscalls[] array is invoked
14. sys_exec runs
15. the header of "/init" ELF file is ready by kernel
16. memory mappings are created for "/init" process
17. Stack is allocated for "/init" process
18. Arguments on setup on process stack for /init
19. name of process "/init" is copied in struct proc
20. page table mappings of 'initcode' are replaced by makpings of 'init'

Question 16

Partially correct

Mark 0.56 out of 1.00

Mark the statements as True or False, w.r.t. mmap()

True	False	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	mmap() can be implemented on both demand paged and non-demand paged systems.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	MAP_FIXED guarantees that the mapping is always done at the specified address
<input checked="" type="radio"/>	<input checked="" type="radio"/>	MAP_PRIVATE leads to a mapping that is copy-on-write
<input checked="" type="radio"/>	<input checked="" type="radio"/>	on failure mmap() returns (void *)-1
<input checked="" type="radio"/>	<input checked="" type="radio"/>	MAP_SHARED leads to a mapping that is copy-on-write
<input checked="" type="radio"/>	<input checked="" type="radio"/>	mmap() results in changes to buffer-cache of the kernel.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	on failure mmap() returns NULL
<input checked="" type="radio"/>	<input checked="" type="radio"/>	mmap() results in changes to page table of a process.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	mmap() is a system call

mmap() can be implemented on both demand paged and non-demand paged systems.: True

MAP_FIXED guarantees that the mapping is always done at the specified address: False

MAP_PRIVATE leads to a mapping that is copy-on-write: True

on failure mmap() returns (void *)-1: True

MAP_SHARED leads to a mapping that is copy-on-write: False

mmap() results in changes to buffer-cache of the kernel.: False

on failure mmap() returns NULL: False

mmap() results in changes to page table of a process.: True

mmap() is a system call: True

Question 17

Incorrect

Mark 0.00 out of 1.00

If one thread opens a file with read privileges then

Select one:

- a. other threads in the same process can also read from that file
- b. none of these
- c. any other thread cannot read from that file
- d. other threads in the another process can also read from that file

Your answer is incorrect.

The correct answer is: other threads in the same process can also read from that file

Question 18

Partially correct

Mark 0.60 out of 1.00

Mark the statements about named and un-named pipes as True or False

True	False	
<input checked="" type="radio"/>	<input type="radio"/> 	Named pipe exists as a file
<input checked="" type="radio"/>	<input type="radio"/> 	Un-named pipes are inherited by a child process from parent.
<input type="radio"/> 	<input checked="" type="radio"/>	The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.
<input checked="" type="radio"/>	<input type="radio"/> 	Both types of pipes are an extension of the idea of "message passing".
<input type="radio"/> 	<input checked="" type="radio"/>	A named pipe has a name decided by the kernel.
<input checked="" type="radio"/>	<input type="radio"/> 	Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.
<input checked="" type="radio"/>	<input type="radio"/> 	Both types of pipes provide FIFO communication.
<input type="radio"/> 	<input checked="" type="radio"/>	Named pipes can be used for communication between only "related" processes.
<input checked="" type="radio"/>	<input type="radio"/> 	Named pipes can exist beyond the life-time of processes using them.
<input type="radio"/> 	<input checked="" type="radio"/>	The pipe() system call can be used to create either a named or un-named pipe.

Named pipe exists as a file.: True

Un-named pipes are inherited by a child process from parent.: True

The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.: False

Both types of pipes are an extension of the idea of "message passing": True

A named pipe has a name decided by the kernel.: False

Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.: True

Both types of pipes provide FIFO communication.: True

Named pipes can be used for communication between only "related" processes.: False

Named pipes can exist beyond the life-time of processes using them.: True

The pipe() system call can be used to create either a named or un-named pipe.: False

Question 19

Partially correct

Mark 0.67 out of 1.00

Select the most common causes of use of IPC by processes

- a. More modular code
- b. Breaking up a large task into small tasks and speeding up computation, on multiple core machines ✓
- c. More security checks
- d. Sharing of information of common interest ✓
- e. Get the kernel performance statistics

The correct answers are: Sharing of information of common interest, Breaking up a large task into small tasks and speeding up computation, on multiple core machines, More modular code

Question 20

Correct

Mark 1.00 out of 1.00

For each function/code-point, select the status of segmentation setup in xv6

after seginit() in main()	gdt setup with 5 entries (0 to 4) on one processor	✓
bootmain()	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	✓
entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
kvmalloc() in main()	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
bootasm.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓

Your answer is correct.

The correct answer is: after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S

Question 21

Partially correct

Mark 0.50 out of 1.00

Mark whether the given sequence of events is possible or not-possible. Also, select the reason for your answer.

For each sequence it's a not-possible sequence if some important event is not mentioned in the sequence.

Assume that the kernel code is non-interruptible and uniprocessor system.

Process P1, user code executing

Timer interrupt

Context changes to kernel context

Generic interrupt handler runs

Generic interrupt handler calls Scheduler

Scheduler selects P2 for execution

After scheduler, Process P2 user code executing

This sequence of events is: not-possible ✓

Because

Generic interrupt handler can not call scheduler ✗

Question 22

Partially correct

Mark 0.63 out of 1.00

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

True	False	
<input checked="" type="radio"/> ✘	<input type="radio"/> ✓	Integer arguments are stored in eax, ebx, ecx, etc. registers
<input type="radio"/> ✓	<input checked="" type="radio"/> ✘	String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	The functions like argint(), argstr() make the system call arguments available in the kernel.
<input checked="" type="radio"/> ✘	<input type="radio"/> ✓	String arguments are first copied to trapframe and then from trapframe to kernel's other variables.
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	The arguments to system call originally reside on process stack.
<input checked="" type="radio"/> ✘	<input type="radio"/> ✓	The arguments to system call are copied to kernel stack in trapasm.S
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	Integer arguments are copied from user memory to kernel memory using argint()
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	The arguments are accessed in the kernel code using esp on the trapframe.

Integer arguments are stored in eax, ebx, ecx, etc. registers: False

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

The functions like argint(), argstr() make the system call arguments available in the kernel.: True

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False

The arguments to system call originally reside on process stack.: True

The arguments to system call are copied to kernel stack in trapasm.S: False

Integer arguments are copied from user memory to kernel memory using argint(): True

The arguments are accessed in the kernel code using esp on the trapframe.: True

Question 23

Not answered

Marked out of 1.00

Given below is a sequence of reference bits on pages before the second chance algorithm runs. Before the algorithm runs, the counter is at the page marked (x). Write the sequence of reference bits after the second chance algorithm has executed once. In the answer write PRECISELY one space BETWEEN each number and do not mention (x).

0 0 1(x) 1 0 1 1

Answer:



The correct answer is: 0 0 0 0 0 1 1

Question 24

Correct

Mark 2.00 out of 2.00

For the reference string

3 4 3 5 2

using FIFO replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.

Select the correct statement.

Select one:

- a. Exhibit Balady's anomaly between 3 and 4 frames
- b. Do not exhibit Balady's anomaly
- c. Exhibit Balady's anomaly between 2 and 3 frames



Your answer is correct.

The correct answer is: Do not exhibit Balady's anomaly

Question 25

Correct

Mark 1.00 out of 1.00

For the reference string

3 4 3 5 2

using LRU replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.

Select the most correct statement.

Select one:

- a. LRU will never exhibit Balady's anomaly ✓
- b. Exhibit Balady's anomaly between 2 and 3 frames
- c. This example does not exhibit Balady's anomaly
- d. Exhibit Balady's anomaly between 3 and 4 frames

Your answer is correct.

The correct answer is: LRU will never exhibit Balady's anomaly

Question 26

Partially correct

Mark 0.55 out of 1.00

Select all the correct statements about process states.

Note that in this question you lose marks for every incorrect choice that you make, proportional to actual number of incorrect choices.

- a. Process state is implemented as a string
- b. Process state is stored in the PCB ✓
- c. A process becomes ZOMBIE when another process bites into its memory
- d. Process state is stored in the processor ✗
- e. The scheduler can change state of a process from RUNNABLE to RUNNING and vice-versa
- f. The scheduler can change state of a process from RUNNABLE to RUNNING ✓
- g. A process becomes ZOMBIE when it calls exit() ✓
- h. Process state is changed only by interrupt handlers
- i. Process state can be implemented as just a number

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: Process state is stored in the PCB, Process state can be implemented as just a number, The scheduler can change state of a process from RUNNABLE to RUNNING, A process becomes ZOMBIE when it calls exit()

Question 27

Partially correct

Mark 0.38 out of 1.00

Consider a demand-paging system with the following time-measured utilizations:

CPU utilization : 20%

Paging disk: 97.7%

Other I/O devices: 5%

For each of the following, indicate whether it will (or is likely to) improve CPU utilization (even if by a small amount). Explain your answers.

a. Install a faster CPU : Yes ✗

b. Install a bigger paging disk. : Yes ✗

c. Increase the degree of multiprogramming. : Yes ✗

d. Decrease the degree of multiprogramming. : Yes ✓

e. Install more main memory.: Yes ✓

f. Install a faster hard disk or multiple controllers with multiple hard disks. : Yes ✓

g. Add prepaging to the page-fetch algorithms. :

May be ✗

h. Increase the page size. : May be ✗

Question 28

Incorrect

Mark 0.00 out of 1.00

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

10011010

Now, there is a request for a chunk of 50 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: 10110010

✗

The correct answer is: 11111010

Question 29

Partially correct

Mark 0.75 out of 1.00

Select the correct points of comparison between POSIX and System V shared memory.

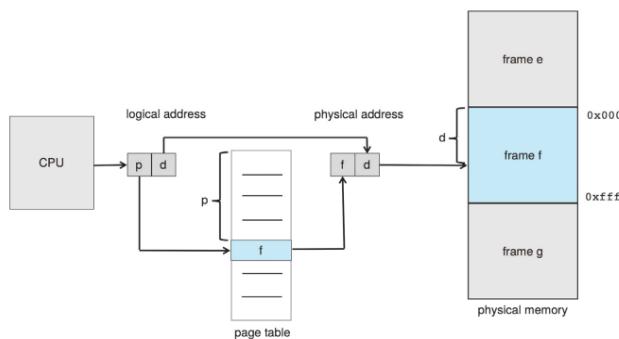
- a. POSIX shared memory is newer than System V shared memory ✓
- b. POSIX shared memory is "thread safe", System V is not ✓
- c. System V is more prevalent than POSIX even today ✓
- d. POSIX allows giving name to shared memory, System V does not

The correct answers are: POSIX shared memory is newer than System V shared memory, POSIX shared memory is "thread safe", System V is not, POSIX allows giving name to shared memory, System V does not, System V is more prevalent than POSIX even today

Question 30

Partially correct

Mark 0.67 out of 1.00

**Figure 9.8** Paging hardware.

Mark the statements as True or False, w.r.t. the above diagram (note that the diagram does not cover all details of what actually happens!)

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The combining of f and d is done by MMU
<input type="radio"/>	<input checked="" type="radio"/>	There are total 3 memory references in this diagram
<input checked="" type="radio"/>	<input type="radio"/>	The split of logical address into p and d is done by MMU
<input checked="" type="radio"/>	<input type="radio"/>	The page table is in physical memory and must be continuous
<input type="radio"/>	<input checked="" type="radio"/>	Using the offset d in the physical page-frame is done by MMU
<input checked="" type="radio"/>	<input type="radio"/>	The logical address issued by CPU is the same one generated by compiler

The combining of f and d is done by MMU: True

There are total 3 memory references in this diagram: False

The split of logical address into p and d is done by MMU: True

The page table is in physical memory and must be continuous: True

Using the offset d in the physical page-frame is done by MMU: False

The logical address issued by CPU is the same one generated by compiler: True

Question 31

Partially correct

Mark 0.50 out of 1.00

Select all the correct statements about signals

Select one or more:

- a. SIGKILL definitely kills a process because its code runs in kernel mode of CPU
- b. Signals are delivered to a process by another process ✗
- c. The signal handler code runs in kernel mode of CPU
- d. SIGKILL definitely kills a process because it can't be caught or ignored, and its default action terminates the process ✓
- e. The signal handler code runs in user mode of CPU ✓
- f. A signal handler can be invoked asynchronously or synchronously depending on signal type ✓
- g. Signal handlers once replaced can't be restored
- h. Signals are delivered to a process by kernel

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: Signals are delivered to a process by kernel, A signal handler can be invoked asynchronously or synchronously depending on signal type, The signal handler code runs in user mode of CPU, SIGKILL definitely kills a process because it can't be caught or ignored, and its default action terminates the process

Question 32

Correct

Mark 1.00 out of 1.00

The data structure used in kalloc() and kfree() in xv6 is

- a. Singly linked circular list
- b. Singly linked NULL terminated list ✓
- c. Double linked NULL terminated list
- d. Doubly linked circular list

Your answer is correct.

The correct answer is: Singly linked NULL terminated list

Question 33

Partially correct

Mark 1.78 out of 2.00

Match the description of a memory management function with the name of the function that provides it, in xv6

Load contents from ELF into existing pages	loaduvm()	✓
Mark the page as in-accessible	clearpteu()	✓
setup the kernel part in the page table	setupkvm()	✓
Switch to kernel page table	switchkvm()	✓
Create a copy of the page table of a process	copyuvm()	✓
Copy the code pages of a process	No such function	✓
Setup and load the user page table for initcode process	inituvm()	✓
Switch to user page table	switchuvm()	✓
Load contents from ELF into pages after allocating the pages first	inituvm()	✗

The correct answer is: Load contents from ELF into existing pages → loaduvm(), Mark the page as in-accessible → clearpteu(), setup the kernel part in the page table → setupkvm(), Switch to kernel page table → switchkvm(), Create a copy of the page table of a process → copyuvm(), Copy the code pages of a process → No such function, Setup and load the user page table for initcode process → inituvm(), Switch to user page table → switchuvm(), Load contents from ELF into pages after allocating the pages first → No such function

Question 34

Partially correct

Mark 0.60 out of 1.00

Mark the statements as True or False, w.r.t. thrashing

True	False	
<input checked="" type="radio"/> ✘	<input type="radio"/> ✓	Thrashing occurs because some process is doing lot of disk I/O.
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.
<input checked="" type="radio"/> ✘	<input checked="" type="radio"/> ✓	mmap() solves the problem of thrashing.
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	The working set model is an attempt at approximating the locality of a process.
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	Thrashing is particular to demand paging systems, and does not apply to pure paging systems.
<input checked="" type="radio"/> ✘	<input type="radio"/> ✓	Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.
<input checked="" type="radio"/> ✘	<input checked="" type="radio"/> ✓	Thrashing can occur even if entire memory is not in use.
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	During thrashing the CPU is under-utilised as most time is spent in I/O
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	Thrashing can be limited if local replacement is used.
<input checked="" type="radio"/> ✓	<input checked="" type="radio"/> ✘	Thrashing occurs when the total size of all processes's locality exceeds total memory size.

Thrashing occurs because some process is doing lot of disk I/O.: False

Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.: True

mmap() solves the problem of thrashing.: False

The working set model is an attempt at approximating the locality of a process.: True

Thrashing is particular to demand paging systems, and does not apply to pure paging systems.: True

Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.: False

Thrashing can occur even if entire memory is not in use.: False

During thrashing the CPU is under-utilised as most time is spent in I/O: True

Thrashing can be limited if local replacement is used.: True

Thrashing occurs when the total size of all processes's locality exceeds total memory size.: True

Question 35

Correct

Mark 1.00 out of 1.00

After virtual memory is implemented

(select T/F for each of the following) One Program's size can be larger than physical memory size

True	False	
<input checked="" type="radio"/>	<input type="radio"/> ✗	Cumulative size of all programs can be larger than physical memory size
<input checked="" type="radio"/>	<input type="radio"/> ✗	Code need not be completely in memory
<input checked="" type="radio"/>	<input type="radio"/> ✗	One Program's size can be larger than physical memory size
<input type="radio"/> ✗	<input checked="" type="radio"/>	Virtual addresses become available to executing process
<input type="radio"/> ✗	<input checked="" type="radio"/>	Virtual access to memory is granted to all processes
<input checked="" type="radio"/>	<input type="radio"/> ✗	Relatively less I/O may be possible during process execution
<input checked="" type="radio"/>	<input type="radio"/> ✗	Logical address space could be larger than physical address space

Cumulative size of all programs can be larger than physical memory size: True

Code need not be completely in memory: True

One Program's size can be larger than physical memory size: True

Virtual addresses become available to executing process: False

Virtual access to memory is granted to all processes: False

Relatively less I/O may be possible during process execution: True

Logical address space could be larger than physical address space: True

◀ (Optional Assignment) lseek system call in xv6

Jump to...

Feedback on Quiz-2 ►

Started on Thursday, 18 March 2021, 2:46 PM

State Finished

Completed on Thursday, 18 March 2021, 3:50 PM

Time taken 1 hour 4 mins

Grade 10.36 out of 20.00 (52%)

Question 1

Partially correct

Mark 0.57 out of 1.00

Mark True, the actions done as part of code of swtch() in swtch.S, in xv6

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Restore new callee saved registers from kernel stack of new context	
<input checked="" type="radio"/>	<input type="radio"/>	Save old callee saved registers on kernel stack of old context	
	<input checked="" type="radio"/>	Save old callee saved registers on user stack of old context	
	<input checked="" type="radio"/>	Switch from old process context to new process context	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Switch from one stack (old) to another(new)	
	<input checked="" type="radio"/>	Restore new callee saved registers from user stack of new context	
	<input checked="" type="radio"/>	Jump to code in new context	

Restore new callee saved registers from kernel stack of new context: True

Save old callee saved registers on kernel stack of old context: True

Save old callee saved registers on user stack of old context: False

Switch from old process context to new process context: False

Switch from one stack (old) to another(new): True

Restore new callee saved registers from user stack of new context: False

Jump to code in new context: False

Question 2

Partially correct

Mark 0.17 out of 0.50

For each function/code-point, select the status of segmentation setup in xv6

bootmain()	gdt setup with 3 entries, right from first line of code of bootloader	✗
kvmalloc() in main()	gdt setup with 5 entries (0 to 4) on one processor	✗
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	✓
after seginit() in main()	gdt setup with 5 entries (0 to 4) on all processors	✗
bootasm.S	gdt setup with 3 entries, right from first line of code of bootloader	✗
entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S

Question 3

Partially correct

Mark 0.38 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- a. The meaning of valid-invalid bit in page table is different in paging and demand-paging. ✓
- b. Demand paging requires additional hardware support, compared to paging. ✓
- c. Paging requires some hardware support in CPU
- d. With paging, it's possible to have user programs bigger than physical memory. ✗
- e. Both demand paging and paging support shared memory pages. ✓
- f. Demand paging always increases effective memory access time.
- g. With demand paging, it's possible to have user programs bigger than physical memory. ✓
- h. Calculations of number of bits for page number and offset are same in paging and demand paging. ✓
- i. TLB hit ration has zero impact in effective memory access time in demand paging.
- j. Paging requires NO hardware support in CPU

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question 4

Partially correct

Mark 0.44 out of 0.50

Suppose a processor supports base(relocation register) + limit scheme of MMU.

Assuming this, mark the statements as True/False

True	False	
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS may terminate the process while handling the interrupt of memory violation
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The hardware detects any memory access beyond the limit value and raises an interrupt
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/>	The hardware may terminate the process while handling the interrupt of memory violation
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS sets up the relocation and limit registers when the process is scheduled
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The process sets up its own relocation and limit registers when the process is scheduled
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The OS detects any memory access beyond the limit value and raises an interrupt
<input type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The compiler generates machine code assuming appropriately sized segments for code, data and stack.

The OS may terminate the process while handling the interrupt of memory violation: True

The hardware detects any memory access beyond the limit value and raises an interrupt: True

The hardware may terminate the process while handling the interrupt of memory violation: False

The OS sets up the relocation and limit registers when the process is scheduled: True

The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;: True

The process sets up its own relocation and limit registers when the process is scheduled: False

The OS detects any memory access beyond the limit value and raises an interrupt: False

The compiler generates machine code assuming appropriately sized segments for code, data and stack.: False

Question 5

Correct

Mark 0.50 out of 0.50

Consider the following list of free chunks, in continuous memory management:

10k, 25k, 12k, 7k, 9k, 13k

Suppose there is a request for chunk of size 9k, then the free chunk selected under each of the following schemes will be

Best fit:

9k



First fit:

10k



Worst fit:

25k

**Question 6**

Partially correct

Mark 0.50 out of 1.00

Select all the correct statements about MMU and its functionality

Select one or more:

- a. MMU is a separate chip outside the processor
- b. MMU is inside the processor ✓
- c. Logical to physical address translations in MMU are done with specific machine instructions
- d. The operating system interacts with MMU for every single address translation ✗
- e. Illegal memory access is detected in hardware by MMU and a trap is raised ✓
- f. The Operating system sets up relevant CPU registers to enable proper MMU translations
- g. Logical to physical address translations in MMU are done in hardware, automatically ✓
- h. Illegal memory access is detected by operating system

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

Question 7

Incorrect

Mark 0.00 out of 0.50

Assuming a 8- KB page size, what is the page numbers for the address 874815 reference in decimal :
(give answer also in decimal)

Answer: 2186



The correct answer is: 107

Question 8

Incorrect

Mark 0.00 out of 0.25

Select the compiler's view of the process's address space, for each of the following MMU schemes:
(Assume that each scheme,e.g. paging/segmentation/etc is effectively utilised)

Segmentation, then paging	Many continuous chunks each of page size	
Relocation + Limit	Many continuous chunks of same size	
Segmentation	one continuous chunk	
Paging	many continuous chunks of variable size	

Your answer is incorrect.

The correct answer is: Segmentation, then paging → many continuous chunks of variable size, Relocation + Limit → one continuous chunk, Segmentation → many continuous chunks of variable size, Paging → one continuous chunk

Question 9

Incorrect

Mark 0.00 out of 0.50

Suppose the memory access time is 180ns and TLB hit ratio is 0.3, then effective memory access time is (in nanoseconds);

Answer: 192



The correct answer is: 306.00

Question 10

Correct

Mark 0.50 out of 0.50

In xv6, The struct context is given as

```
struct context {
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};
```

Select all the reasons that explain why only these 5 registers are included in the struct context.

- a. The segment registers are same across all contexts, hence they need not be saved ✓
- b. esp is not saved in context, because context{} is on stack and it's address is always argument to swtch() ✓
- c. xv6 tries to minimize the size of context to save memory space
- d. esp is not saved in context, because it's not part of the context
- e. eax, ecx, edx are caller save, hence no need to save ✓

Your answer is correct.

The correct answers are: The segment registers are same across all contexts, hence they need not be saved, eax, ecx, edx are caller save, hence no need to save, esp is not saved in context, because context{} is on stack and it's address is always argument to swtch()

Question 11

Partially correct

Mark 0.83 out of 1.50

Arrange the following events in order, in page fault handling:

Disk interrupt wakes up the process

7	✓
---	---

The reference bit is found to be invalid by MMU

1	✓
---	---

OS makes available an empty frame

6	✗
---	---

Restart the instruction that caused the page fault

9	✓
---	---

A hardware interrupt is issued

3	✗
---	---

OS schedules a disk read for the page (from backing store)

5	✓
---	---

Process is kept in wait state

4	✗
---	---

Page tables are updated for the process

8	✓
---	---

Operating system decides that the page was not in memory

2	✗
---	---

Your answer is partially correct.

You have correctly selected 5.

The correct answer is: Disk interrupt wakes up the process → 7, The reference bit is found to be invalid by MMU → 1, OS makes available an empty frame → 4, Restart the instruction that caused the page fault → 9, A hardware interrupt is issued → 2, OS schedules a disk read for the page (from backing store) → 5, Process is kept in wait state → 6, Page tables are updated for the process → 8, Operating system decides that the page was not in memory → 3

Question 12

Incorrect

Mark 0.00 out of 0.50

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

00001010

Now, there is a request for a chunk of 70 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: 11101010



The correct answer is: 11111010

Question 13

Incorrect

Mark 0.00 out of 0.25

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- a. end, 4MB
- b. P2V(end), P2V(PHYSTOP)
- c. end, P2V(4MB + PHYSTOP)
- d. P2V(end), PHYSTOP ✗
- e. end, (4MB + PHYSTOP)
- f. end, PHYSTOP
- g. end, P2V(PHYSTOP)

Your answer is incorrect.

The correct answer is: end, P2V(PHYSTOP)

Question 14

Partially correct

Mark 0.33 out of 0.50

Match the pair

Hashed page table	Linear search on collision done by OS (e.g. SPARC Solaris) typically	✓
Inverted Page table	Linear/Parallel search using frame number in page table	✗
Hierarchical Paging	More memory access time per hierarchy	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Hashed page table → Linear search on collision done by OS (e.g. SPARC Solaris) typically, Inverted Page table → Linear/Parallel search using page number in page table, Hierarchical Paging → More memory access time per hierarchy

Question 15

Partially correct

Mark 0.29 out of 0.50

After virtual memory is implemented

(select T/F for each of the following) One Program's size can be larger than physical memory size

True	False	
<input checked="" type="radio"/>	<input type="radio"/> ✗	Code need not be completely in memory
<input checked="" type="radio"/>	<input type="radio"/> ✗	Cumulative size of all programs can be larger than physical memory size
<input type="radio"/> ✗	<input checked="" type="radio"/>	Virtual access to memory is granted
<input checked="" type="radio"/>	<input type="radio"/> ✗	Logical address space could be larger than physical address space
<input type="radio"/> ✗	<input checked="" type="radio"/>	Virtual addresses are available
<input checked="" type="radio"/>	<input checked="" type="radio"/> ✗	Relatively less I/O may be possible during process execution
<input checked="" type="radio"/>	<input type="radio"/> ✗	One Program's size can be larger than physical memory size

Code need not be completely in memory: True

Cumulative size of all programs can be larger than physical memory size: True

Virtual access to memory is granted: False

Logical address space could be larger than physical address space: True

Virtual addresses are available: False

Relatively less I/O may be possible during process execution: True

One Program's size can be larger than physical memory size: True

Question 16

Partially correct

Mark 0.64 out of 1.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only
Mark statements True or False**True****False**

<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The stack allocated in entry.S is used as stack for scheduler's context for first processor	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The free page-frame are created out of nearly 222 MB	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The kernel code and data take up less than 2 MB space	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	PHYSTOP can be increased to some extent, simply by editing memlayout.h	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	xv6 uses physical memory upto 224 MB only	✗
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The process's address space gets mapped on frames, obtained from ~2MB:224MB range	✓
<input checked="" type="radio"/>	<input type="radio"/> <input checked="" type="checkbox"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context	✗

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

The free page-frame are created out of nearly 222 MB: True

The kernel code and data take up less than 2 MB space: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

xv6 uses physical memory upto 224 MB only: True

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

Question 17

Incorrect

Mark 0.00 out of 1.50

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using LRU replacement is:

Answer: ✖

#6# 6,4# 6,4,2 # 0,4,2#0,1,2#6,1,2#6,9,2#0,9,2#0,5,2

The correct answer is: 9

Question 18

Partially correct

Mark 0.31 out of 0.50

Consider the image given below, which explains how paging works.

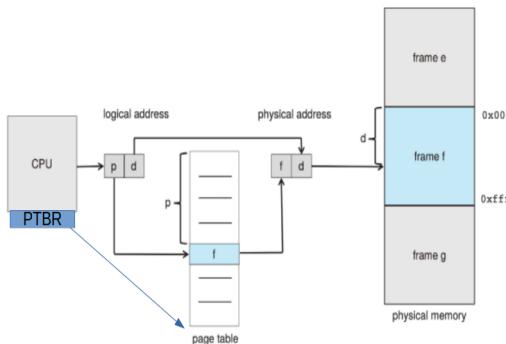


Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The PTBR is present in the CPU as a register
<input type="radio"/>	<input checked="" type="radio"/>	The page table is indexed using frame number
<input checked="" type="radio"/>	<input type="radio"/>	The page table is indexed using page number
<input type="radio"/>	<input checked="" type="radio"/>	The locating of the page table using PTBR also involves paging translation
<input type="radio"/>	<input checked="" type="radio"/>	Size of page table is always determined by the size of RAM
<input checked="" type="radio"/>	<input type="radio"/>	The page table is itself present in Physical memory
<input checked="" type="radio"/>	<input type="radio"/>	Maximum Size of page table is determined by number of bits used for page number
<input checked="" type="radio"/>	<input type="radio"/>	The physical address may not be of the same size (in bits) as the logical address

The PTBR is present in the CPU as a register: True

The page table is indexed using frame number: False

The page table is indexed using page number: True

The locating of the page table using PTBR also involves paging translation: False

Size of page table is always determined by the size of RAM: False

The page table is itself present in Physical memory: True

Maximum Size of page table is determined by number of bits used for page number: True

The physical address may not be of the same size (in bits) as the logical address: True

Question 19

Correct

Mark 2.00 out of 2.00

Given below is shared memory code with two processes sharing a memory segment.

The first process sends a user input string to second process. The second capitalizes the string. Then the first process prints the capitalized version.

Fill in the blanks to complete the code.

// First process

```
#define SHMSZ 27

int main()
{
    char c;
    int shmid;
    key_t key;
    char *shm, *s, string[128];
    key = 5679;
    if ((shmid =
        shmget
        ✓ (key, SHMSZ, IPC_CREAT | 0666)) < 0) {
        perror("shmget");
        exit(1);
    }
    if ((shm =
        shmat
        ✓ (shmid, NULL, 0)) == (char *) -1) {
        perror("shmat");
        exit(1);
    }
    s = shm;
    *s = '$';
    scanf("%s", string);
    strcpy(s + 1, string);
    *s =
        @
        ✓ ';' //note the quotes
    while(*s != '
        $
        ')
        sleep(1);
        printf("%s\n", s + 1);
        exit(0);
}
```

//Second process

```
#define SHMSZ 27

int main()
{
    int shmid;
    key_t key;
    char *shm, *s;
    int i;
    char string[128];
    key =
        5679
```

```

✓ ;
if ((shmid = shmget(key, SHMSZ, 0666)) < 0) {
    perror("shmget");
    exit(1);
}
if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
    perror("shmat");
    exit(1);
}
s =

✓ ;
while(*s != '@')
    sleep(1);
for(i = 0; i < strlen(s + 1); i++)
    s[i + 1] = toupper(s[i + 1]);
*s = '$';
exit(0);
}

```

Question 20

Partially correct

Mark 0.25 out of 0.50

Map the functionality/use with function/variable in xv6 code.

return a free page, if available; 0, otherwise

Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed

Array listing the kernel memory mappings, to be used by setupkvm()

Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

Setup kernel part of a page table, and switch to that page table

 kinit1()

 mappages()

 kmap[]

 kvmalloc()

 walkpgdir()

 setupkvm()

Your answer is partially correct.

You have correctly selected 3.

The correct answer is: return a free page, if available; 0, otherwise → kalloc(), Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → mappages(), Array listing the kernel memory mappings, to be used by setupkvm() → kmap[], Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → setupkvm(), Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → walkpgdir(), Setup kernel part of a page table, and switch to that page table → kvmalloc()

Question 21

Partially correct

Mark 1.53 out of 2.50

Order events in xv6 timer interrupt code

(Transition from process P1 to P2's code.)

P2 is selected and marked RUNNING

12 ✓

Change of stack from user stack to kernel stack of P1

3 ✓

Timer interrupt occurs

2 ✓

alltraps() will call iret

17 ✗

change to context of P2, P2's kernel stack in use now

13 ✓

P2's trap() will return to alltraps

16 ✗

jump in vector.S

4 ✓

P2 will return from sched() in yield()

14 ✗

yield() is called

8 ✓

trap() is called

7 ✓

Process P2 is executing

18 ✗

P1 is marked as RUNNABLE

9 ✓

P2's yield() will return in trap()

15 ✗

Process P1 is executing

1 ✓

sched() is called,

11 ✗

change to context of the scheduler, scheduler's stack in use now

10 ✗

jump to alltraps

5 ✓

Trapframe is built on kernel stack of P1

6 ✓

Your answer is partially correct.

You have correctly selected 11.

The correct answer is: P2 is selected and marked RUNNING → 12, Change of stack from user stack to kernel stack of P1 → 3, Timer interrupt occurs → 2, alltraps() will call iret → 18, change to context of P2, P2's kernel stack in use now → 13, P2's trap() will return to alltraps → 17, jump in vector.S → 4, P2 will return from sched() in yield() → 15, yield() is called → 8, trap() is called → 7, Process P2 is executing → 14, P1 is marked as RUNNABLE → 9, P2's yield() will return in trap() → 16, Process P1 is executing → 1, sched() is called, → 10, change to context of the scheduler, scheduler's stack in use now → 11, jump to alltraps → 5, Trapframe is built on kernel stack of P1 → 6

Question 22

Incorrect

Mark 0.00 out of 1.00

Given that the memory access time is 200 ns, probability of a page fault is 0.7 and page fault handling time is 8 ms,
The effective memory access time in nanoseconds is:

Answer: ✖

The correct answer is: 5600060.00

Question 23

Correct

Mark 0.25 out of 0.25

Select the state that is not possible after the given state, for a process:

- New: Running ✓
- Ready : Waiting ✓
- Running: None of these ✓
- Waiting: Running ✓

Question 24

Partially correct

Mark 0.63 out of 1.00

Select the correct statements about sched() and scheduler() in xv6 code

- a. scheduler() switches to the selected process's context ✓
- b. When either sched() or scheduler() is called, it does not return immediately to caller ✓
- c. After call to swtch() in sched(), the control moves to code in scheduler()
- d. Each call to sched() or scheduler() involves change of one stack inside swtch() ✓
- e. After call to swtch() in scheduler(), the control moves to code in sched()
- f. When either sched() or scheduler() is called, it results in a context switch ✓
- g. sched() switches to the scheduler's context ✓
- h. sched() and scheduler() are co-routines

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: sched() and scheduler() are co-routines, When either sched() or scheduler() is called, it does not return immediately to caller, When either sched() or scheduler() is called, it results in a context switch, sched() switches to the scheduler's context, scheduler() switches to the selected process's context, After call to swtch() in scheduler(), the control moves to code in sched(), After call to swtch() in sched(), the control moves to code in scheduler(), Each call to sched() or scheduler() involves change of one stack inside swtch()

Question 25

Correct

Mark 0.25 out of 0.25

The data structure used in kalloc() and kfree() in xv6 is

- a. Doubly linked circular list
- b. Singly linked circular list
- c. Double linked NULL terminated list
- d. Singly linked NULL terminated list



Your answer is correct.

The correct answer is: Singly linked NULL terminated list

[◀ \(Assignment\) lseek system call in xv6](#)

Jump to...

Dashboard / My courses / Computer Engineering & IT / CEIT-Even-sem-20-21 / QS-Even-sem-2020-21 / 16 May - 22 May / End Sem Exam OS-2021

Started on Saturday, 22 May 2021, 8:00 AM

State Finished

Completed on Saturday, 22 May 2021, 9:30 AM

Time taken 1 hour 30 mins

Grade 26.12 out of 40.00 (65%)

Question 1

Incorrect

Mark 0.00 out of 1.00

A 4 GB disk with 1 KB of block size would require these many number of **blocks** for its free block bitmap:

Answer: 4096 ✖

The correct answer is: 512

Question 2

Correct

Mark 1.00 out of 1.00

Given that the memory access time is 110 ns, probability of a page fault is 0.5 and page fault handling time is 12 ms,

The effective memory access time in nanoseconds is:

Answer: 6000165 ✓

The correct answer is: 6000055.00

Question 3

Incorrect

Mark 0.00 out of 1.00

The maximum size of a file in number of blocks of BSIZE in xv6 code is

(write a number only)

Answer: 268 ✖

The correct answer is: 138

Question 4

Incorrect

Mark 0.00 out of 1.00

Calculate the average waiting time using

Round Robin scheduling with time quantum of 5 time units
for the following workload

assuming that they arrive in the order written below.

Process Burst Time

P1	5
P2	7
P3	6
P4	2

Write only a number in the answer upto two decimal points.

Answer: 40.75 ✖

The correct answer is: 10.25



Question 5

Correct

Mark 1.00 out of 1.00

For the reference string

4 2 5 1 0 1 2 5 4 1 2

the number of page faults, including initial ones,
with FIFO replacement and 2 frames are :

Answer: 10 ✓

4 -

4 2

5 2

5 1

0 1

-

2 1

2 5

4 5

4 1

2 1

The correct answer is: 10

Question 6

Correct

Mark 1.00 out of 1.00

Assuming a 16- KB page size, what is the page number for the address 428517 reference in decimal :

(give answer also in decimal)

Answer: 27 ✓

The correct answer is: 26



Question 7

Correct

Mark 1.00 out of 1.00

In the code below assume that each function can be executed concurrently by many threads/processes.
Ignore syntactical issues, and focus on the semantics.

This program is an example of

```
spinlock a, b; // assume initialized
thread1() {
    spinlock(b);
    //some code;
    spinlock(a);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
thread2() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
```

- a. Deadlock ✓
- b. Self Deadlock
- c. None of these
- d. Deadlock or livelock depending on actual race
- e. Livelock

Your answer is correct.

The correct answer is: Deadlock



Question 8

Partially correct

Mark 1.33 out of 2.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.
"..." means some code.

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

    // The + in "+m" denotes a read-modify-write operand.
    asm volatile("lock; xchgl %0, %1" :
        "+m" ("addr"), "=a" (result) :
        "1" (newval) :
        "cc");
    return result;
}
```

Atomic compare and swap instruction (to be expanded inline into code)



```
void
sleep(void *chan, struct spinlock *lk)
{
    ...
    if(lk != &ptable.lock){
        acquire(&ptable.lock);
        release(lk);
    }
}
```

If you don't do this, a process may be running on two processors parallelly



```
void
acquire(struct spinlock *lk)
{
    ...
    __sync_synchronize();
}
```

Tell compiler not to reorder memory access beyond this line



Your answer is partially correct.

You have correctly selected 2.

The correct answer is: static inline uint
xchg(volatile uint *addr, uint newval)

```
{
    uint result;
```

```
// The + in "+m" denotes a read-modify-write operand.
asm volatile("lock; xchgl %0, %1" :
    "+m" ("addr"), "=a" (result) :
    "1" (newval) :
    "cc");
return result;
} → Atomic compare and swap instruction (to be expanded inline into code), void
sleep(void *chan, struct spinlock *lk)
{
    ...
    if(lk != &ptable.lock){
        acquire(&ptable.lock);
        release(lk);
    } → Avoid a self-deadlock, void
    acquire(struct spinlock *lk)
{
    ...
    __sync_synchronize(); → Tell compiler not to reorder memory access beyond this line
```

Question 9

Correct

Mark 1.00 out of 1.00

Predict the output of the program given here.

Assume that all the path names for the programs are correct. For example "/usr/bin/echo" will actually run echo command.

Assume that there is no mixing of print output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good

output

should be written as good output

--

```
main() {  
    int i;  
    i = fork();  
    if(i == 0)  
        execl("/usr/bin/echo", "/usr/bin/echo", "hi", 0);  
    else  
        wait(0);  
    fork();  
    execl("/usr/bin/echo", "/usr/bin/echo", "one", 0);  
}
```

Answer: hi one one 

The correct answer is: hi one one

Question 10

Partially correct

Mark 1.67 out of 2.00

Select all the blocks that may need to be written back to disk (if updated, of-course), as "Yes", when an operation of deleting a file is carried out on ext2 file system.

An option has to be correct entirely to be marked "Yes"

Superblock

Yes 

One or multiple data blocks of the parent directory

No 

One or more data bitmap blocks for the parent directory

No 

Block bitmap(s) for all the blocks of the file

No 

Possibly one block bitmap corresponding to the parent directory

Yes 

Data blocks of the file

No 

Your answer is partially correct.

only one data block of parent directory. multiple blocks not possible. an entry is always contained within one single block

You have correctly selected 5.

The correct answer is: Superblock → Yes, One or multiple data blocks of the parent directory → No, One or more data bitmap blocks for the parent directory → No, Block bitmap(s) for all the blocks of the file → Yes, Possibly one block bitmap corresponding to the parent directory → Yes, Data blocks of the file → No

Question 11

Correct

Mark 1.00 out of 1.00

Select all the correct statements about bootloader.

Every wrong selection will deduct marks proportional to $1/n$ where n is total wrong choices in the question.

You will get minimum a zero.

- a. Modern Bootloaders often allow configuring the way an OS boots
- b. Bootloaders allow selection of OS to boot from
- c. Bootloader must be one sector in length
- d. The bootloader loads the BIOS
- e. LILO is a bootloader



Your answer is correct.

The correct answers are: LILO is a bootloader, Modern Bootloaders often allow configuring the way an OS boots, Bootloaders allow selection of OS to boot from



Question 12

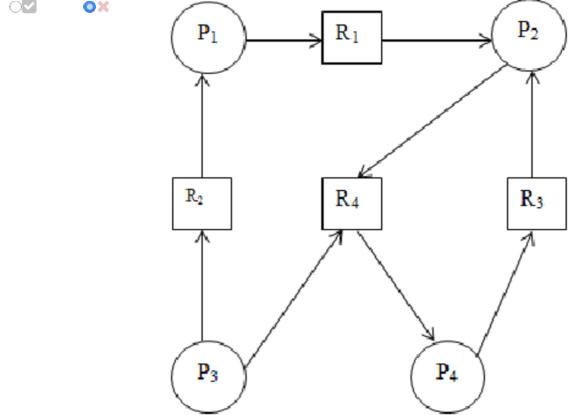
Incorrect

Mark 0.00 out of 1.00

For each of the resource allocation diagram shown,
infer whether the graph contains at least one deadlock or not.

Yes

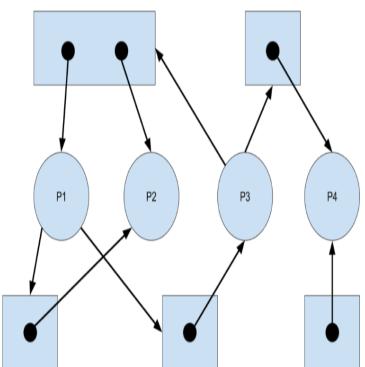
No



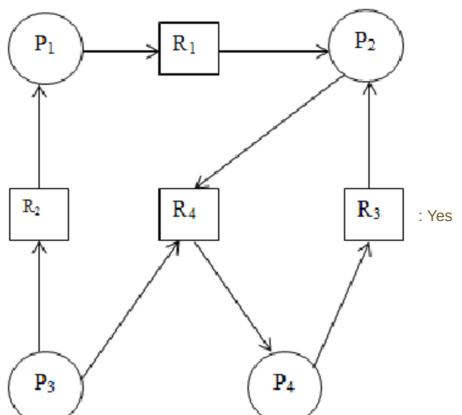
✗

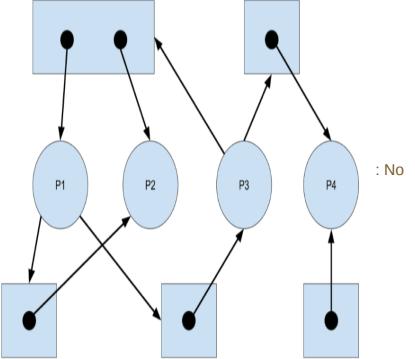
✗

✓



✗



**Question 13**

Partially correct

Mark 0.71 out of 1.00

Mark the statements about device drivers by marking as True or False.

True	False	
<input checked="" type="radio"/>	<input type="radio"/> ✘	It's possible that a particular hardware has multiple device drivers available for it.
<input checked="" type="radio"/>	<input type="radio"/> ✘	xv6 has device drivers for IDE disk and console.
<input checked="" type="radio"/>	<input type="radio"/> ✘	A disk driver converts OS's logical view of disk into physical locations on disk.
<input checked="" type="radio"/>	<input type="radio"/> ✘	A device driver code is specific to a hardware device
<input checked="" type="radio"/>	<input type="radio"/> ✘	All devices of the same type (e.g. 2 hard disks) can typically use the same device driver
<input checked="" type="radio"/>	<input type="radio"/> ✘	Writing a device driver mandatorily demands reading the technical documentation about the hardware.
<input type="radio"/> ✘	<input checked="" type="radio"/>	Device driver is an intermediary between the end-user and OS

It's possible that a particular hardware has multiple device drivers available for it.: True

xv6 has device drivers for IDE disk and console.: True

A disk driver converts OS's logical view of disk into physical locations on disk.: True

A device driver code is specific to a hardware device: True

All devices of the same type (e.g. 2 hard disks) can typically use the same device driver: True

Writing a device driver mandatorily demands reading the technical documentation about the hardware.: True

Device driver is an intermediary between the end-user and OS: False

Question 14

Partially correct

Mark 0.33 out of 1.00

Consider this program.

Some statements are identified using the // comment at the end.

Assume that `=` is an atomic operation.

```
#include <stdio.h>
#include <pthread.h>
long c = 0, c1 = 0, c2 = 0, run = 1;
void *thread1(void *arg) {
    while(run == 1) { //E
        c = 10; //A
        c1 = c2 + 5; //B
    }
}
void *thread2(void *arg) {
    while(run == 1) { //F
        c = 20; //C
        c2 = c1 + 3; //D
    }
}
int main() {
    pthread_t th1, th2;
    pthread_create(&th1, NULL, thread1, NULL);
    pthread_create(&th2, NULL, thread2, NULL);
    sleep(2);
    run = 0;
    printf(stdout, "c = %ld c1+c2 = %ld c1 = %ld c2 = %ld \n", c, c1+c2, c1, c2);
    fflush(stdout);
}
```

Which statements are part of the critical Section?

Yes	No	
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	F
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/> <input type="checkbox"/>	D
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	C
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	A
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/> <input type="checkbox"/>	B
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	E

F: No

D: Yes

C: No

A: No

B: Yes

E: No

Question 15

Partially correct

Mark 1.43 out of 2.00

Mark statements as T/F

All statements are in the context of preventing deadlocks.

True**False**

<input checked="" type="radio"/>	<input type="radio"/>	A process holding one resources and waiting for just one more resource can also be involved in a deadlock.	✓
<input type="radio"/>	<input checked="" type="radio"/>	If a resource allocation graph contains a cycle then there is a guarantee of a deadlock	✗
<input type="radio"/>	<input checked="" type="radio"/>	The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order	✗
<input checked="" type="radio"/>	<input type="radio"/>	Circular wait is avoided by enforcing a lock ordering	✓
<input checked="" type="radio"/>	<input type="radio"/>	Hold and wait means a thread/process holding some locks and waiting for acquiring some.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens	✓

A process holding one resources and waiting for just one more resource can also be involved in a deadlock.: True

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

Circular wait is avoided by enforcing a lock ordering: True

Hold and wait means a thread/process holding some locks and waiting for acquiring some.: True

Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

Question 16

Correct

Mark 1.00 out of 1.00

Match the left side use(or non-use) of a synchronization primitive with the best option on the right side.

This is the smallest primitive made available in software, using the hardware provided atomic instructions

 spinlock ✓

This tool is useful for event-wait scenarios

 semaphore ✓

This tool is more useful on multiprocessor systems

 spinlock ✓

This tool is quite attractive in solving the main bounded buffer problem

 semaphore ✓

This tool is very useful for waiting for 'something'

 condition variables ✓

Your answer is correct.

The correct answer is: This is the smallest primitive made available in software, using the hardware provided atomic instructions → spinlock, This tool is useful for event-wait scenarios → semaphore, This tool is more useful on multiprocessor systems → spinlock, This tool is quite attractive in solving the main bounded buffer problem → semaphore, This tool is very useful for waiting for 'something' → condition variables

Question 17

Correct

Mark 1.00 out of 1.00

The permissions -rwx--x--x on a file mean

- a. The file can be read only by the owner
- b. 'cat' on the file by owner will not work
- c. 'cat' on the file by any user will work
- d. 'rm' on the file by any user will work
- e. The file can be executed by anyone
- f. The file can be written only by the owner



Your answer is correct.

The correct answers are: The file can be executed by anyone, The file can be read only by the owner, The file can be written only by the owner, 'rm' on the file by any user will work

Question 18

Incorrect

Mark 0.00 out of 1.00

Note: for this question you get full marks if you select all and only correct options, you get ZERO if at least one option is wrong or not selected.

Select all the correct statements about log structured file systems.

- a. a transaction is said to be committed when all operations are written to file system
- b. log may be kept on same block device or another block device
- c. file system recovery may end up losing data
- d. even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery
- e. file system recovery recovers all the lost data



Your answer is incorrect.

The correct answers are: file system recovery may end up losing data, log may be kept on same block device or another block device, even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery

Question 19

Incorrect

Mark 0.00 out of 1.00

Consider the structure of directory entry in ext2, as shown in this diagram.

	inode	rec_len	file_type	name_len	name
0	21	12	1	2	.
12	22	12	2	2	.
24	53	16	5	2	h o m e
40	67	28	3	2	u s r
52	0	16	7	1	o l d f i l e
68	34	12	4	2	s b i n

Select the correct statements about the directory entry in ext2 file system.

The correct formula for rec_len is (when entries are continuously stored)

- a. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + (-1) * (\text{strlen(name)} \% 4))$
- b. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + (\text{strlen(name)} - 4) \% 4)$
- c. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + 4 - (\text{strlen(name)} \% 4))$ ✗
- d. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + (-1) * (\text{strlen(name)} - 4))$
- e. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} \% 4)$
- f. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + \text{strlen(name)}$

Your answer is incorrect.

The correct answer is: $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + (-1) * (\text{strlen(name)} - 4))$

Question 20

Partially correct

Mark 0.50 out of 1.00

Mark whether the given sequence of events is possible or not-possible. Also, select the reason for your answer.

For each sequence it's a not-possible sequence if some important event is not mentioned in the sequence.

Assume that the kernel code is non-interruptible and uniprocessor system.

Process P1 executing a system call
 Timer interrupt
 Generic interrupt handler runs
 Scheduler runs
 Scheduler selects P2 for execution
 P2 returns from timer interrupt handler
 Process p2, user code executing

This sequence of events is: ✓

Because

✗

Question 21

Incorrect

Mark 0.00 out of 1.00

The given semaphore implementation faces which problem?

Assume any suitable code for signal()

Note: blocks means waits in a wait queue.

```
struct semaphore {  
    int val;  
    spinlock lk;  
};  
sem_init(semaphore *s, int initval) {  
    s->val = initval;  
    s->sl = 0;  
}  
wait(semaphore *s) {  
    spinlock(&(s->sl));  
    while(s->val <=0)  
        ;  
    (s->val)--;  
    spinunlock(&(s->sl));  
}
```

- a. blocks holding a spinlock
- b. deadlock
- c. too much spinning, bounded wait not guaranteed
- d. not holding lock after unblock



Your answer is incorrect.

The correct answer is: deadlock



Question 22

Partially correct

Mark 0.80 out of 1.00

Mark statements True/False w.r.t. change of states of a process.

Reference: The process state diagram (and your understanding of how kernel code works)

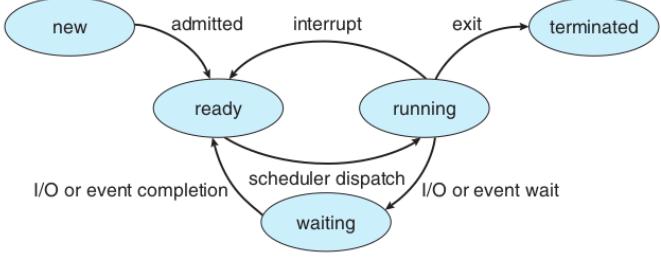


Figure 3.2 Diagram of process state.

True

False

<input type="radio"/> ✗	<input checked="" type="checkbox"/> ✗	A process in RUNNING state only can become TERMINATED because scheduler moves it to ZOMBIE state	✓
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.	✗
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred	✓
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	Every process has to go through ZOMBIE state, at least for a small duration.	✓
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	Only a process in READY state is considered by scheduler	✓

A process in RUNNING state only can become TERMINATED because scheduler moves it to ZOMBIE state: False

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred: True

Every process has to go through ZOMBIE state, at least for a small duration.: True

Only a process in READY state is considered by scheduler: True

Question 23

Correct

Mark 1.00 out of 1.00

Select T/F for statements about Volume Managers.

Do pay attention to the use of the words physical partition and physical volume.

True**False**

<input checked="" type="radio"/>	<input type="radio"/> ✗	The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A logical volume can be extended in size but upto the size of volume group	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A logical volume may span across multiple physical volumes	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	The volume manager stores additional metadata on the physical disk partitions	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A physical partition should be initialized as a physical volume, before it can be used by volume manager.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A volume group consists of multiple physical volumes	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A logical volume may span across multiple physical partitions	<input checked="" type="checkbox"/> since a physical volume is made up of physical partitions, and a volume can span across multiple PVs, it can also span across multiple PP

The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.: True

A logical volume can be extended in size but upto the size of volume group: True

A logical volume may span across multiple physical volumes: True

The volume manager stores additional metadata on the physical disk partitions: True

A physical partition should be initialized as a physical volume, before it can be used by volume manager.: True

A volume group consists of multiple physical volumes: True

A logical volume may span across multiple physical partitions: True

Question 24

Correct

Mark 1.00 out of 1.00

Map the block allocation scheme with the problem it suffers from

(Match pairs 1-1, match a scheme with the problem that it suffers from relatively the most, compared to others)

Continuous allocation	need for compaction	<input checked="" type="checkbox"/>
Linked allocation	Too many seeks	<input checked="" type="checkbox"/>
Indexed Allocation	Overhead of reading metadata blocks	<input checked="" type="checkbox"/>

Your answer is correct.

The correct answer is: Continuous allocation → need for compaction, Linked allocation → Too many seeks, Indexed Allocation → Overhead of reading metadata blocks

Question 25

Correct

Mark 1.00 out of 1.00

This one is not a system call:

- a. open
- b. read
- c. write
- d. scheduler



Your answer is correct.

The correct answer is: scheduler



Question 26

Correct

Mark 1.00 out of 1.00

Match the pairs.

This question is based on your general knowledge about operating systems/related concepts and their features.

Java threads	monitors,re-entrant locks, semaphores	✓
Linux threads	atomic-instructions, spinlocks, etc.	✓
POSIX threads	semaphore, mutex, condition variables	✓

Your answer is correct.

The correct answer is: Java threads → monitors,re-entrant locks, semaphores, Linux threads → atomic-instructions, spinlocks, etc., POSIX threads → semaphore, mutex, condition variables

Question 27

Correct

Mark 1.00 out of 1.00

Consider the following list of free chunks, in continuous memory management:

7k, 15k, 21k, 14k, 19k, 6k

Suppose there is a request for chunk of size 5k, then the free chunk selected under each of the following schemes will be

Best fit:	6k	✓
First fit:	7k	✓
Worst fit:	21k	✓

Question 28

Correct

Mark 1.00 out of 1.00

This one is not a scheduling algorithm

- a. Round Robin
- b. SJF
- c. Mergesort
- d. FCFS



Your answer is correct.

The correct answer is: Mergesort

Question 29

Correct

Mark 1.00 out of 1.00

Mark whether the concept is related to scheduling or not.

Yes	No	
<input checked="" type="radio"/>	<input type="radio"/>	timer interrupt
<input checked="" type="radio"/>	<input type="radio"/>	context-switch
<input checked="" type="radio"/>	<input type="radio"/>	ready-queue
<input type="radio"/>	<input checked="" type="radio"/>	file-table
<input checked="" type="radio"/>	<input type="radio"/>	runnable process

timer interrupt: Yes

context-switch: Yes

ready-queue: Yes

file-table: No

runnable process: Yes



Question 30

Partially correct

Mark 1.00 out of 2.00

Map ext2 data structure features with their purpose

Many copies of Superblock Choose...**Free blocks count in superblock and group descriptor**

Redundancy to ensure the most crucial data structure is not lost

**Used directories count in group descriptor**

is redundant and helps do calculations of directory entries faster

**Combining file type and access rights in one variable**

saves 1 byte of space

**rec_len field in directory entry**

Try to keep all the data of a directory and its file close together in a group

**File Name is padded**

aligns all memory accesses on word boundary, improving performance

**Inode bitmap is one block**

limits total number of files that can belong to a group

**Block bitmap is one block**

Limits the size of a block group, thus improvising on purpose of a group

**Mount count in superblock**

to enforce file check after certain amount of mounts at boot time

**Inode table location in Group Descriptor**

is redundant and helps do calculations of directory entries faster

**Inode table**

All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk

**A group**

Redundancy to ensure the most crucial data structure is not lost



Your answer is partially correct.

You have correctly selected 6.

The correct answer is: **Many copies of Superblock** → Redundancy to ensure the most crucial data structure is not lost, **Free blocks count in superblock and group descriptor** → Redundancy to help fsck restore consistency, **Used directories count in group descriptor** → attempt is made to evenly spread the first-level directories, this count is used there, **Combining file type and access rights in one variable** → saves 1 byte of space, **rec_len field in directory entry** → allows holes and linking of entries in directory, File Name is padded → aligns all memory accesses on word boundary, improving performance, **Inode bitmap is one block** → limits total number of files that can belong to a group, **Block bitmap is one block** → Limits the size of a block group, thus improvising on purpose of a group, **Mount count in superblock** → to enforce file check after certain amount of mounts at boot time, **Inode table location in Group Descriptor** → Obvious, as it's per group and not per file-system, **Inode table** → All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk, **A group** → Try to keep all the data of a directory and its file close together in a group

Question 31

Partially correct

Mark 1.85 out of 2.00

Mark True/False

Statements about scheduling and scheduling algorithms

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The nice() system call is used to set priorities for processes
<input checked="" type="radio"/>	<input type="radio"/>	Aging is used to ensure that low-priority processes do not starve in priority scheduling.
<input type="radio"/>	<input checked="" type="radio"/>	In non-pre-emptive priority scheduling, the highest priority process is scheduled and runs until it gives up CPU.
<input checked="" type="radio"/>	<input type="radio"/>	xv6 code does not care about Processor Affinity
<input checked="" type="radio"/>	<input type="radio"/>	In pre-emptive priority scheduling, priority is implemented by assigning more time quantum to higher priority process.
<input checked="" type="radio"/>	<input type="radio"/>	A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.
<input checked="" type="radio"/>	<input type="radio"/>	Processor Affinity refers to memory accesses of a process being stored on cache of that processor
<input checked="" type="radio"/>	<input type="radio"/>	Response time will be quite poor on non-interruptible kernels
<input checked="" type="radio"/>	<input type="radio"/>	Shortest Remaining Time First algorithm is nothing but pre-emptive Shortest Job First algorithm
<input checked="" type="radio"/>	<input type="radio"/>	On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread
<input checked="" type="radio"/>	<input type="radio"/>	Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.
<input checked="" type="radio"/>	<input type="radio"/>	Pre-emptive scheduling leads to many race conditions in kernel code.
<input checked="" type="radio"/>	<input type="radio"/>	Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.

The nice() system call is used to set priorities for processes.: True

Aging is used to ensure that low-priority processes do not starve in priority scheduling.: True

In non-pre-emptive priority scheduling, the highest priority process is scheduled and runs until it gives up CPU.: True

xv6 code does not care about Processor Affinity: True

In pre-emptive priority scheduling, priority is implemented by assigning more time quantum to higher priority process.: True

A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.: True

Processor Affinity refers to memory accesses of a process being stored on cache of that processor: True

Response time will be quite poor on non-interruptible kernels: True

Shortest Remaining Time First algorithm is nothing but pre-emptive Shortest Job First algorithm: True

On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread: True

Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.: True

Pre-emptive scheduling leads to many race conditions in kernel code.: True

Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.: True

Question 32

Partially correct

Mark 1.17 out of 2.00

The unix file semantics demand that changes to any open file are visible immediately to any other processes accessing that file at that point in time.

Select the data-structure/programmatic features that ensure the implementation of unix semantics. (Assume there is no mmap())

Yes	No	
<input type="radio"/> <input checked="" type="checkbox"/>	All processes accessing the same file share the file descriptor among themselves	✓
<input type="radio"/> <input checked="" type="checkbox"/>	The pointer entry in the file descriptor array entry points to the data of the file directly	✓
<input checked="" type="checkbox"/> <input type="radio"/>	There is only one global file structure per on-disk file.	✗
<input type="radio"/> <input checked="" type="checkbox"/>	All file accesses are made using only global variables	✓
<input checked="" type="checkbox"/> <input type="radio"/>	The 'file offset' is shared among all the processes that access the file.	✗
<input type="radio"/> <input checked="" type="checkbox"/>	No synchronization is implemented so that changes are made available immediately.	✓
<input checked="" type="checkbox"/> <input type="radio"/>	A single spinlock is to be used to protect the unique global 'file structure' representing the file, thus synchronizing access, and making other processes wait for earlier process to finish writing so that writes get visible immediately.	✗
<input checked="" type="checkbox"/> <input type="radio"/>	There is only one in-memory copy of the on disk file's contents in kernel memory/buffers	✓
<input checked="" type="checkbox"/> <input type="radio"/>	The file descriptors in every PCB are pointers to the same global file structure.	✗
<input type="radio"/> <input checked="" type="checkbox"/>	The file descriptor array is external to PCB and all processes that share a file, have pointers to same file-descriptors' array	✓
<input checked="" type="checkbox"/> <input type="radio"/>	All file structures representing any open file, give access to the same in-memory copy of the file's contents	✓
<input checked="" type="checkbox"/> <input type="radio"/>	The 'file offset' index is stored outside the file-structure to which file-descriptor array points	✗

All processes accessing the same file share the file descriptor among themselves: No

The pointer entry in the file descriptor array entry points to the data of the file directly: No

There is only one global file structure per on-disk file.: No

All file accesses are made using only global variables: No

The 'file offset' is shared among all the processes that access the file.: No

No synchronization is implemented so that changes are made available immediately.: No

A single spinlock is to be used to protect the unique global 'file structure' representing the file, thus synchronizing access, and making other processes wait for earlier process to finish writing so that writes get visible immediately.: No

There is only one in-memory copy of the on disk file's contents in kernel memory/buffers: Yes

The file descriptors in every PCB are pointers to the same global file structure.: No

The file descriptor array is external to PCB and all processes that share a file, have pointers to same file-descriptors' array: No

All file structures representing any open file, give access to the same in-memory copy of the file's contents: Yes

The 'file offset' index is stored outside the file-structure to which file-descriptor array points: No

Question 33

Partially correct

Mark 0.33 out of 2.00

Map the function in xv6's file system code, to its perceived logical layer.

namei	inode	✗
filestat()	Choose...	
dirlookup	directory	✓
ialloc	file descriptor	✗
stati	Choose...	
ideintr	buffer cache	✗
bread	Choose...	
balloc	file descriptor	✗
sys_chdir()	system call	✓
skipelem	system call	✗
commit	system call	✗
bmap	system call	✗

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: namei → pathname lookup, filestat() → file descriptor, dirlookup → directory, ialloc → inode, stati → inode, ideintr → disk driver, bread → buffer cache, balloc → block allocation on disk, sys_chdir() → system call, skipelem → pathname lookup, commit → logging, bmap → inode

[◀ Course Exit Feedback](#)[Jump to...](#)[xv6-public-master ►](#)

Started on Monday, 24 January 2022, 7:07:42 PM

State Finished

Completed on Monday, 24 January 2022, 8:08:11 PM

Time taken 1 hour

Grade 8.90 out of 20.00 (45%)

Question 1

Complete

Mark 0.80 out of 1.00

Match the register with the segment used with it.

ebp	ss
eip	cs
edi	ds
esp	ss
esi	ds

The correct answer is: ebp → ss, eip → cs, edi → es, esp → ss, esi → ds

Question 2

Complete

Mark 1.00 out of 1.00

```
int value = 5;
int main()
{
    pid_t pid;
    pid = fork();
    if (pid == 0) { /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("%d", value); /* LINE A */
    }
    return 0;
}
```

What's the value printed here at LINE A?

Answer:

The correct answer is: 5

Question 3

Complete

Mark 0.50 out of 0.50

Is the command "cat README > done &" possible on xv6? (Note the & in the end)

- a. no
- b. yes

The correct answer is: yes

Question 4

Complete

Mark 0.00 out of 2.00

xv6.img: bootblock kernel

```
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- a. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file.
- b. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk.
- c. The size of xv6.img is exactly = (size of bootblock) + (size of kernel)
- d. xv6.img is the virtual processor used by the qemu emulator
- e. The size of the kernel file is nearly 5 MB
- f. Blocks in xv6.img after kernel may be all zeroes.
- g. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk.
- h. The kernel is located at block-1 of the xv6.img
- i. The bootblock is located on block-0 of the xv6.img
- j. The bootblock may be 512 bytes or less (looking at the Makefile instruction)
- k. The size of the xv6.img is nearly 5 MB

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question 5

Complete

Mark 0.43 out of 1.00

Rank the following storage systems from slowest (first) to fastest(last)

You can drag and drop the items below/above each other.

Registers
Cache
Main memory
Nonvolatile memory
Magnetic tapes
Optical disk
Hard-disk drives

The correct order for these items is as follows:

1. Magnetic tapes
2. Optical disk
3. Hard-disk drives
4. Nonvolatile memory
5. Main memory
6. Cache
7. Registers

Question 6

Complete

Mark 1.00 out of 1.00

How does the distinction between kernel mode and user mode function as a rudimentary form of protection (security) ?

Select one:

- a. It disallows hardware interrupts when a process is running
- b. It prohibits one process from accessing other process's memory
- c. It prohibits a user mode process from running privileged instructions
- d. It prohibits invocation of kernel code completely, if a user program is running

The correct answer is: It prohibits a user mode process from running privileged instructions

Question 7

Complete

Mark 0.00 out of 2.00

Which of the following are NOT a part of job of a typical compiler?

- a. Suggest alternative pieces of code that can be written
- b. Check the program for syntactical errors
- c. Check the program for logical errors
- d. Convert high level language code to machine code
- e. Process the # directives in a C program
- f. Invoke the linker to link the function calls with their code, extern globals with their declaration

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question 8

Complete

Mark 0.00 out of 2.00

Match the program with its output (ignore newlines in the output. Just focus on the count of the number of 'hi')

main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi hi

main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi hi hi

main() { int i = NULL; fork(); printf("hi\n"); }

No output

main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }

hi

The correct answer is: main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi, main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi, main() { int i = NULL; fork(); printf("hi\n"); } → hi hi, main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); } → hi hi

Question 9

Complete

Mark 0.83 out of 2.00

Select all statements that correctly explain the use/purpose of system calls.

Select one or more:

- a. Provide an environment for process creation
- b. Switch from user mode to kernel mode
- c. Handle ALL types of interrupts
- d. Handle exceptions like division by zero
- e. Run each instruction of an application program
- f. Allow I/O device access to user processes
- g. Provide services for accessing files

The correct answers are: Switch from user mode to kernel mode, Provide services for accessing files, Allow I/O device access to user processes, Provide an environment for process creation

Question 10

Complete

Mark 0.50 out of 0.50

Compare multiprogramming with multitasking

- a. A multiprogramming system is not necessarily multitasking
- b. A multitasking system is not necessarily multiprogramming

The correct answer is: A multiprogramming system is not necessarily multitasking

Question 11

Complete

Mark 0.60 out of 1.00

Select all the correct statements about two modes of CPU operation

Select one or more:

- a. The two modes are essential for a multitasking system
- b. Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode
- c. The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously
- d. The two modes are essential for a multiprogramming system
- e. There is an instruction like 'iret' to return from kernel mode to user mode

The correct answers are: The two modes are essential for a multiprogramming system, The two modes are essential for a multitasking system, There is an instruction like 'iret' to return from kernel mode to user mode, The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously, Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode

Question 12

Complete

Mark 0.50 out of 0.50

Is the terminal a part of the kernel on GNU/Linux systems?

- a. no
- b. yes

The correct answer is: no

Question 13

Complete

Mark 1.00 out of 1.00

Why should a program exist in memory before it starts executing ?

- a. Because the hard disk is a slow medium
- b. Because the processor can run instructions and access data only from memory
- c. Because the variables of the program are stored in memory
- d. Because the memory is volatile

The correct answer is: Because the processor can run instructions and access data only from memory

Question 14

Complete

Mark 1.33 out of 2.00

Which of the following instructions should be privileged?

Select one or more:

- a. Read the clock.
- b. Access memory management unit of the processor
- c. Access I/O device.
- d. Turn off interrupts.
- e. Set value of a memory location
- f. Set value of timer.
- g. Access a general purpose register
- h. Modify entries in device-status table
- i. Switch from user to kernel mode.

The correct answers are: Set value of timer., Access memory management unit of the processor, Turn off interrupts., Modify entries in device-status table, Access I/O device., Switch from user to kernel mode.

Question 15

Complete

Mark 0.07 out of 2.00

Select all the correct statements about calling convention on x86 32-bit.

- a. The ebp pointers saved on the stack constitute a chain of activation records
- b. The return value is either stored on the stack or returned in the eax register
- c. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables
- d. Parameters may be passed in registers or on stack
- e. Return address is one location above the ebp
- f. Parameters may be passed in registers or on stack
- g. Compiler may allocate more memory on stack than needed
- h. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function
- i. Parameters are pushed on the stack in left-right order
- j. during execution of a function, ebp is pointing to the old ebp
- k. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

Question 16

Complete

Mark 0.33 out of 0.50

Order the following events in boot process (from 1 onwards)

Shell	3
BIOS	1
Init	4
OS	6
Login interface	5
Boot loader	2

The correct answer is: Shell → 6, BIOS → 1, Init → 4, OS → 3, Login interface → 5, Boot loader → 2

[◀ \(Task\) Compulsory xv6 task](#)

Jump to...

[\(Optional Assignment\) Shell Programming\(Conformance tests\) ▶](#)

Started on Wednesday, 9 February 2022, 7:00:12 PM

State Finished

Completed on Wednesday, 9 February 2022, 7:46:38 PM

Time taken 46 mins 26 secs

Grade 3.00 out of 11.00 (27%)

Question 1

Complete

Mark 0.00 out of 1.00

The number of GDT entries setup during boot process of xv6 is

- a. 2
- b. 3
- c. 0
- d. 256
- e. 4
- f. 255

The correct answer is: 3

Question 2

Complete

Mark 0.00 out of 1.00

x86 provides which of the following type of memory management options?

- a. segmentation and one level paging
- b. segmentation or one or two level paging
- c. segmentation and two level paging
- d. segmentation or paging
- e. segmentation and one or two level paging
- f. segmentation only

The correct answer is: segmentation and one or two level paging

Question 3

Complete

Mark 0.00 out of 1.00

which of the following is not a difference between real mode and protected mode

- a. in real mode general purpose registers are 16 bit, in protected mode they are 32 bit
- b. in real mode the addressable memory is more than in protected mode
- c. in real mode the addressable memory is less than in protected mode
- d. in real mode the segment is multiplied by 16, in protected mode segment is used as index in GDT
- e. processor starts in real mode

The correct answer is: in real mode the addressable memory is more than in protected mode

Question 4

Complete

Mark 0.00 out of 1.00

The kernel ELF file contains how many Program headers?

- a. 4
- b. 2
- c. 3
- d. 9
- e. 10

The correct answer is: 3

Question 5

Not answered

Marked out of 0.50

code line, MMU setting: Match the line of xv6 code with the MMU setup employed

Answer:

The correct answer is: inb \$0x64,%al

Question 6

Complete

Mark 1.00 out of 1.00

The kernel is loaded at Physical Address

- a. 0x000100000
- b. 0x00010000
- c. 0x80100000
- d. 0x800000000

The correct answer is: 0x000100000

Question 7

Complete

Mark 0.00 out of 1.00

Why is the code of entry() in Assembly and not in C?

- a. Because the symbol entry() is inside the ELF file
- b. Because the kernel code must begin in assembly
- c. Because it needs to setup paging
- d. There is no particular reason, it could also be in C

The correct answer is: Because it needs to setup paging

Question 8

Complete

Mark 1.00 out of 1.00

The ljmp instruction in general does

- a. change the CS and EIP to 32 bit mode, and jumps to next line of code
- b. change the CS and EIP to 32 bit mode, and jumps to new value of EIP
- c. change the CS and EIP to 32 bit mode
- d. change the CS and EIP to 32 bit mode, and jumps to kernel code

The correct answer is: change the CS and EIP to 32 bit mode, and jumps to new value of EIP

Question 9

Complete

Mark 0.00 out of 1.00

The variable \$stack in entry.S is

- a. located at the value given by %esp as setup by bootmain()
- b. located at 0x7c00
- c. a memory region allocated as a part of entry.S
- d. located at less than 0x7c00
- e. located at 0

The correct answer is: a memory region allocated as a part of entry.S

Question 10

Not answered

Marked out of 0.50

Match the pairs of which action is taken by whom

Answer:

The correct answer is: kernel

Question 11

Complete

Mark 0.00 out of 1.00

ELF Magic number is

- a. 0xFFFFFFFF
- b. 0
- c. 0xELFELFELF
- d. 0xELF
- e. 0x0x464CELF
- f. 0x464C457FL
- g. 0x464C457FU

The correct answer is: 0x464C457FU

Question 12

Complete

Mark 1.00 out of 1.00

The right side of line of code "entry = (void(*)(void))(elf->entry)" means

- a. Convert the "entry" in ELF structure into void
- b. Get the "entry" in ELF structure and convert it into a function pointer accepting no arguments and returning nothing
- c. Get the "entry" in ELF structure and convert it into a function void pointer
- d. Get the "entry" in ELF structure and convert it into a void pointer

The correct answer is: Get the "entry" in ELF structure and convert it into a function pointer accepting no arguments and returning nothing

[◀ Homework questions: Basics of MM, xv6 booting](#)

Jump to...

[\(Code\) Files, redirection, dup, \(IPC\)pipe ▶](#)

Started on Monday, 21 February 2022, 7:00:28 PM

State Finished

Completed on Monday, 21 February 2022, 7:49:24 PM

Time taken 48 mins 56 secs

Grade 8.30 out of 10.00 (83%)

Question 1

Complete

Mark 0.80 out of 1.00

Match the elements of C program to their place in memory

Local Static variables	Data
Global variables	Data
Code of main()	Code
Function code	Code
Arguments	Stack
Mallocoed Memory	Heap
#include files	No Memory needed
#define MACROS	No memory needed
Local Variables	Stack
Global Static variables	Data

The correct answer is: Local Static variables → Data, Global variables → Data, Code of main() → Code, Function code → Code, Arguments → Stack, Mallocoed Memory → Heap, #include files → No memory needed, #define MACROS → No Memory needed, Local Variables → Stack, Global Static variables → Data

Question 2

Complete

Mark 1.00 out of 1.00

What will be the output of this program

```
int main() {
    int fd;
    printf("%d ", open("/etc/passwd", O_RDONLY));
    close(1);
    fd = printf("%d ", open("/etc/passwd", O_RDONLY));
    close(fd);
    fd = printf("%d ", open("/etc/passwd", O_RDONLY));
}
```

- a. 3 1 2
- b. 3 3 3
- c. 3 1 1
- d. 1 1 1
- e. 3 4 5
- f. 2 2 2

The correct answer is: 3 1 1

Question 3

Complete

Mark 1.00 out of 1.00

Arrange in correct order, the files involved in execution of system call

vectors.S	2
trap.c	4
usys.S	1
trapasm.S	3

The correct answer is: vectors.S → 2, trap.c → 4, usys.S → 1, trapasm.S → 3

Question 4

Complete

Mark 1.00 out of 1.00

The "push 0" in vectors.S is

- a. A placeholder to match the size of struct trapframe
- b. Place for the error number value
- c. To indicate that it's a system call and not a hardware interrupt
- d. To be filled in as the return value of the system call

The correct answer is: Place for the error number value

Question 5

Complete

Mark 0.00 out of 1.00

A process blocks itself means

- a. The kernel code of an interrupt handler, moves the process to a waiting queue and calls scheduler
- b. The application code calls the scheduler
- c. The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler
- d. The kernel code of system call calls scheduler

The correct answer is: The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler

Question 6

Complete

Mark 1.00 out of 1.00

Select the odd one out

- a. Kernel stack of new process to Process stack of new process
- b. Process stack of running process to kernel stack of running process
- c. Kernel stack of running process to kernel stack of scheduler
- d. Kernel stack of scheduler to kernel stack of new process
- e. Kernel stack of new process to kernel stack of scheduler

The correct answer is: Kernel stack of new process to kernel stack of scheduler

Question 7

Complete

Mark 0.50 out of 0.50

Match the File descriptors to their meaning

0 Standard Input

2 Standard error

1 Standard output

The correct answer is: 0 → Standard Input, 2 → Standard error, 1 → Standard output

Question 8

Complete

Mark 1.00 out of 1.00

Which of the following is not a task of the code of swtch() function

- a. Switch stacks
- b. Change the kernel stack location
- c. Load the new context
- d. Jump to next context EIP
- e. Save the return value of the old context code
- f. Save the old context

The correct answers are: Save the return value of the old context code, Change the kernel stack location

Question 9

Complete

Mark 0.50 out of 0.50

Match the names of PCB structures with kernel

linux struct task_struct

xv6 struct proc

The correct answer is: linux → struct task_struct, xv6 → struct proc

Question 10

Complete

Mark 0.50 out of 0.50

Match the MACRO with its meaning

KERNBASE	2 GB
KERNLINK	2.224 GB
PHYSTOP	224 MB

The correct answer is: KERNBASE → 2 GB, KERNLINK → 2.224 GB, PHYSTOP → 224 MB

Question 11

Complete

Mark 1.00 out of 1.00

The trapframe, in xv6, is built by the

- a. hardware, trapasm.S
- b. hardware, vectors.S
- c. hardware, vectors.S, trapasm.S, trap()
- d. hardware, vectors.S, trapasm.S
- e. vectors.S, trapasm.S

The correct answer is: hardware, vectors.S, trapasm.S

Question 12

Complete

Mark 0.00 out of 0.50

Which of the following state transitions are not possible?

- a. Running -> Waiting
- b. Ready -> Waiting
- c. Waiting -> Terminated
- d. Ready -> Terminated

The correct answers are: Ready -> Terminated, Waiting -> Terminated, Ready -> Waiting

[◀ Description of some possible course mini projects](#)

Jump to...

[\(Code\) mmap related programs ▶](#)

Started on Saturday, 26 February 2022, 5:18:30 PM

State Finished

Completed on Saturday, 26 February 2022, 6:30:44 PM

Time taken 1 hour 12 mins

Grade 8.55 out of 15.00 (57%)

Question 1

Complete

Mark 0.50 out of 0.50

Map the technique with its feature/problem

static linking	large executable file
dynamic loading	allocate memory only if needed
dynamic linking	small executable file
static loading	wastage of physical memory

The correct answer is: static linking → large executable file, dynamic loading → allocate memory only if needed, dynamic linking → small executable file, static loading → wastage of physical memory

Question 2

Complete

Mark 0.00 out of 1.00

Calculate the EAT in NANO-seconds (upto 2 decimal points) w.r.t. a page fault, given

Memory access time = 299 ns

Average page fault service time = 6 ms

Page fault rate = 0.8

Answer:

The correct answer is: 4800059.80

Question 3

Complete

Mark 1.00 out of 1.00

Given six memory partitions of 300 KB , 600 KB , 350 KB , 200 KB , 750 KB , and 125 KB (in order), how would the first-fit, best-fit, and worst-fit algorithms place processes of size 115 KB and 500 KB (in order)?

best fit 115 KB	125 KB
best fit 500 KB	600 KB
worst fit 500 KB	635 KB
worst fit 115 KB	750 KB
first fit 500 KB	600 KB
first fit 115 KB	300 KB

The correct answer is: best fit 115 KB → 125 KB, best fit 500 KB → 600 KB, worst fit 500 KB → 635 KB, worst fit 115 KB → 750 KB, first fit 500 KB → 600 KB, first fit 115 KB → 300 KB

Question 4

Complete

Mark 0.29 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- a. TLB hit ration has zero impact in effective memory access time in demand paging.
- b. Both demand paging and paging support shared memory pages.
- c. Calculations of number of bits for page number and offset are same in paging and demand paging.
- d. Demand paging requires additional hardware support, compared to paging.
- e. The meaning of valid-invalid bit in page table is different in paging and demand-paging.
- f. Paging requires NO hardware support in CPU
- g. With paging, it's possible to have user programs bigger than physical memory.
- h. With demand paging, it's possible to have user programs bigger than physical memory.
- i. Paging requires some hardware support in CPU
- j. Demand paging always increases effective memory access time.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question 5

Complete

Mark 0.36 out of 0.50

Map the parts of a C code to the memory regions they are related to

local variables	stack
global un-initialized variables	bss
static variables	data
global initialized variables	data
function arguments	stack
malloced memory	stack
functions	stack

The correct answer is: local variables → stack, global un-initialized variables → bss, static variables → data, global initialized variables → data, function arguments → stack, malloced memory → heap, functions → code

Question 6

Complete

Mark 0.75 out of 1.00

which of the following, do you think, are valid concerns for making the kernel pageable?

- a. The kernel must have some dedicated frames for it's own work
- b. No data structure of kernel should be pageable
- c. The kernel's own page tables should not be pageable
- d. The disk driver and disk interrupt handler should not be pageable
- e. No part of kernel code should be pageable.
- f. The page fault handler should not be pageable

The correct answers are: The kernel's own page tables should not be pageable, The page fault handler should not be pageable, The kernel must have some dedicated frames for it's own work, The disk driver and disk interrupt handler should not be pageable

Question 7

Complete

Mark 0.75 out of 1.00

W.r.t the figure given below, mark the given statements as True or False.

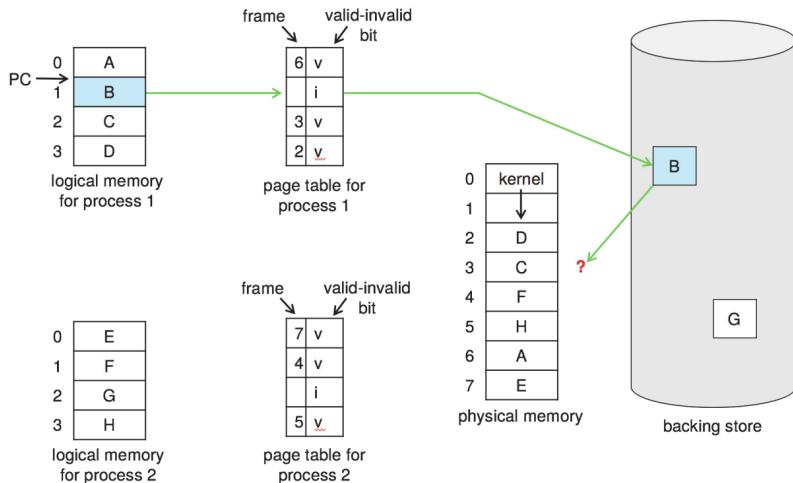


Figure 10.9 Need for page replacement.

True False

- Kernel occupies two page frames
- Handling this scenario demands two disk I/Os
- Local replacement means chose any of the frames 2, 3, 6
- Local replacement means chose any of the frame from 2 to 7
- Global replacement means chose any of the frame from 0 to 7
- Global replacement means chose any of the frame from 2 to 7
- The kernel's pages can not used for replacement if kernel is not pageable.
- Page 1 of process 1 needs a replacement

Kernel occupies two page frames: True

Handling this scenario demands two disk I/Os: True

Local replacement means chose any of the frames 2, 3, 6: True

Local replacement means chose any of the frame from 2 to 7: False

Global replacement means chose any of the frame from 0 to 7: False

Global replacement means chose any of the frame from 2 to 7: True

The kernel's pages can not used for replacement if kernel is not pageable.: True

Page 1 of process 1 needs a replacement: True

Question 8

Complete

Mark 0.67 out of 1.00

Shared memory is possible with which of the following memory management schemes ?

Select one or more:

- a. paging
- b. segmentation
- c. continuous memory management
- d. demand paging

The correct answers are: paging, segmentation, demand paging

Question 9

Complete

Mark 0.60 out of 1.00

Given below is the "maps" file for a particular instance of "vim.basic" process.

Mark the given statements as True or False, w.r.t. the contents of the map file.

55a43501b000-55a435049000 r--p 00000000 103:05 917529	/usr/bin/vim.basic
55a435049000-55a435248000 r-xp 0002e000 103:05 917529	/usr/bin/vim.basic
55a435248000-55a4352b6000 r--p 0022d000 103:05 917529	/usr/bin/vim.basic
55a4352b7000-55a4352c5000 r--p 0029b000 103:05 917529	/usr/bin/vim.basic
55a4352c5000-55a4352e2000 rw-p 002a9000 103:05 917529	/usr/bin/vim.basic
55a4352e2000-55a4352f0000 rw-p 00000000 00:00 0	[heap]
55a436bc9000-55a436e5b000 rw-p 00000000 00:00 0	/usr/lib/x86_64-linux-
7f275b0a3000-7f275b0a6000 r--p 00000000 103:05 917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so	gnu/libnss_files-2.31.so
7f275b0a6000-7f275b0ad000 r-xp 00003000 103:05 917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so	gnu/libnss_files-2.31.so
7f275b0ad000-7f275b0af000 r--p 0000a000 103:05 917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so	gnu/libnss_files-2.31.so
7f275b0af000-7f275b0b0000 r--p 0000b000 103:05 917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so	gnu/libnss_files-2.31.so
7f275b0b0000-7f275b0b1000 rw-p 0000c000 103:05 917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so	gnu/libnss_files-2.31.so
7f275b0b1000-7f275b0b7000 rw-p 00000000 00:00 0	
7f275b0b7000-7f275b8f5000 r--p 00000000 103:05 925247	/usr/lib/locale/locale-archive
7f275b8f5000-7f275b8fa000 rw-p 00000000 00:00 0	
7f275b8fa000-7f275b8fc000 r--p 00000000 103:05 924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4	gnu/libogg.so.0.8.4
7f275b8fc000-7f275b901000 r-xp 00002000 103:05 924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4	gnu/libogg.so.0.8.4
7f275b901000-7f275b904000 r--p 00007000 103:05 924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4	gnu/libogg.so.0.8.4
7f275b904000-7f275b905000 ---p 0000a000 103:05 924216	
gnu/libogg.so.0.8.4	gnu/libogg.so.0.8.4
7f275b905000-7f275b906000 r--p 0000a000 103:05 924216	
gnu/libogg.so.0.8.4	gnu/libogg.so.0.8.4
7f275b906000-7f275b907000 rw-p 0000b000 103:05 924216	
gnu/libogg.so.0.8.4	gnu/libogg.so.0.8.4
7f275b907000-7f275b90a000 r--p 00000000 103:05 924627	
gnu/libvorbis.so.0.4.8	gnu/libvorbis.so.0.4.8
7f275b90a000-7f275b921000 r-xp 00003000 103:05 924627	
gnu/libvorbis.so.0.4.8	gnu/libvorbis.so.0.4.8
7f275b921000-7f275b932000 r--p 0001a000 103:05 924627	
gnu/libvorbis.so.0.4.8	gnu/libvorbis.so.0.4.8
7f275b932000-7f275b933000 ---p 0002b000 103:05 924627	
gnu/libvorbis.so.0.4.8	gnu/libvorbis.so.0.4.8
7f275b933000-7f275b934000 r--p 0002b000 103:05 924627	
gnu/libvorbis.so.0.4.8	gnu/libvorbis.so.0.4.8
7f275b934000-7f275b935000 rw-p 0002c000 103:05 924627	
gnu/libvorbis.so.0.4.8	gnu/libvorbis.so.0.4.8
7f275b935000-7f275b937000 rw-p 00000000 00:00 0	
7f275b937000-7f275b938000 r--p 00000000 103:05 917914	/usr/lib/x86_64-linux-
gnu/libutil-2.31.so	gnu/libutil-2.31.so
7f275b938000-7f275b939000 r-xp 00001000 103:05 917914	/usr/lib/x86_64-linux-
gnu/libutil-2.31.so	gnu/libutil-2.31.so
7f275b939000-7f275b93a000 r--p 00002000 103:05 917914	/usr/lib/x86_64-linux-
gnu/libutil-2.31.so	gnu/libutil-2.31.so
7f275b93a000-7f275b93b000 r--p 00002000 103:05 917914	/usr/lib/x86_64-linux-
gnu/libutil-2.31.so	gnu/libutil-2.31.so
7f275b93b000-7f275b93c000 rw-p 00003000 103:05 917914	/usr/lib/x86_64-linux-

```

gnu/libutil-2.31.so
7f275b93c000-7f275b93e000 r--p 00000000 103:05 915906      /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b93e000-7f275b94f000 r-xp 00002000 103:05 915906      /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b94f000-7f275b955000 r--p 00013000 103:05 915906      /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b955000-7f275b956000 ---p 00019000 103:05 915906      /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b956000-7f275b957000 r--p 00019000 103:05 915906      /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b957000-7f275b958000 rw-p 0001a000 103:05 915906      /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b958000-7f275b95c000 r--p 00000000 103:05 923645      /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b95c000-7f275b978000 r-xp 00004000 103:05 923645      /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b978000-7f275b982000 r--p 00020000 103:05 923645      /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b982000-7f275b983000 ---p 0002a000 103:05 923645      /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b983000-7f275b985000 r--p 0002a000 103:05 923645      /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b985000-7f275b986000 rw-p 0002c000 103:05 923645      /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b986000-7f275b988000 r--p 00000000 103:05 924057      /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b988000-7f275b98d000 r-xp 00002000 103:05 924057      /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b98d000-7f275b98f000 r--p 00007000 103:05 924057      /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b98f000-7f275b990000 r--p 00008000 103:05 924057      /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b990000-7f275b991000 rw-p 00009000 103:05 924057      /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b991000-7f275b995000 r--p 00000000 103:05 921934      /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b995000-7f275b9a3000 r-xp 00004000 103:05 921934      /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9a3000-7f275b9a9000 r--p 00012000 103:05 921934      /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9a9000-7f275b9aa000 r--p 00017000 103:05 921934      /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9aa000-7f275b9ab000 rw-p 00018000 103:05 921934      /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9ab000-7f275b9ad000 rw-p 00000000 00:00 0
7f275b9ad000-7f275b9af000 r--p 00000000 103:05 924631      /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9af000-7f275b9b4000 r-xp 00002000 103:05 924631      /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b4000-7f275b9b5000 r--p 00007000 103:05 924631      /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b5000-7f275b9b6000 ---p 00008000 103:05 924631      /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b6000-7f275b9b7000 r--p 00008000 103:05 924631      /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b7000-7f275b9b8000 rw-p 00009000 103:05 924631      /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b8000-7f275b9ba000 r--p 00000000 103:05 924277      /usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0
7f275b9ba000-7f275ba1e000 r-xp 00002000 103:05 924277      /usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0

```

7f275ba1e000-7f275ba46000 r--p 00066000 103:05 924277	/usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0	
7f275ba46000-7f275ba47000 r--p 0008d000 103:05 924277	/usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0	
7f275ba47000-7f275ba48000 rw-p 0008e000 103:05 924277	/usr/lib/x86_64-linux-
gnu/libpcre2-8.so.0.9.0	
7f275ba48000-7f275ba6d000 r--p 00000000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275ba6d000-7f275bbe5000 r-xp 00025000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bbe5000-7f275bc2f000 r--p 0019d000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bc2f000-7f275bc30000 ---p 001e7000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bc30000-7f275bc33000 r--p 001e7000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bc33000-7f275bc36000 rw-p 001ea000 103:05 917893	/usr/lib/x86_64-linux-
gnu/libc-2.31.so	
7f275bc36000-7f275bc3a000 rw-p 00000000 00:00 0	
7f275bc3a000-7f275bc41000 r--p 00000000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc41000-7f275bc52000 r-xp 00007000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc52000-7f275bc57000 r--p 00018000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc57000-7f275bc58000 r--p 0001c000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc58000-7f275bc59000 rw-p 0001d000 103:05 917906	/usr/lib/x86_64-linux-
gnu/libpthread-2.31.so	
7f275bc59000-7f275bc5d000 rw-p 00000000 00:00 0	
7f275bc5d000-7f275bcce000 r--p 00000000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275bcce000-7f275bf29000 r-xp 00071000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275bf29000-7f275c142000 r--p 002cc000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275c142000-7f275c143000 ---p 004e5000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275c143000-7f275c149000 r--p 004e5000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275c149000-7f275c190000 rw-p 004eb000 103:05 917016	/usr/lib/x86_64-linux-
gnu/libpython3.8.so.1.0	
7f275c190000-7f275c1b3000 rw-p 00000000 00:00 0	
7f275c1b3000-7f275c1b4000 r--p 00000000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b4000-7f275c1b6000 r-xp 00001000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b6000-7f275c1b7000 r--p 00003000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b7000-7f275c1b8000 r--p 00003000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b8000-7f275c1b9000 rw-p 00004000 103:05 917894	/usr/lib/x86_64-linux-
gnu/libdl-2.31.so	
7f275c1b9000-7f275c1bb000 rw-p 00000000 00:00 0	
7f275c1bb000-7f275c1c0000 r-xp 00000000 103:05 923815	/usr/lib/x86_64-linux-
gnu/libgpm.so.2	
7f275c1c0000-7f275c3bf000 ---p 00005000 103:05 923815	/usr/lib/x86_64-linux-
gnu/libgpm.so.2	
7f275c3bf000-7f275c3c0000 r--p 00004000 103:05 923815	/usr/lib/x86_64-linux-
gnu/libgpm.so.2	
7f275c3c0000-7f275c3c1000 rw-p 00005000 103:05 923815	/usr/lib/x86_64-linux-
gnu/libgpm.so.2	

```

7f275c3c1000-7f275c3c3000 r--p 00000000 103:05 923315          /usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253
7f275c3c3000-7f275c3c8000 r-xp 00002000 103:05 923315          /usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253
7f275c3c8000-7f275c3ca000 r--p 00007000 103:05 923315          /usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253
7f275c3ca000-7f275c3cb000 r--p 00008000 103:05 923315          /usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253
7f275c3cb000-7f275c3cc000 rw-p 00009000 103:05 923315          /usr/lib/x86_64-linux-
gnu/libacl.so.1.1.2253
7f275c3cc000-7f275c3cf000 r--p 00000000 103:05 923446          /usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5
7f275c3cf000-7f275c3d9000 r-xp 00003000 103:05 923446          /usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5
7f275c3d9000-7f275c3dd000 r--p 0000d000 103:05 923446          /usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5
7f275c3dd000-7f275c3de000 r--p 00010000 103:05 923446          /usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5
7f275c3de000-7f275c3df000 rw-p 00011000 103:05 923446          /usr/lib/x86_64-linux-
gnu/libcanberra.so.0.2.5
7f275c3df000-7f275c3e5000 r--p 00000000 103:05 924431          /usr/lib/x86_64-linux-
gnu/libselinux.so.1
7f275c3e5000-7f275c3fe000 r-xp 00006000 103:05 924431          /usr/lib/x86_64-linux-
gnu/libselinux.so.1
7f275c3fe000-7f275c405000 r--p 0001f000 103:05 924431          /usr/lib/x86_64-linux-
gnu/libselinux.so.1
7f275c405000-7f275c406000 ---p 00026000 103:05 924431          /usr/lib/x86_64-linux-
gnu/libselinux.so.1
7f275c406000-7f275c407000 r--p 00026000 103:05 924431          /usr/lib/x86_64-linux-
gnu/libselinux.so.1
7f275c407000-7f275c408000 rw-p 00027000 103:05 924431          /usr/lib/x86_64-linux-
gnu/libselinux.so.1
7f275c408000-7f275c40a000 rw-p 00000000 00:00 0
7f275c40a000-7f275c418000 r--p 00000000 103:05 924540          /usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2
7f275c418000-7f275c427000 r-xp 0000e000 103:05 924540          /usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2
7f275c427000-7f275c435000 r--p 0001d000 103:05 924540          /usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2
7f275c435000-7f275c439000 r--p 0002a000 103:05 924540          /usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2
7f275c439000-7f275c43a000 rw-p 0002e000 103:05 924540          /usr/lib/x86_64-linux-
gnu/libtinfo.so.6.2
7f275c43a000-7f275c449000 r--p 00000000 103:05 917895          /usr/lib/x86_64-linux-
gnu/libm-2.31.so
7f275c449000-7f275c4f0000 r-xp 0000f000 103:05 917895          /usr/lib/x86_64-linux-
gnu/libm-2.31.so
7f275c4f0000-7f275c587000 r--p 000b6000 103:05 917895          /usr/lib/x86_64-linux-
gnu/libm-2.31.so
7f275c587000-7f275c588000 r--p 0014c000 103:05 917895          /usr/lib/x86_64-linux-
gnu/libm-2.31.so
7f275c588000-7f275c589000 rw-p 0014d000 103:05 917895          /usr/lib/x86_64-linux-
gnu/libm-2.31.so
7f275c589000-7f275c58b000 rw-p 00000000 00:00 0
7f275c5ae000-7f275c5af000 r--p 00000000 103:05 917889          /usr/lib/x86_64-linux-gnu/ld-
2.31.so
7f275c5af000-7f275c5d2000 r-xp 00001000 103:05 917889          /usr/lib/x86_64-linux-gnu/ld-
2.31.so
7f275c5d2000-7f275c5da000 r--p 00024000 103:05 917889          /usr/lib/x86_64-linux-gnu/ld-
2.31.so
7f275c5db000-7f275c5dc000 r--p 0002c000 103:05 917889          /usr/lib/x86_64-linux-gnu/ld-
2.31.so

```

```
7f275c5dc000-7f275c5dd000 rw-p 0002d000 103:05 917889          /usr/lib/x86_64-linux-gnu/ld-
2.31.so
7f275c5dd000-7f275c5de000 rw-p 00000000 00:00 0
7ffd22d2f000-7ffd22d50000 rw-p 00000000 00:00 0          [stack]
7ffd22db0000-7ffd22db4000 r--p 00000000 00:00 0          [vvar]
7ffd22db4000-7ffd22db6000 r-xp 00000000 00:00 0          [vdso]
ffffffff600000-ffffffff601000 --xp 00000000 00:00 0          [vsyscall]
```

True False

- The size of the heap is one page
- This is a virtual memory map (not physical memory map)
- vim.basic uses the math library
- The 5th entry 55a4352c5000-55a4352e2000 may correspond to "data" of the vim.basic
- The size of the stack is one page

The size of the heap is one page: False

This is a virtual memory map (not physical memory map): True

vim.basic uses the math library: True

The 5th entry 55a4352c5000-55a4352e2000 may correspond to "data" of the vim.basic: True

The size of the stack is one page: False

Question 10

Complete

Mark 0.00 out of 1.00

Select all the correct statements, w.r.t. Copy on Write

- a. Fork() used COW technique to improve performance of new process creation.
- b. Vfork() assumes that there will be no write, but rather exec()
- c. COW helps us save memory
- d. If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated
- e. use of COW during fork() is useless if child called exit()
- f. use of COW during fork() is useless if exec() is called by the child

The correct answers are: Fork() used COW technique to improve performance of new process creation., If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated, COW helps us save memory, Vfork() assumes that there will be no write, but rather exec()

Question 11

Complete

Mark 1.00 out of 1.00

Assuming a 8- KB page size, what is the page numbers for the address 1093943 reference in decimal :
 (give answer also in decimal)

Answer:

The correct answer is: 134

Question 12

Complete

Mark 0.00 out of 1.00

Order the following events, related to page fault handling, in correct order

1. Page fault handler detects that it's a page fault and not illegal memory access
2. Disk interrupt handler runs
3. MMU detects that a page table entry is marked "invalid"
4. Page faulted process is moved to ready-queue
5. Empty frame is found
6. Page fault interrupt is generated
7. Other processes scheduled by scheduler
8. Page table of page faulted process is updated
9. Disk Interrupt occurs
10. Disk read is issued
11. Page fault handler in kernel starts executing
12. Page faulting process is made to wait in a queue

The correct order for these items is as follows:

1. MMU detects that a page table entry is marked "invalid"
2. Page fault interrupt is generated
3. Page fault handler in kernel starts executing
4. Page fault handler detects that it's a page fault and not illegal memory access
5. Empty frame is found
6. Disk read is issued
7. Page faulting process is made to wait in a queue
8. Other processes scheduled by scheduler
9. Disk Interrupt occurs
10. Disk interrupt handler runs
11. Page table of page faulted process is updated
12. Page faulted process is moved to ready-queue

Question 13

Complete

Mark 1.00 out of 1.00

Page sizes are a power of 2 because

Select one:

- a. Certain bits are reserved for offset in logical address. Hence page size = $2^{(\text{no.of offset bits})}$
- b. Power of 2 calculations are highly efficient
- c. Certain bits are reserved for offset in logical address. Hence page size = $2^{(32 - \text{no.of offset bits})}$
- d. operating system calculations happen using power of 2
- e. MMU only understands numbers that are power of 2

The correct answer is: Certain bits are reserved for offset in logical address. Hence page size = $2^{(\text{no.of offset bits})}$

Question 14

Complete

Mark 1.00 out of 1.00

Given below is the output of the command "ps -eo min_flt,maj_flt,cmd" on a Linux Desktop system. Select the statements that are consistent with the output

```
626729 482768 /usr/lib/firefox/firefox -contentproc -parentBuildID 20220202182137 -prefsLen 9256 -prefMapSize 264738 -appDir /usr/lib/firefox/browser 6094 true rdd
2167 687 /usr/sbin/apache2 -k start
1265185 222 /usr/bin/gnome-shell
102648 111 /usr/sbin/mysqld
9813 0 bash
15497 370 /usr/bin/gedit --gapplication-service
```

- a. All of the processes here exhibit some good locality of reference
- b. Firefox has likely been running for a large amount of time
- c. The bash shell is mostly busy doing work within a particular locality
- d. Apache web-server has not been doing much work

The correct answers are: Firefox has likely been running for a large amount of time, Apache web-server has not been doing much work, The bash shell is mostly busy doing work within a particular locality, All of the processes here exhibit some good locality of reference

Question 15

Complete

Mark 0.14 out of 1.00

Suppose two processes share a library between them. The library consists of 5 pages, and these 5 pages are mapped to frames 9, 15, 23, 4, 7 respectively. Process P1 has got 6 pages, first 3 of which consist of process's own code/data and 3 correspond to library's pages 0, 2, 4. Process P2 has got 7 pages, first 3 of which consist of process's own code/data and remaining 4 correspond to library's pages 0, 1, 3, 4. Fill in the blanks for page table entries of P1 and P2.

Page table of P1, Page 5	<input type="text" value="23"/>
Page table of P1, Page 3	<input type="text" value="9"/>
Page table of P2, Page 0	<input type="text" value="6"/>
Page table of P2, Page 3	<input type="text" value="7"/>
Page table of P2, Page 4	<input type="text" value="9"/>
Page table of P2, Page 1	<input type="text" value="5"/>
Page table of P1, Page 4	<input type="text" value="9"/>

The correct answer is: Page table of P1, Page 5 → 7, Page table of P1, Page 3 → 9, Page table of P2, Page 0 → 9, Page table of P2, Page 3 → 4, Page table of P2, Page 4 → 7, Page table of P2, Page 1 → 15, Page table of P1, Page 4 → 23

Question 16

Complete

Mark 0.50 out of 1.00

For the reference string

3 4 3 5 2

the number of page faults (including initial ones) using

FIFO replacement and 2 page frames is :

FIFO replacement and 3 page frames is :

◀ (Code) mmap related programs

Jump to...

Points from Mid-term feedback ►

Started on Monday, 7 March 2022, 7:00:12 PM

State Finished

Completed on Monday, 7 March 2022, 8:00:04 PM

Time taken 59 mins 52 secs

Grade 9.78 out of 15.00 (65%)

Question 1

Complete

Mark 1.00 out of 1.00

Why is there a call to kinit2? Why is it not merged with knit1?

- a. call to seginit() makes it possible to actually use PHYSTOP in argument to kinit2()
- b. When kinit1() is called there is a need for few page frames, but later kinit2() is called to serve need of more page frames
- c. Because there is a limit on the values that the arguments to kinit1() can take.
- d. kinit2 refers to virtual addresses beyond 4MB, which are not mapped before kalloc() is called

The correct answer is: kinit2 refers to virtual addresses beyond 4MB, which are not mapped before kalloc() is called

Question 2

Complete

Mark 1.50 out of 1.50

Arrange the following in the correct order of execution (w.r.t. 'init')

initcode() returns in forkret()

6

'initcode' process is marked RUNNABLE

3

'initcode' struct proc is created

2

userinit() is called

1

mpmain() calls scheduler()

4

initcode() calls exec("/init", ...)

8

initcode() returns from trapret()

7

scheduler() schedules initcode() process

5

The correct answer is: initcode() returns in forkret() → 6, 'initcode' process is marked RUNNABLE → 3, 'initcode' struct proc is created → 2, userinit() is called → 1, mpmain() calls scheduler() → 4, initcode() calls exec("/init", ...) → 8, initcode() returns from trapret() → 7, scheduler() schedules initcode() process → 5

Question 3

Complete

Mark 0.00 out of 2.00

exec() does this: curproc->tf->eip = elf.entry, but userinit() does this: p->tf->eip = 0; Select all the statements from below, that collectively explain this

- a. the 'entry' in initcode is anyways 0
- b. exec() loads from ELF file and the address of first instruction to be executed is given by 'entry'
- c. the initcode is created using objcopy, which discards all relocation information and symbols (like entry)
- d. elf.entry is anyways 0, so both statements mean the same
- e. In userinit() the function inituvm() has mapped the code of 'initcode' to be starting at virtual address 0
- f. the code of 'initcode' is loaded at physical address 0

The correct answers are: exec() loads from ELF file and the address of first instruction to be executed is given by 'entry', In userinit() the function inituvm() has mapped the code of 'initcode' to be starting at virtual address 0, the initcode is created using objcopy, which discards all relocation information and symbols (like entry)

Question 4

Complete

Mark 1.00 out of 1.00

What does userinit() do ?

- a. initializes the users
- b. sets up the 'initcode' process to start execution in forkret()
- c. sets up the 'initcode' process to start execution in forkret()
- d. sets up the 'initcode' process to start execution in trapret()
- e. sets up the 'init' process to start execution in forkret()
- f. initializes the process 'init' and starts executing it

The correct answer is: sets up the 'initcode' process to start execution in forkret()

Question 5

Complete

Mark 0.67 out of 1.00

Which of the following is done by mappages()?

- a. allocate page directory if required
- b. allocate page frame if required
- c. allocate page table if required
- d. create page table mappings for the range given by "va" and "va + size"
- e. create page table mappings to the range given by "pa" and "pa + size"

The correct answers are: create page table mappings for the range given by "va" and "va + size", allocate page table if required, create page table mappings to the range given by "pa" and "pa + size"

Question 6

Complete

Mark 0.00 out of 1.00

Select the statement that most correctly describes what setupkvm() does

- a. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array
- b. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array and makes kpgdir point to it
- c. creates a 1-level page table for the use by the kernel, as specified in kmap[] global array
- d. creates a 2-level page table for the use of the kernel, as specified in gdtdesc

The correct answer is: creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array

Question 7

Complete

Mark 0.00 out of 1.00

The approximate number of page frames created by kinit1 is

- a. 4
- b. 16
- c. 4000
- d. 3000
- e. 2000
- f. 1000
- g. 10

The correct answer is: 3000

Question 8

Complete

Mark 1.20 out of 1.50

Which of the following is DONE by allocproc() ?

- a. setup kernel memory mappings for the process
- b. Select an UNUSED struct proc for use
- c. ensure that the process starts in trapret()
- d. allocate kernel stack for the process
- e. allocate PID to the process
- f. ensure that the process starts in forkret()
- g. setup the trapframe and context pointers appropriately
- h. setup the contents of the trapframe of the process properly

The correct answers are: Select an UNUSED struct proc for use, allocate PID to the process, allocate kernel stack for the process, setup the trapframe and context pointers appropriately, ensure that the process starts in forkret()

Question 9

Complete

Mark 1.00 out of 1.00

Map the virtual address to physical address in xv6

KERNLINK	0x100000
0xFE000000	0xFE000000
80108000	0x108000
KERNBASE	0

The correct answer is: KERNLINK → 0x100000, 0xFE000000 → 0xFE000000, 80108000 → 0x108000, KERNBASE → 0

Question 10

Complete

Mark 0.42 out of 1.00

Select all the correct statements about initcode

- a. code of initcode is loaded at virtual address 0
- b. The data and stack of initcode is mapped to one single page in userinit()
- c. code of 'initcode' is loaded along with the kernel during booting
- d. the size of 'initcode' is 2c
- e. code of initcode is loaded in memory by the kernel during userinit()
- f. initcode is the 'init' process
- g. initcode essentially calls exec("/init",...)

The correct answers are: code of 'initcode' is loaded along with the kernel during booting, the size of 'initcode' is 2c, The data and stack of initcode is mapped to one single page in userinit(), initcode essentially calls exec("/init",...)

Question 11

Complete

Mark 1.00 out of 1.00

The variable 'end' used as argument to kinit1 has the value

- a. 801154a8
- b. 81000000
- c. 80110000
- d. 80102da0
- e. 8010a48c
- f. 80000000

The correct answer is: 801154a8

Question 12

Complete

Mark 1.00 out of 1.00

What does seginit() do?

- a. Nothing significant, just repetition of earlier GDT setup but with 2-level paging setup done
- b. Nothing significant, just repetition of earlier GDT setup but with free frames list created now
- c. Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 0
- d. Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 3
- e. Nothing significant, just repetition of earlier GDT setup but with kernel page table allocated now

The correct answer is: Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 3

Question 13

Complete

Mark 1.00 out of 1.00

Does exec() code around clearptau() lead to wastage of one page frame?

- a. no
- b. yes

The correct answer is: yes

[◀ Questions for test on kalloc/kfree/kvmalloc, etc.](#)

Jump to...

[\(Optional Assignment\) Slab allocator in xv6 ►](#)

Dashboard / My courses / Computer Engineering & IT / CEIT-Even-sem-20-21 / QS-Even-sem-2020-21 / 16 May - 22 May / End Sem Exam OS-2021

Started on Saturday, 22 May 2021, 8:00 AM

State Finished

Completed on Saturday, 22 May 2021, 9:30 AM

Time taken 1 hour 30 mins

Grade 26.12 out of 40.00 (65%)

Question 1

Incorrect

Mark 0.00 out of 1.00

A 4 GB disk with 1 KB of block size would require these many number of **blocks** for its free block bitmap:

Answer: 4096 ✖

The correct answer is: 512

Question 2

Correct

Mark 1.00 out of 1.00

Given that the memory access time is 110 ns, probability of a page fault is 0.5 and page fault handling time is 12 ms,

The effective memory access time in nanoseconds is:

Answer: 6000165 ✓

The correct answer is: 6000055.00

Question 3

Incorrect

Mark 0.00 out of 1.00

The maximum size of a file in number of blocks of BSIZE in xv6 code is

(write a number only)

Answer: 268 ✖

The correct answer is: 138

Question 4

Incorrect

Mark 0.00 out of 1.00

Calculate the average waiting time using

Round Robin scheduling with time quantum of 5 time units
for the following workload

assuming that they arrive in the order written below.

Process Burst Time

P1	5
P2	7
P3	6
P4	2

Write only a number in the answer upto two decimal points.

Answer: 40.75 ✖

The correct answer is: 10.25



Question 5

Correct

Mark 1.00 out of 1.00

For the reference string

4 2 5 1 0 1 2 5 4 1 2

the number of page faults, including initial ones,
with FIFO replacement and 2 frames are :

Answer: 10 ✓

4 -

4 2

5 2

5 1

0 1

-

2 1

2 5

4 5

4 1

2 1

The correct answer is: 10

Question 6

Correct

Mark 1.00 out of 1.00

Assuming a 16- KB page size, what is the page number for the address 428517 reference in decimal :

(give answer also in decimal)

Answer: 27 ✓

The correct answer is: 26



Question 7

Correct

Mark 1.00 out of 1.00

In the code below assume that each function can be executed concurrently by many threads/processes.
Ignore syntactical issues, and focus on the semantics.

This program is an example of

```
spinlock a, b; // assume initialized
thread1() {
    spinlock(b);
    //some code;
    spinlock(a);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
thread2() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
```

- a. Deadlock ✓
- b. Self Deadlock
- c. None of these
- d. Deadlock or livelock depending on actual race
- e. Livelock

Your answer is correct.

The correct answer is: Deadlock



Question 8

Partially correct

Mark 1.33 out of 2.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.
"..." means some code.

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

    // The + in "+m" denotes a read-modify-write operand.
    asm volatile("lock; xchgl %0, %1" :
        "+m" ("*addr), "=a" (result) :
        "1" (newval) :
        "cc");
    return result;
}
```

Atomic compare and swap instruction (to be expanded inline into code)



```
void
sleep(void *chan, struct spinlock *lk)
{
    ...
    if(lk != &ptable.lock){
        acquire(&ptable.lock);
        release(lk);
    }
}
```

If you don't do this, a process may be running on two processors parallelly



```
void
acquire(struct spinlock *lk)
{
    ...
    __sync_synchronize();
}
```

Tell compiler not to reorder memory access beyond this line



Your answer is partially correct.

You have correctly selected 2.

The correct answer is: static inline uint
xchg(volatile uint *addr, uint newval)

```
{
    uint result;
```

// The + in "+m" denotes a read-modify-write operand.

```
asm volatile("lock; xchgl %0, %1" :
    "+m" ("*addr), "=a" (result) :
    "1" (newval) :
    "cc");
return result;
} → Atomic compare and swap instruction (to be expanded inline into code), void
sleep(void *chan, struct spinlock *lk)
{
    ...
    if(lk != &ptable.lock){
        acquire(&ptable.lock);
        release(lk);
    } → Avoid a self-deadlock, void
    acquire(struct spinlock *lk)
{
    ...
    __sync_synchronize(); → Tell compiler not to reorder memory access beyond this line
}
```

Question 9

Correct

Mark 1.00 out of 1.00

Predict the output of the program given here.

Assume that all the path names for the programs are correct. For example "/usr/bin/echo" will actually run echo command.

Assume that there is no mixing of print output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good

output

should be written as good output

--

```
main() {  
    int i;  
    i = fork();  
    if(i == 0)  
        execl("/usr/bin/echo", "/usr/bin/echo", "hi", 0);  
    else  
        wait(0);  
    fork();  
    execl("/usr/bin/echo", "/usr/bin/echo", "one", 0);  
}
```

Answer: hi one one 

The correct answer is: hi one one

Question 10

Partially correct

Mark 1.67 out of 2.00

Select all the blocks that may need to be written back to disk (if updated, of-course), as "Yes", when an operation of deleting a file is carried out on ext2 file system.

An option has to be correct entirely to be marked "Yes"

Superblock

Yes 

One or multiple data blocks of the parent directory

No 

One or more data bitmap blocks for the parent directory

No 

Block bitmap(s) for all the blocks of the file

No 

Possibly one block bitmap corresponding to the parent directory

Yes 

Data blocks of the file

No 

Your answer is partially correct.

only one data block of parent directory. multiple blocks not possible. an entry is always contained within one single block

You have correctly selected 5.

The correct answer is: Superblock → Yes, One or multiple data blocks of the parent directory → No, One or more data bitmap blocks for the parent directory → No, Block bitmap(s) for all the blocks of the file → Yes, Possibly one block bitmap corresponding to the parent directory → Yes, Data blocks of the file → No

Question 11

Correct

Mark 1.00 out of 1.00

Select all the correct statements about bootloader.

Every wrong selection will deduct marks proportional to $1/n$ where n is total wrong choices in the question.

You will get minimum a zero.

- a. Modern Bootloaders often allow configuring the way an OS boots
- b. Bootloaders allow selection of OS to boot from
- c. Bootloader must be one sector in length
- d. The bootloader loads the BIOS
- e. LILO is a bootloader



Your answer is correct.

The correct answers are: LILO is a bootloader, Modern Bootloaders often allow configuring the way an OS boots, Bootloaders allow selection of OS to boot from



Question 12

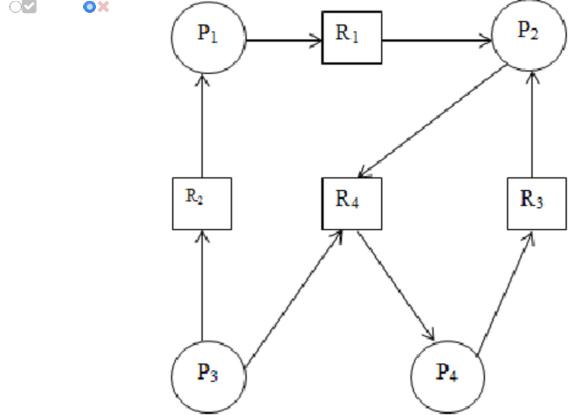
Incorrect

Mark 0.00 out of 1.00

For each of the resource allocation diagram shown,
infer whether the graph contains at least one deadlock or not.

Yes

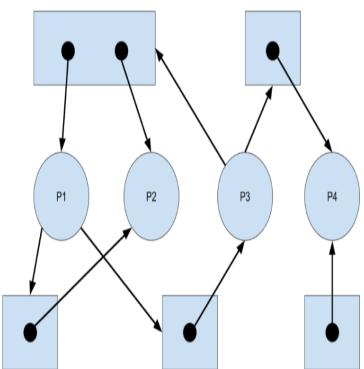
No



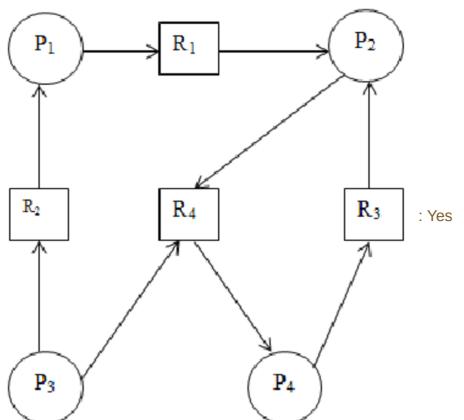
✗

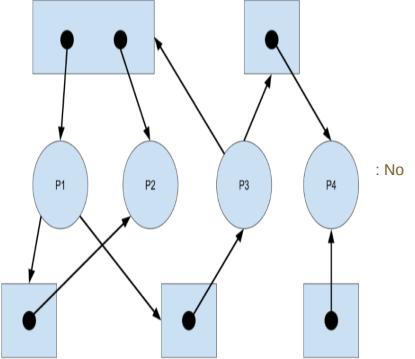
✗

✓



✗





Question 13

Partially correct

Mark 0.71 out of 1.00

Mark the statements about device drivers by marking as True or False.

True	False	
<input checked="" type="radio"/>	<input type="radio"/> ✘	It's possible that a particular hardware has multiple device drivers available for it.
<input checked="" type="radio"/>	<input type="radio"/> ✘	xv6 has device drivers for IDE disk and console.
<input checked="" type="radio"/>	<input type="radio"/> ✘	A disk driver converts OS's logical view of disk into physical locations on disk.
<input checked="" type="radio"/>	<input type="radio"/> ✘	A device driver code is specific to a hardware device
<input checked="" type="radio"/>	<input type="radio"/> ✘	All devices of the same type (e.g. 2 hard disks) can typically use the same device driver
<input checked="" type="radio"/>	<input type="radio"/> ✘	Writing a device driver mandatorily demands reading the technical documentation about the hardware.
<input type="radio"/> ✘	<input checked="" type="radio"/>	Device driver is an intermediary between the end-user and OS

It's possible that a particular hardware has multiple device drivers available for it.: True

xv6 has device drivers for IDE disk and console.: True

A disk driver converts OS's logical view of disk into physical locations on disk.: True

A device driver code is specific to a hardware device: True

All devices of the same type (e.g. 2 hard disks) can typically use the same device driver: True

Writing a device driver mandatorily demands reading the technical documentation about the hardware.: True

Device driver is an intermediary between the end-user and OS: False

Question 14

Partially correct

Mark 0.33 out of 1.00

Consider this program.

Some statements are identified using the // comment at the end.

Assume that `=` is an atomic operation.

```
#include <stdio.h>
#include <pthread.h>
long c = 0, c1 = 0, c2 = 0, run = 1;
void *thread1(void *arg) {
    while(run == 1) { //E
        c = 10; //A
        c1 = c2 + 5; //B
    }
}
void *thread2(void *arg) {
    while(run == 1) { //F
        c = 20; //C
        c2 = c1 + 3; //D
    }
}
int main() {
    pthread_t th1, th2;
    pthread_create(&th1, NULL, thread1, NULL);
    pthread_create(&th2, NULL, thread2, NULL);
    sleep(2);
    run = 0;
    printf(stdout, "c = %ld c1+c2 = %ld c1 = %ld c2 = %ld \n", c, c1+c2, c1, c2);
    fflush(stdout);
}
```

Which statements are part of the critical Section?

Yes	No	
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	F
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/> <input type="checkbox"/>	D
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	C
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	A
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input checked="" type="radio"/> <input type="checkbox"/>	B
<input checked="" type="radio"/> <input checked="" type="checkbox"/>	<input type="radio"/> <input checked="" type="checkbox"/>	E

F: No

D: Yes

C: No

A: No

B: Yes

E: No

Question 15

Partially correct

Mark 1.43 out of 2.00

Mark statements as T/F

All statements are in the context of preventing deadlocks.

True**False**

<input checked="" type="radio"/>	<input type="radio"/>	A process holding one resources and waiting for just one more resource can also be involved in a deadlock.	✓
<input type="radio"/>	<input checked="" type="radio"/>	If a resource allocation graph contains a cycle then there is a guarantee of a deadlock	✗
<input type="radio"/>	<input checked="" type="radio"/>	The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order	✗
<input checked="" type="radio"/>	<input type="radio"/>	Circular wait is avoided by enforcing a lock ordering	✓
<input checked="" type="radio"/>	<input type="radio"/>	Hold and wait means a thread/process holding some locks and waiting for acquiring some.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens	✓

A process holding one resources and waiting for just one more resource can also be involved in a deadlock.: True

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

Circular wait is avoided by enforcing a lock ordering: True

Hold and wait means a thread/process holding some locks and waiting for acquiring some.: True

Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

Question 16

Correct

Mark 1.00 out of 1.00

Match the left side use(or non-use) of a synchronization primitive with the best option on the right side.

This is the smallest primitive made available in software, using the hardware provided atomic instructions

 spinlock ✓

This tool is useful for event-wait scenarios

 semaphore ✓

This tool is more useful on multiprocessor systems

 spinlock ✓

This tool is quite attractive in solving the main bounded buffer problem

 semaphore ✓

This tool is very useful for waiting for 'something'

 condition variables ✓

Your answer is correct.

The correct answer is: This is the smallest primitive made available in software, using the hardware provided atomic instructions → spinlock, This tool is useful for event-wait scenarios → semaphore, This tool is more useful on multiprocessor systems → spinlock, This tool is quite attractive in solving the main bounded buffer problem → semaphore, This tool is very useful for waiting for 'something' → condition variables

Question 17

Correct

Mark 1.00 out of 1.00

The permissions -rwx--x--x on a file mean

- a. The file can be read only by the owner
- b. 'cat' on the file by owner will not work
- c. 'cat' on the file by any user will work
- d. 'rm' on the file by any user will work
- e. The file can be executed by anyone
- f. The file can be written only by the owner



Your answer is correct.

The correct answers are: The file can be executed by anyone, The file can be read only by the owner, The file can be written only by the owner, 'rm' on the file by any user will work

Question 18

Incorrect

Mark 0.00 out of 1.00

Note: for this question you get full marks if you select all and only correct options, you get ZERO if at least one option is wrong or not selected.

Select all the correct statements about log structured file systems.

- a. a transaction is said to be committed when all operations are written to file system
- b. log may be kept on same block device or another block device
- c. file system recovery may end up losing data
- d. even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery
- e. file system recovery recovers all the lost data



Your answer is incorrect.

The correct answers are: file system recovery may end up losing data, log may be kept on same block device or another block device, even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery

Question 19

Incorrect

Mark 0.00 out of 1.00

Consider the structure of directory entry in ext2, as shown in this diagram.

	inode	rec_len	file_type	name_len	name
0	21	12	1	2	.
12	22	12	2	2	.
24	53	16	5	2	h o m e
40	67	28	3	2	u s r
52	0	16	7	1	o l d f i l e
68	34	12	4	2	s b i n

Select the correct statements about the directory entry in ext2 file system.

The correct formula for rec_len is (when entries are continuously stored)

- a. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + (-1) * (\text{strlen(name)} \% 4))$
- b. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + (\text{strlen(name)} - 4) \% 4)$
- c. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + 4 - (\text{strlen(name)} \% 4))$ ✗
- d. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + (-1) * (\text{strlen(name)} - 4))$
- e. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} \% 4)$
- f. $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + \text{strlen(name)}$

Your answer is incorrect.

The correct answer is: $\text{rec_len} = \text{sizeof(inode entry)} + \text{sizeof(name len entry)} + \text{sizeof(file type entry)} + (\text{strlen(name)} + (-1) * (\text{strlen(name)} - 4))$

Question 20

Partially correct

Mark 0.50 out of 1.00

Mark whether the given sequence of events is possible or not-possible. Also, select the reason for your answer.

For each sequence it's a not-possible sequence if some important event is not mentioned in the sequence.

Assume that the kernel code is non-interruptible and uniprocessor system.

Process P1 executing a system call
 Timer interrupt
 Generic interrupt handler runs
 Scheduler runs
 Scheduler selects P2 for execution
 P2 returns from timer interrupt handler
 Process p2, user code executing

This sequence of events is: ✓

Because

✗

Question 21

Incorrect

Mark 0.00 out of 1.00

The given semaphore implementation faces which problem?

Assume any suitable code for signal()

Note: blocks means waits in a wait queue.

```
struct semaphore {  
    int val;  
    spinlock lk;  
};  
sem_init(semaphore *s, int initval) {  
    s->val = initval;  
    s->sl = 0;  
}  
wait(semaphore *s) {  
    spinlock(&(s->sl));  
    while(s->val <=0)  
        ;  
    (s->val)--;  
    spinunlock(&(s->sl));  
}
```

- a. blocks holding a spinlock
- b. deadlock
- c. too much spinning, bounded wait not guaranteed
- d. not holding lock after unblock



Your answer is incorrect.

The correct answer is: deadlock



Question 22

Partially correct

Mark 0.80 out of 1.00

Mark statements True/False w.r.t. change of states of a process.

Reference: The process state diagram (and your understanding of how kernel code works)

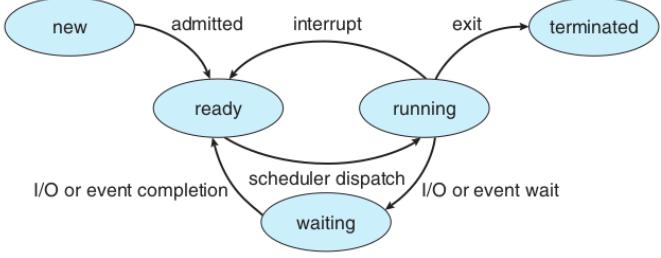


Figure 3.2 Diagram of process state.

True

False

<input type="radio"/> ✗	<input checked="" type="checkbox"/> ✗	A process in RUNNING state only can become TERMINATED because scheduler moves it to ZOMBIE state	✓
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.	✗
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred	✓
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	Every process has to go through ZOMBIE state, at least for a small duration.	✓
<input checked="" type="checkbox"/> ✗	<input type="radio"/> ✗	Only a process in READY state is considered by scheduler	✓

A process in RUNNING state only can become TERMINATED because scheduler moves it to ZOMBIE state: False

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred: True

Every process has to go through ZOMBIE state, at least for a small duration.: True

Only a process in READY state is considered by scheduler: True

Question 23

Correct

Mark 1.00 out of 1.00

Select T/F for statements about Volume Managers.

Do pay attention to the use of the words physical partition and physical volume.

True**False**

<input checked="" type="radio"/>	<input type="radio"/> ✗	The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A logical volume can be extended in size but upto the size of volume group	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A logical volume may span across multiple physical volumes	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	The volume manager stores additional metadata on the physical disk partitions	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A physical partition should be initialized as a physical volume, before it can be used by volume manager.	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A volume group consists of multiple physical volumes	<input checked="" type="checkbox"/>
<input checked="" type="radio"/>	<input type="radio"/> ✗	A logical volume may span across multiple physical partitions	<input checked="" type="checkbox"/> since a physical volume is made up of physical partitions, and a volume can span across multiple PVs, it can also span across multiple PP

The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.: True

A logical volume can be extended in size but upto the size of volume group: True

A logical volume may span across multiple physical volumes: True

The volume manager stores additional metadata on the physical disk partitions: True

A physical partition should be initialized as a physical volume, before it can be used by volume manager.: True

A volume group consists of multiple physical volumes: True

A logical volume may span across multiple physical partitions: True

Question 24

Correct

Mark 1.00 out of 1.00

Map the block allocation scheme with the problem it suffers from

(Match pairs 1-1, match a scheme with the problem that it suffers from relatively the most, compared to others)

Continuous allocation	need for compaction	<input checked="" type="checkbox"/>
Linked allocation	Too many seeks	<input checked="" type="checkbox"/>
Indexed Allocation	Overhead of reading metadata blocks	<input checked="" type="checkbox"/>

Your answer is correct.

The correct answer is: Continuous allocation → need for compaction, Linked allocation → Too many seeks, Indexed Allocation → Overhead of reading metadata blocks

Question 25

Correct

Mark 1.00 out of 1.00

This one is not a system call:

- a. open
- b. read
- c. write
- d. scheduler



Your answer is correct.

The correct answer is: scheduler



Question 26

Correct

Mark 1.00 out of 1.00

Match the pairs.

This question is based on your general knowledge about operating systems/related concepts and their features.

Java threads	monitors,re-entrant locks, semaphores	✓
Linux threads	atomic-instructions, spinlocks, etc.	✓
POSIX threads	semaphore, mutex, condition variables	✓

Your answer is correct.

The correct answer is: Java threads → monitors,re-entrant locks, semaphores, Linux threads → atomic-instructions, spinlocks, etc., POSIX threads → semaphore, mutex, condition variables

Question 27

Correct

Mark 1.00 out of 1.00

Consider the following list of free chunks, in continuous memory management:

7k, 15k, 21k, 14k, 19k, 6k

Suppose there is a request for chunk of size 5k, then the free chunk selected under each of the following schemes will be

Best fit:	6k	✓
First fit:	7k	✓
Worst fit:	21k	✓

Question 28

Correct

Mark 1.00 out of 1.00

This one is not a scheduling algorithm

- a. Round Robin
- b. SJF
- c. Mergesort
- d. FCFS



Your answer is correct.

The correct answer is: Mergesort

Question 29

Correct

Mark 1.00 out of 1.00

Mark whether the concept is related to scheduling or not.

Yes	No	
<input checked="" type="radio"/>	<input type="radio"/>	timer interrupt
<input checked="" type="radio"/>	<input type="radio"/>	context-switch
<input checked="" type="radio"/>	<input type="radio"/>	ready-queue
<input type="radio"/>	<input checked="" type="radio"/>	file-table
<input checked="" type="radio"/>	<input type="radio"/>	runnable process

timer interrupt: Yes

context-switch: Yes

ready-queue: Yes

file-table: No

runnable process: Yes



Question 30

Partially correct

Mark 1.00 out of 2.00

Map ext2 data structure features with their purpose

Many copies of Superblock Choose...**Free blocks count in superblock and group descriptor**

Redundancy to ensure the most crucial data structure is not lost

**Used directories count in group descriptor**

is redundant and helps do calculations of directory entries faster

**Combining file type and access rights in one variable**

saves 1 byte of space

**rec_len field in directory entry**

Try to keep all the data of a directory and its file close together in a group

**File Name is padded**

aligns all memory accesses on word boundary, improving performance

**Inode bitmap is one block**

limits total number of files that can belong to a group

**Block bitmap is one block**

Limits the size of a block group, thus improvising on purpose of a group

**Mount count in superblock**

to enforce file check after certain amount of mounts at boot time

**Inode table location in Group Descriptor**

is redundant and helps do calculations of directory entries faster

**Inode table**

All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk

**A group**

Redundancy to ensure the most crucial data structure is not lost



Your answer is partially correct.

You have correctly selected 6.

The correct answer is: **Many copies of Superblock** → Redundancy to ensure the most crucial data structure is not lost, **Free blocks count in superblock and group descriptor** → Redundancy to help fsck restore consistency, **Used directories count in group descriptor** → attempt is made to evenly spread the first-level directories, this count is used there, **Combining file type and access rights in one variable** → saves 1 byte of space, **rec_len field in directory entry** → allows holes and linking of entries in directory, File Name is padded → aligns all memory accesses on word boundary, improving performance, **Inode bitmap is one block** → limits total number of files that can belong to a group, **Block bitmap is one block** → Limits the size of a block group, thus improvising on purpose of a group, **Mount count in superblock** → to enforce file check after certain amount of mounts at boot time, **Inode table location in Group Descriptor** → Obvious, as it's per group and not per file-system, **Inode table** → All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk, **A group** → Try to keep all the data of a directory and its file close together in a group

Question 31

Partially correct

Mark 1.85 out of 2.00

Mark True/False

Statements about scheduling and scheduling algorithms

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The nice() system call is used to set priorities for processes
<input checked="" type="radio"/>	<input type="radio"/>	Aging is used to ensure that low-priority processes do not starve in priority scheduling.
<input type="radio"/>	<input checked="" type="radio"/>	In non-pre-emptive priority scheduling, the highest priority process is scheduled and runs until it gives up CPU.
<input checked="" type="radio"/>	<input type="radio"/>	xv6 code does not care about Processor Affinity
<input checked="" type="radio"/>	<input type="radio"/>	In pre-emptive priority scheduling, priority is implemented by assigning more time quantum to higher priority process.
<input checked="" type="radio"/>	<input type="radio"/>	A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.
<input checked="" type="radio"/>	<input type="radio"/>	Processor Affinity refers to memory accesses of a process being stored on cache of that processor
<input checked="" type="radio"/>	<input type="radio"/>	Response time will be quite poor on non-interruptible kernels
<input checked="" type="radio"/>	<input type="radio"/>	Shortest Remaining Time First algorithm is nothing but pre-emptive Shortest Job First algorithm
<input checked="" type="radio"/>	<input type="radio"/>	On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread
<input checked="" type="radio"/>	<input type="radio"/>	Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.
<input checked="" type="radio"/>	<input type="radio"/>	Pre-emptive scheduling leads to many race conditions in kernel code.
<input checked="" type="radio"/>	<input type="radio"/>	Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.

The nice() system call is used to set priorities for processes.: True

Aging is used to ensure that low-priority processes do not starve in priority scheduling.: True

In non-pre-emptive priority scheduling, the highest priority process is scheduled and runs until it gives up CPU.: True

xv6 code does not care about Processor Affinity: True

In pre-emptive priority scheduling, priority is implemented by assigning more time quantum to higher priority process.: True

A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.: True

Processor Affinity refers to memory accesses of a process being stored on cache of that processor: True

Response time will be quite poor on non-interruptible kernels: True

Shortest Remaining Time First algorithm is nothing but pre-emptive Shortest Job First algorithm: True

On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread: True

Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.: True

Pre-emptive scheduling leads to many race conditions in kernel code.: True

Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.: True

Question 32

Partially correct

Mark 1.17 out of 2.00

The unix file semantics demand that changes to any open file are visible immediately to any other processes accessing that file at that point in time.

Select the data-structure/programmatic features that ensure the implementation of unix semantics. (Assume there is no mmap())

Yes	No	
<input type="radio"/> <input checked="" type="checkbox"/>	All processes accessing the same file share the file descriptor among themselves	✓
<input type="radio"/> <input checked="" type="checkbox"/>	The pointer entry in the file descriptor array entry points to the data of the file directly	✓
<input checked="" type="checkbox"/> <input type="radio"/>	There is only one global file structure per on-disk file.	✗
<input type="radio"/> <input checked="" type="checkbox"/>	All file accesses are made using only global variables	✓
<input checked="" type="checkbox"/> <input type="radio"/>	The 'file offset' is shared among all the processes that access the file.	✗
<input type="radio"/> <input checked="" type="checkbox"/>	No synchronization is implemented so that changes are made available immediately.	✓
<input checked="" type="checkbox"/> <input type="radio"/>	A single spinlock is to be used to protect the unique global 'file structure' representing the file, thus synchronizing access, and making other processes wait for earlier process to finish writing so that writes get visible immediately.	✗
<input checked="" type="checkbox"/> <input type="radio"/>	There is only one in-memory copy of the on disk file's contents in kernel memory/buffers	✓
<input checked="" type="checkbox"/> <input type="radio"/>	The file descriptors in every PCB are pointers to the same global file structure.	✗
<input type="radio"/> <input checked="" type="checkbox"/>	The file descriptor array is external to PCB and all processes that share a file, have pointers to same file-descriptors' array	✓
<input checked="" type="checkbox"/> <input type="radio"/>	All file structures representing any open file, give access to the same in-memory copy of the file's contents	✓
<input checked="" type="checkbox"/> <input type="radio"/>	The 'file offset' index is stored outside the file-structure to which file-descriptor array points	✗

All processes accessing the same file share the file descriptor among themselves: No

The pointer entry in the file descriptor array entry points to the data of the file directly: No

There is only one global file structure per on-disk file.: No

All file accesses are made using only global variables: No

The 'file offset' is shared among all the processes that access the file.: No

No synchronization is implemented so that changes are made available immediately.: No

A single spinlock is to be used to protect the unique global 'file structure' representing the file, thus synchronizing access, and making other processes wait for earlier process to finish writing so that writes get visible immediately.: No

There is only one in-memory copy of the on disk file's contents in kernel memory/buffers: Yes

The file descriptors in every PCB are pointers to the same global file structure.: No

The file descriptor array is external to PCB and all processes that share a file, have pointers to same file-descriptors' array: No

All file structures representing any open file, give access to the same in-memory copy of the file's contents: Yes

The 'file offset' index is stored outside the file-structure to which file-descriptor array points: No

Question 33

Partially correct

Mark 0.33 out of 2.00

Map the function in xv6's file system code, to its perceived logical layer.

namei	inode	✗
filestat()	Choose...	
dirlookup	directory	✓
ialloc	file descriptor	✗
stati	Choose...	
ideintr	buffer cache	✗
bread	Choose...	
balloc	file descriptor	✗
sys_chdir()	system call	✓
skipelem	system call	✗
commit	system call	✗
bmap	system call	✗

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: namei → pathname lookup, filestat() → file descriptor, dirlookup → directory, ialloc → inode, stati → inode, ideintr → disk driver, bread → buffer cache, balloc → block allocation on disk, sys_chdir() → system call, skipelem → pathname lookup, commit → logging, bmap → inode

[◀ Course Exit Feedback](#)[Jump to...](#)[xv6-public-master ►](#)

Started on Saturday, 30 April 2022, 2:05:41 PM

State Finished

Completed on Saturday, 30 April 2022, 6:47:56 PM

Time taken 4 hours 42 mins

Grade **34.99** out of 40.00 (87%)

Question 1

Correct

Mark 1.00 out of 1.00

which of the following scheduling algorithms discriminate either in favor of or against short processes, from the perspective of minimizing waiting time.

For	Against
<input type="radio"/> ✗	<input checked="" type="radio"/> FCFS
<input checked="" type="radio"/> ✗	<input type="radio"/> Round Robin
<input checked="" type="radio"/> ✗	<input type="radio"/> Multilevel feedback queues

FCFS: Against

Round Robin: For

Multilevel feedback queues: For

Question 2

Correct

Mark 1.00 out of 1.00

Consider a computer system with a 32-bit logical address and 8- KB page size. The system supports up to 1 GB of physical memory. How many entries are there in each of the following?

a. A conventional, single-level page table (write a decimal number):

524288



b. An inverted page table (number of entries only, write a decimal number):

131072



Question 3

Correct

Mark 1.00 out of 1.00

Select T/F for the disk block allocation scheme related statements

True	False	
<input checked="" type="radio"/>	<input type="radio"/> X	Continuous allocation leads to faster file access
<input checked="" type="radio"/>	<input type="radio"/> X	FAT uses linked allocation
<input checked="" type="radio"/>	<input type="radio"/> X	The unix inode is based on indexed allocation
<input checked="" type="radio"/>	<input type="radio"/> X	Continuous allocation allows to fetch any block with just one seek
<input checked="" type="radio"/>	<input type="radio"/> X	NVM storage devices are pushing for search for new block allocation schemes
<input checked="" type="radio"/>	<input type="radio"/> X	Continuous allocation may involve a costly search for free space
<input checked="" type="radio"/>	<input type="radio"/> X	Maximum file size limit is determined by disk block allocation scheme, as one of the factors.
<input checked="" type="radio"/>	<input type="radio"/> X	Linked allocation does away with file size limitation to a large extent

Continuous allocation leads to faster file access: True

FAT uses linked allocation: True

The unix inode is based on indexed allocation: True

Continuous allocation allows to fetch any block with just one seek: True

NVM storage devices are pushing for search for new block allocation schemes: True

Continuous allocation may involve a costly search for free space: True

Maximum file size limit is determined by disk block allocation scheme, as one of the factors.: True

Linked allocation does away with file size limitation to a large extent: True

Question 4

Correct

Mark 1.00 out of 1.00

Which one of the following is not a system call

- a. lock
- b. open
- c. lseek
- d. mount

The correct answer is: lock

Question 5

Partially correct

Mark 1.13 out of 1.50

The following processes are being scheduled using a pre-emptive, priority-based, round-robin scheduling algorithm.

<u>Process</u>	<u>Priority</u>	<u>Burst</u>	<u>Arrival</u>
P_1	8	15	0
P_2	3	20	0
P_3	4	20	20
P_4	4	20	25
P_5	5	5	45
P_6	5	15	55

Each process is assigned a numerical priority, with a higher number indicating a higher relative priority. The scheduler will execute the currently highest priority process for its full duration, unless it gets pre-empted by newly arriving higher priority process. For processes with the same priority, a round-robin scheduler will be used with a time quantum of 10 units. If a process is pre-empted by a higher-priority process, the pre-empted process is placed at the front of the queue of same priority processes (so that its turn continues when higher priority process is over).

The order in which the processes get scheduled is (write your answer without a space, e.g. P1,P2,P3,P4,P5) :

P1,P2,P3,P4,P3,P5,P3,P6,P4,P2



The turn-around time for the process P2 is:

95



The turn-around time for the process P6 is:

15



The process P5 finishes at time unit:

50



0-15 P1

15-20 P2

20-40 P3

40-45 P4

45-50 P5

50-55 P4

55-70 P6

70-80 P4

80-95 P2

--

P2 turnaround time = 95 - 15 = 80

Question 6

Partially correct

Mark 1.62 out of 3.00

Suppose it is required to add the chown() (without notion of a group, just the notion of owner and others) system call to xv6. Select the changes, from the options given below, which are an absolute must to effect the addition of this system call.

Changes w.r.t. other system calls should be selected if those system calls are necessary for the implementation of the chown() system call.

Must	Not-Must	
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Add a EUID field in the struct proc representing Effective user ID
<input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Add a mode field representing file permissions, inside dinode
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Add a UID field in the struct proc represenging the USER-ID of the user who started this process
<input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Add a uid field representing the owner, inside dinode
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Add few extra files belonging to different users, using mkfs.c
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Modify exec() to check SUID bit on the executable file and set EUID of the process to UID of the executable
<input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Add the chown() system call which checks if the user with id "0" is calling this system call and then change the ownership of the on-disk and in-memory inode.
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Add a setuid(), seteuid() system call, callable only by the user with ID or EUID equal to "0" to set the new user id of the process
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Create an application program su.c which itself is a SUID program, and calls setuid() and setuid() with specified user-ID and then does exec() of a shell.
<input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Modify mkfs.c to add owner and permissions to each file created
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Inherit the UID and EUID of the parent in fork()
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Mandatorily create a file like "passwd" which maps user-names to user-IDs and modify other utilities (like "ls") to show user-name instead of user-ID for the files.
<input checked="" type="radio"/>	<input checked="" type="checkbox"/>	Add the username field to the on disk inode, that is dinode.

Add a EUID field in the struct proc representing Effective user ID: Not-Must

Add a mode field representing file permissions, inside dinode: Must

Add a UID field in the struct proc represenging the USER-ID of the user who started this process: Not-Must

Add a uid field representing the owner, inside dinode: Must

Add few extra files belonging to different users, using mkfs.c: Not-Must

Modify exec() to check SUID bit on the executable file and set EUID of the process to UID of the executable: Not-Must

Add the chown() system call which checks if the user with id "0" is calling this system call and then change the ownership of the on-disk and in-memory inode.: Must

Add a setuid(), seteuid() system call, callable only by the user with ID or EUID equal to "0" to set the new user id of the process: Not-Must

Create an application program su.c which itself is a SUID program, and calls setuid() and setuid() with specified user-ID and then does exec() of a shell.: Not-Must

Modify mkfs.c to add owner and permissions to each file created: Must
 Inherit the UID and EUID of the parent in fork(): Not-Must
 Mandatorily create a file like "passwd" which maps user-names to user-IDs and modify other utilities (like "ls") to show user-name instead of user-ID for the files.: Not-Must
 Add the username field to the on disk inode, that is dinode.: Not-Must

Question 7

Correct

Mark 1.00 out of 1.00

The following diagram that explains the concept of multi-threading

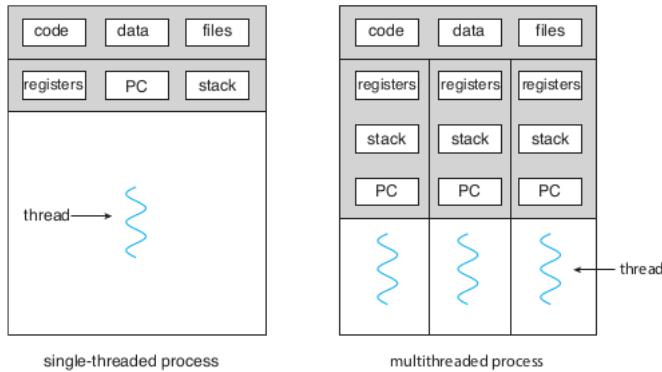


Figure 4.1 Single-threaded and multithreaded processes.

Leads to the following conclusions about implementation of threads

True	False	
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	A multi-threaded program can be split in multiple ELF files.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The memory management for the process in the kernel should not (most typically) change due to creation of a thread.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	Switching between threads requires a "context switch" with change of registers involved in it
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	Each thread should be represented by a function that starts execution on a separate stack.

A multi-threaded program can be split in multiple ELF files.: False

The memory management for the process in the kernel should not (most typically) change due to creation of a thread.: True

Switching between threads requires a "context switch" with change of registers involved in it: True

Each thread should be represented by a function that starts execution on a separate stack.: True

Question 8

Correct

Mark 1.00 out of 1.00

Consider this program.

Some statements are identified using the // comment at the end.

Assume that = is an atomic operation.

```
#include <stdio.h>
#include <pthread.h>
long c = 0, c1 = 0, c2 = 0, run = 1;
void *thread1(void *arg) {
    while(run == 1) { //E
        c = c1; //A
        c1++; //B
    }
}
void *thread2(void *arg) {
    while(run == 1) { //F
        c = c2; //C
        c2++; //D
    }
}
int main() {
    pthread_t th1, th2;
    pthread_create(&th1, NULL, thread1, NULL);
    pthread_create(&th2, NULL, thread2, NULL);
    sleep(2);
    run = 0;
    fprintf(stdout, "c = %ld c1+c2 = %ld c1 = %ld c2 = %ld \n", c, c1+c2, c1, c2);
    fflush(stdout);
}
```

Which statements are part of the critical Section?

Yes	No
<input type="radio"/> ✗	<input checked="" type="checkbox"/> E
<input type="radio"/> ✗	<input checked="" type="checkbox"/> F
<input type="radio"/> ✗	<input checked="" type="checkbox"/> B
<input type="radio"/> ✗	<input checked="" type="checkbox"/> D
<input checked="" type="checkbox"/>	<input type="radio"/> A
<input checked="" type="checkbox"/>	<input type="radio"/> C

As per Remzi's book: "A critical section is a piece of code that accesses a shared variable (or more generally, a shared resource and must not be concurrently executed by more than one thread."

As per Galvin's Book: a critical section, in which the process may be accessing — and updating — data that is shared with at least one other process. The important feature of the system is that when one process is executing in its critical section, no other process is allowed to execute in its critical section.

Since A and C refer to a variable that is shared, it's critical section. Here the hardware guarantees (as = is atomic) that the critical section is accessed by only one thread at a time.

E: No

F: No

B: No

D: No

A: Yes

C: Yes

Question 9

Partially correct

Mark 0.38 out of 1.00

It is proposed that when a process does an illegal memory access, xv6 terminate the process by printing the error message "Illegal Memory Access". Select all the changes that need to be done to xv6 for this as True (Note that the changes proposed here may not cover the exhaustive list of all changes required) and the un-necessary/wrong changes as False.

Required	Un-necessary/Wrong	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change in the Makefile and instruct cc/ld to start the code of each program at some address other than 0
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Ensure that the address 0 is mapped to invalid
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Mark each page as readonly in the page table mappings
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change allocuvm() to call mappages() with proper permissions on each page table entry
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Add code that checks if the illegal memory access trap was due to an actual illegal memory access.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Handle the Illegal memory acceess trap in trap() function, and terminate the currently running process.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change exec to treat text/data sections separately and call allocuvm() with proper flags for page table entries
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change mappages() to set specified permissions on each page table entry

Change in the Makefile and instruct cc/ld to start the code of each program at some address other than 0: Required

Ensure that the address 0 is mapped to invalid: Required

Mark each page as readonly in the page table mappings: Un-necessary/Wrong

Change allocuvm() to call mappages() with proper permissions on each page table entry: Required

Add code that checks if the illegal memory access trap was due to an actual illegal memory access.: Un-necessary/Wrong

Handle the Illegal memory acceess trap in trap() function, and terminate the currently running process.: Required

Change exec to treat text/data sections separately and call allocuvm() with proper flags for page table entries: Required

Change mappages() to set specified permissions on each page table entry: Un-necessary/Wrong

Question 10

Partially correct

Mark 0.63 out of 1.00

Select which of the following data structures may need an update, in ext2 file system, when 5 bytes get removed from the end of an existing file

Yes	No	
<input type="radio"/> ✘	<input checked="" type="radio"/> ✓	The inode of the parent directory
<input checked="" type="radio"/> ✓	<input type="radio"/> ✘	The superblock
<input checked="" type="radio"/> ✓	<input type="radio"/> ✘	Inode of the file
<input type="radio"/> ✘	<input checked="" type="radio"/> ✓	The inode bitmap
<input checked="" type="radio"/> ✓	<input type="radio"/> ✘	The group descriptor
<input type="radio"/> ✘	<input type="radio"/> ✓	A free block on the file-system
<input checked="" type="radio"/> ✓	<input type="radio"/> ✘	One of the block bitmaps
<input type="radio"/> ✘	<input checked="" type="radio"/> ✓	The boot sector

The inode of the parent directory: No

The superblock: Yes

Inode of the file: Yes

The inode bitmap: No

The group descriptor: Yes

A free block on the file-system: No

One of the block bitmaps: Yes

The boot sector: No

Question 11

Partially correct

Mark 1.75 out of 2.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.

"..." means some code.

```
void
sleep(void *chan, struct spinlock *lk)
{
...
if(lk != &ptable.lock){
    acquire(&ptable.lock);
    release(lk);
}

void
panic(char *s)
{
...
panicked = 1;
```

If you don't do this, a process may be running on two processors parallelly ✖

```
struct proc*
myproc(void) {
...
pushcli();
c = mycpu();
p = c->proc;
popcli();
...
}
```

Disable interrupts to avoid another process's pointer being returned ✓

```
void
acquire(struct spinlock *lk)
{
...
__sync_synchronize();
```

Tell compiler not to reorder memory access beyond this line ✓

```
void
acquire(struct spinlock *lk)
{
...
getcallerpcs(&lk, lk->pcs);
```

Traverse ebp chain to get sequence of instructions followed in functions calls ✓

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

// The + in "+m" denotes a read-modify-write operand.
asm volatile("lock; xchgl %0, %1":
    "+m" (*addr), "=a" (result) :
    "1" (newval) :
    "cc");
    return result;
}
```

Atomic compare and swap instruction (to be expanded inline into code) ✓

```
void
yield(void)
{
...
release(&ptable.lock);
}
```

Release the lock held by some another process

✓

```
void
acquire(struct spinlock *lk)
{
    pushcli();
```

Disable interrupts to avoid deadlocks

✓

Your answer is partially correct.

You have correctly selected 7.

The correct answer is: void
sleep(void *chan, struct spinlock *lk)
{
...
if(lk != &ptable.lock){
 acquire(&ptable.lock);
 release(lk);
} → Avoid a self-deadlock, void
panic(char *s)
{
...
panicked = 1; → Ensure that no printing happens on other processors, struct proc*
myproc(void) {
...
 pushcli();
 c = mycpu0;
 p = c->proc;
 popcli();
...
}

→ Disable interrupts to avoid another process's pointer being returned, void
acquire(struct spinlock *lk)
{
...
__sync_synchronize();

→ Tell compiler not to reorder memory access beyond this line, void
acquire(struct spinlock *lk)
{
...
getcallerpcs(&lk, lk->pcs);

→ Traverse ebp chain to get sequence of instructions followed in functions calls, static inline uint
xchg(volatile uint *addr, uint newval)
{
 uint result;

 // The + in "+m" denotes a read-modify-write operand.
 asm volatile("lock; xchgl %0, %1" :
 "+m" (*addr), "=a" (result) :
 "1" (newval) :
 "cc");
 return result;
} → Atomic compare and swap instruction (to be expanded inline into code), void

```

yield(void)
{
...
release(&ptable.lock);
}

```

→ Release the lock held by some another process, **void**

acquire(struct spinlock *lk)

{

pushcli();

→ Disable interrupts to avoid deadlocks

Question 12

Correct

Mark 1.00 out of 1.00

Mark statements as T/F

All statements are in the context of preventing deadlocks.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens
<input type="radio"/>	<input checked="" type="radio"/>	The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order
<input checked="" type="radio"/>	<input type="radio"/>	The lock ordering to be followed to avoid circular wait is a protocol to be followed by programmers.
<input checked="" type="radio"/>	<input type="radio"/>	Circular wait is avoided by enforcing a lock ordering
<input checked="" type="radio"/>	<input type="radio"/>	If a resource allocation graph contains a cycle then there is a possibility of a deadlock
<input type="radio"/>	<input checked="" type="radio"/>	If a resource allocation graph contains a cycle then there is a guarantee of a deadlock
<input checked="" type="radio"/>	<input type="radio"/>	Deadlock is not possible if any of these conditions is not met: Mutual exclusion, hold and wait, no pre-emption, circular wait.

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

The lock ordering to be followed to avoid circular wait is a protocol to be followed by programmers.: True

Circular wait is avoided by enforcing a lock ordering: True

If a resource allocation graph contains a cycle then there is a possibility of a deadlock: True

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

Deadlock is not possible if any of these conditions is not met: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

Question 13

Correct

Mark 1.00 out of 1.00

Predict the output of the program given here.

Assume that there is no mixing of printf output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good

output

should be written as good output

--

```
int main() {  
    int pid;  
    printf("hi\n");  
    pid = fork();  
    if(pid == 0) {  
        exit(0);  
    }  
    printf("bye\n");  
    fork();  
    printf("ok\n");  
}
```

Answer: hi bye ok ok



The correct answer is: hi bye ok ok

Question 14

Partially correct

Mark 0.63 out of 1.00

For Virtual File System to work, which of the following changes are required to be done to an existing OS code (e.g. xv6)?

- a. A mount() system call should be provided to mount a partition onto some directory in existing namespace rooted at "/"
- b. The filesystem related system calls (e.g. read, write) need to invoke the file system specific functions (e.g. ext2_read, ext2_write, ntfs_read, ntfs_write) using function pointers. ✓
- c. The file system specific function pointers, for file system system-calls, need to be setup in the generic inode during lookup. ✓
- d. The generic inode needs to have a field representing if this inode is a mount point and also to refer/point to the root of the mounted file system's inode.
- e. Each open() needs to copy the function pointers from the inode of the parent directory into the inode of the child (if not already done), unless it's traversing a mount point. (This may be done as part of lookup() which is called by open()) ✓
- f. Each file-system writer needs to provide the set of function pointers for VFS, and these function pointers need to be setup in generic inode of "/" of that file system during mount() ✓
- g. The operating system in-memory inode needs to be a generic-inode representing "inode" like data structure across multiple file systems.
- h. The lookup() operation needs to check if it's crossing a mount point and call FS specific operations to read inodes/directories ✓

The correct answers are: A mount() system call should be provided to mount a partition onto some directory in existing namespace rooted at "/", The filesystem related system calls (e.g. read, write) need to invoke the file system specific functions (e.g. ext2_read, ext2_write, ntfs_read, ntfs_write) using function pointers., The file system specific function pointers, for file system system-calls, need to be setup in the generic inode during lookup., The operating system in-memory inode needs to be a generic-inode representing "inode" like data structure across multiple file systems., The generic inode needs to have a field representing if this inode is a mount point and also to refer/point to the root of the mounted file system's inode., The lookup() operation needs to check if it's crossing a mount point and call FS specific operations to read inodes/directories, Each file-system writer needs to provide the set of function pointers for VFS, and these function pointers need to be setup in generic inode of "/" of that file system during mount(), Each open() needs to copy the function pointers from the inode of the parent directory into the inode of the child (if not already done), unless it's traversing a mount point. (This may be done as part of lookup() which is called by open())

Question 15

Partially correct

Mark 0.40 out of 0.50

Which of the following statements are correct about xv6 and multiprocessor support ?

- a. Each processor on xv6 starts code execution in mpenter() when first processor sets "started" to 1.
- b. xv6 supports only SMP ✓
- c. In xv6 on x86, the first processor configures other processors in mpinit() ✓
- d. xv6 supports a variable number of processors (upto 8) and adjusts its data structures according to number of processors ✓
- e. At any point in time, after main(), the kernel may be parallelly executing on any of the processors. ✓

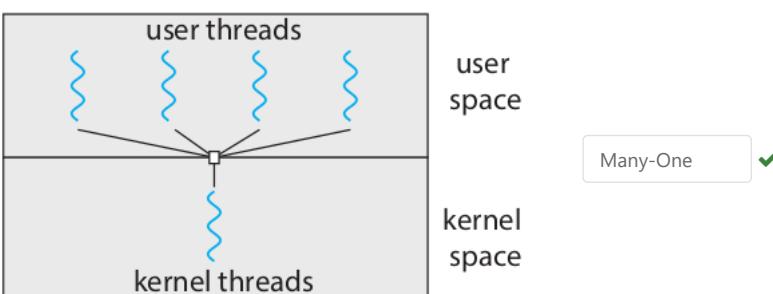
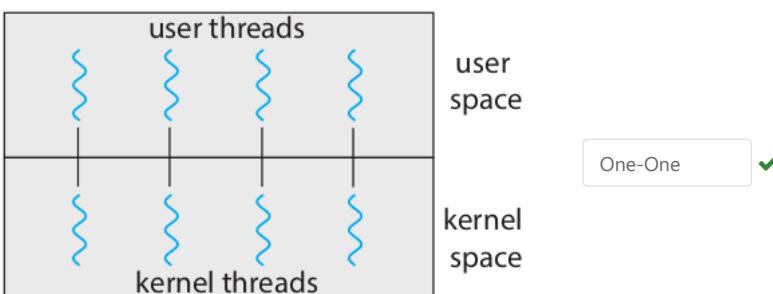
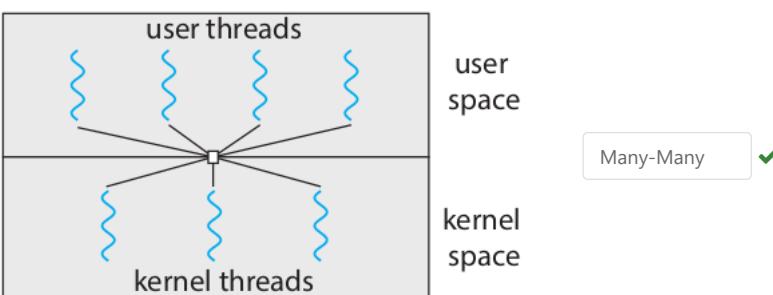
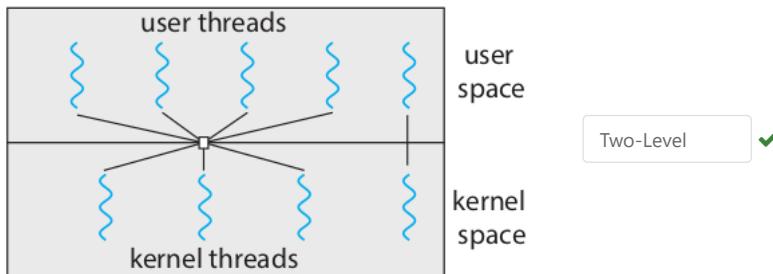
The correct answers are: xv6 supports only SMP, xv6 supports a variable number of processors (upto 8) and adjusts its data structures according to number of processors, Each processor on xv6 starts code execution in mpenter() when first processor sets "started" to 1., In xv6 on x86, the first processor configures other processors in mpinit(), At any point in time, after main(), the kernel may be parallelly executing on any of the processors.

Question 16

Correct

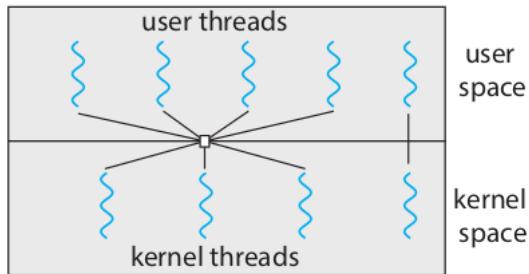
Mark 1.00 out of 1.00

Match the diagram with the threading model

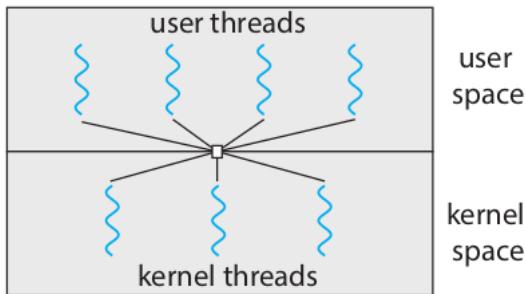


Your answer is correct.

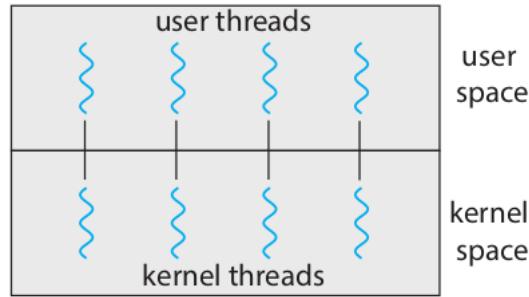
The correct answer is:



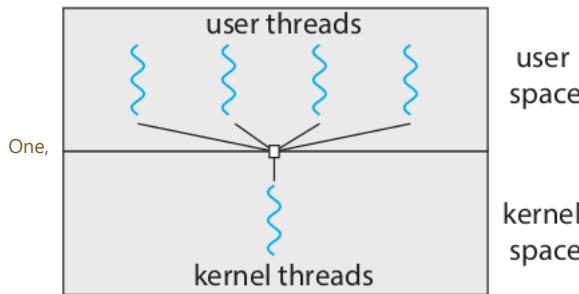
→ Two-Level,



→ Many-Many,



→ One-



→ Many-One

Question 17

Correct

Mark 1.00 out of 1.00

Select all the correct statements about Shell

- a. The essential job of the shell is to fork-exec the specified application ✓
- b. Examples of shell are: bash, ksh, csh, msh, ooosh, nosh, etc. ✗
- c. Shell is a layer on top of hardware ✗
- d. Shell converts the application code into system calls. ✗
- e. The default shell for a user is specified in /etc/passwd on typical GNU/Linux systems. ✓
- f. Examples of shell are: bash, ksh, csh, zsh, sh, etc. ✓

The correct answers are: The essential job of the shell is to fork-exec the specified application, Examples of shell are: bash, ksh, csh, zsh, sh, etc., The default shell for a user is specified in /etc/passwd on typical GNU/Linux systems.

Question 18

Correct

Mark 1.00 out of 1.00

Which one of the following is not a scheduling algorithm

- a. Multilevel Feedback Queue
- b. Priority
- c. FCFS
- d. paging



The correct answer is: paging

Question 19

Correct

Mark 1.00 out of 1.00

Mark the statements as True/False w.r.t. the basic concepts of memory management.

True	False	
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The compiler interacts with the kernel continuously while compiling a program and obtains the correct set of memory addresses for code/stack/heap/data and then generates the machine code file.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The compiler generates the address references for code/data/stack/heap in the executable file as per the memory management schema chosen by the compiler itself, and then the kernel ensures that program is executed with this schema.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	When a process is executing, each virtual address is converted into physical address by the CPU hardware directly.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	When a process is executing, each virtual address is converted into physical address by the kernel directly.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The compiler generates address references for code/data/stack/heap in the executable file, depending on the MM architecture provided by CPU and kernel.
<input type="radio"/> ✗	<input checked="" type="radio"/> ✗	The kernel refers to the page table for converting each virtual address to physical address.
<input checked="" type="radio"/> ✗	<input type="radio"/> ✗	The kernel ensures that the MMU is setup before scheduling a process and then the CPU/MMU ensures that the address translation takes place.

The compiler interacts with the kernel continuously while compiling a program and obtains the correct set of memory addresses for code/stack/heap/data and then generates the machine code file.: False

The compiler generates the address references for code/data/stack/heap in the executable file as per the memory management schema chosen by the compiler itself, and then the kernel ensures that program is executed with this schema.: False

When a process is executing, each virtual address is converted into physical address by the CPU hardware directly.: True

When a process is executing, each virtual address is converted into physical address by the kernel directly.: False

The compiler generates address references for code/data/stack/heap in the executable file, depending on the MM architecture provided by CPU and kernel.: True

The kernel refers to the page table for converting each virtual address to physical address.: False

The kernel ensures that the MMU is setup before scheduling a process and then the CPU/MMU ensures that the address translation takes place.: True

Question 20

Partially correct

Mark 0.83 out of 1.00

Which of the following instructions should be privileged?

- a. Modify entries in device-status table. ✓
- b. Read the clock.
- c. Switch from user to kernel mode. ✓
- d. Set value of timer. ✓
- e. Turn off interrupts. ✓
- f. Access I/O device. ✓
- g. Issue a trap instruction.

The correct answers are: Set value of timer., Access I/O device., Issue a trap instruction., Switch from user to kernel mode., Modify entries in device-status table., Turn off interrupts.

Question 21

Correct

Mark 1.00 out of 1.00

Calculate the average waiting time using
FCFS scheduling
for the following workload
assuming that they arrive in this order during the first time unit:

Process Burst Time

P1	2
P2	6
P3	2
P4	3

Write only a number in the answer upto two decimal points.

Answer: 5.00



P2 waits for 2 units

P3 waits for 2+6 units

P4 waits for 2 + 6 +2 units of time

Total waiting = $2 + 2 + 6 + 2 + 6 + 2 = 20$ units

Average waiting time = $20/4 = 5$

The correct answer is: 5

Question 22

Partially correct

Mark 0.25 out of 0.50

Select all the correct statements related to implementation of a Hypervisor like VirtualBox

- a. All Traps in hardware are handled by HOST OS, but Host OS may hand it over to Guest OS if trap was from within guest OS
- b. Typically they run using CPU privilege level 1 or 2 (lower than kernel's but higher than applications)
- c. When an application runs a privileged instruction it traps into the actual hardware ✓
- d. The HOST OS determines whether a process (in Guest) or the guest OS itself was doing privileged instruction depending on the privilege level. ✓

The correct answers are: Typically they run using CPU privilege level 1 or 2 (lower than kernel's but higher than applications), When an application runs a privileged instruction it traps into the actual hardware, All Traps in hardware are handled by HOST OS, but Host OS may hand it over to Guest OS if trap was from within guest OS, The HOST OS determines whether a process (in Guest) or the guest OS itself was doing privileged instruction depending on the privilege level.

Question 23

Partially correct

Mark 0.75 out of 1.00

Match each suggested semaphore implementation (discussed in class)

with the problems that it faces

```
struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->s1 = 0;
}

wait(semaphore *s) {
    spinlock(&(s->s1));
    while(s->val <=0) {
        spinunlock(&(s->s1));
        spinlock(&(s->s1));
    }
    (s->val)--;
    spinunlock(&(s->s1));
}
```

too much spinning, bounded wait not guaranteed



```
struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->s1 = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->s1));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->s1));
}
```

blocks holding a spinlock



```
struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->s1 = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->s1));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->s1));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->s1));
}

signal(semaphore *s) {
    spinlock(*s->s1);
    (s->val)++;
    x = dequeue(s->s1) and enqueue(readyq, x);
    spinunlock(*s->s1);
}
```

not holding lock after unblock



```

struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0)
    ;
    (s->val)--;
    spinunlock(&(s->sl));
}

```

livelock



Your answer is partially correct.

You have correctly selected 3.

The correct answer is:

```

struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        spinunlock(&(s->sl));
        spinlock(&(s->sl));
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

```

→ too much spinning, bounded wait not guaranteed,

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

```

→ blocks holding a spinlock,

```

struct semaphore {
    int val;
    spinlock lk;
    list l;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

block(semaphore *s) {
    listappend(s->l, current);
    spinunlock(&(s->sl));
    schedule();
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0) {
        block(s);
    }
    (s->val)--;
    spinunlock(&(s->sl));
}

signal(semaphore *s) {
    spinlock(*(&s->sl));
    (s->val)++;
    x = dequeue(s->sl) and enqueue(readyq, x);
    spinunlock(*(&s->sl));
}

```

→ not holding lock after unblock,

```

struct semaphore {
    int val;
    spinlock lk;
};

sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}

wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <=0)
    ;
    (s->val)--;
    spinunlock(&(s->sl));
}

```

→ deadlock

Question 24

Correct

Mark 1.00 out of 1.00

In an ext2 file system, if the block size is 4KB and partition size is 2 GB, then the number of block groups will be:

Answer: 

$\text{size} * 1024 * 1024 / 4 \rightarrow \text{no of blocks}$

each group = $8 * 4 * 1024$ blocks = 32768 blocks

so $\text{size} * 1024 * 1024 / (4 * 32768)$ number of groups

The correct answer is: 16.00

Question 25

Correct

Mark 1.00 out of 1.00

Match the pair

Inverted Page table	Linear/Parallel search using page number in page table	✓
Hierarchical Paging	More memory access time per hierarchy	✓
Hashed page table	Linear search on collision done by OS (e.g. SPARC Solaris) typically	✓

Your answer is correct.

The correct answer is: Inverted Page table → Linear/Parallel search using page number in page table, Hierarchical Paging → More memory access time per hierarchy, Hashed page table → Linear search on collision done by OS (e.g. SPARC Solaris) typically

Question 26

Correct

Mark 1.00 out of 1.00

Write the possible contents of the file /tmp/xyz after this program.

In the answer if you want to mention any non-text character, then write \0 For example abc\0\0 means abc followed by any two non-text characters

```
int main(int argc, char *argv[]) {
    int fd1, fd2, n, i;
    char buf[128];

    fd1 = open("/tmp/xyz", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);
    write(fd1, "hello", 5);
    fd2 = open("/tmp/xyz", O_WRONLY, S_IRUSR|S_IWUSR);
    write(fd2, "bye", 3);
    close(fd1);
    close(fd2);
    return 0;
}
```

Answer: byelo ✓

The correct answer is: byelo

Question 27

Correct

Mark 1.00 out of 1.00

Select the correct statements about hard and soft links

Select one or more:

- a. Soft links increase the link count of the actual file inode
- b. Hard links enforce separation of filename from its metadata in on-disk data structures. ✓
- c. Hard links can span across partitions while soft links can't
- d. Deleting a soft link deletes only the actual file
- e. Soft links can span across partitions while hard links can't ✓
- f. Deleting a soft link deletes the link, not the actual file ✓
- g. Soft link shares the inode of actual file
- h. Deleting a hard link deletes the file, only if link count was 1 ✓
- i. Deleting a soft link deletes both the link and the actual file
- j. Deleting a hard link always deletes the file
- k. Hard links share the inode ✓
- l. Hard links increase the link count of the actual file inode ✓

Your answer is correct.

The correct answers are: Soft links can span across partitions while hard links can't, Hard links increase the link count of the actual file inode, Deleting a soft link deletes the link, not the actual file, Deleting a hard link deletes the file, only if link count was 1, Hard links share the inode, Hard links enforce separation of filename from its metadata in on-disk data structures.

Question 28

Partially correct

Mark 2.54 out of 3.00

Suppose you are required to implement the priority scheduling algorithm in xv6. Select all the options that correctly reflect the changes that are required to be done in code:

Needed/Optional	Not Needed	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Set the priority of the new process to the default(using inheritance) during fork()
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Add a system call (like nice()) that allows to set priority of a process
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change sched() to set the priority of the process.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Add a priority field to the struct proc
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change scheduler() to calculate the priority of the process using user specified value
<input checked="" type="radio"/>	<input checked="" type="radio"/>	The nice() system call needs to acquire the ptable.lock for setting the priority.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change swtch() to set the timer value for the process
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change yield() to re-set the timer value to zero for the process
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Optionally remove the default timer value of 10000000 in lapicinit()
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change exec() to add a priority to the process.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Set the priority of the process as per information in ELF file
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Insert code in scheduler() after if(p->state != RUNNABLE); continue; by a code that selects the highest priority process.
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Set a new value for the timer before you call swtch() in scheduler()

Set the priority of the new process to the default(using inheritance) during fork(): Needed/Optional

Add a system call (like nice()) that allows to set priority of a process: Needed/Optional

Change sched() to set the priority of the process.: Not Needed

Add a priority field to the struct proc: Needed/Optional

Change scheduler() to calculate the priority of the process using user specified value: Not Needed

The nice() system call needs to acquire the ptable.lock for setting the priority.: Needed/Optional

Change swtch() to set the timer value for the process: Not Needed

Change yield() to re-set the timer value to zero for the process: Not Needed

Optionally remove the default timer value of 10000000 in lapicinit(): Needed/Optional

Change exec() to add a priority to the process.: Not Needed

Set the priority of the process as per information in ELF file: Not Needed

Insert code in scheduler() after if(p->state != RUNNABLE); continue; by a code that selects the highest priority process.: Needed/Optional

Set a new value for the timer before you call swtch() in scheduler(): Needed/Optional

Question **29**

Correct

Mark 0.50 out of 0.50

Doing a lookup on the pathname /a/b/b/c/d for opening the file "d" requires reading no. of inodes. Assume that there are no hard/soft links on the path.

Write the answer as a number.

The correct answer is: 6

Question 30

Partially correct

Mark 1.86 out of 2.00

Given below is an incomplete code for reader-writer lock with preference for readers. Fill in the blanks to complete the code.

Note1:

In your answer, if an expression is to be written, then please separate all tokens by exactly one space in between. For example

i = i + 1 is correct while

i=i+1 is not correct or

i= i +1 is also not correct.

Note2:

Correct: a->b++ (no spaces here!)

Any other notation is incorrect.

Note3: The code of downgrade() and upgrade() has proportional to 4/7 marks.

Code of rwlock:

```
typedef struct rwlock {
```

```
    int nActive;  
    int nPendingReads;  
    int nPendingWrites;  
    spinlock_t
```

```
    sl
```

```
    ✓ ;
```

```
    condition canRead;  
    condition canWrite;
```

```
}rwlock;
```

```
void lockShared(rwlock *r) {
```

```
    spin_lock(&r->sl);
```

```
    r->nPendingReads++
```

```
    ✓ ;
```

```
    while(r->nActive < 0)  
        wait(
```

```
        &r->canRead
```

```
    ✓ , &r->sl);
```

```
    r->nActive++
```

```
    ✓ ;
```

```
    r->nPendingReads--;  
    spin_unlock(&r->sl);
```

```
}
```

```
void unlockShared(rwlock *r) {
```

```
    spin_lock(&r->sl);
```

```
    r->nActive--
```

```
    ✓ ;
```

```
    if(r->nActive == 0) {  
        spin_unlock(&r->sl);  
        do_signal(
```

```
        &r->canWrite
```

```
    ✓ );
```

```
    } else  
        spin_unlock(&r->sl);
```

```
}
```

```
void lockExclusive(rwlock *r) {
```

```
    spin_lock(&r->sl);  
    r->nPendingWrites++;
```

```
    while(r->nActive || r->nPendingReads)  
        wait(
```

```

    &r->canWrite
    ✓ ,
    &r->sl
    ✓ );
    r->nPendingWrites--;

    r->nActive = -1
    ✓ ;
    spin_unlock(&r->sl);
}
void unlockExclusive(rwlock *r) {
    boolean_t wakeReaders;
    int i;
    spin_lock(&r->sl);

    r->nActive = 0
    ✓ ;
    if(r->nPendingReads != 0) {
        for(i = 0; i <
            r->nPendingReads
        ✓ ; i++)
            do_signal(&r->canRead);
    }
    else
        do_signal(
    &r->canWrite
    ✓ );
    spin_unlock(&r->sl);
}

void downgrade(rwlock *r) {
    boolean_t wakeReaders;
    int i;
    spin_lock(&r->sl);
    r->nActive =
    1
    ✓ ;
    if(r->nPendingReads != 0)
        for(i = 0; i <
            nPendingReads
        ✗ ; i++)
            do_signal(
    &r->canRead
    ✓ );
    spin_unlock(&r->sl);
}

void upgrade(rwlock *r) {
    spin_lock(&r->sl);
    if(r->nActive == 1) {
        r->nActive =
        -1
    }
    ✓ ;
    } else {
        r->nPendingWrites++;
        r->nActive--;
        while(r->nActive !=
        0
    ✓ )
        wait(
    &r->canWrite

```

```
✓ , &r->sl);
```

```
r->nPendingWrites--
```

```
✓ ;  
    r->nActive =  
    -1  
✓ ;  
}  
    spin_unlock(&r->sl);  
}
```

Question 31

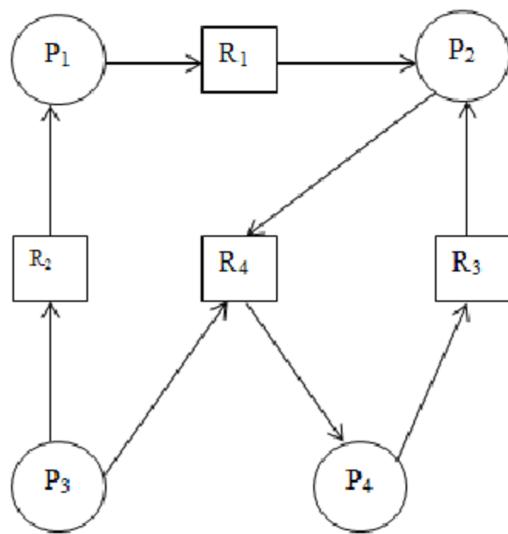
Correct

Mark 1.00 out of 1.00

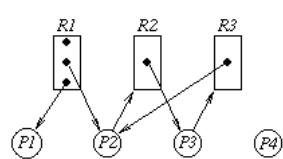
For each of the resource allocation diagram shown,
infer whether the graph contains at least one deadlock or not.

Yes

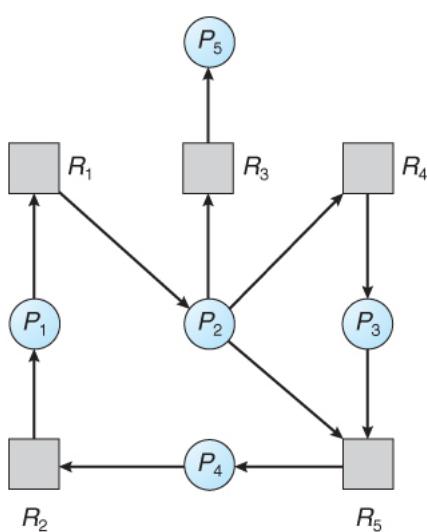
No



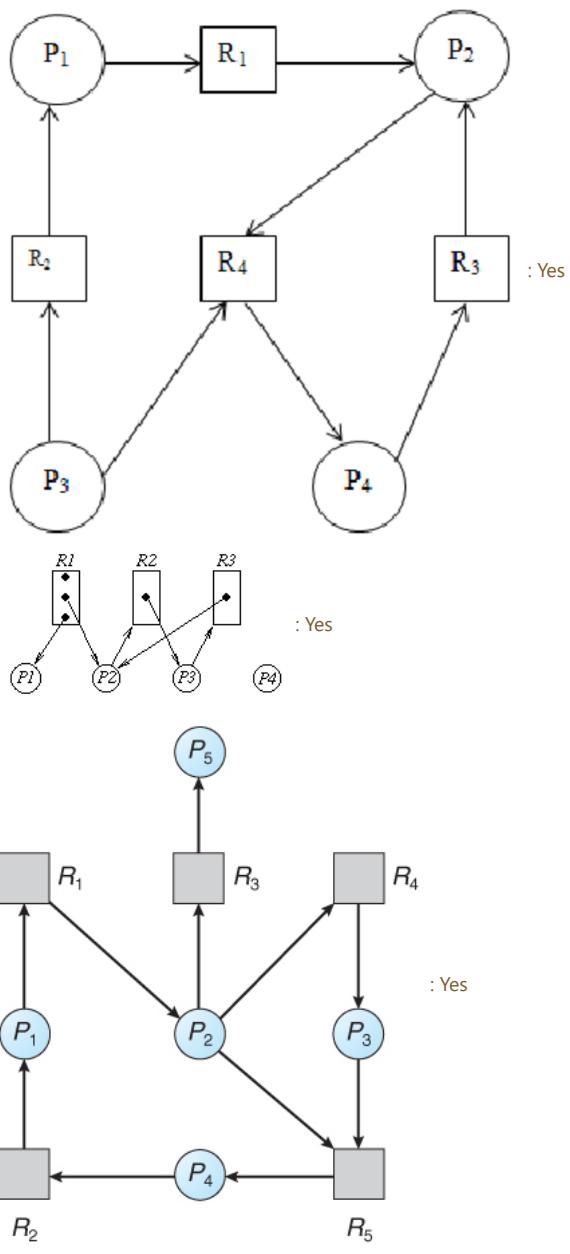
✓



✓



✓



Question 32

Correct

Mark 1.00 out of 1.00

Map each signal with its meaning

SIGPIPE	Broken Pipe	✓
SIGUSR1	User Defined Signal	✓
SIGSEGV	Invalid Memory Reference	✓
SIGALRM	Timer Signal from alarm()	✓
SIGCHLD	Child Stopped or Terminated	✓

The correct answer is: SIGPIPE → Broken Pipe, SIGUSR1 → User Defined Signal, SIGSEGV → Invalid Memory Reference, SIGALRM → Timer Signal from alarm(), SIGCHLD → Child Stopped or Terminated

Question 33

Correct

Mark 1.00 out of 1.00

Not a FOSS: Which of the following is/are not a FOSS operating system ?

- a. Minix
- b. Open Solaris
- c. xv6
- d. GNU/Linux
- e. Darwin (core kernel of Mac-OS)
- f. FreeBSD
- g. Windows
- h. BSD Unix



The correct answer is: Windows

Question 34

Correct

Mark 1.00 out of 1.00

Will this code work for a spinlock() operation? The intention here is to call compare-and-swap() only if the lock is not held (the if condition checks for the same).

```
void spinlock(int *lock) {  
{  
    while (true) {  
        if (*lock == 0) {  
            /* lock appears to be available */  
            if (!compare_and_swap(lock, 0, 1))  
                break  
        }  
    }  
}
```

- a. No, because this breaks the atomicity requirement of compare-and-test.
- b. Yes, because there is no race to update the lock variable
- c. No, because in the case of both processes succeeding in the "if" condition, both may end up acquiring the lock.
- d. Yes, because no matter in which order the if-check and compare-and-swap run in multiple processes, only one process will succeed in compare-and-swap() and others will keep looping in while-loop.



Your answer is correct.

The correct answer is: Yes, because no matter in which order the if-check and compare-and-swap run in multiple processes, only one process will succeed in compare-and-swap() and others will keep looping in while-loop.

Question 35

Partially correct

Mark 0.75 out of 1.00

Mark the statements as True/False with respect to Mobile systems and swapping.

True	False	
<input checked="" type="radio"/>	<input type="radio"/> 	Mobile systems generally do not support swapping
<input checked="" type="radio"/>	<input checked="" type="radio"/> 	Size of flash memory is one reason for restriction on use of swap on mobile devices.
<input checked="" type="radio"/>	<input type="radio"/> 	Limited number of write operations that flash memory can tolerate, is one reason for limitations of swapping on mobile devices.
<input checked="" type="radio"/>	<input type="radio"/> 	Poor throughput between main memory and flash memory is one reason for restriction on use of swap on mobile devices.

Mobile systems generally do not support swapping.: True

Size of flash memory is one reason for restriction on use of swap on mobile devices.: True

Limited number of write operations that flash memory can tolerate, is one reason for limitations of swapping on mobile devices.: True

Poor throughput between main memory and flash memory is one reason for restriction on use of swap on mobile devices.: True

[◀ Course Exit Feedback](#)

Jump to...